




Deterministic Concurrency

Edward Ashford Lee

Copyright ©2026
Edward Ashford Lee
All rights reserved

Work in Progress, Version 0.05
July 15, 2026

Contents

Preface	vii
Acknowledgements	xiii
Acronyms	xv
Notation	xvii
 I Background	 1
1 Determinism	3
1.1 Determinism as a Property of Models	4
1.2 What is Determinism?	13
1.3 Determinism vs. Predictability	22
1.4 Determinism in physics	26
1.5 Incompleteness of Determinism	41
1.6 Conclusion	50

2	Concurrency	53
2.1	Concurrent, Parallel, and Distributed	54
2.2	Consistency	57
2.3	The CAL Theorem	74
2.4	Conclusion	99
3	Nondeterministic Concurrency	101
3.1	Threads	101
3.2	Actors	101
3.3	Publish and Subscribe	101
3.4	Remote Procedure Calls	101
3.5	Addressing the Challenges	101
II	Untimed Models	103
4	Process Networks	105
4.1	Kahn Process Networks	106
4.2	Semantics of Process Networks	110
4.3	Execution of Process Networks	122
4.4	Convergence to the Denotational Semantics	134
4.5	Beyond Kahn-MacQueen	140
4.6	Dynamic Kahn Networks	144
4.7	Nondeterministic Extensions of Kahn Networks	146
4.8	Rendezvous	150
4.9	Conclusion	150
5	Dataflow	153
5.1	Firing and Firing Rules	155
5.2	Synchronous Dataflow	172
5.3	Boolean Dataflow	189

5.4	Modal Dataflow	196
5.5	Conclusion	202
III Timed Models		205
6	Synchronous Reactive Models	207
6.1	The Global Clock Model	208
6.2	Alternative Clock Models	239
6.3	Conclusion	252
7	Discrete-Event Systems	255
7.1	Discrete-Event Signals and Components	256
7.2	Metric Space Semantics	262
7.3	Generalized Ultrametric Semantics	278
7.4	Conclusion	293
8	Reactors	297
8.1	TBD	297
IV Appendices		299
A	Partially Ordered Sets	301
A.1	Sets	302
A.2	Relations and Functions	303
A.3	Sequences	307
A.4	Partial Orders	307
A.5	Functions on Posets	318
A.6	Fixed Points	322
A.7	Cardinality	323
A.8	Discrete Sets	328

A.9	Computability	329
A.10	Conclusion	330
B	Metric Spaces	331
B.1	Varieties of Metric Spaces	332
B.2	Metrics on Functions	334
B.3	Fixed-Point Theorems	340
B.4	Conclusion	344
	Bibliography	347
	Index	368

Preface

To many traditional computer science theorists, the title of this book may appear to be an oxymoron. It is not uncommon, particularly in the early literature, to equate concurrency with nondeterminism and to treat the two together as a profound difficulty encountered when trying to understand programs. For example, [Hennessy and Milner \(1980\)](#) say, “what exactly is meant by the behaviour of non-deterministic or concurrent programs is far from clear.” They go on:

Any satisfactory comparison of the behaviour of concurrent programs must take into account their intermediate states as they progress through a computation, because differing intermediate states can be exploited in larger programming environments to produce different overall behaviour (e.g. deadlock).

Indeed, nearly all concurrent programming models at the time (and to this day, see [Chapter 3](#)) have the property that they can exhibit observably different behavior given the same inputs.

But we must be careful. What do we mean by “observably different behavior”? Whether a program deadlocks or not is typically quite important. Deadlocking is an observable behavior. But so is timing, and for many programs, the exact time at which an output is produced is observable but not important. When considering a

simple, single-threaded imperative program, we are likely to declare the program to be deterministic. This means that given the same inputs, it always produces the same observable behavior, where what we mean by “observable behavior” is the printed outputs. We are unlikely to include in “observable behavior” the time at which each output is printed. A concurrent program may be similarly deterministic, as long as we restrict what we mean by “observable behavior.”

And when we say, “given the same inputs,” what do we mean by “inputs”? Is the time at which an input is presented part of the input? If inputs come from multiple sources, like the keyboard and the network, is the order in which they are seen part of the input? In a distributed system, inputs may arrive at physically separated nodes, in which case, it may not be possible to assign an order to their arrival. How we define “inputs” becomes important when deciding whether a particular system exhibits observably different behavior given the same inputs.

The first chapter of this book confronts the meaning of the word “determinism.” In this book, determinism is not a property of some physical system (a computer running a computer program, for example). Instead, it is a property of a model of that system. And the structure of the model matters. The same physical system may have multiple useful models, some deterministic and some nondeterministic. The model must include a clear definition of “inputs” and “behavior” before it becomes sensible to talk about determinism.

Any given physical realization of a system may be modeled in many different ways. Some of these models may be deterministic while some are nondeterministic. All of these models are approximations. By nature of being models, they omit some details of the physical realization. Hence, having a deterministic model for a system does not mean that the physical realization will always exhibit the same behavior given the same inputs. It just means that when the assumptions that make the model valid are met, the behavior, as defined by the model, will be the same whenever the inputs, as defined by the model, are the same. These are important caveats.

Deterministic concurrency should be thought of as a way to construct a clear definition of a “correct” or “desired” behavior for a system. It is a model of what we want. But we won’t always get what we want. The computer will most likely not behave like the model while it is being unplugged, for example.

Having a clear specification of a “correct” or “desired” behavior has many advantages. One of the key ones is that it becomes far easier to detect incorrect behaviors. In other words, deterministic concurrency makes it easier to detect and react to faults. Faults *will* occur. Murphy’s Law guarantees it. No amount of determinism in your models will prevent it. But being able to *recognize* faults is a prerequisite for handling them reasonably.

About the Cover

Are schooling fish deterministic? If you believe that everything that occurs in the physical world is a consequence of the preconditions, which I do not, then yes, schooling fish are deterministic. However, this is hardly a useful perspective. Their behavior is certainly not predictable. But they rarely collide with one another, they generally agree on whether they are being pursued by a predator, and they negotiate and then agree on a direction in which to go. They communicate and act more as a whole than as individual parts. And their coordinated behavior persists despite the inevitable loss of some of the components.

These are the hallmarks of the most interesting concurrent systems. Even the most carefully engineered such systems are rarely, if ever, predictable. Yet their concurrent parts communicate and coordinate to act as a whole. They exhibit observable properties reliably despite the vagaries in their context. And they tolerate, up to a point, failures of the individual components.

This book is not about some illusion of creating concurrent systems that behave like flawless clockwork. Instead, this book is about leveraging the power of deterministic models to make systems that are more reliable and more robust. It is the *models*, not the physical systems, that are deterministic.

For the fish, a reasonable deterministic model might be one where any individual fish can make a decision to turn, and, with reasonable time delays, using feedback control, the surrounding fish will react by turning as well. All the fish in the school are simultaneously making such decisions and reacting to the movements of the surrounding fish, making small adjustments to the course. The whole reacts with useful dynamics when presented with, for example, a predator.

Such a deterministic model could take the form of a set of complicated nonlinear partial differential equations. These will be hard to solve, hard to simulate, and

impossible to predict because of chaos. But these equations are nevertheless deterministic. And their determinism has value towards yielding understanding of the system as a whole and how its holistic behavior is affected by the behaviors of the parts.

A property of concurrent systems that is sometimes conflated with determinism is consistency. As we will see in Chapter 2, consistency is a subtle property with several definitions. Loosely, it refers to agreement about something across the concurrent parts of a system. For schooling fish, for example, the direction of travel may be a feature requiring agreement. But for schooling fish, such agreement is never perfect. Each individual is constantly adjusting its direction with input from its neighbors.

Often, the property that is sought is eventual consistency, which means that if inputs stop arriving, all parts of the system will eventually agree on something. Imagine schooling fish where all external stimulation stops. Will they all converge to a single linear direction of travel? Is this even a useful question?

Many modern techniques that achieve eventual consistency are nondeterministic. What is agreed upon is not determined by the inputs. The same inputs could lead to many different agreements. For some applications, this is acceptable, but it comes at a cost. For one, testing such a system is much more difficult. There is no single “known good” final value to be agreed upon, and hence, tests that compare outcomes against “known good” results are harder to design. It is often difficult to ascertain the totality of possible outcomes. Hence, such systems could lead to agreed-upon values that are unexpected, surprising, or undesirable.

Neither determinism nor consistency, by itself, is a valid objective. It is easy, for example, to design a deterministic and consistent distributed banking system by keeping all copies of all bank balances at zero. This is not useful. We almost certainly want the bank balances to be functions of the inputs, the deposits and withdrawals. But in a distributed system, the inputs are not just a set of numbers representing deposits and withdrawals. Instead, we have a dynamic system, where deposits and withdrawals occur in time at different locations. It is impossible, from first principles, to totally order such deposits and withdrawals, so how should we define the inputs? At best, it is a partially ordered set of deposits and withdrawals. But if the notion of “inputs” is poorly defined, then the partial order will be different for each node in the system. Consequently, given the same inputs, one node may reject a withdrawal due to insufficient balance while an-

other node accepts it, leading to divergent opinions about the resulting balance. A more practical design is more like the schooling fish, where small temporary discrepancies in the local representations of the balance are tolerated as long as the discrepancies are continually corrected over time. The fish do not diverge despite never perfectly agreeing on the direction of travel.

For many applications, stronger forms of consistency are required. For interactive systems that do not terminate, there is no point at which inputs stop arriving. Consequently, an eventually consistent system may never be consistent. It may not be acceptable to be forever in a state of disagreement. When a bank balance is reported to a customer, for example, all branches of the bank should agree on the balance. In Chapter 2, I offer a stronger notion of consistency that uses the notion of logical time. This form of consistency is quantifiable, in the sense that we can design systems that tolerate limited inconsistency.

Conversely, for many other applications, perfect consistency is not required, eventual or not. It is never really necessary for all fish in a school to be simultaneously traveling in exactly the same direction. If one fish turns, the surrounding fish will turn similarly within bounded time. This, of course, can give rise to interesting, even chaotic dynamics. All individual fish are simultaneously making decisions, and making course corrections that blend observations of the surrounding fish with local decisions. The goal becomes coherent behavior of the whole, not perfect agreement of the parts.

This book is not trying to enshrine determinism as a holy grail, an idealized and never perfectly achievable objective. It is instead trying to highlight the value of deterministic models. A big emphasis is on sophisticated mathematical treatments of a variety of deterministic concurrent models. I start with untimed models, specifically process networks and dataflow, then turn to timed models, specifically synchronous-reactive models and discrete-event systems. Until Chapter 8, the exposition is independent of any programming language. In Chapter 8, which might be viewed as the crescendo of the book, I show how the reactor model of computation realizes the best features of all four of the previously considered models. This is illustrated using the *Lingua Franca* coordination language, a free and open-source framework with surprising versatility. The most important principle is that *Lingua Franca* enables the construction of concurrent and distributed programs that are deterministic under clearly stated assumptions, and that it makes it possible to detect when the assumptions are violated and to treat these violations as faults.

Acknowledgements

I would like to thank my current and former students and postdocs for their contributions to my thinking on this subject. Over the years, I have learned a great deal from them. I would like particularly single out the ones who contributed directly to the formalisms presented in this book, including Shuvra S. Bhattacharyya, Joe Buck, Adam Cataldo, Stephen Edwards, Soonhoi Ha, Ben Lickly, Shaokai Lin, Xiaojun Liu, Marten Lohstroh, and Eleftherios Matsikoudis.

In addition, a number of my colleagues and friends significantly influenced my thinking on this subject, including most notably, Gerard Berry, Tom Henzinger, Dave Messerschmitt, Alberto Sangiovanni-Vincentelli, and Marjan Sirjani.

I would also like to thank Brian Ward and his AI tool Proofreader (<https://github.com/binarybison/proofreader> and Ward (2026)) for checking the proofs in this book. Proofreader found two serious errors that invalidated two of the proofs and offered suggestions for improvement in a number of other cases.

Acronyms

APG	acyclic precedence graph	185
API	application programming interface	71
BDF	boolean dataflow	192
CAL	consistency, availability, latency	60
CAP	consistency, availability, partitioning	59
CCS	calculus of communicating systems	150
CET	causally and eventually consistent with timestamps	67
CRC	cyclic redundancy check	25
CSDF	cyclo-static dataflow	196
DDF	dynamic dataflow	195
FIFO	first in first out	107
GLB	greatest lower bound	312
HDF	heterochronous dataflow	199
IDF	integer dataflow	192
KPN	Kahn process network	106
IP	internet protocol	8
LET	logical execution time	93
LEX	network penalty	79
LUB	least upper bound	311

NTP	network time protocol	67
OS	operating system	55
PASS	periodic admissible sequential schedule	182
PN	process network	106
PSDF	parameterized SDF	202
PTP	precision time protocol	71
SDF	synchronous dataflow	172
TCP	transmission control protocol	8
TSN	time sensitive network	71
UDP	user datagram protocol	12
UTC	coordinated universal time	67
WCET	worst-case execution time	86

Notation

$\mathbb{B} = \{0, 1\}$	binary digits	302
$\mathbb{T} = \{\text{false}, \text{true}\}$	truth values	302
$\mathbb{N} = \{0, 1, 2, \dots\}$	natural numbers	302
$\mathbb{Z} = \{\dots, -1, 0, 1, 2, \dots\}$	integers	302
$\mathbb{R}$	real numbers	302
$\mathbb{R}_+$	non-negative real numbers	302
$A \subseteq B$	subset	302
$\wp(A)$	powerset	302
2^A	powerset	302
$\emptyset$	empty set	302
$A \setminus B$	set subtraction	302
$A \times B$	cartesian product	303
$(a, b) \in A \times B$	tuple	303
A^0	singleton set	303
$f: A \rightarrow B$	function	303
$f: A \rightharpoonup B$	partial function	303
$g \circ f$	function composition	304
$f^n: A \rightarrow A$	function to a power	304
$(A \rightarrow B)$	set of all functions from A to B	304
B^A	set of all functions from A to B	304

$f^0(a)$	identity function	304
$\hat{f}: 2^A \rightarrow 2^B$	image function	304
$f \upharpoonright C$	restriction	282
π_I	projection	306
$\hat{\pi}_I$	lifted projection	307
$\omega = \{\emptyset, \{\emptyset\}, \{\emptyset, \{\emptyset\}\}, \dots\}$	von Neumann numbers	308
A^ω	infinite sequences	308
A^*	finite sequences	307
$A^\mathbb{N}$	infinite sequences	307
A^{**}	finite and infinite sequences	307
λ	empty sequence	307
Λ_m^n	empty function	156
ϕ	dataflow actor functional	157
$\leq$	partial order	309
$(A, \leq)$	poset	309
$\sqsubseteq$	prefix order	309
$\sqcap$	common prefix	282
$(A, <)$	strict poset	311
$\mathbb{Q}$	rational numbers	312
$\vee B$	join, least upper bound, LUB	312
$a \vee a'$	join, least upper bound, LUB	312
$\wedge B$	meet, greatest lower bound, GLB	312
$a \wedge a'$	meet, greatest lower bound, GLB	312
$\perp$	bottom element	310
$[A \rightarrow B]$	set of all continuous functions	320
$\bar{C}$	inconsistency	76
$\bar{A}$	unavailability	77
L	network latency	79
$\mathbb{G}$	set of tags	73
$\mathbb{I}$	time interval	70
$\mathbb{T}$	time values	70
$\mathcal{T}: \mathbb{G} \rightarrow \mathbb{T}$	physical time interpretation for a tag	73
$a.b$	concatenation	108
$\perp$	unknown	212
ϵ	absent	212

Part I

Background

Determinism

1.1	Determinism as a Property of Models	4
1.1.1	Engineering Models vs. Scientific Models	6
1.1.2	Determinism, Causation, and Randomness	7
1.1.3	The Practical Value of Determinism	8
1.1.4	The Practical Value of Nondeterminism	10
1.1.5	The Cost of Determinism	12
1.2	What is Determinism?	13
1.2.1	Behavior, State, and Observation	14
1.2.2	Input-Output vs. State-Trajectory Behavior	15
1.2.3	Time	19
1.2.4	Observers	20
1.3	Determinism vs. Predictability	22
1.3.1	Murphy's Law and Faults	24
1.4	Determinism in physics	26
1.4.1	Nondeterminism in Newtonian Physics	27
1.4.2	Metastable States	33
1.4.3	Relativity	34
1.4.4	Quantum Physics	37
1.5	Incompleteness of Determinism	41
1.5.1	Nondeterministic Collisions	42

1.5.2 Incompleteness	48
1.6 Conclusion	50

1.1 Determinism as a Property of Models

For most of my professional research career, I have sought more deterministic mechanisms for solving various engineering problems. My focus has always been on systems that combine the clean and neat world of computation with the messy and unpredictable physical world. Why the obsession with deterministic mechanisms? My wife, who is an expert in stochastic models, gives me a hard time about this obsession, observing, correctly, that deterministic models are just a special case. Why not, then, focus on the more general set of nondeterministic models?

Today, our society relies heavily on deterministic engineered systems. The balances in our bank accounts are a consequence of the inputs to the bank’s computing systems. The email received is the email that was sent. The files on our computers contain the data that we put there. Our cars start when we push the start button. None of these is perfect, of course. Files and communication can get corrupted and machinery can fail. But for many systems, we easily see trillions of bits of information before encountering an error and years of reliable operation before encountering a failure.

None of this reliability comes for free. A great deal of engineering effort has gone into making the behavior of these systems repeatable. The underlying physics, a sea of electrons sloshing around, for example, offers no such repeatability. We should be in awe of these systems, but we are not because we have gotten used to their dependability. In the early days of computing, vacuum tubes would fail frequently enough that programs had to be kept short so that they would complete before a failure occurred. Computer memories were nowhere near reliable enough to replace paper ledgers. And cars frequently would not start. But today, computers operate continuously for months and years, performing billions of operations per second without error. When errors do occur, a second layer kicks in, using error-correcting codes or redundant systems to self-correct. As a result, paper ledgers have vanished. And cars start.

There are many reasons for this dependability, but for the purposes of this book, I would like to focus on my obsession: the principle of determinism. I am not talking about the philosophical view that everything that happens is an inevitable consequence of pre-existing conditions, although that concept is related. I am instead talking about an engineering principle, where a model defines exactly one correct behavior for a system in response to its inputs. Like the philosophical concept of determinism, the engineering version has many fine points, subtleties, and consequences that are subject to debate. My purpose in this chapter is to address these fine points. I will start informally, but I will not shy away from subtleties.

The first issue is that I will treat determinism as a property of a model, not as a property of a physical artifact.

Example 1.1: Consider a logic gate, an AND gate for example. Is it deterministic? This is, of course, a trick question. If what we mean by “an AND gate” is a piece of silicon, then the answer to this question depends on whether the underlying physics of electrons sloshing through semiconductor atoms buffeted by thermal noise under the influence of electric fields is deterministic. I will consider later in the chapter the question of whether modern physics can answer that question, but suffice it to say, for now, that there is a different interpretation of the question that makes it much easier to answer.

If what we mean by “an AND gate” is a model, an abstraction, that given two boolean inputs, produces the output “true” if the two boolean inputs are “true” and otherwise produces the output “false,” then this AND gate is deterministic. It defines exactly one correct behavior for each pattern of inputs. This model, however, leaves a great deal unsaid. When must the inputs be provided? When is the output provided? How are “true” and “false” to be represented in the physical world? Who or what is the observer of the output? What can that observer perceive about the AND gate? Determinism does not require that *all* aspects of a system be prescribed. It only requires that what we construe as “behavior” be prescribed.

Implicitly, I have assumed that what we mean by “inputs” is a pair of boolean values from the set $\{true, false\}$ and what we mean by “behavior” is producing a boolean output chosen from the same set. This is a mathemat-

ical abstraction, not a physical artifact. If an observer is able to perceive, for example, the temperature of the AND gate, then this could be construed as part of its behavior, and the AND gate is no longer deterministic. Many possible temperatures are consistent with correct behavior of the AND gate. The abstraction also says nothing about the time at which these inputs are presented and the output is produced, so if the timing is observable, then again the AND gate is no longer deterministic.

Whether something is deterministic or not, therefore, depends on what can be observed about it and whether we consider these observations part of the “behavior.”

1.1.1 Engineering Models vs. Scientific Models

A mathematical abstraction is a model. What makes this model useful? In my book, *Plato and the Nerd* (Lee, 2017), I observe that the purpose of a **scientific model** is to emulate a physical system. For an **engineering model**, the purpose of the physical system is to emulate the model. The mathematical abstraction of the AND gate that I have given is a rather poor model of electrons sloshing in silicon, so it is not likely to be used as a scientific model. The model is valuable because we are able to construct silicon that behaves like the model, not because it accurately describes what electrons do as they move through silicon atoms. Hence, it is more useful as an engineering model than as a scientific model.

In this book, I will focus primarily on deterministic engineering models, and hence, whether the physical world is in fact deterministic is not nearly as important as whether we can coerce the physical world to behave like a deterministic model. The reality today is that humanity has figured out how to get silicon electronics to behave like deterministic models with astonishing fidelity.

Of course, no model is perfect, and no physical realization is perfect. The famous quote from George Box, “all models are wrong, but some are useful” (Box and Draper, 1987), applies to scientific models, where it is incumbent on the model to match the physical world. The mirror image, “all physical realizations are wrong, but some are useful” (Lee and Sirjani, 2018), applies to engineering models, where it is incumbent on the physical world to match the model.

In this book, determinism is a property of models, not of physical realizations. I will analyze why this property is so valuable, and why it is worth a great deal of engineering effort to build physical systems that emulate deterministic models. I will argue that using deterministic models is not at odds with dealing with uncertainty nor with building fault-tolerant systems. On the contrary, it makes these things easier. I will discuss the relationship between deterministic models and determinism in the physical world. And I will review results that show that determinism is incomplete in that sets of deterministic models that are rich enough to model the physical world do not contain their own limit points. From a practical perspective, this means that nondeterminism is unavoidable in a broad class of models.

1.1.2 Determinism, Causation, and Randomness

Determinism in philosophy is the principle that every action is a consequence of its preconditions and that fixed rules uniquely determine these consequences. The closely related principle of **causation** (or **causality**) is that the preconditions and the rules *cause* the consequences. Nondeterminism, under this principle, can be thought of as **uncaused action**. Many people find it difficult to accept that there can be uncaused action and will go to great lengths to find some cause when none is obvious, resorting eventually to God or some other supernatural cause when all else fails.

Causation is an explanation of *why* something happens. Instead of focusing on why, we can focus on *what* happens. In this view, it is the *uniqueness* of the consequences that is the core of the concept of determinism, not causation. There is only one possible consequence of each precondition. This uniqueness is the property I focus on in this book.

Nondeterminism is closely related to the concept of **randomness**, but there are important distinctions. In vernacular use, a random event is an unpredictable one. But even deterministic processes can be unpredictable, so under this view, randomness does not imply nondeterminism. Nondeterministic processes are always unpredictable, however, so nondeterminism does imply randomness in this vernacular understanding of the word.

There is a more technical interpretation to the word “random,” however, which seeks to quantify the *likelihood* of various possible outcomes with **probabilities**.

Nondeterminism, in contrast, is about *possibilities*, not probabilities. It says nothing about likelihoods.

This more technical interpretation of randomness leads to a deeper split with nondeterminism, making these concepts almost orthogonal. Consider what we mean by “likelihood” and “probability.” In the classical **frequentist** interpretation of these concepts, a probability is a prediction of the proportion of repeated experiments that will turn out some way. For example, what fraction of die tosses will produce snake eyes? But no repeat of any experiment begins with the same preconditions as the previous experiment, so this frequentist interpretation has no need for nondeterminism.

Some experiments *cannot* be repeated, and yet we consider them random and assign them probabilities. For example, what is the probability of a major earthquake in San Francisco in the next 30 years? This probability has no valid frequentist interpretation. In the **Bayesian** interpretation, a probability is a measure of how much we *know* about the outcomes, not a measure of how frequently they occur (see Chapter 11 of [Lee \(2017\)](#) for a discussion of the frequentists vs. the Bayesians). Once again, randomness has no need for nondeterminism. If we know little about some deterministic system, we can consider its outcomes to be random and assign them probabilities. Whether the underlying system is nondeterministic becomes irrelevant.

1.1.3 The Practical Value of Determinism

Let me begin by pointing out that our engineering toolkits are full of extremely useful deterministic models. Differential equations are deterministic (with some important exceptions that I will discuss below). Most computer programming languages, if we exclude their mechanisms for concurrency, are deterministic. Synchronous digital logic, the most widely used electronic circuit design paradigm, is deterministic (and concurrent). Instruction set architectures are mostly deterministic. **TCP/IP (transmission control protocol/Internet protocol)**, the central protocol in the Internet, is deterministic in the sense that a stream of bits in yields the same stream of bits out, even though packets may be dropped (even deliberately to shape traffic) or arrive out of order. Computer memory is deterministic, even when it is built on unreliable components, as demonstrated for example by the RAID project (redundant arrays of inexpensive disks) ([Patterson et al., 1988](#)).

Today, it is even possible to build deterministic cloud services on top of farms of unreliable servers.

All of the above are deterministic *models*. The underlying physical realization is almost certainly not deterministic, and considerable effort and cost has gone into building physical systems that are *sufficiently* faithful to these deterministic models. Perfect faithfulness is not achievable (“all physical realizations are wrong...”). But we would not be putting in all that effort if determinism were not quite valuable.

There are many reasons why determinism is valuable. Here are a few:

1. **Repeatability.** Every engineer tests a design while developing it. Such testing is valuable only under the assumption that if a system works once, it will work again in largely the same way. The behavior is repeatable given the same inputs. This principle is systematized in test-driven design, where collections of regression tests have well-defined correct behaviors. When a change is made to a design, the regression tests will reveal whether the change has resulted in some unexpected behavior.
2. **Consistency.** When two agents come to a conclusion, it is often valuable that the conclusions be the same. Your bank and you usually agree on your bank balance. You both keep track of that balance using deterministic computations. Given the same input information, both will produce the same results. Consistency is discussed in much more detail in Chapter 2.
3. **Predictability.** If behaviors in response to certain inputs are predictable, then they do not need to be discovered through testing or by surprise after deployment. As I will explain below, determinism does not assure predictability, but nondeterminism assures unpredictability. Some deterministic models are predictable.
4. **Fault detection.** A deterministic model gives an unambiguous definition of “correct behavior.” This enables fault detection. Any behavior that deviates from the correct behavior is faulty behavior, meaning that some assumption has not been satisfied. An AND gate, for example, will not behave in conformance with the model if the temperature is too high. A CRC check on a value read from memory (something computed using a deterministic algorithm) reveals whether the memory has experienced a

fault that caused a bit to flip. Both of these faults would be harder to detect without a deterministic model defining correct behavior.

5. **Simplicity.** In a deterministic design, one input implies one behavior. In a nondeterministic design, it is easy to get an exponentially growing number of allowed behaviors. This makes comprehensive testing much more challenging. It can also compromise the ability to analyze a model using, for example, formal methods.
6. **Unsurprising behavior.** Often, we want engineered systems to be boring. They should not surprise us with unexpected behaviors. When a computer programmer gets unexpected behaviors from a multithreaded program, for example, it can be extremely disruptive, costly, and difficult to fix (Lee, 2006b).
7. **Composability.** When building large systems out of smaller components, having a clear understanding of the possible behaviors of those smaller components becomes essential. Deterministic models of those components makes this far easier.

When engineers are forced to move beyond these deterministic mechanisms, building correct designs becomes far more difficult. Unfortunately, many of today's most popular programming frameworks supporting parallel and distributed computing make it quite difficult to achieve deterministic behaviors (see Chapter 3).

1.1.4 The Practical Value of Nondeterminism

None of what I have said implies that nondeterminism has no value. Nondeterministic models can also be valuable. Here are a few reasons:

1. **Abstraction.** A nondeterministic model may provide a much simpler abstraction of a deterministic model (Cousot and Cousot, 1977). Note that this does not contradict Property 5 above. The nondeterministic *model* is simpler, but if it is a sound abstraction, the number of *behaviors* allowed by the model cannot possibly be smaller than those of the model it abstracts.

A smaller model may be easier to understand. It may also be easier to formally analyze, as long as the formal analysis does not require exhaustively exploring all possible behaviors.

2. **Uncertainty.** Nondeterministic models are useful as *scientific* models of systems where our knowledge of their behavior is incomplete. For example, modeling a human operator of a car is probably not a reasonable task for deterministic modeling. Such a model of a human operator, however, is almost certainly a scientific model, not an engineering model.
3. **Deferred design decisions.** For *engineering* models, nondeterminism can be a useful way to capture deferred design decisions. Deferred design decisions represent a different form of uncertainty. It is uncertainty about the model, not uncertainty about what is being modeled.
4. **Security.** Many techniques for securing software systems rely on good random number generators. Pseudo-random number generators are deterministic, a weakness that opens a vulnerability. Seeding a pseudo-random number generator with the result of a nondeterministic process can help.
5. **Don't care.** A model may have many acceptable behaviors in response to a given input. Serverless architectures in the Internet, for example, are useful because computations in response to inputs can be mapped to any available server. It is irrelevant to correctness in what order and where these computations are performed.
6. **Surprising behavior.** We don't always want engineered systems to be boring. An automated musical accompanist, for example, might be enriched by unpredictable behavior. Any artificial intelligence attempting to exhibit human-like behavior will also need to at least appear to be nondeterministic (see my book (Lee, 2020, Chapter 10) for a discussion of the connection between creativity and nondeterminism).

Notice that valid uses of determinism and nondeterminism are quite different from one another, which suggests that engineers should consciously choose between them. This is rarely easy. For example, when writing multithreaded or distributed software, almost all available languages, frameworks, and middleware are nondeterministic by default. Achieving determinism is left to the designer and is sometimes impossibly difficult (Lee, 2006b; Lohstroh and Lee, 2019).

Probabilistic, stochastic, or random models (I will treat these words as synonyms) are also useful, but as I pointed out in Section 1.1.2, the concept represented by these words is largely orthogonal to nondeterminism. In computer security, for example, random numbers are usually generated using a deterministic algorithm that ensures a desirable empirical probability distribution on observations. A non-deterministic source of random numbers may, in fact, be hopelessly inadequate for computer security without some characterization of its probability distribution. In the Bayesian interpretation of probability, it is irrelevant whether the thing being observed is deterministic. What is relevant is whether the observation can be anticipated given prior knowledge.

1.1.5 The Cost of Determinism

Building physical systems that behave like deterministic models usually comes at a price. A synchronous digital circuit, for example, has to be clocked slowly enough to provide comfortable margins to accommodate delay variability due to manufacturing tolerances and temperature. These margins are a cost in performance. Some power-users of computers have discovered that they can often “overclock” their CPU without introducing too many errors. This may be acceptable for gaming, for example, but it would probably not be wise for a bank to overclock their CPUs.

For distributed systems, determinism comes at the cost of latency (see Chapter 2). For applications where latency is a key performance metric, nondeterministic solutions may be a better choice. For example, some distributed applications choose to use **UDP (user datagram protocol)** rather than TCP for network communication. The UDP protocol is simpler and faster than TCP, but it sacrifices the guarantee of eventual in-order delivery of messages. If occasional packet losses are acceptable for a particular application, then this may be a reasonable choice.

As with all engineering, there are tradeoffs. To evaluate these tradeoffs, it is helpful to have a deeper understanding of what determinism really is. We look at that next.

1.2 What is Determinism?

John Earman, in his *Primer on Determinism*, states that “determinism is a doctrine about the nature of the world” and concludes that “a real understanding of determinism cannot be achieved without simultaneously constructing a comprehensive philosophy of science” (Earman, 1986, p. 21). If instead of a “doctrine about the nature of the world” we view determinism as a property of models, then no such philosophy is needed. We can focus instead on the usefulness of the concept of determinism.

As a property of models, determinism is easy to define:

Definition 1.1. A model is **deterministic** if given all the **inputs** that are provided to the model, the model defines exactly one possible **behavior**.

In other words, a model is deterministic if it is not possible for it to react in two or more ways to the same inputs.² Only one reaction is possible in the model. More precisely, only one reaction is *correct*; any other reaction is not one given by the model. In this definition, I have emphasized words that must be defined within the modeling paradigm to complete the definition, specifically, “inputs” and “behavior.”

Example 1.2: If the behavior of a particle is its position $x(t)$ in a Euclidean space as a function of time t , where both time and space are continuums, and if the input $F(t)$ is a force applied to the particle with mass m at each instant t , then Newton’s second law,

$$F(t) = m \frac{d^2}{dt^2} x(t) \quad (1.1)$$

is a deterministic model (mostly, see Section 1.4).

The *same particle*, however, may have an equally valid nondeterministic model. For the same definitions of behavior and input, the statement, “The particle accelerates in the direction of the net force,” provides a nondeterministic model. It admits many more behaviors than the deterministic

²For a nice formalization of this concept, see Edwards (2018).

model. This latter model is a sound abstraction because the behavior of the deterministic model is among the behaviors of the nondeterministic one.

When considering a *physical* particle, the thing-in-itself rather than its model,³ then Earman's hesitation comes to the foreground. A definitive answer to whether the actual particle is deterministic may never be possible. Any discussion of the determinism of a particle is necessarily a discussion of some model, not of the thing-in-itself, even if that fact goes unsaid. Some models of the particle are deterministic and some are not. I will address the physics of determinism in Section 1.4. But first, let us address some of the subtleties with models before we complicate the picture with the thing-in-itself.

1.2.1 Behavior, State, and Observation

When determining whether a model is deterministic, we need to define "behavior." Many models tie behavior to the notion of "state." Newtonian physics, for example, does this by defining a time continuum and modeling the state of a system as positions and momentums at a shared "instant" in time. "Behavior" can then be defined as the evolution of this state in the time continuum in response to the inputs, which are forces. Modern physics complicates this simple picture (see Section 1.4), but the Newtonian models nevertheless remain useful.

In their classic book on automata theory, Hopcroft and Ullman define "state" as follows:

Definition 1.2. *The **state** of the system summarizes the information concerning past inputs that is needed to determine the behavior of the system on subsequent inputs. (Hopcroft and Ullman, 1979, p. 13)*

The state is all the information about its past that can affect its future behavior. In other words, the state of a system is the information about the past such that any additional information tells us nothing about its future behavior. This definition requires defining "behavior," but it also requires a notion of "past" and "future" as well as the boundary between these, the "present." This notion

³The philosopher Immanuel Kant made the distinction between the world as-it-is, what he called the thing-in-itself (*das Ding an sich* in German), and the phenomenal world, or the world as it appears to us.

turns out to be problematic for very fundamental reasons. It is problematic in modern physics, but also in practical realizations of parallel, distributed, and cyber-physical systems (Lee, 2008). These systems have no well-defined “present” separating past and future. The notion of state, therefore, useful as it is, sits on shaky foundations.

In computing, automata theory is built around the notion of state (or equivalent notions working with sequences of symbols). Variants of automata theory use either a Newtonian time continuum or a discrete, countable model of time, both of which provide a “present” that separates past from future. Automata theories share with Newtonian physics the idea that time advances uniformly throughout the system and that there is a shared notion of an “instant” of time, a “present,” at which the system is in some state.

In automata theory, unlike Newtonian physics, even if time is a continuum, the *state* evolves in discrete jumps. At an instant, the state of the system changes from some value s to some other value s' . Again, “behavior” can be defined as the evolution of this state in time. When the system is distributed or concurrent, where it consists of multiple interacting components, possibly spread out over space, nondeterminism proves to be a useful way to model the uncertainty about the order in which state changes occur in the distinct components.

1.2.2 Input-Output vs. State-Trajectory Behavior

An alternative way to define behavior is to introduce the notion of **outputs**, data or symbols that some observer interprets as the behavior of the system in response to some input. Instead of state-trajectory behavior, we have input-output behavior. The relationship between input-output behavior and state-trajectory behavior is fascinating, subtle, and complex. The essential question is what can an observer observe? The possible observations are the “output” of the system. But this is true even if we define “behavior” to be a state trajectory. In that case, we have implicitly defined an observer that can observe the state of the system at an instant in time. This definition will require us to define “an instant in time,” which, as we will see, becomes difficult for a distributed system.

Consider automata theory, where a state transition system may be endowed with an explicit notion of inputs and outputs. Inputs trigger state transitions and outputs result from state transitions. A transition system evolves as a *sequence*




Figure 1.1: Three component system that is either deterministic or not depending on how you define the observer.

of state transitions, and hence, both the inputs and the outputs are sequences of symbols from some alphabet. A sequence of symbols forms a **sentence**, and the set of all sentences that are possible defines a **language**. In an input-output interpretation of behavior, a deterministic model is one for which, given any input sentence, there is only one possible output sentence.

The fact that inputs and outputs are defined in automata theory to be *sequences* of symbols is important. An observer only ever sees sequences of symbols. These sequences cause no end of difficulties with modeling concurrent systems, where the order in which symbols occur may not be relevant or even well defined. Starting in the 1970s, Robin Milner and Tony Hoare pioneered methods for formally modeling such systems (see Chapter 14 of Winskel (1993) for a nice summary of these results). I will illustrate these difficulties with some key observations due to Milner.

In 1980, Milner published a series of lecture notes giving an elegant formalism that he called a Calculus of Communicating Systems (CCS) (Milner, 1980). This formalism made it abundantly clear that looking only at the sentences of inputs and outputs is insufficient.

Example 1.3: For a simple illustration of this, consider a distributed system consisting of three components C_1 , C_2 , and C_3 running on three computers, illustrated in Figure 1.1. C_1 sends a message to each of the other two to initiate the computation. In response, the other two perform some deterministic computation and output a sequence of results (a sentence).

Suppose C_2 always outputs AB and C_3 always outputs CD . Is the overall system deterministic?

As described, this system has no input, so, to be deterministic it should have exactly one “behavior.” If by “behavior” we mean a single output sentence, as Milner does, then we are forced to combine the output sentences produced by C_2 and C_3 . How should we combine them? If we define the “observer” to be some entity that simply observes the symbols A, B, C , and D as they are produced, then that observer will see some arbitrary interleaving of the two sentences AB and CD . There are six such interleavings, as shown at the upper right in the figure. Hence, the model is nondeterministic.

Arguably, the nondeterminism that arises here is a side effect of our insistence that an observer can only see a single output sequence. Indeed, Milner found this sort of nondeterminism unsatisfying and introduced a notion that he called **confluence**. In his formalism, as long as C_2 and C_3 cannot interfere with each other’s ability to produce their output, this system is deemed confluent, presented as a useful replacement for the concept of determinism.

There is another solution, however, which does not require replacing the notion of determinism. Instead, we can change what we mean by behavior by changing what an observer sees. If, instead of a single output sequence, the behavior is defined to be a *pair* of output sequences, then the model immediately becomes deterministic, as shown at the lower right in the figure. The one and only behavior is (AB, CD) , a pair of sentences, each containing a sequence of two symbols.

Redefining behavior may seem like a sleight of hand, a trick. It is not. *Any* definition of behavior depends on a notion of an **observer**, and fundamentally, for any system, different observers see different things. Let me make this crystal clear with a trivial example.

Example 1.4: Consider the following C program:

```
1 int main(int argc, char* argv[]) {
2     printf("Hello World.\n");
3 }
```

Is this program deterministic? If the observer is a human sitting at a computer screen, then for this program to be deterministic, that observer should

be able to see only exactly one possible observation from running this program. Is that the case? What color will the characters “Hello World” be rendered in? How long will it take before the observer sees “Hello World”? Neither of these observable properties are specified by the program, so, if we include these properties in the notion of behavior, then the program is nondeterministic. Almost certainly, however, this is not what we intended. If we carefully define “observer” and restrict that observer to observing the sequence of symbols produced by the program, then the program becomes deterministic.

Returning to our example with C_1 , C_2 , and C_3 , a different notion of deterministic programs was introduced by Gilles Kahn in 1974, a class of models that are now called Kahn process networks (KPN) (Kahn, 1974). These are closely related to dataflow models (Lee and Parks, 1995; Lee and Matsikoudis, 2009), which share with them a different notion of observer from Milner’s that leads to different conclusions about determinism. A KPN is a network of processes that send messages to each other along defined channels, where each channel is assumed to preserve the order of the messages and deliver them reliably (like TCP). Kahn gave an elegant construction using the mathematics of partial orders to give conditions on the processes such that the overall behavior of the program is deterministic (see Chapter 4 and Kahn (1974)). Kahn and MacQueen later showed that a simple constraint on the processes in the network, blocking reads, is sufficient to ensure determinism (Kahn and MacQueen, 1977).

The definition of behavior here, however, is not the sentences of Milner. Kahn defined behavior to be a collection of (possibly infinite) sequences of messages, each recording the messages traversed on one channel. “Behavior” in a KPN, therefore, is a tuple of possibly infinite sequences. Under the KPN model, C_1 , C_2 , and C_3 together form a deterministic system. No notion of confluence is needed.

Kahn networks also do not require any notion of state. Each process can be usefully modeled as a state machine, but there is no need to ever talk about a global state, some combination of the states of all the processes at some instant in time. The approach Kahn took towards defining the semantics of programs without appealing to a notion of state can be generalized to many other kinds of concurrent systems (Lee, 2006a), as done for example in the tagged signal model (Lee and Sangiovanni-Vincentelli, 1998). This is explained in Chapter 4.

An issue with Kahn networks is that coordinated decision making in a distributed system can become difficult (Lohstroh and Lee, 2019). There is no notion of the state of the system, so writing code like “if the state of the system is X do Y ” becomes challenging.

A number of alternative concurrency models have emerged that preserve determinism and reintroduce a semantic notion of state without having to appeal to confluence. One such alternative concurrency model is embodied in synchronous-reactive languages (Benveniste and Berry, 1991), explained in Chapter 6. These languages introduce the notion of a global “clock” that ticks discretely, much like that found in synchronous digital logic design. At each tick of this now conceptual (rather than physical) clock, many computations are performed, possibly in parallel, until the entire system settles to a well-defined state. With certain constraints on the component computations and on their scheduling (Edwards and Lee, 2003a), the resulting state is a unique function of the inputs. The resulting model can be conceptualized as lying somewhere between Milner’s transition systems and Kahn’s asynchronous process networks. During the computation at a tick, the behavior is more like a Kahn network, where data precedences constrain the order in which things happen, but no notion of a single global state trajectory is needed and computations can proceed in parallel and asynchronously. At the conclusion of a tick, things settle to a well-defined state, enabling a higher-level state-trajectory model that treats all the computation at a tick as a single atomic state transition. These models are deterministic as long as the observer is constrained to observe the state only at the conclusion of each tick. The execution of synchronous-reactive languages may be thought of as “punctuated chaos,” where periods of chaotic, parallel, asynchronous computation are marked by isolated points of stable, well-defined state.

1.2.3 Time

Although synchronous-reactive languages have “clocks,” they do not really have a notion of time. For cyber-physical systems (CPSs), which combine computation and networking with physical components (Lee, 2008), some notion of time becomes essential (Lee, 2009). It is not necessary (nor is it physically possible) to insist on the Newtonian notion, where time is a continuum with well-defined instants t that are shared by all components in the system. Instead, leveraging the punctuated chaos of synchronous-reactive languages, we can assign a semantic measure

of time elapsed between ticks ([von Hanxleden et al., 2017](#); [Schulz-Rosengarten et al., 2018](#)). This makes it possible for models to combine Newtonian models of physical components with computational models, thereby offering a rigorous approach to CPS design that does not sacrifice determinism ([Cremona et al., 2017](#)).

We can make an even bigger commitment to a notion of time by explicitly time-stamping events, as done in discrete-event systems (see Chapter 8). In such models, the concept of a global “tick” is replaced by timestamped messages with the constraint that every component processes messages in timestamp order. With an additional constraint that messages with identical timestamps be processed in a well-defined order, the model becomes deterministic. This is the principle behind the recently introduced reactor model ([Lohstroh et al., 2019](#)) as realized in the Lingua Franca language ([Lohstroh et al., 2020](#)). Semantically, such discrete-event models can be viewed as a generalization of synchronous-reactive models ([Lee and Zheng, 2007](#)) and have even been called “sparse synchronous” models by [Edwards and Hui \(2020\)](#). Reactors are described in Chapter 8.

1.2.4 Observers

Another major issue that emerges from Milner’s calculus is the tension between state transitions and output sentences. Which are observed? It is easy to construct an automaton that makes nondeterministic state transitions and yet produces a deterministic output in response to inputs. Should this automaton be deemed deterministic or not? This tension has prompted some researchers to attempt to codify this distinction, sometimes using two distinct words, “determinate” and “deterministic”; see for example [von Hanxleden et al. \(2014\)](#). The word “determinate” is meant to capture the idea that observable outputs are uniquely defined by the inputs, whereas the word “deterministic” attempts to capture the idea that there is a unique way in which the outputs are determined. Using this distinction, our constructed automaton (the one that makes nondeterministic state transitions and yet produces a deterministic output) is determinate but not deterministic.⁴

⁴A similar distinction is given by [Wisniewski et al. \(2020\)](#), who define “strong” and “weak” determinism in terms of Petri nets augmented with inputs and outputs. They make a distinction between “stable markings” (ones where no transitions are enabled, given the inputs) and “unstable markings” (which can be viewed as transitory markings towards stable markings). A “weakly deterministic” Petri net is one where for each possible stable marking and input, there is exactly one successor stable marking, and a “strongly deterministic” Petri net is one where for any marking

However, this distinction is specious. An automaton is a model, not an implementation. An implementation might be realized by a computer program, in which case, the automaton is a model of the behavior of that program. The program itself is a model of the computations to be performed by one or more instruction set architectures (ISAs) communicating over some network fabric using some protocols. The ISA and the network protocols are themselves models of a physical system with electrons sloshing around. At which level should we determine whether the system is deterministic? How many words for determinism will we need to cover all these levels of models? If we are clear about what we mean by an observer, then no such distinction is needed and one word is sufficient. Determinism becomes a property of the combination of the model and the observer.

Another context in which this tension comes to the foreground is with Alonzo Church's lambda calculus ([Church, 1932](#)), a model of computation in which "lambda expressions" are subjected to syntactic rewriting following a set of rules. The Church-Rosser theorem shows if such an expression can be reduced to the point where none of the rewriting rules can be applied any more, then the same final expression results regardless of the order in which the rewriting rules are applied ([Church and Rosser, 1936](#)). This model of computation can be construed as nondeterministic if the observer can see the intermediate expressions or as deterministic if only the final expression is visible.

The notion of an observer, it turns out, has its own subtleties. Should an observer be passive and objective, or can the observer interact with the system? This distinction turns out to be important. In computer security, for example, any discussion of the security of the system requires a threat model that is explicit about the capabilities of an attacker. It is important whether the attacker can interact with the system or is restricted to passively observing it.

Milner's own definition of determinism exposes such subtleties in the concept of an observer. Milner's definition of determinism depends on a relation between automata that Milner called **bisimulation**. It is beyond the scope of this chapter to explain bisimulation, but suffice it to say that this concept depends on an observer that is not just passive and objective, but rather can interact with the system being

and input, there is exactly one successor marking. Another similar distinction is given by [Khomenko et al. \(2008\)](#), who define "output-determinacy" as a relaxation of determinacy.

observed.⁵ In 1980, David Park found a gap in Milner’s prior and simpler notion of **simulation**, in which a passive and objective observer automaton emulates the behavior of an observed automaton (Park, 1980). Park noticed that even if two automata simulate each other, they can exhibit significant differences in behavior. These differences are not observable by any passive, objective observer, but if the observer can interact with the observed automaton, providing inputs that depend on its observations, then the differences become visible. Park’s observation led Milner to develop the notion of bisimulation (and the closely related notion of **observational equivalence**), an interactive form of simulation that ensures that two automata are indistinguishable even through interaction. He then based his notion of determinism on bisimulation (Milner, 1989).⁶

1.3 Determinism vs. Predictability

Determinism does not imply predictability. For a model to be predictable, we must be able to anticipate its behavior by examining the model rather than by just watching what it does. Once again, we must be careful to define behavior. No computer program is predictable, for example, if behavior includes generating heat. The computer program alone is not sufficient to anticipate how much heat will be generated by executing the program.

Turing machines are deterministic models. The “input” to a Turing machine is a binary bit sequence, which Alan Turing described as the initial sequence of marks on a tape. The “behavior” can be defined to be a final bit sequence, the final sequence of marks on the tape when the machine halts, or a special result, often called **bottom** and written with the symbol $\perp$, that indicates failure to halt. A Turing machine, therefore, is a model whose operation computes a function that maps an input bit sequence to either an output bit sequence or $\perp$.

A Turing machine is also deterministic to an observer that can observe the sequence of operations that lead to this final output, but this stronger form of determinism is less important to the notion of computation. In fact, one of the

⁵See my book (Lee, 2020, Chapter 12) for an in depth discussion of the distinction between observation and interaction. That chapter includes a gentle introduction to the concept of bisimulation.

⁶Sangiorgi (2009) gives a nice overview of the historical development of this idea. He notes that essentially the same concept of bisimulation had also been developed in the fields of philosophical logic and set theory.

most interesting things about the foundational theory of computation is that many different mechanisms, including some nondeterministic ones, can be used to compute the same set of functions. A human using pencil and paper and following well-defined rules, given enough paper and time, can compute the same set of functions. Your laptop, which uses low level operations that are quite different from those of a Turing machine and involve no tape, given enough time and memory, can also compute the same set of functions. Church's lambda calculus, which has expressiveness equivalent to Turing machines, is deterministic in the weaker sense, where "behavior" is the final irreducible expression, but its mechanisms for finding that expression need not be deterministic. Reduction rules may be applied in any order. The low-level mechanisms are not as important as the function that is computed, so we will stick with the definition of "behavior" that restricts the observer to observe the final binary result.

The question now becomes, is the deterministic behavior of a Turing machine predictable? Alan Turing showed that it is not. He showed that there is no mechanism, no systematic procedure that can predict, for all Turing machines, whether an execution will halt for a particular input (Turing, 1936). Similarly, there is no such procedure for determining whether an execution uses a finite amount of memory, which translates to a finite amount of tape in the classic Turing machine. This result carries over to all the similarly expressive mechanisms (so called **Turing complete** mechanisms), including Church's lambda calculus, a human with paper and pencil, and your laptop. This result decisively decouples determinism from predictability. Determinism does not imply predictability.

Turing machines and lambda calculus underly modern imperative and functional programming languages, respectively. Programmers, therefore, face the possibility that they will not be able to predict whether their programs will terminate. Usually, however, a programmer needs assurances that the program will terminate (e.g. a program that converts a text file into a PDF file) or will not terminate (e.g. a web server or an operating system). Fortunately, most programs are predictable in this sense, even if, in theory, it is impossible for all programs to be predictable in this sense.

Many computer programs are unpredictable in a more informal sense. In this more informal sense, the question is, by examining the program but not executing it, can we anticipate its behavior? Some programs are designed to be unpredictable in exactly this sense. Pseudo-random number generators, for example, fall in this category. You cannot tell from looking at the program what numbers it will

generate. You must execute the program instead. Machine-learning algorithms turn out also to be unpredictable in this sense. Cellular automata, a class of deterministic computational machines that are also unpredictable in this sense, are capable of surprising and complex behaviors, so complex as to prompt some thinkers to conclude that they underly all of physics (Wolfram, 2002).

Unpredictable deterministic models also arise in Newtonian physics. Nonlinear differential equation models, such as those modeling the thermodynamics of weather, exhibit behaviors that are so unpredictable that they are called “chaotic” (Lorenz, 1963). A chaotic model is one where arbitrarily small perturbations in the inputs or the initial conditions have arbitrarily large effects in the future.

Deterministic chaotic systems are, *not even in principle*, predictable. Lewis and MacGregor (2006) propose a thought experiment involving two spheres colliding with each other within a contained space. They calculate the precision with which the initial conditions must be known in order to predict the behavior after a certain number of collisions. They then assume that the initial positions of the spheres are to be determined optically and derive the wavelength of light that will achieve the required precision. They then show that a single photon of such light would have “more energy than is currently posited for the entire universe in order to resolve the initial state of the system with precision sufficient to predict its behavior after just 35 collisions” (Lewis and MacGregor, 2006, p. 10-11).

The “system” that Lewis and MacGregor analyze is, however, a model, not a physical system. The question we should be asking, therefore, is whether the *model* is predictable. The model, given by Newton’s laws with discrete collisions, admits no closed-form solution, and therefore would have to be solved numerically to predict its behavior. A similar analysis could be done to determine the arithmetic precision and computational load required to accurately predict the behavior after 35 collisions. I have not done this analysis, but I suspect it would show that this is equally impossible.

1.3.1 Murphy’s Law and Faults

An argument that I hear frequently against deterministic models goes something like this: In the real world, things will go wrong. Nothing is really predictable. Even deterministic models will be violated in practice, so why bother with de-

terministic models? Why not, instead, assume everything is random and design your system to tolerate this randomness?

This argument, if carried too far, suggests we should not bother with the reliable in-order packet delivery of TCP, the linchpin of the Internet. We should not bother with **CRC (cyclic redundancy check)** to ensure the integrity of computer memories. We should not bother with synchronous digital logic design, and instead use circuits that may or may not produce expected results. But we do bother with all these things. Why?

Today, computer memories have replaced ledgers in finance and law. Digital signatures have become an acceptable way to finalize legal contracts. Stock trades execute without human intervention or paper records. None of these would be possible without deterministic models and our ability to build physical systems that are highly faithful to these models. Electronic circuits perform billions of arithmetic operations per second and go for years without errors.

I repeat the core principles because they are so important. Determinism is a property of models, not of physical systems. An *engineering* model is a specification of how a physical system *should* behave, not a model of how the physical system *does* behave (the latter is a scientific model). When you have a model that defines how a system *should* behave, then you get, for free, the notion of a **fault**. A fault is a behavior that deviates from the specification.

The existence of faults does not undermine the value of deterministic models. In fact, the very notion of a fault is strengthened by deterministic models because they define more clearly what behavior a physical system *should* have. Detecting faults, therefore, is easier. A CRC, for example, detects at least some violations of a simple deterministic model of a computer memory. This enables fault-tolerant design, where the system reacts in predictable ways to faults.

Every realization of an engineering model can exhibit faults. When we successfully build a physical system that reliably behaves like a model, it does so only under certain assumptions. No computer will correctly execute a program if it overheats, is crushed, or is submerged in salt water. The model is faithfully emulated only under the assumption that none of these things has happened.

Making the assumptions clear also has value. The Ptides model (Zhao et al., 2007) for deterministic distributed execution of discrete-event programs, for example, which is realized in Google Spanner (Corbett et al., 2012) and Lingua

Franca ([Lohstroh et al., 2020](#)), assumes a bound on clock synchronization error and a bound on network latency to achieve extremely efficient distributed and fault-tolerant coordination of program components (see Chapter 8). Violations of these assumptions will occur in practice, but, in a well-designed system, they will be rare. Moreover, such violations are detectable because they manifest as software components seeing events out of timestamp order. A deterministic model, together with clearly stated and quantified assumptions under which a physical realization emulates the model, enable efficient designs that can react in predictable ways to faults. Hence, despite Murphy’s Law, deterministic models are useful, even in the face of unpredictable failures.

1.4 Determinism in physics

The question of whether the physical world is deterministic has been controversial for a long time. In the early 1800s, Pierre-Simon Laplace argued that if someone (a “great intellect,” later known as “Laplace’s demon”) were to know the precise location and velocity of every particle in the universe, then the past and future locations and velocities for each particle would be completely determined and could be calculated from the laws of classical mechanics ([Laplace, 1901](#)). Is this true?

In 2008, David Wolpert proved that Laplace’s demon cannot exist ([Wolpert, 2008](#)). No such calculation is possible. Wolpert’s proof relies on the observation that such a demon, were it to exist, would have to exist in the very physical world that it predicts. This results in a self-referentiality that yields contradictions, not unlike Turing’s undecidability and Gödel’s incompleteness theorems.

But Laplace’s demon is about *prediction*, not just determinism. We already know that determinism does not imply predictability. Even if prediction is known to be impossible, we cannot conclude that the world is nondeterministic. We will look at the question of whether Newtonian physics, the state-of-the-art when Laplace lived, is a deterministic model. You may be surprised by the answer.

More recently than Laplace, Karl Popper, high priest of scientific positivism, also insists on a deterministic universe:

One sometimes hears it said that the movements of the planets obey strict laws, whilst the fall of a die is fortuitous, or subject to chance. In

my view the difference lies in the fact that we have so far been able to predict the movement of the planets successfully, but not the individual results of throwing dice. In order to deduce predictions one needs laws and initial conditions; if no suitable laws are available or if the initial conditions cannot be ascertained, the scientific way of predicting breaks down. In throwing dice, what we lack is, clearly, sufficient knowledge of initial conditions. With sufficiently precise measurements of initial conditions it would be possible to make predictions in this case also. (Popper, 1959, p. 198)

This quotation reflects a conventional wisdom, which dictates that Newton's laws provide a deterministic model of the universe. Also conventional wisdom is that quantum physics dashed that determinism. Neither of these is strictly true. A concise summary of the ways that determinism in these physics models have been interpreted is given by Hoefer (2016). A more in-depth study is given by Earman's *Primer on Determinism* (Earman, 1986). Here, I will relate some of these interpretations to the above discussion of determinism in engineering models and give my own perspective on the subject.

1.4.1 Nondeterminism in Newtonian Physics

A (rather controversial) example of nondeterminism in Newton's laws is due to the philosopher of science John Norton (Norton, 2007).⁷ Norton considers a point mass precariously balanced on top of a smooth frictionless dome. Norton shows that, without violating any of Newton's laws, the mass can spontaneously begin sliding down the side of the dome in an arbitrary direction at an arbitrary time without anything causing it to start sliding. His argument is carefully constructed and surprised me when I first heard it. I was sure his argument was wrong, but I finally concluded that my certainty was based on circular reasoning.

Newton's second law, given in equation (1.1), states that at any time instant, the force imposed on an object equals its mass times its acceleration. If there is no force, the acceleration must be zero. If the acceleration is zero, then the velocity is not changing. Hence, it would seem that if the mass is not moving, balanced on the top of the dome, and no force is applied, then it should remain still, with velocity equal to zero. But Norton points out that it is possible for the mass to

⁷Another example with similar properties is given by Dhar (1993).




Figure 1.2: Norton's dome (from <http://www.pitt.edu/~jdnorton/Goodies/Dome/index.html>). From <http://www.pitt.edu/~jdnorton/Goodies/Dome/index.html>.

start sliding down the dome at any arbitrary time T without violating this law and without any force initiating the slide. At the instant T , the mass will have velocity zero and acceleration zero, so it is not moving. But at any time greater than T , say at $T + \epsilon$, no matter how small ϵ is, the mass may be no longer centered on the top of the dome. It will now be sitting on a slope, which means that gravity will exert a nonzero force in the downhill direction, and the mass will have a nonzero acceleration.

Specifically, Norton proposes a dome shape where the dome drops by a distance $h = (2/3g)r^{3/2}$, where r is the distance along the surface of the dome from the center of the dome and g is the force of gravity (see Figure 1.2).⁸ There is nothing particularly special about this shape; it just makes the math work out simply. With this choice, the force on the mass tangent to the dome, as a function of the distance r from the center of the dome, is $F(r) = \sqrt{r}$. With this function, equation (1.1) admits many solutions. In particular, if we assume unit mass $m = 1$, then the following is a solution for any T :

$$r(t) = \begin{cases} (1/144)(t - T)^4 & t \geq T \\ 0 & \text{otherwise} \end{cases} \quad (1.2)$$

At the instant $t = T$, the mass is not moving, the net force on the mass is zero, and the mass is not accelerating. At any time larger than T , the mass is moving, the net force is not zero, and the mass is accelerating down the dome. T can have

⁸Gareth Davies points out in his blog that this is a rather odd specification of the dome shape and that the extent of the dome has to be limited to $r \leq g^2$. See <https://blog.gruffd Davies.com/2017/12/24/newtonian-physics-is-deterministic-sorry-norton/>.

any value without violating Newton's second law. Equation (1.1) holds at every instant t .

There are many subtleties around Norton's example. First, it may be helpful to realize that there is no *first* instant at which the mass accelerates. Instead, the time $t = T$ is the *last* instant at which the mass is *not* accelerating. At all instants greater than T , there is a nonzero net downward force, gravity on a slope, so the mass accelerates. Zinkernagel (2010) claims that this property is the key flaw that would allow us to declare Norton's dome to not be a "Newtonian system." He says the force lacks a "first cause." But requiring a first cause would force us to reject many other innocuous systems that do not exhibit nondeterminism, including many reasonable models where force appears gradually (Fletcher, 2012).

Malament (2008) offers a fascinating analysis of the mathematics behind Norton's dome. First he addresses a preconception that was part of what made me initially skeptical about Norton's claim. I had always assumed that Newton's second law (1.1) would have a unique solution for any input force function F that could be generated by a reasonable Newtonian system. Norton's system seems reasonable in this sense because the force is just Newtonian gravity and the dome is a reasonably simple geometric shape. But Malament points out that uniqueness is not guaranteed if F is not continuously differentiable. Norton's function is not continuously differentiable at $r = 0$, and (1.1) admits many solutions. We could restrict "Newtonian Systems" to include only forces that are continuously differentiable.⁹ This is not satisfactory to me, however. We would have to rule out a large number of shapes, including shapes that do not lead to nondeterminism, such as a table with an edge where a mass slides off the edge.

Another possible objection to Norton's dome is that such a shape could never be constructed perfectly. This argument is specious, however, because *every* shape that we can describe mathematically will have flaws when constructed in the physical world. This objection would effectively eviscerate Newtonian physics.

Yet another approach to rejecting Norton's claim is to impose constraints on the *solutions* to (1.1) rather than the force function. For example, Davies suggests that higher-order derivatives of the solution need to exist and be continuous for the solution to be a reasonable model of physical behavior.¹⁰ The family of solutions

⁹ A weaker but sufficient constraint would be to require the force to be locally Lipschitz.

¹⁰ See Gareth Davies' blog: <https://blog.gruffdavies.com/2017/12/24/newtonian-physics-is-deterministic-sorry-norton/>

given by (1.2) does not meet this requirement. However, this too is not satisfactory to me. Again, it would rule out a mass sliding off the edge of a table. Moreover, this amounts to saying something like, “the model is valid if we choose from among its behaviors the one that seems reasonable.”

It seems that to regain determinism, we need to augment Newtonian physics with additional axioms that are not derivable from the core concepts. Newton could have given us a fourth law of motion going something like this: “a mass can only have one possible motion that conforms with the previous three laws.” This would assume away nondeterminism in order to obtain a deterministic model of the physical world.

Fletcher (2012) points out, however, even this restriction has several possible versions, and any choice between these seems arbitrary. For example, we could declare Norton’s dome to not be a “Newtonian system” altogether, or we could declare it to not be a “Newtonian system” only when the mass starts or ends at the peak of the dome. If the initial position of the mass is somewhere else on the dome, Newton’s second law gives us a deterministic model. What if it starts somewhere else on the dome with some momentum towards the peak and crosses the peak? At the moment of crossing, does the system suddenly and instantaneously cease to be a Newtonian system?

Are Newton’s other laws violated when the mass spontaneously starts sliding? Newton’s first law states that an object will remain at rest or in uniform motion in a straight line unless acted upon by an external force. It may seem that having the mass spontaneously start to move violates this law, but actually it does not, at least under one reasonable interpretation of this law. Since an external force may vary in time, this first law needs also to be interpreted as a statement that holds at each instant of time. Under this interpretation, the first law becomes a special case of the second law where $F(t) = 0$.

There is another interpretation of Newton’s first law, however, that restores determinism for this example, but this interpretation requires more than Newtonian physics and is ultimately based on circular reasoning. The first law can be interpreted to mean that there can be no *uncaused* changes in momentum, where the cause is some external effect. If the mass slides off the dome as I have described, then there must be some such external effect causing this. That effect must be something other than a Newtonian force, however, because any force would lead to a trajectory different from (1.2), which satisfies the second law. So what could

that required external non-force effect be? The force of will of a conscious mind? The force of God? There is nothing in Newtonian physics that qualifies.

One could take this absence of a model for such an external non-force as an argument that the mass will not move. But the *absence* of a model cannot be construed as evidence. According to Norton, there is no *need* for such an external non-force for the mass to slide down the side of the dome, so there is no need for a model for this external non-force. Newton's laws are still satisfied. It is equally valid to demand a model for whatever keeps the mass perched on the dome. What non-force is that?

It will still disturb many readers that the mass can start moving with no provocation. Newton's laws are *also* satisfied by a mass that behaves itself and remains quietly perched at the top of the dome for all eternity. Isn't it reasonable to assume it will do that? An empirical approach would actually find this interpretation *unreasonable*, since any practical realization of Norton's dome will result in the mass sliding off the peak. The apparent reasonableness of this interpretation is due to a distaste for uncaused action, i.e., a distaste for nondeterminism. It is not due to empirical evidence and indeed flies in the face of empirical evidence. Hence, the argument that the mass will remain at the top of the dome is circular. It is not supported by the mathematics of Newton's laws alone and instead depends on the assumption that nothing happens without provocation. In other words, it concludes determinism based on an assumption of determinism.

The presupposition that every behavior has a cause is a difficult one to give up. In 1913, Bertrand Russell challenged the scientific world to give it up:

All philosophers, of every school, imagine that causation is one of the fundamental axioms or postulates of science, yet, oddly enough, in advanced sciences such as gravitational astronomy, the word "cause" never occurs ... The law of causality, I believe, like much that passes muster among philosophers, is a relic of a bygone age, surviving, like the monarchy, only because it is erroneously supposed to do no harm. (Russell, 1913)

Objective discussion of causality is difficult because the notion of causality lurks in every aspect of natural language.¹¹ In my recent book, *The Coevolution* (Lee,

¹¹A nice collection of essays on the deep influences of the notion of causality on language is found in Copley and Martin (2014).

2020), I examine this presupposition of causality in much more depth, leveraging the arguments of Judea Pearl in (Pearl and Mackenzie, 2018) (who argues that reasoning about causality requires subjective involvement) and evolutionary biologists (who argue that the notion of causality may have arisen because of its evolutionary survival value rather than because it is a fact about the world). But for our purposes here, it is sufficient to simply observe that Newton’s laws do not imply causality.

What about Newton’s third law, which states that every action has an equal and opposite reaction? When gravity exerts a force on a mass, causing the mass to fall, the mass exerts an equal and opposite force on the earth, causing the earth to rise. Since the mass also gains momentum in some lateral direction, the earth must acquire an equal and opposite lateral momentum. The mass of the earth, however, is so much larger than the masses in Norton’s example that, to the earth, the force exerted on it is negligible. At all times $t > T$, the nonzero net force downward and laterally on the mass will be balanced by an equal and opposite nonzero net force pulling the earth up and laterally. The effect of that force will not be measurable, but it is there, so the third law is also not violated.

Newtonian physics is often assumed to be time reversible, but reversing time on Norton’s example has curious effects. Consider the scenario where the mass starts on the outskirts of the dome and we push it with just enough force that it reaches the top of the dome and stops. If we push too hard, it will go over the top of the dome. If we don’t push hard enough, it will not reach the top of the dome and will fall back down. But if we push it with the Goldilocks force, just right, it will stop at the top, stay there for an arbitrary amount of time, and then spontaneously slide down the dome again sometime in the future. As soon as the mass perches at the top, its history is lost. Nothing in its state reveals when the mass arrived at the top, thereby foiling Laplace’s demon.

For this scenario to work, Norton points out that the shape of the dome is important. If the dome is a perfect hemisphere, then with the Goldilocks force, it will take infinite time for the mass to reach the top of the dome. It will keep slowing down as it approaches the top, but it will never actually reach the top. But there are many other dome shapes where the mass reaches the top in finite time. Norton’s example, where the dome drops by a distance $h = (2/3g)r^{3/2}$, is one such dome shape.

In conclusion, either Newtonian physics admits nondeterminism, or Newtonian physics needs to be augmented with additional axioms that preclude nondeterminism. As I will show in Section 1.5, however, any set of additional axioms that preclude nondeterminism will also preclude many useful models of physical phenomena.

1.4.2 Metastable States

It turns out that many systems are vulnerable to similarly uncaused action under Newtonian physics. When the mass is perched on the top of the dome, it is in a **metastable state**. A metastable state is marginally stable, where infinitesimal disruptions throw the system out of its precarious state. Norton's mass is vulnerable to falling out of its metastable state with no provocation.

Electronic circuits, particularly ones at the boundary between the continuous physical world and the discrete world of digital electronics, have long been known to be vulnerable to lingering for unbounded periods of time in a metastable state (Marino, 1981; Kinniment, 2007; Mandler et al., 2012). So the problem is not limited to cute examples of masses on domes. It is a fundamental problem at the boundary between the discrete, computational world of computers and the continuous physical world.

A particular kind of metastable system is a **bistable** system, one that has exactly two stable states and can persist for an indeterminate period of time in a metastable state between the two stable states. A digital circuit can be thought of as a piece of electronics that wants to be in one of two states, and in principle, situations where it can linger indefinitely between these two states are unavoidable.

Designers of such circuits go to great lengths to make sure that the probability of lingering gets extremely small as time advances. Consequently, it is rare for circuits to persist in a metastable state for very long. Such situations can occur, however, but are difficult to reproduce in the lab. They have occasionally been implicated in otherwise inexplicable crashes of computers. Is the basic operation of such circuits deterministic? If not, then any electronic implementation of a Turing machine, a deterministic model, is actually nondeterministic.

Bistable behavior has also been observed by biologists in nerve axons. Under certain circumstances, these axons can linger for an indefinite period of time

before settling into one of two resting potentials (Ditlevsen and Samson, 2013). It is likely that such metastability plays a role in brain function.

Building a physical realization of Norton's dome is impossible because no physical dome is perfectly smooth and frictionless and no mass is a point mass. Nevertheless, it is common to deliberately build systems that come as close as possible to such metastability. Sensitive instruments depend on metastable states. The instrument hovers in its (nearly) metastable state until the slightest nudge from the thing being measured pushes it off in one direction or the other. The circuit that reads the contents of a DRAM, a commonly used computer memory, for example, makes use of such metastability to read a tiny stored charge.

1.4.3 Relativity

The special and general theories of relativity are mostly deterministic in a similar sense that Newtonian physics is mostly deterministic. There are only a few corner cases, specifically singularities, that result in many possible futures given a specific past. The event horizons of black holes present such singularities, but these are unobservable in the rest of spacetime, and hence arguably pose no problem to the determinism of the theory. A conceivable class of singularities, called "naked singularities," however, have no event horizon and become observable. In 1969, Sir Roger Penrose posited that such naked singularities do not exist in the universe, a principle called "cosmic censorship." If this principle holds, then general relativity adds no sources of nondeterminism over and above any already present in classical physics.

Nevertheless, there is one aspect in which relativity complicates the notion of a deterministic model of physics. Specifically, it undermines the notion of "the state of the system," as used by Milner and Newton, and consequently, we can no longer talk about determinism in terms of the evolution of the state of the system in time. In relativity, there is no "instant in time." This Newtonian concept is replaced by a Cauchy surface, a hypersurface in four-dimensional space time. But the Cauchy surface is different for each observer, so two distinct observers can disagree about the state of the system.

Example 1.5: The difficulty can be illustrated with the example in Figure 1.3. Consider a distributed system with two physically separated subsystems,


one of which makes an instantaneous transition from state A to B , and the other of which makes a transition from C to D . Under a Newtonian model of time, one of the following transition sequences must be the true one, in the sense that it is actually what happens when the system transitions from (A, C) to (B, D) :

$$\begin{aligned}(A, C) &\rightarrow (B, C) \rightarrow (B, D) \\ (A, C) &\rightarrow (A, D) \rightarrow (B, D) \\ (A, C) &\rightarrow (B, D)\end{aligned}$$


However, according to relativity, it is possible for none of these to be “true” for all observers. The order of physically separated events may depend on the observer. It is not that two different observers just *see* things delayed in different ways, it is that the true order for the different observers is different. There is no ground truth, in the sense that we cannot make a statement like, “at time t , the system was in state (B, C) .” That statement can be true for one observer and false for another.

This problematic notion of time (and hence state) follows easily from the fact that the speed of light is the same for all observers.

Example 1.6: Consider a simple thought experiment, a variant of the famous Einstein train example. Suppose that Randy is seated at the center of a rapidly moving train car, and Jane is standing on a platform at a station where the train does not stop. Suppose that just at the instant that Randy passes Jane, he hits a button that emits a brief pulse of bright light. In



Physically separated state machines.



Transition system model.

Figure 1.3: Model of a distributed system that relies on a notion of “system state.”

Randy's frame of reference, the light will hit the front and back ends of the train car at the same time. However, in Jane's frame of reference, the light will hit the back of the train car before it hits the front. In Randy's frame of reference, the light has to travel the same distance in either direction, and since the speed of light is constant, it will strike both ends of the car simultaneously. But in Jane's frame of reference, the distance the light has to travel to the back of the train is less because the back of the train is moving toward where the light was emitted while the light is traveling toward the back of the train. Hence, the two events of light striking the ends of the train are simultaneous for Randy but not for Jane.

Note this example is not just a cute toy that is only realizable if you have trains that can travel close to the speed of light. The phenomenon is called the Sagnac effect, and it is used in ring laser gyroscopes, a key component of many inertial guidance systems, including many used in commercial aircraft. The Sagnac effect also has to be taken into account in the design of GPS because of the rotation of the earth.

Instead of an evolving state in time, a relativistic system is a collection of events where, given a pair of events, one may precede the other. This is analogous to the contrast between Milner's transition systems (analogous to classical physics) and Kahn's precedence order constraints (analogous to relativity). Of course, if one event *causes* another, then relativity is careful to ensure that all observers see the one before the other, so nothing magical happens. To ensure this consistency, relativity posits that the effects of any event cannot propagate through space faster than the speed of light.

Interestingly, a similar dichotomy exists in the philosophical study of time, dubbed an "A-Series" (classical) model of time or a "B-Series" model by [Gale \(1966\)](#). An A-Series model of time is built on the tensed notions of past, present, and future, while a B-Series model of time is built on a partial order, a "precedes" relation between events. For example, the statement "I will complete this chapter today" is an A-Series statement and will have a different meaning if uttered tomorrow. In contrast, "Completion of this chapter precedes the 2026 New Year" is a tenseless B-Series statement. It has the same meaning whenever it is uttered. Relativity requires a B-Series model of time.

The illusion of state is a powerful illusion, however, because no single observer can observe a system to be in two states at once. For example, no observer of the example of Figure 1.3 will see both (B, C) and (A, D) . This becomes important when considering quantum models, which I do next.

1.4.4 Quantum Physics

It is common to assume that quantum physics immediately undermines any notion of the world being deterministic, but the story is more subtle. There are many interpretations, not all of which lead to nondeterminism.

First, under quantum mechanics, the evolution of the particle's wave function is deterministic, following the Schrödinger equation. If we redefine "behavior" to be the evolution of the wave function in time, rather than position and momentum, then the model is deterministic. The fact that the Schrödinger equation is deterministic prompted Stephen Hawking to colorfully proclaim that determinism can be salvaged:

At first, it seemed that these hopes for a complete determinism would be dashed by the discovery early in the 20th century that events like the decay of radioactive atoms seemed to take place at random. It was as if God was playing dice, in Einstein's phrase. But science snatched victory from the jaws of defeat by moving the goal posts and redefining what is meant by a complete knowledge of the universe. ... In quantum theory, it turns out one doesn't need to know both the positions and the velocities [of the particles]. ([Hawking, 2002](#))

It is enough to know how the wave function evolves in time.

But the wave function is not directly observable. It cannot be measured by instruments. If deterministic behavior has no observer, does it lose its value? If, on the other hand, we define "behavior" in terms of measurable quantities, such as position and momentum, then the quantum model becomes nondeterministic, albeit in a subtle way. The wave function is commonly interpreted as giving the probabilities of the various possible inherently random measurement outcomes.

However, what could "probability" mean here? It turns out that it cannot be the usual notion of probability that you probably learned in an undergraduate

class. Quantum probability is not an indicator of the relative frequency of various outcomes when performing repeated experiments. The so-called “no cloning theorem” in quantum mechanics says that such repeated experiments are impossible! The frequentist interpretation of probability has to be replaced by the Bayesian notion, which gives a different meaning to the word “probability” (see my book (Lee, 2017, Chapter 11) for a discussion of these two interpretations). In the Bayesian interpretation, a probability is a measure of what is unknown. Quantum probability is even stronger; it is a measure of what is *unknowable*.

It turns out that the mechanics of probability theory does not depend much on which interpretation you adopt. The math is the same. Interestingly, much of modern probability theory was developed by Laplace, who was convinced the world was completely deterministic. There is no contradiction here, however, because Laplace was a Bayesian. His probabilities measure uncertainty, or lack of knowledge, not intrinsic randomness.

Even if we take “behavior” to be the evolution of the wave function in time, despite our inability to measure it, subtleties remain. The Schrödinger equation gives the evolution of the wave function in classical Newtonian time, not relativistic time. Hence, even if we could measure the wave function, it would become possible for two distinct observers to see two different wave function evolutions for the same physical system! If the system includes metastable components, these distinct evolutions could be quite different indeed.

The role of an observer has always been important in quantum physics, albeit not in this relativistic sense. In what is now called the Copenhagen interpretation, originally proposed in the years 1925 to 1927 by Niels Bohr and Werner Heisenberg, the state of a system continues to be defined by probabilities until an external observer observes the state, and only at that point do the probabilities influence the outcome. Prior to being observed, all possible outcomes represented by the probabilities continue to remain possible. This requires an “observer” who is somehow separate from the system and measures the position of the particle. A “collapse” of the wave function occurs at the instant that a “measurement” is made, converting possibilities into certitude. This interpretation leaves unspecified what a “measurement” or “observation” is and typically puts the measurement apparatus outside the domain of quantum mechanics. This has led to sometimes bizarre interpretations, for example that conscious minds play a central role, which presumably means that the physical world didn’t exist or played by different rules before conscious minds formed.

The Copenhagen interpretation, however, is firmly rooted in an insistence on imposing a notion of “state” on the system. The outcome of a measurement is taken as a definitive answer, a fact about the system. A more fully quantum interpretation would continue to model the system, even after measurement, using a wave function.

In the 1950s, the physicist Hugh Everett III dispensed with the distinct observer, instead bundling observer and observed under a single wave function that evolves deterministically under the Schrödinger equation. The measurement apparatus and measured system entangle and evolve together in a single wave function. This view is a straightforward, simple, and direct interpretation of quantum mechanics until one insists on the same sort of certitude that one gets from the collapse of the wave function posited by the Copenhagen interpretation. With this insistence, the theory gets rather extravagant.

Consider a particle, say, an electron, moving through space. In classical physics, at each instant t , it has a definite position and momentum. In quantum physics, it has a wave function. If we measure its position, say by putting a phosphorescent screen in its way, then the position of the electron will be revealed by a flash on the screen. Under the Copenhagen interpretation, at the time the electron hits the screen, its position is drawn randomly from a probability density function that puts weights equal to their likelihood on all regions of possible positions. Under Everett’s interpretation, the photon, the screen, and the human observer become entangled, and a single wave function covering all of them continues to evolve deterministically.

But, under Everett’s interpretation, where is the electron? Physicists seem to continue to insist that it must be somewhere, that it has “state,” which leads to the most bizarre part of Everett’s thesis. In his thesis, the electron is everywhere that its wave function permits it to be, but now in an uncountably infinite number of split-off universes. In each such universe, the electron is at one of the infinitely many possible positions. Each of these universes is somehow weighted by the probability dictated by the wave function.¹² The insistence on state leads to uncountably many universes being spawned anytime there is an interaction between components of a system.

¹²This is a rather mysterious part of the many worlds interpretation because the wave function defines a probability *density* for position, not a probability, so each of the uncountably many universes would have to have weight equal to zero. This stretches the notion of “existence” to the breaking point.

Because of the proliferation of universes, Everett's thesis is often called the "many worlds" interpretation of quantum mechanics. It means that every time there is an interaction between an observer and a subject, the universe splits. Depending on what is being measured, it could split in two or into an infinite number of possibilities. Since interactions are occurring all the time, a wildly extravagant proliferation of universes occurs at every instant in time. The model is deterministic, but the split universes cannot interact with one another, and hence, from the perspective of any entity in any one of those universes, the outcome of the experiment appears to be random.

Many physicists consider this consequence of Everett's thesis to be a *reductio ad absurdum* proof of the invalidity of the thesis. But others, including many highly respected physicists, accept the thesis as the best available explanation of quantum mechanics (see [Becker \(2018\)](#) for a very readable overview of the alternative explanations).

There is, however, a simpler interpretation that is consistent with Everett's core idea that there is a single wave function that continues to evolve deterministically. The simpler interpretation rejects the notion that there is a ground truth about the state of a system at an instant in time. The electron is never at a definitive location. Despite seeing the flash of light, even a human observer is mistaken to interpret this as definitive evidence of a true location. No experimental apparatus is perfect, and it takes a great deal of effort to build an experimental apparatus that delivers even reasonable confidence, much less certainty. Moreover, no human perceptual system, which sees the flash, is perfect, and no brain has perfect memory. It is simply wrong to conclude that the flash is an indicator of a definitive truth. It is no more than compelling evidence of a high likelihood. This simpler interpretation has no need for a proliferation of universes. In exchange, it sacrifices certainty. Specifically, it sacrifices the notion of "state."

Einstein famously resisted aspects of quantum physics, a theory he helped to build, insisting that the theory was incomplete. He spent much of his later career searching for the "hidden variables," hitherto unknown aspects of the state of a physical system that would reconcile quantum theory with the principles of locality implied by relativity (that no event's consequences can propagate at faster than the speed of light). The conflict posed by quantum theory is between locality and a principle that some physicists call "reality." Here, **reality** is the principle that objects have real properties that determine the outcome of measurements, i.e., that they have what I have been calling a "state." **Locality** is the principle that

reality at one location in space is not influenced instantaneously by measurements performed at a distant location (or more specifically, that influences propagate no faster than the speed of light).

In 1964, John Stewart Bell proved that quantum theory is incompatible with the principles of reality and locality (Bell, 1964). That is, if quantum theory is a faithful scientific model of the physical world, something that is well verified experimentally, then one of the two principles must be false. Most physicists, including Bell himself, seem to prefer to sacrifice locality over reality.

However, this makes quantum theory incompatible with relativity, which requires locality. Relativity, however, does not require *reality* in this sense. In fact, as I argued in the previous section, relativity is inconsistent with reality in this sense. Relativity shows that a physical system, at least one that is spread out over space, cannot have “state” that, at an “instant” (a boundary between the past and the future), determines all of its properties. Any notion of state is dependent on the observer. The so-called “relational interpretation” of quantum mechanics, first attributed to the theoretical physicist Carlo Rovelli, takes this perspective. The state of a quantum system is a relation between the observer and the system. There is no notion of state that is independent of an observer (see Rovelli (2017, 2018) for readable explanations of this approach).

In summary, the world is not predictable under either classical or modern physics. Is it deterministic? That depends on what model we use, and all deterministic models have logical problems. A deterministic Newtonian model requires a presupposition of determinism. A deterministic relativistic model requires giving up the notion of state. A deterministic quantum model requires either an extravagant proliferation of universes or a similar forsaking of the notion of state. I will argue next that these logical problems are inevitable. Any set of deterministic models rich enough to say interesting things about the world is demonstrably incomplete, so these logical problems are unavoidable.

1.5 Incompleteness of Determinism

In 2016, I published a paper on the limits of modeling for cyber-physical systems in which I showed a sense in which determinism is incomplete (Lee, 2016) (see also my book (Lee, 2017, Chapter 10)). I review that result here. The short summary is that any set of deterministic models that is rich enough to encompass Newton’s

laws and also admits discrete behaviors does not contain its own limit points. Thus, any approach to modeling that relies on such a set of deterministic models has corner cases that exhibit nondeterminism. Nondeterminism, therefore, is inescapable in physical models.

1.5.1 Nondeterministic Collisions

Consider a model of the collisions of two billiard balls, as shown in Figure 1.4. Suppose that we model a collision as a discrete event, where the collision occurs in an instant, having no duration in time. Assume that the balls are ideally elastic, meaning that no kinetic energy is lost when they collide. In this case, Newton's laws require that both energy and momentum be conserved; the total momentum and energy must be the same after the collision as before.

Let v_1 and v'_1 be the velocity of the first ball before and after the collision, respectively. Let v_2 and v'_2 similarly represent the velocity of the second ball before and after the collision. Conservation of momentum requires

$$m_1v'_1 + m_2v'_2 = m_1v_1 + m_2v_2, \quad (1.3)$$




Figure 1.4: Collision of ideal billiard balls on a frictionless surface.

where m_1 and m_2 are the masses of the two balls, respectively. Conservation of kinetic energy requires

$$\frac{m_1(v'_1)^2}{2} + \frac{m_2(v'_2)^2}{2} = \frac{m_1(v_1)^2}{2} + \frac{m_2(v_2)^2}{2}. \quad (1.4)$$

Assuming we know the starting speeds v_1 and v_2 and the masses, then we have two equations and two unknowns, v'_1 and v'_2 . This is a quadratic problem with two solutions.

Solution 1: Ignore the collision:

$$v'_1 = v_1, \quad v'_2 = v_2.$$

Solution 2:

$$\begin{aligned} v'_1 &= \frac{v_1(m_1 - m_2) + 2m_2v_2}{m_1 + m_2} \\ v'_2 &= \frac{v_2(m_2 - m_1) + 2m_1v_1}{m_1 + m_2}. \end{aligned}$$

Note that solution 1 looks like tunneling, where the balls pass through one another without affecting each other, as a neutrino might. In solution 2, if $m_1 = m_2$, then the two masses simply exchange velocities, as suggested in Figure 1.4. If we rule out solution 1, then the model is deterministic.

Now consider a more elaborate scenario, shown in Figure 1.5. Two (ideal) billiard balls approach a third stationary ball from opposite sides on a frictionless surface and collide with the stationary ball simultaneously. How should they react? If we consider only conservation of momentum and energy, we still have two equations, but now there are *three* unknowns.




Figure 1.5: Collision of three billiard balls on a frictionless surface.

At the time of the collision, there are *two* collisions. We could attempt to treat these as a pair of two-ball collisions, each of which has two possible outcomes. If we reject the tunneling solutions, then it would seem that only one outcome remains for each of the two collisions. But it is not obvious how to combine the two non-tunneling outcomes.

A first (naive) solution, using what is known as **Newton's hypothesis**, just superimposes the resulting momentums of the two non-tunneling outcomes. If the balls all have the same mass, then the left ball will transfer its momentum to the middle ball, the right ball will also transfer its momentum to the middle ball, and the equal and opposite momentums will cancel. All balls stop. Momentum is conserved, but not energy. This solution is shown at the top of Figure 1.6. Since there is no mechanism for energy dissipation in this model, this resolution is not satisfactory.

An alternative solution, using what is known in the literature as **Poisson's hypothesis**, introduces a form of superdense time (Lee, 2014). At the time of the collision, the kinetic energy of the balls is instantly converted to potential energy by compressing the middle ball. Then, without time elapsing, in a second microstep, the potential energy is reconverted to kinetic energy by the middle ball expanding. But how should this ball apportion the kinetic energy to the two outer balls? An intuitively appealing solution is shown at the bottom of Figure 1.6, which assumes the masses are equal and the two balls bounce off the center ball and end up with equal and opposite velocities. In general, however, there are many solutions that conserve both momentum and energy! We seem to have no basis for picking one solution over another, so the model appears to become intrinsically nondeterministic.

This thought experiment asked us to consider that the two collisions occur *simultaneously*. But, as we saw in the previous section, we have to choose an observer before simultaneity has any meaning. One observer may see the balls colliding simultaneously, another may see the left collision occurring first, and a third may see the right collision occurring first. It seems that any solution we come up with needs to have some sort of consistency across the experiences of these three observers.

Let's consider what happens when we treat the collisions as occurring in some order, but without any time elapsing between the collisions (this again relies on a superdense model of time and can be seen as a limiting case where an infinitesimal




Figure 1.6: Newton's hypothesis vs. Poisson's hypothesis.

amount of time elapses between the collisions). As shown in Figure 1.7, when the collisions occur, we arbitrarily pick either the left collision or the right one, temporarily ignoring the other one. Rejecting the tunneling solution, we get a deterministic exchange of momentum. Without time elapsing, we find ourselves in state (b) in the figure, at which point we must handle the second collision. Again we get a deterministic solution, leaving us in state (c). Again, without time elapsing, we handle a third collision, which leaves us in state (d). After time elapses, we find ourselves in state (e).

I first studied this problem while developing the hybrid system simulator in Ptolemy II (Cardoso et al., 2014). I experimented with different masses and with handling the collisions in different orders. I kept seeing plots like those shown in Figure 1.8. I was convinced that these plots should have been the same, that the order in which the software handled the collisions should not matter. Time did not elapse between collisions, so according to Newton's laws, the state of the system should not change. I spent weeks looking for the bug in the software before I finally realized that there was no bug in the software. More than one final state conserves both momentum and energy.




Figure 1.7: One of two orderings for handling collisions.

Inevitably, when I present this example, someone asks me how the “real world” behaves in this situation. This is a difficult question. One of the consequences of quantum mechanics is the Heisenberg uncertainty principle, which states that we cannot simultaneously know the position and momentum of an object to arbitrary precision. But discrete modeling of these collisions depends on modeling position and momentum precisely.

To many readers, it may seem odd to be invoking quantum mechanics on macro-scale physics problems, where Newtonian mechanics usually works just fine. But when the *order* of events affects the outcome, we find ourselves inevitably at quantum scales. The order of events can change with arbitrarily small differences in time or space.

Both relativity and quantum physics expose difficulties with this mix of time and space continuums with discrete events. Perhaps it is the mix that is problematic and we should instead reject either continuums or discreteness. Rejecting con-




Figure 1.8: If the masses are different, the behavior depends on which collision is handled first.

tinuums amounts to accepting a hypothesis sometimes called “digital physics,” a position that I challenge in my book *Plato and the Nerd* (Lee, 2017, Chapter 8).

Rejecting discreteness is a bit harder to debunk, but it has the same flavor as the assumptions that Malament and Fletcher point out are required to rule out Norton’s dome as a valid Newtonian system (Malament, 2008; Fletcher, 2012). Discrete collisions are singularities, not unlike the singularity that makes Norton’s dome result in forces that are not continuously differentiable. But ruling them out simply to preserve determinism seems rather arbitrary. More disturbingly, I have developed an example (a flyback diode circuit) (Lee, 2016) that shows that rejecting discreteness has the consequence of also having to reject causality.

Nevertheless, for the billiard balls example, rejecting discreteness is relatively easy. Using only classical mechanics, we can model the balls as stiff springs, yielding a model where there are no discontinuous changes in momentum. I have shown that this makes the model deterministic, but chaotic (Lee, 2016). Hence, arbitrarily small perturbations in initial conditions lead to arbitrarily large changes in final state. Combine this observation with quantum uncertainty, and the result is indistinguishable from nondeterminism in any practical sense.

1.5.2 Incompleteness

The three-billiard-ball example has the odd property that if the collisions are not simultaneous, then no matter how small the time difference between collisions is, the resulting behavior is deterministic. As the time between collisions approaches zero, as long as it remains nonzero, we have a deterministic model. But in the limit, the model becomes nondeterministic. This suggests that the set of deterministic models is incomplete, in that it does not contain its own limit points.

I review here a construction from my previous paper (Lee, 2016) of such an incomplete set of deterministic models. Consider the set M of models describing one-dimensional motion of $N = 3$ ideal elastic balls subject to Newton's second law, where collisions are handled with impulsive forces, and all behaviors that conserve momentum and energy are allowed except for tunneling. Every model in M is closed, in that there are no external forces, so all behaviors are a consequence of the initial conditions. To keep things simple, the set M includes three balls with specific masses $m_1 = 0.2$, $m_2 = 1.0$, and $m_3 = 5$ (these are the masses that generate the behaviors shown in Figure 1.8). The units do not matter, as long as they are consistent. Assume initial positions $x_i(0) \in \mathbb{R}$, and initial velocities $v_i(0) \in \mathbb{R}$ given as follows:

$$\begin{aligned} x_1(0) &= -1 \\ v_1(0) &= 1 \\ x_2(0) &= 0 \\ v_2(0) &= 0 \\ x_3(0) &= 1 + \Delta \\ v_3(0) &= -1 \end{aligned}$$

where $\Delta \in \mathbb{R}$ is a real number. Δ is the only parameter that distinguishes models in M , but this one variable is sufficient to give us an uncountably infinite set of models. When $\Delta \neq 0$, the collisions will not be simultaneous, as illustrated in Figure 1.9.

Following Lee and Sangiovanni-Vincentelli (1998), we formally define a **model** as a set of behaviors. It is sufficient to consider the behavior of each model over the time interval $[0, 2]$ only. Let B be the set of all functions of the form $b: [0, 2] \rightarrow \mathbb{R}^3$. For a particular model $A \in M$ with parameter Δ_A , we say that $x \in B$ is a behavior of A if for all $t \in [0, 2]$, $x_i(t)$ is the position in our one-dimensional space of ball i at




Figure 1.9: Non-simultaneous collisions.

time t , for $i \in \{1, 2, 3\}$. These x_i are functions of time that satisfy the equations of motion and conservation laws. In other words, a behavior of a model $A \in M$ is a function x giving the positions of the three balls as a function of time.

If a particular model $A \in M$ has exactly one behavior, then A is deterministic. As shown above, M contains only one nondeterministic model, let's call it $N \in M$, the one where $\Delta_N = 0$. We can now construct a sequence of deterministic models in D that should converge to N but does not. To do that, we need some notion of convergence.

Consider the subset $D \subset M$ of deterministic models. We can define a metric on D that allows us to talk about models being "close." Consider two models $A, A' \in D$. Because they are in D , these models are deterministic. Each admits one behavior, $x, x' \in B$, respectively. Because each has exactly one behavior, we can define a distance function d as follows,

$$d(A, A') = \frac{1}{2} \int_0^2 \|x(t) - x'(t)\| dt, \quad (1.5)$$

where $\|x\|$ is the L1 norm of a real vector x . This distance function measures the difference between two ball trajectories. It is easy to show that d is a metric, and hence (D, d) is a metric space.

Consider a sequence of models $A_i \in D$, $i \in \{1, 2, \dots\}$ where

$$\Delta_{A_i} = 1/i^2.$$

As i gets larger, the time between collisions gets smaller, so in some sense, these models approximate the nondeterministic case N where $\Delta_N = 0$ ever more closely as i gets larger. It is easy to show that the sequence A_i is Cauchy, which means that for any $\epsilon > 0$, we can find a positive integer N such that for all positive integers $i, j > N$,

$$d(A_i, A_j) < \epsilon.$$

This means that as i and j get large, the trajectories of the balls in A_i get ever closer to the trajectories of the balls in A_j . With sufficiently large i and j , A_i becomes nearly indistinguishable from A_j . Nevertheless, this sequence has no limit in D .

The one nondeterministic model $N \in M$, where the collisions are exactly simultaneous, $\Delta_N = 0$, admits behaviors that are very distant from the behaviors of the models in the sequence A_i using the same metric. Some of the behaviors of N will be close (in this metric space sense) to behaviors of a very different Cauchy sequence, $C_i \in D$, $i \in \{1, 2, \dots\}$ where

$$\Delta_{C_i} = -1/i^2.$$

Here, the time between collisions is approaching zero from the other side.

A metric space that has Cauchy sequences that have no limit in the space is said to be “incomplete.” The metric space (D, d) of deterministic models is incomplete. It does not contain all its limit points.

Every model in the sequence A_i is deterministic, and the models in the sequence get arbitrarily close to one another. Moreover, any set of deterministic models rich enough to encompass Newton’s laws that allows discrete collisions must be rich enough to include this sequence of models, and therefore will be incomplete.

This has profound consequences that I explore in my book, *Plato and the Nerd* (Lee, 2017). Specifically, many people assume without proof that if some modeling technique can be shown to be able to arbitrarily closely approximate any model of interest, then this technique is “good enough” for all practical purposes. This example shows that this assumption is untrue unless the modeling technique defines a complete set of models. Our Cauchy sequence A_i arguably represents an arbitrarily close approximation to N , and yet it fails to capture a hugely important property of N , that it is nondeterministic.

1.6 Conclusion

I hope I have convinced you that determinism is a deep subject. In reality, I have only scratched the surface here. In my book, *The Coevolution*, I explore an even harder aspect, the connection between determinism and free will (Lee, 2020, Chapter 12). (Hint: observers become important again.) The focus there is to try to get a handle on the question of whether humans will ever build machines that

have, in any sense, free will. More practically, the question is whether we humans will ever build machines that can and should be held accountable for their actions. That question is beyond the scope of this book.

This book focuses instead on the practical question of the usefulness and completeness of deterministic models. I have argued that they are extremely useful, and that even though nondeterministic models also have their uses, those uses are disjoint. The choice of whether to use deterministic or nondeterministic models should be front-and-center in any engineering design effort, and yet it rarely is. Whether the models are deterministic or not typically depends on the choice of modeling and design frameworks, and many of these are nondeterministic by accident, not by intent.

I have shown that whether a model is deterministic or not depends on how one defines the inputs and behavior of the model. To define behavior, one has to define an observer. I compared and contrasted two classes of ways to do it, one based on the notion of “state” and another that more flexibly defines the observables. A state-based model further requires a restrictive model of time, one that is known to be an inaccurate model of physical time. Specifically, it requires an unambiguous simultaneity.

I examine determinism in models of the physical world. In what may surprise many readers, I show that Newtonian physics admits nondeterminism (unless one presupposes determinism) and that quantum physics may be interpreted as a deterministic model. Moreover, I show that both relativity and quantum mechanics undermine the notion of “state” and therefore require the more flexible ways of defining observables. Finally, I show that sufficiently rich sets of deterministic models are incomplete. Hence, no matter how much we value them, they will not solve all our problems.

In engineering practice, whether to use nondeterministic models becomes a central question. For scientific models, where it is incumbent on the model to match the thing being modeled, nondeterminism can be useful if it reflects inherent randomness or uncertainty about the thing being modeled. In this case, nondeterminism enables exploration of the range of possible behaviors. However, nondeterminism says nothing about the likelihood of any behavior, only about its possibility. Hence, a probabilistic model may be more useful than a nondeterministic one. A probabilistic model may usefully model a deterministic system (by interpreting probabilities as a measure of uncertainty) and may even, in some cir-

cumstances, be interpreted as a deterministic model (as we can do with quantum physics).

For engineering models, where it is incumbent on the thing being modeled to match the model, nondeterminism can be a useful abstraction mechanism, a way to get simpler models for analysis. It can also be useful for deferring design decisions. However, it comes at a steep price. The model becomes less useful for testing, for evaluating the degree to which the thing being modeled matches the model. Nondeterminism that arises from sloppiness in the modeling framework or language, however, is rarely useful and should not be excused by the fact that the physical world is unpredictable in practice.

My own focus as an engineer is on cyber-physical systems, which combine the neat world of computation with the messy and unpredictable physical world. Deterministic models play an important role on both sides of this divide. For example, on the physical side, deterministic differential equation models can be useful descriptions of how a robot arm *should* behave, particularly if coupled with probabilistic models of how it may actually behave. On the cyber side, for distributed software, clear specifications of how the components *should* coordinate are useful, particularly if coupled with probabilistic models of network behavior that may compromise these specifications. In both cases, we are talking about a combination of engineering and scientific models.

Concurrency

2.1	Concurrent, Parallel, and Distributed	54
2.2	Consistency	57
2.2.1	The CAP Theorem	59
2.2.2	Consistency and Time	61
2.2.3	Causal Consistency	63
2.2.4	Causal Consistency is Insufficient	66
2.2.5	Using Physical Time to Assign Logical Times	67
2.2.6	Tags and Physical Time	68
	<i>Sidebar: Clock Synchronization</i>	69
	<i>Sidebar: PTP and TSN</i>	71
2.3	The CAL Theorem	74
2.3.1	Inconsistency	74
2.3.2	Unavailability	77
2.3.3	Apparent Latency	78
2.3.4	The CAL Theorem	80
	<i>Sidebar: Real-Time Processes</i>	81
2.3.5	Self Updates	83
	<i>Sidebar: Constant Network Approximation</i>	84

This chapter draws from [Lee et al. \(2021\)](#) and hence has some overlap with subsequently published papers [Lee et al. \(2023a,b\)](#). This chapter reformulates the CAL theorem without processing offsets.

2.3.6	Max-Plus Formulation	88
2.4	Conclusion	99

2.1 Concurrent, Parallel, and Distributed

The word **concurrency** comes from the Latin *concurrere*, which means “run together.” The verb “to run” has an amazing variety of meanings. A human might run in a marathon, run before the wind (sail downwind), run a machine, run for office, run illegal drugs, or run a company. A rumor can run through a crowd, and a seminar can run for two days. A faucet might be left running, and a street might run all the way to the harbor. Colors can run when clothing is washed. Eggs can run six dollars a dozen. Machines can run even if they don’t move. It is really only this last sense of the word that we care about in this book.

In computing, when talking about concurrency, it is programs that are running. In this context, a **program** is a specification for operations to be performed by one or more computers. The program may execute entirely within a process or on several processes spread across several machines. A program is concurrent if it specifies operations that may be performed simultaneously. Whether they are actually performed simultaneously is irrelevant to whether it is a concurrent program.

Traditionally, a program gives an **algorithm**, a step-by-step procedure that operates on inputs and terminates to produce some desired output. An algorithm, in this sense, is not concurrent because it is a step-by-step procedure. It is, fundamentally, a *sequence* of discrete steps. A program that gives an algorithm is a **sequential program**. Concurrent programs will typically have sequential sub-tasks, but overall, they are not sequential.

Modern computers are able to do more than one thing at a time. Even at the lowest level, operations that appear to be performed in sequence are often performed simultaneously on distinct hardware. These operations are said to be performed **in parallel**. At this low level, however, the hardware is carefully designed to ensure that the steps *appear* to be carried out in sequence. For example, a processor may speculatively execute instructions after a branch instruction while it simultaneously evaluates the condition that determines whether the branch will be taken. This kind of concurrency is not relevant to us as programmers because

it is hidden from us. It just makes our programs run a bit faster, on average. But they are still sequential programs.

At a higher level, most modern microprocessors contain more than one **core**, where each core is a distinct computer able to execute its own programs. Such a machine is a **parallel computer** but, unlike the parallel low-level operations, no attempt is made to hide this parallelism from the programmer. If programmers wish to take advantage of the multiple available cores, they must write concurrent programs.

A **symmetric multicore** machine is a computer with multiple cores that share a main memory. The cores typically have a local **cache memory** that mirrors some portion of the main memory and is quicker to access but, conceptually, the machine is designed to look as if there is a single memory. Thus, programs running simultaneously on different cores see the same memory and can read and write the same memory locations. As discussed in the next chapter, this can lead to no end of headaches for programmers unless it is managed very carefully, for example by using the deterministic concurrency mechanisms covered in this book.

A **distributed system** is a collection of computers that do not directly share memory but instead communicate over networks. The computers may be housed in the same box, on the same rack, in the same datacenter, or in different cities. The essential feature that makes them “distributed” is that they do not share memory and must send each other messages over networks to communicate.

A modern operating system (**OS**), such as Linux, macOS, Windows, iOS, and Android, runs programs within a **process**. A process, conceptually, has its own memory which is invisible to other processes. Physically, the machine may be concurrently running several processes, and these will share the same physical memory, but, at least in theory, one process is not able to read or write memory locations that belong to another process. It is the job of the operating system, with some support from the hardware, to preserve this isolation. If a program consists of multiple processes, those processes must send each other messages,² as in a distributed system.

²Modern operating systems do contain mechanisms for processes that are running on the same multicore machine to define regions of memory that they share, but this just makes them more like threads than processes.

A process may include more than one **thread**. A thread is a step-by-step procedure, like an algorithm except that it may not terminate. Threads within a process share memory. They can read and write the same memory locations. In a modern programming language, threads read and write **shared variables**. The next chapter discusses some of the pitfalls of such a model.

A concurrent program may be written in many different ways. We will explore some of the more commonly used ways in the next chapter, but for now, it is sufficient to note that most of these ways of writing concurrent programs are nondeterministic. The program specifies more than one behavior in response to any given inputs. A key objective of this book is to provide the foundations for programming models that are both concurrent and deterministic.

A concurrent program may be executed on hardware in many different ways. Classically, computers were capable of only executing one step-by-step procedure at a time. **Multitasking** operating systems were invented to make a single such sequential machine appear to be multiple machines, each of which could execute a set of programs simultaneously. This was done by executing a thread for some period of time, then preempting it, storing its state in memory, restoring the state of another executing thread, and resuming with the execution of the other thread. This gives the illusion of executing in parallel without actually executing in parallel.

Concurrent programs are, in theory, agnostic to whether they are being executed in a multitasking manner or in parallel on multiple machines. The distinction affects their timing properties, but since their semantics usually ignores timing, it has no effect on their defined functional behaviors. With nondeterministic programs, however, timing will typically affect which of the nondeterministic choices are made. As a consequence, observed behavior will be different on a multitasking system vs. a multicore system. One of the goals of deterministic concurrency is to eliminate this discrepancy.

Modern multicore systems typically use the same multitasking model, except that now they have several cores available and hence can actually handle more than one thread in parallel. If the threads belong to the same process, then the shared memory of the multicore machine provides a convenient way for the threads within the same process to access shared variables.

Concurrent programs can also be mapped onto distributed machines. It is even possible to create the illusion of shared memory across distributed machines, but it is considerably more costly than on a symmetric multicore machine. As a consequence, programs intended to run on a distributed system are better written using some message-passing abstraction than a multithreading abstraction (see the next chapter for some such abstractions).

In summary, concurrency is a conceptual property of a program. If it does not specify a single sequence of operations, it is a concurrent program. A concurrent program may be executed on a single sequential machine, in which case its concurrent parts will be executed sequentially in some order determined by the OS or the hardware. The order in which these concurrent parts are executed is nondeterministic in the sense that it is not specified by the program, but the order may or may not affect the observable outcomes of the program, so the program itself may or may not be deterministic. Hence, concurrency and determinism are separate concepts. Programs may be concurrent and deterministic, concurrent and nondeterministic, sequential and deterministic, or sequential and nondeterministic.³ A program, concurrent or not, is deterministic only if *every* correct execution yields the same outcomes.

2.2 Consistency

Besides “running,” the other part of concurrency is “together.” If you are out for a run with a friend, you each modulate your speed to stay together, and you constantly adjust the distance between you. Maintaining togetherness requires coordination, and coordination requires communication. Togetherness is agreement, not on everything, but on some crucial things, such as speed. If your average speeds differ, even by a small amount, you won’t be together for long.

In the software world, the word **consistency** is often used to represent agreement. Ironically, however, in the software literature, there is little agreement on what consistency actually means (Helland, 2021). Nevertheless, among the reasons cited in Section 1.1.3 of the previous chapter that deterministic models are valuable, **consistency** looms large. There, consistency is described qualitatively as two

³The last of these is not as rare as you might think. Languages used for formal verification often model programs as finite state machines, which are sequential and can allow nondeterministic state transitions. See for example Chapter 3 of Lee and Seshia (2017).

agents agreeing on something. The speed with which you and your partner are running is just such a thing.

A naive view of consistency in software is that two concurrent parts of a program agree on the value of some shared information at every instant in time. If the shared information is stored in a single place in a shared memory, then such consistency is tautological. A memory cell is consistent with itself.

The problem with consistency arises when the shared information is stored in multiple memories. You and your running partner are adjusting your own pace of running and, by observing each other (a form of communication), continually adjusting your shared speed information. Under the naive notion of consistency, the multiple memories will have to have the same contents at all times. This means that if the contents change, all copies must change simultaneously. If you speed up your running, your partner will speed up precisely in synchrony with you rather than temporarily lagging behind and then catching up. However, such instantaneous communication is not possible, even in theory. Even more fundamentally, in a distributed system, we do not even have a well-defined notion of simultaneity. Time is a local phenomenon.

The fundamental problems with this naive view are that for any system spread across hardware resources, there is no shared instant in time, and communication takes time. This is not just a theoretical truth, following from the physical theory of relativity, but also a practical truth. No two clocks can be perfectly synchronized, and no communication can occur faster than the speed of light. Hence, there is no well-founded way to talk about two processes on two separate machines sharing an instant in time, much less agreeing on something at that instant in time. If it seems odd to invoke relativity at the scale of parallel machines, consider that at the speed of light, a nanosecond is about one foot long, and modern machines execute many instructions per nanosecond.

So what is consistency, if not agreement at an instant? That is the main topic of this chapter. We will see that consistency has more to do with the order of events than the instant at which they occur, and that achieving agreement takes time. Schooling fish agree on the direction of travel, not at an instant, but with a temporal dynamics that leads to undulations and fluctuations.

The relationship between consistency and determinism is similarly fraught. It should be clear that deterministic programs can help ensure consistency under

any reasonable definition of consistency. If two deterministic programs operate on the same inputs in the same way, then, by the definition of determinism, they will have the same outcomes. Therefore, they will agree on these outcomes. But when in time will they agree? And how do we ensure that they see the same inputs and operate on them in the same way? These are the questions we address here.

2.2.1 The CAP Theorem

In the year 2000, Eric Brewer gave a keynote talk ([Brewer, 2000](#)) at the Symposium on Principles of Distributed Computing (PODC) in which he introduced the **CAP** theorem, which states that you can have only two of the following three properties in a distributed system:

- **Consistency:** distributed components agree on the value of shared state;
- **Availability:** ability to respond to user requests for that value; and
- **Partition tolerance:** the ability to keep operating when communication failures divide the system into islands.

This keynote is credited by many in the distributed computing community with opening up the field, enabling innovative approaches that offer differing tradeoffs between these properties.

[Gilbert and Lynch \(2002\)](#) proved a couple of variants of this theorem, one rather strong result for “asynchronous networks” (which have no clocks) ([Lynch, 1996](#), chapter 8) and one weaker result for “partially synchronous networks” (which have unsynchronized clocks that measure the passage of time at the same rate). These results reinforced Brewer’s original interpretation that you could obtain any two of the three properties, but not all three.

[Brewer \(2012\)](#) provided a retrospective in which he observes that the tradeoff is more subtle than this. First, the “P” property is not really one you can trade off against the others. You cannot choose whether to have network partitions because these are the result of system failures. Brewer clarified the design problem as one of how to trade off consistency against availability when network partitions occur.

In the same year, [Abadi \(2012\)](#) reacted that CAP is irrelevant when there are no network partitions, but [Brewer \(2012\)](#) points out that network partitions are not a binary property; all networks have latency, and a complete communication failure is just the limiting case when the latency goes to infinity. [Abadi \(2012\)](#), in fact, emphasizes the tradeoff in database design between consistency and latency experienced by the user, arguing that it has had more impact on actual database system designs than CAP. He then observes that “availability and latency are arguably the same thing” and calls for unifying the consistency/latency tradeoff with the CAP tradeoffs.

There are (at least) two latencies of interest: communication latency and user latency. Whereas communication latency is the time that it takes messages to propagate from one node to another, user latency, which we will refer to as unavailability, is the time it takes to respond to a user. Specifically, unavailability is the physical time that elapses between when a system first receives a request from a user and when the system sends a response to the user.

Network partitioning is the limiting case of communication latency. When the communication latency goes to infinity, the network is partitioned. Hence, we derive the **CAL** theorem, replacing “Partition tolerance” with “Latency.” This theorem will establish a numerical relationship between inconsistency, unavailability, and communication latency, thereby unifying Brewer’s and Abadi’s positions and having the CAP theorem as a special case. Our resulting formulas make clear the tradeoff between these quantities and enable a range of designs that vary in their emphasis on one or more of these quantities.

To establish the relationship, we have to carefully define the three variables: **inconsistency**, **unavailability**, and network **latency**. Most designers have strong intuitive understandings of the meanings of these terms, but there are pitfalls. For example, an intuitive statement about a consistent view of a shared variable x goes something like this: If component 1 writes to x “before” component 2 reads the value of x , then component 2 will read the updated value. Another such statement is that if component 1 writes to x “before” component 2 also writes to x , then all observers of x will “subsequently” see the value written by component 2. These statements are difficult to formalize, however, because the words “before” and “subsequently” are problematic. The order in which two physically separated events occur may be not only unknown, but also, at a fundamental level, **unknowable**. From special relativity, the order in which physically separated events occur can depend on the frame of reference. In

modern distributed systems—some of which span the globe and all of which have communication that is limited by the speed of light—if many updates and reads of shared variables are occurring, some of those are sure to have ambiguous ordering. So we will have to do better than this. We need a more solid foundation for the notion of consistency.

[Yu and Vahdat \(2006\)](#) have also quantified availability and consistency and shown a tradeoff between them. Their quantifications, however, are in terms of fractions of satisfied accesses (availability) and fractions of out-of-order writes (inconsistency), and they show that finding the availability as a function of consistency is NP-hard. In contrast, we define these quantities as time intervals and derive a strikingly simple relationship.

2.2.2 Consistency and Time

Consider a shared variable x that is read and written to by various sequential processes running on nodes distributed across a network. An intuitive form of consistency requires that all nodes agree on the order of operations on x (reads and writes). [Gilbert and Lynch \(2002\)](#) say “The most natural way of formalizing the idea of a consistent service is as an atomic data object,” where **atomic** comes from [Lamport \(1986\)](#). Atomic here is interchangeable with the term **linearizable** from [Herlihy and Wing \(1990\)](#). This interpretation defines “consistent” to mean that all observers see the same sequence of atomic operations (reads and writes) on the data object.

For most applications, however, this is too strong. It is sufficient to agree on the order in which *writes* occur, and to agree on the order in which reads occur relative to those writes, but otherwise, there is usually no need to agree on the order in which the reads occur relative to each other. But even this weaker agreement may be too strong. Consider two independent writes to a shared variable x that occur nearly simultaneously on two distinct nodes. Depending on how these writes are to be merged, it may not be necessary to define the order in which they occur.

There is also the question of what we mean by **agreement**. When should the nodes agree? And what is agreement?

These questions have as a backdrop some notion of physical time. Intuitively, physical time has to play a major role. It would be strange to design a system

that considers a write to x that occurs three weeks after another write to x to have occurred “before.” But physical time is a slippery, poorly understood thing (Rovelli, 2018; Muller, 2016). Newtonian time, with all its beauty and simplicity, has limited applicability, and there is no perfect measurement of it. Fundamentally, the ordering and agreement that we need are logical (or semantic) properties, not physical properties. But to make them useful in applications, we will need physical realizations that are reasonably faithful to our semantic models using only practical, realizable measurements of physical time. We require, therefore, formal semantic models and physical realizations that, with high confidence, conform to these models.

The events we will consider in this chapter are reading and writing values to or from some local copy of a shared variable x . The value of x may be a text, a number, or some data structure, and the semantics of a “write” will matter. In general, a “write” performs a **merge operation** to update the old value using the new value. For example, x may be a number and a “write” may **replace** the old value or simply **add** to it. The latter merge operation is associative and commutative, assuming perfect arithmetic, where there are no rounding errors, and overflows wrap. The replacement merge operation is associative but not commutative. If x is a text, the “write” may **append** to the text, again giving a merge operation that is associative but not commutative.

If the merge operation is associative and commutative, then we have two of the three properties of the **ACID 2.0** principle, proposed by Helland and Campbell (2009). The letters stand for Associative, Commutative, Idempotent, and Distributed. **Idempotence** means that operations that are applied more than once have the same effect as operations that are applied exactly once. We will assume that idempotence is realized at a layer below, not in the application code; for example, using **TCP**, we can assume a model where messages are delivered reliably exactly once. Hence, the application code can assume that writes are applied exactly once (in the absence of failures, at which point the *model* becomes invalid).

If we are to have availability in the presence of network partitioning, then some amount of **replication** is necessary (Abadi, 2012). That is, the value of x will need to be stored at multiple nodes. It is this replication that creates the consistency problem, but without it, we cannot assure availability. We will assume in this chapter that reads and writes to local copies of x can occur at any of the nodes, and that we want these values of x to agree within some bounded amount of time.

2.2.3 Causal Consistency

An elegant formal model of consistency, called **causal consistency**, is analyzed in depth by Schwarz and Mattern (1994). They define a causality relation, written $e_1 \rightarrow e_2$, between events e_1 and e_2 to mean that e_1 can causally affect e_2 . The phrase “causally affect” is rather difficult to pin down (see Chapter 11 of Lee (2020) for the subtleties around the notion of causation). Intuitively, $e_1 \rightarrow e_2$ means e_2 cannot behave as if e_1 had not occurred. Put another way, if the effect of an event is reflected in the state of a local replica of a variable x , then any cause of the event must also be reflected. Put yet another way, an observer must never observe an effect before observing its cause.

For example, if event e_1 writes $x = 1$ and e_2 sends the value of x to another process, then $e_1 \rightarrow e_2$ means that whatever value e_2 sends is either the $x = 1$ value written by e_1 or some subsequently written value.

Following Schwarz and Mattern, we will impose a bit more structure on our distributed applications.

Definition 2.1. A *sequential process* is a totally ordered sequence of events occurring on a single machine. A sequential process, therefore, is the **trace** (the record of events) of a single *thread* of execution.

We will simplify our model to sequential processes that maintain a local replica of just one shared variable x . Our abstraction will consider only processes that are sequences of events of the following types:

Definition 2.2. *Event types:*

1. w_x : **write event**: Merge a value with the local replica of x .
2. d_x : **read event**: Read the value of the local replica of x .
3. s_x : **send event**: Send the value of the local replica of x to some set of other processes.
4. r_x : **receive event**: Receive a new value for x and merge it with the local replica.

When the type of event is important, we will use one of the four symbols above to denote the event. Otherwise, we will use the symbol e . If the variable x is clear from context, we will omit the subscript.

Typical **merge operations** are **replace**, **append**, and **add**, but the merge operation can be arbitrarily sophisticated. To maintain consistency, we will find it necessary for all replicas of x to use the same merge operation.

Definition 2.3. The *causality relation* of [Schwarz and Mattern \(1994\)](#) is the smallest transitive relation such that $e_1 \rightarrow e_2$ if e_1 precedes e_2 in a process, or e_1 is the sending of a value in one process (event type s_x) and e_2 is the receipt of the value in another process (event type r_x). If neither $e_1 \rightarrow e_2$ nor $e_2 \rightarrow e_1$ holds, then we write $e_1 || e_2$ or $e_2 || e_1$ and say that e_1 and e_2 are *incomparable*.

This causality relation is identical to the **happened before** relation of [Lamport \(1978\)](#), but Schwarz and Mattern prefer the term “causality relation” because even if e_1 occurs unambiguously earlier than e_2 in physical time, they may nevertheless be incomparable, $e_1 || e_2$.

The causality relation is a **strict partial order** (see Appendix A). Schwarz and Mattern use their causality relation to define a “consistent global snapshot” of a distributed computation to be a subset S of all the events E in the execution that is a **downset**, meaning that if $e' \in S$ and $e \rightarrow e'$, then $e \in S$. A consistent global snapshot had been previously called a “consistent cut” by [Mattern \(1988\)](#).

Example 2.1: Consider a shared online bulletin board, like a Facebook Timeline. Consider the following events, adapted from [Bailis et al. \(2013\)](#):

- w_1 : Joe posts a picture of Sally at a recent party by writing to a local copy.
- w_2 : Sally posts that her son Billy is missing, writing to a local copy.
- s_1 : Sally’s message is sent to Joe’s process.
- r_1 : Joe’s machine receives the message and updates his local copy.
- w_3 : Sally posts that her son has been found (on the local copy).
- s_2 : Sally’s message is sent to Joe’s process.
- r_2 : Joe’s machine receives the second message and updates his local copy.




Figure 2.1: Causal ordering of events in a bulletin board example.




Figure 2.2: Observer of the exchange between Joe and Sally.

- d_1 : Joe reads Sally's messages.
- w_4 : Joe posts "That's good news, a relief."

The immediate causal relations are depicted by arrows in Figure 2.1. The complete causal relation is the transitive closure of the depicted causal relations.

Suppose that in this application the merge operation is **append**. Each write to the local copy of the bulletin board simply appends the message to the end of the board. This operation is associative but not commutative.

After the writes w_1 and w_2 , Sally and Joe have inconsistent views of the shared variable. A third observer, as indicated in Figure 2.2, may have yet a third inconsistent view. In the figure, after the event r_4 , the p_3 observer sees both messages, albeit in an arbitrary order that has been determined by the happenstances of communication latency. However, in this case, there is no violation of *causal* consistency because $w_1 \parallel w_2$. Semantically, Joe's and Sally's first writes are unordered, and different observers may see none, one, or both messages, and in the latter case, they may see them in either order. Causal consistency tolerates this form of nondeterminism.

A worse problem, however, is illustrated in Figure 2.2. After event r_5 , the p_3 observer has seen Sally's first post and Joe's second ("That's good news") but has not seen Sally's second post. At the read event d_2 , this observer will think that Joe has responded "that's good news" to Sally's message that Billy is missing! This observer would have to conclude that Joe really does not like Sally's son. After event r_5 , the state of p_3 's local replica is *causally inconsistent* because p_3 sees the effects of w_4 , but not the effects of w_3 , and $w_3 \rightarrow w_4$.

Leveraging the results of [Charron-Bost \(1991\)](#), Schwarz and Mattern argue that there is no simple solution that ensures causal consistency that does not also impose additional unnecessary constraints. They give a non-simple solution using **vector clocks** ([Mattern, 1988](#)). It is non-simple in that it requires, in the worst case, message sizes that are order N , where N is the number of communicating processes. Chapter 8 discusses an alternative solution using timestamps.

2.2.4 Causal Consistency is Insufficient

In practice, causal consistency is necessary but not sufficient for most applications. It may be perfectly reasonable to impose additional constraints. In the bulletin board application, for example, it may not be acceptable for different observers to see Sally and Joe's first posts in arbitrary order. One common way to prevent that is to timestamp the posts using some clock measuring physical time on the machine where the post originates. All observers could then arrange whatever posts they see in timestamp order, breaking ties with some consistent algorithm (e.g. alphabetical ordering). These timestamps give a semantics to the problematic "before" and "subsequently" that we encountered in Section 2.2.1. A general form of this solution is often called "last writer wins," but that phrase only makes sense if the merge policy is **replace**. So we will call the policy **timestamp-ordered**.

Some care is required here because unsynchronized clocks could result in rather unexpected orderings. Hence, this solution assumes some measure of clock synchronization, a topic we discuss further in Section 2.2.5. Nevertheless, this solution is fairly common, found in many distributed database systems, such as Google Spanner ([Corbett et al., 2012, 2013](#)).

By itself, however, time-stamping events is not enough to ensure causal consistency. After r_5 in Fig. 2.2, w_3 will still be missing even after the posts seen are arranged in timestamp order. Ensuring causal consistency in a distributed system will require more work. All of the techniques discussed in Part II offer solutions.

An additional requirement in many applications is **eventual consistency**. Eventual consistency, informally, means that if writes stop occurring at some point in time, then, eventually, all replicas of x will have the same value. Fig. 2.2 would become eventually consistent if we add send-receive pairs so that p_1 sends its writes to p_2 . However, even though it will be eventually consistent, it will still not be causally consistent unless something is done about the mis-ordered r_5 and r_6 .

We will henceforth in this chapter assume that our distributed system is **causally** and **eventually consistent** with **timestamps**, which we shorten to **CET**. The first and last properties are absolute in that they hold at all times during an execution. Eventual consistency, however, is rather odd in that an eventually consistent system may, in fact, never be consistent! It may become consistent only if all processes stop writing to shared variables, and even then, it may take an unbounded amount of time to become consistent.

We will formally define a time interval $\bar{C}$ (a timestamp interval) over which a system is allowed to be inconsistent. $\bar{C}$ will become a measure of **inconsistency** in the system, and we will give relations between $\bar{C}$, a measure $\bar{A}$ of **unavailability**, and a measure $\mathcal{L}$ of (apparent) **communication latency**. If $\bar{C}$ is bounded, then the system is eventually consistent. If $\bar{C}$ is zero, we call the system **strongly consistent**. A strongly consistent system is obviously eventually consistent.

2.2.5 Using Physical Time to Assign Logical Times

To replace the naive “earlier” with something realizable, we borrow an idea from [Lamport \(1984\)](#), who proposed that each update of a shared variable be assigned a timestamp. The timestamp is based on a local clock, and the local clocks throughout the system are sufficiently synchronized to provide an application-dependent reasonable notion of “before” or “earlier.” Once a timestamp is assigned, it is treated as a **logical time** and defines the semantic order of events. For an explanation of how clock synchronization works, see the sidebar on page 69.

In the bulletin board example, the timestamps could be based on **coordinated universal time (UTC)** and synchronized via the **network time protocol (NTP)** ([Mills, 2003, 2006](#)), which is designed to synchronize computer clocks to within a few milliseconds of UTC. NTP is ubiquitous in modern computers, so this is not likely to present a practical problem, and much higher precision clock synchronization mechanisms are available when needed by an application (see sidebar on page 71). If a post in London gets a timestamp 10 milliseconds less than a post in Singapore, the London post will be deemed to have occurred before. Whether it physically occurred before the post in Singapore is irrelevant.⁴

⁴The physical order of these two events is unknowable. At the speed of light, Singapore is about 32 milliseconds away from London, which means that even with perfect clock synchronization,

Given such timestamped events, the behavior of the distributed system can then be defined by the numerical order of those timestamps. Once the timestamps are assigned, their numerical order is a semantic property, and the goal of a **CET** system design is to ensure that every component sees events in timestamp order. This is the semantics adopted by Google Spanner (Corbett et al., 2012, 2013), a globally distributed replicated database operating today.

Because the timestamps are numbers, there is always the possibility of two events having the same timestamp. These will be deemed to be **logically simultaneous**. How to handle logically simultaneous events is application dependent. If the merge operation is associative and commutative, then the simultaneous operations can be applied in different orders without adverse consequences. If, further, no user can observe the intermediate results of these operations, then the whole system will be strongly consistent. Without associative and commutative operations, however, the order of operations will need to be the same at all nodes in order to achieve consistency. For simultaneous events, consistent ordering requires a consistent mechanism for assigning distinct priorities to each event.

Note that even sloppy clock synchronization can be useful because it defines a semantic ordering. In the bulletin board application, does it really matter if the clocks are off even by a few seconds (which would be unusually poor performance for NTP)? We will see that clock synchronization error is indistinguishable from communication latency, and so the sum of the two quantities is the only meaningful measure. Consequently, a large clock synchronization error coupled with a requirement for strong consistency will result in high unavailability. This would be true with any mechanism for assigning a total order to distributed updates because any such order, in effect, assigns some kind of logical “time” to each event. The mere act of putting them in order is the same as assigning them a time.

2.2.6 Tags and Physical Time

We are interested in times of events and time intervals between events, but we have to be careful because these cannot simply be Newtonian time. Newtonian time is inaccessible to distributed software. We will instead use two distinct

these two events are outside each other’s light cones. This implies that the true physical order of the two events will be different in different frames of reference.

Sidebar: Clock Synchronization

The figure at the right shows how PTP works. The clock master A initiates an exchange of messages. The first message is sent at time t_1 (by the master clock) and contains the value of t_1 . That message is received by the slave B at time t_2 (by the master clock), but since the slave does not have access to the master clock, the slave records the time t'_2 that it receives the message according to its own clock. If its clock is off by e vs. the master clock, then $t'_2 = t_2 + e$. The slave responds by sending a message back to the master at time t'_3 according to its clock, or t_3 according to the master's clock, so

$$t'_3 = t_3 + e.$$


The master receives this second message at t_4 , and replies with a third message containing the value of t_4 . The slave now has t_1, t'_2, t'_3 , and t_4 . Now notice that the round trip communication latency (the time that a message from A to B and a reply message spend in transport) is

$$r = (t_2 - t_1) + (t_4 - t_3) = (t_4 - t_1) - (t_3 - t_2).$$

At B , this value can be calculated even though t_2 and t_3 are not known, because $(t_3 - t_2) = (t'_3 - t'_2)$, and slave B has t'_2 , and t'_3 . If the communication latencies are symmetric, then the one way latency is $r/2$. To correct the clock of B , we simply need to estimate e . This will tell us whether the clock is ahead or behind, so we can slow it down or speed it up, respectively. If the communication channel has symmetric delays (i.e. $t_2 - t_1 = t_4 - t_3$), then a very good estimate is given by

$$\tilde{e} = t'_2 - t_1 = r/2.$$

If the communication latency is exactly symmetric, then $\tilde{e} = e$, the exact clock error. B can now adjust its local clock by $\tilde{e}$.



notions of time, logical and physical. A **physical time** $T \in \mathbb{T}$ is an imperfect measurement of time taken from some physical clock somewhere in the system.

Let the set $\mathbb{T}$ contain all the possible times that a physical clock can report. We assume that $\mathbb{T}$ is totally ordered and includes two special members: $\infty \in \mathbb{T}$ is larger than any time any clock can report, and $-\infty \in \mathbb{T}$ is smaller than any time any clock can report. For example, $\mathbb{T}$ could be the set of integers $\mathbb{Z}$ augmented with the two infinite members.

Given any $T_1, T_2 \in \mathbb{T}$, the **physical time interval** (or just **time interval** if there is no ambiguity) between the two times is written $i = T_1 - T_2$. Time intervals are assumed to be members of a group $\mathbb{I}$ with a largest member ∞ and smallest member $-\infty$ and a commutative and associative addition operation. For example, $\mathbb{I}$ could again be $\mathbb{Z}$ augmented with the two infinite members. Addition involving the infinite members behaves in the expected way in that for any $i \in \mathbb{I} \setminus \{\infty, -\infty\}$,

$$\begin{aligned} i + \infty &= \infty \\ i + (-\infty) &= -\infty \end{aligned}$$

We also assume that addition of infinite intervals saturates, as in

$$\begin{aligned} \infty + \infty &= \infty \\ (-\infty) + (-\infty) &= -\infty \\ \infty + (-\infty) &\text{ is undefined} \end{aligned}$$

Note that we use the same symbols ∞ and $-\infty$ for the special members of both the set of physical times $\mathbb{T}$ and the set of intervals $\mathbb{I}$. We hope this will not create confusion.

Intervals can be added to a physical time value, and we assume that this addition is associative. I.e., for any $T \in \mathbb{T}$ and any $i_1, i_2 \in \mathbb{I}$,

$$T + (i_1 + i_2) = (T + i_1) + i_2 \in \mathbb{T}. \quad (2.1)$$

Addition of infinite intervals to a time value saturates in a manner similar to addition of infinite intervals.

These idealized requirements for physical times and time intervals can be efficiently approximated in practical implementations. First, it is convenient to have

the set $\mathbb{T}$ represent a common definition of physical time, such as [UTC](#) because, otherwise, comparisons between times will not correlate with physical reality. UTC, however, has the significant complexity of **leap seconds**, which are seconds inserted or removed from a measurement of time to keep time aligned with the rotation of the earth.

Sidebar: PTP and TSN

Many distributed systems rely on a common notion of time. On local area networks and in data centers, clock synchronization with much higher precision than [NTP](#) is available with **PTP (precision time protocol)** ([Eidson, 2006](#); [Eidson and Stanton, 2015](#)), a centerpiece of the **TSN (time sensitive network)** family of protocols being developed by the IEEE 802.1 working group ([Lo Bello and Steiner, 2019](#)). The clock synchronization standard part of TSN, IEEE 802.1AS-2020, a subset of the older standard IEEE 1588, is widely supported in networking hardware, even if it has not yet appeared in standard programming APIs (application programming interfaces). This standard is capable of synchronizing clocks to within less than a nanosecond on local area networks, and on the order of microseconds on wide area networks. The White Rabbit project at CERN, for example, claims to be able to synchronize clocks on a network spanning several kilometers to under 100 picoseconds ([Gaderer et al., 2009](#)). This means that if you simultaneously ask two clocks separated by, say, 10 kilometers of networking cable what time it is, their response will differ by less than 100 picoseconds.

GPS (the global positioning system) provides a complementary alternative that is sometimes combined with PTP. With GPS, clocks can be synchronized to sub-nanosecond precision globally, in principle ([Jefferson et al., 1996](#)). Google Spanner synchronizes master clocks at data centers using GPS combined with atomic clocks, and then synchronizes clocks within the data center using PTP ([Corbett et al., 2012](#)). GPS, however, is not always available to systems (particularly indoor systems), and it is vulnerable to spoofing and jamming. In addition, avoiding brittleness in such systems, where the source of the clock becomes a single point of failure that can bring down the entire system, may be expensive and a significant engineering challenge ([Kopetz, 1997](#); [Kopetz and Bauer, 2003](#)).

Example 2.2: One possible concrete representation for time is **Unix time**, where $T \in \mathbb{T}$ represents the number of seconds (ignoring leap seconds) that have elapsed since the **epoch**, defined to be midnight, January 1, 1970, Greenwich mean time. If 32-bit signed integers are used (a common choice), then the minimum representable time is sometime on Friday, December 13, 1901, and the maximum representable time is on Tuesday, January 19, 2038. To conform with the abstraction given above, the largest and smallest times can represent ∞ and $-\infty$, respectively, in which case addition and subtraction should be implemented to respect the above saturation requirements. The time intervals $\mathbb{I}$ could be represented with the same set of 32-bit integers.

Note that the set of 32-bit integers is not the same as the set $\mathbb{Z}$ of integers because it is finite. As a consequence, addition can overflow. If such overflow saturates at ∞ or $-\infty$, addition is no longer associative. For example, $T + (i_1 + i_2)$ may not overflow while $(T + i_1) + i_2$ does overflow. Such overflow will become a problem in the year 2038, creating what is known as the **year 2038 problem**.

If 64-bit signed integers are used instead, then the time interval extends 292.3 billion years in both directions, which is unlikely to create a problem.

Example 2.3: Another representation for time common in Unix-like systems represents time as a C struct called **timespec** consisting of two integers, one storing the Unix time as in the previous example, and another storing fractional seconds counted in nanoseconds. For the second number, a 32-bit signed or unsigned integer is a suitable choice because it does not overflow within the one billion nanoseconds in a second.

Example 2.4: A third representation for time that is convenient for timestamps and efficient on 64-bit architectures is to use a 64-bit signed or unsigned integer to represent the number of nanoseconds elapsed since the **epoch**. This is the representation used in the Ptolemy II system (Ptolemaeus, 2014) and the Lingua Franca coordination language (Lohstroh et al., 2021), for example. Using a signed 64-bit number, overflow occurs approximately 292 years after the epoch or sometime in the year 2262.

Example 2.5: A commonly used representation for time, particularly in simulation systems, is floating-point numbers. This is a poor choice because, even in the absence of overflow, addition is not associative and the **resolution** (the difference between two representable times) is not independent of the magnitude (Broman et al., 2015). Later times have less resolution than earlier times.

Logical time may or may not be represented identically to physical time. Let G represent the set of possible logical time values. G is required to be a **totally ordered set**.

In this book, logical time will take many forms. In all of these forms, logical time can be represented as an element $g \in G$ for some totally ordered set G . When g resembles a **physical time** value, we call it a **timestamp**. More generally, following Lee and Sangiovanni-Vincentelli (1998), we call g a **tag**.

In a **CET** system, each event is associated with a tag $g \in G$. From the perspective of any component of a distributed system, the order in which events occur is defined by the order of their tags. If two distinct events have the same tag, they are **logically simultaneous**. We assume the tag set G has an element ∞ that is larger than any other tag and another $-\infty$ that is smaller than any other tag.

Example 2.6: In Ptolemy II and Lingua Franca, $G = \mathbb{T} \times \mathbb{U}$, where $\mathbb{U}$ is the set of 32-bit unsigned integers representing the **microstep** of a **superdense time** system (Maler et al., 1992; Cataldo et al., 2006; Cremona et al., 2017).

We use the term **tag** instead of timestamp to allow for such a richer model of logical time. For the purposes of this chapter, however, microsteps will not matter, and hence you can think of a tag as a timestamp and ignore the microstep. Microsteps will become important in later chapters. We will consistently denote tags with a lower case $g \in G$ and measurements of physical time $T \in \mathbb{T}$ with upper case.

We will need operations that combine tags and physical times. To do this, we assume a monotonically nondecreasing (order-preserving) function $\mathcal{T}: G \rightarrow \mathbb{T}$ that gives a **physical time interpretation** to any tag. For any tag g , we call $\mathcal{T}(g)$ its **timestamp**.

Example 2.7: In Ptolemy II and Lingua Franca, for any tag $g = (t, m) \in \mathbb{G}$, $\mathcal{T}(g) = t$. Hence, to get a timestamp from a tag, you just have to ignore the microstep.

The set $\mathbb{G}$ also includes infinite elements such that $\mathcal{T}(\infty_{\mathbb{G}}) = \infty_{\mathbb{T}}$ and $\mathcal{T}(-\infty_{\mathbb{G}}) = -\infty_{\mathbb{T}}$, where the subscripts disambiguate which infinity we are referring to.

In a **CET** system, an external input from outside the system, such as a user input or query that wishes to read the value of a shared quantity x , will be assigned a tag g such that $\mathcal{T}(g) = T$, where T is a measurement of physical time taken from the local clock where the input first enters the system. If microsteps are used, this tag is normally given microstep 0, $g = (T, 0)$.

2.3 The CAL Theorem⁵

We can now (finally!) define the terms in the CAL theorem, **inconsistency**, **unavailability**, and (apparent) **latency**. As we will see, communication latency is only one of several factors that can affect inconsistency and unavailability, so we drop the reference to the “network” and simply consider latency, trusting that the reader will not apply the word “latency” to unavailability, as [Abadi \(2012\)](#) does. Specifically, we will see that communication latency is indistinguishable from clock synchronization error, and that computation times have the same detrimental effects as communication latency.

2.3.1 Inconsistency

In this section, we will define **inconsistency**, $\bar{C}$, to be a non-negative time interval in $\mathbb{I}$ that has the value zero for strongly consistent systems. When it has a value greater than zero, it will represent a **logical time** interval over which components in a distributed system can disagree about the values assigned to a shared data object. If this number is bounded, then we have **eventual consistency**.

⁵This section draws from [Lee et al. \(2021\)](#) and hence has some overlap with subsequently published papers ([Lee et al., 2023a,b](#)). This section reformulates the CAL theorem without processing offsets.

Assume we are given a **trace** of the execution of N **sequential processes** in a distributed system, where the processes consist of the **event types** in Definition 2.2. Assume the k -th event in each process has a tag g_k and a physical time T_k . The physical time T_k is the reading on a local clock at the time where the event starts being processed.

Requirement 2.1. *We require that events in a process have nondecreasing tags and strictly increasing physical times. That is, if g_k is the tag and T_k is the physical time of the k -th event, then $g_k \leq g_{k+1}$ and $T_k < T_{k+1}$.*

The constraint that a process have nondecreasing tags is central to maintaining causal consistency, but it comes with significant costs. Chapters 4 through 8 all provide mechanisms that can associate such nondecreasing tags with such events. For now, we just assume it is enforced somehow.

The constraint that physical times be increasing seems natural for a sequential process, but it may not be trivial to ensure. The physical time of an event is the reading of a local clock when the processing of the event starts. But local clocks are not always assured of increasing monotonically. A clock that is synchronized using **NTP**, for example, can move backward when updated. Moreover, a local clock will have some granularity, so it may not be *strictly* increasing. It could report the same time for two successive events.

Our formalism, however, requires that these physical times be strictly increasing. This can be ensured in software by providing an intermediate API function for accessing a local physical clock that always returns a value strictly greater than any previous value it has returned. Such an interface can be found in *Lingua Franca*, for example.⁶

A **write event** is an update to a local copy of a shared variable. A **read event** is a reading of the local copy of a shared variable. A **send event** is the launching into the network of a notification of an updated value. A **receive event** is the receipt of such a notification; we assume that every receive event also updates the local copy of the shared variable. All updates are mediated by some application-specific **merge operation** that combines the update value with the previous value in some specified way.

⁶<https://lf-lang.org>

Within each process, every read or send event with tag g yields the value of a shared variable x that was merged into the local copy of x by the write or receive event in the same process with the largest tag g' where $g' \leq g$. If $g' = g$, we require that $T' < T$, where T' is the physical time of the write or receive event and T is the physical time of the read event. This requirement ensures that a read event reads a value that was written at an earlier physical time. This requirement is the principal reason for the requirement that physical times in a process be monotonically increasing.

We have additional requirements on tags:

Requirement 2.2.

1. Every write or receive event in a process with the same tag as a read or send event must precede that read or send event in the process.
2. Every receive event on process i with tag g_i has a corresponding send event on some process j with tag $g_j \leq g_i$.

Note that the physical time of a receive event relative to the physical time of its corresponding send event is unconstrained. This may be confusing because it seems to allow the physical time of a receive event to be less than that of the corresponding send event. However, this can occur because the physical times of these two events come from different physical clocks. These clocks are not (and cannot be) perfectly synchronized.

Definition 2.4. For each receive event on process i with tag g_i , there is a send event that sends this update on some process j . The **minimum inconsistency** $\bar{C}_{ij} \in \mathbb{I}$ from j to i for x is defined to be

$$\bar{C}_{ij} = \min(\mathcal{T}(g_i) - \mathcal{T}(g_j)), \quad (2.2)$$

where the minimization is over all receive events on process i that receive from process j . If there is no such send event, let $\bar{C}_{ij} = \infty$.

If there is a send event on j but no corresponding receive event on i , then $\bar{C}_{ij}$ is defined as above, where g_i is the tag of the receive event that i would have processed had the communication not failed.

With requirement 2.2(2), it is clear that $\bar{C}_{ij} \geq 0$. If $\bar{C} = 0$, we have **strong consistency**. We will see that this strong consistency comes at a price in availability, and that network failures can result in unbounded unavailability. If $\bar{C}$ is bounded, we have **eventual consistency**, and $\bar{C}$ puts a lower bound on “eventual.”

If there is no send from i to j , we interpret this to mean that process i , if it has made any updates to x , has chosen not to share those updates with j . This means that the distributed program explicitly tolerates unbounded inconsistency between these two nodes. This motivates the choice $\bar{C}_{ij} = \infty$.

For the second special case, where there is a send event on j but no corresponding receive event on i , we assume that the intended tag to be seen by i is specified by the program and is therefore known.

We are interested in the *minimum* inconsistency because, as we will see, it is the minimum that will determine the unavailability.

2.3.2 Unavailability

In database systems, unavailability is a measure of the time it takes for a system to respond to user requests (Kleppmann, 2015). If it takes a long time (or it never responds), then the system is unavailable, whereas if responses are instantaneous, then the system is highly available. A robust database system will duplicate records, so x in our system could be a record where each node in the system maintains a local copy. Unavailability, therefore, is the time it takes to read or update x .

First, note that, from Requirements 2.1 and 2.2, the physical time at which a local copy of a shared variable x can be read for tag g cannot be earlier than the physical time at which it is updated with tags g or less. Hence, we can focus on delays incurred performing those updates.

Definition 2.5. For each write or receive event on process i , let g_i be its tag and T_i be the physical time at which it is processed. The **unavailability** $\bar{A}_i \in \mathbb{I}$ at process i is defined to be

$$\bar{A}_i = \max(T_i - \mathcal{T}(g_i)), \quad (2.3)$$

where the maximization is over all write or receive events on process i . If there are no such write or receive events on process i , then $\bar{A}_i = -\infty$. If there is a send event on some process j that is never received by i , then we take the physical time at which it is received to be $T_i = \infty$, giving $\bar{A}_i = \infty$.

The $-\infty$ is because a process with no write or receive events consists only of read and send events, and there is no lower bound on the physical time at which these

read or send events can occur. This is not actually an interesting case because if there are no write or receive events updating x on process i then x is a constant and is always available. For more interesting scenarios, we can often assume that $\mathcal{T}(g_i) \leq T_i$ because g_i is typically assigned using the local physical clock. In such scenarios, $\bar{A}_i \geq 0$. We will have more to say about this later.

Definition 2.6. *The **lag** of an event with tag g that is processed at physical time T is defined to be $T - \mathcal{T}(g)$.*

With this definition, the unavailability is the maximum lag of write or receive events on a process.

2.3.3 Apparent Latency

When a send from process j to process i occurs, some time will elapse before a corresponding receive event on process i performs a corresponding update to its local copy of the shared variable. The worst-case delay for this update is the **apparent latency**, defined as follows:

Definition 2.7. *Let T_i be the physical time at which process i processes a receive event with tag g_i from process j . The **apparent latency** $\mathcal{L}_{ij} \in \mathbb{I}$ for communication from j to i is*

$$\mathcal{L}_{ij} = \max(T_i - \mathcal{T}(g_i)), \quad (2.4)$$

where maximization is over all such receive events in process i . If there is a send event on j to process i but no corresponding receive event on i , then $\mathcal{L}_{ij} = \infty$ (a message was sent but never received). If there are no send or receive events at all from j to i , then $\mathcal{L}_{ij} = -\infty$ (there is no communication).

The apparent latency $\mathcal{L}_{ij}$ is the maximum **lag** among receive events on process i that receive messages sent from process j . Hence, we have $\bar{A}_i \geq \mathcal{L}_{ij}$ for every j .

Note that this apparent latency is trivially easy to measure, unlike an actual physical one-way latency. Note further that the tag g_j of the sent event may differ from the tag g_i of the received event. This difference has the **inconsistency** $\bar{C}_{ij}$ as a lower bound. We will say more below about the case where $i = j$.

In Definition 2.7, apparent latency is defined as a maximum over a trace. Consider a *single* communication from j to i (rather than an entire trace). Let g_j and

T_j be the tag and physical time of the send event, and let g_i and T_i be the tag and physical time of the receive event. The apparent latency for this one communication is simply $T_i - \mathcal{T}(g_i)$. This latency can be modeled as a combination of five components,

$$T_i - \mathcal{T}(g_i) = (T_j - \mathcal{T}(g_j)) + L_{ij} + E_{ij} + X_{ij} - (\mathcal{T}(g_i) - \mathcal{T}(g_j)) \quad (2.5)$$

The first term is the unavailability $\bar{A}_j$ at j for this one communication. Its value is determined by the fact that the message cannot be sent before the value of x at tag g_j can be read, and it can't be read until the previous write or receive has happened.

The last term is the **inconsistency** $\bar{C}_{ij}$ for this particular communication. A larger value here expresses a tolerance in the program for a larger inconsistency between j and i . It offsets the other latencies; from the perspective of the receiving process, if $\mathcal{T}(g_i) - \mathcal{T}(g_j) > 0$, then the message appears to have been received earlier than expected.

The three middle terms are physical time offsets. L_{ij} is the **communication latency** or **network latency** from j to i . Note that such a one-way latency is not a measurable quantity on its own, but, nevertheless, the apparent latency given by (2.5) is measurable. E_{ij} is the **clock synchronization error**. X_{ij} is **execution time** overhead at node j for sending a message to node i plus whatever overhead at i might be incurred actually receiving the message. These three middle terms are collectively called the **network penalty** from j to i and can be referred to by the shorthand **LEX**.

The clock synchronization error E_{ij} can be positive or negative, whereas X_{ij} , and L_{ij} are always nonnegative. If E_{ij} or $\bar{A}_j$ is a sufficiently large negative number, or if $\bar{C}_{ij}$ is a sufficiently large positive number, then the apparent latency can be negative. In such a case, the receive event will appear to have been processed before any remote physical event such as user input that triggered it. Imperfect clocks make this unavoidable, but, more interestingly, the inconsistency makes it a feature we can exploit. By tolerating inconsistency, as we will see, we can offset the effects of network latencies and computation times.

Equation (2.5) can be simplified to

$$T_i = T_j + L_{ij} + E_{ij} + X_{ij}. \quad (2.6)$$

This is a reassuringly simple statement: The physical time at which a message is received is equal to the physical time at which it is sent (as measured by a

different clock) plus the network penalty **LEX**. But the form in (2.5) is more useful because it tells us that unavailability increases as **LEX** increases and decreases as the tolerance for inconsistency increases. This leads to the **CAL** theorem.

2.3.4 The CAL Theorem

We can now put everything together and identify the sources of the **unavailability** of Definition 2.5. Comparing Definitions 2.5 and 2.7, we see that

$$\bar{A}_i \geq \max_{j \in \{1 \dots N\}} (\mathcal{L}_{ij}), \quad (2.7)$$

where $\mathcal{L}_{ij}$ is the **apparent latency** from j to i . Comparing the definitions carefully, we can see that this inequality accounts for all the constraints on $\bar{A}_i$ except those due to **write events**.

To complete the picture we can convert all write events on process i to receive events where process i sends a message to itself to perform the write. Adapting (2.5) to let $j = i$, we get

$$T_i - \mathcal{T}(g_i) = (T'_i - \mathcal{T}(g'_i)) + L_{ii} + E_{ii} + X_{ii} - (\mathcal{T}(g_i) - \mathcal{T}(g'_i)), \quad (2.8)$$

where g'_i and T'_i are the tag and physical time at which process i sends itself a message, and g_i and T_i are the tag and physical time at which it receives this message. For example, a process may read x at tag g'_i , perform some computation, and write a new value for x at a larger tag g_i .

As before, we can simplify (2.8). Let's assume $L_{ii} = E_{ii} = 0$ because the sending and receiving are on the same host. Then (2.8) reduces to

$$T_i = T'_i + X_{ii}, \quad (2.9)$$

which simply tells us that the update to x occurs at a physical time equal to the sum of the physical time at which the previous value of x can be read plus the execution time. This is what we would expect.

In many applications, however, (2.9) is too optimistic. There may be additional constraints on the physical time at which this update can occur, as illustrated by the following example.

Example 2.8: Suppose that the write to the local copy of the shared variable is because of reading a local sensor value, for example in a distributed control system. The sensor could be, for example, a camera observing the local environment. In this case, the tag g_i assigned to the event would naturally be such that $\mathcal{T}(g_i)$ is equal to the physical time, according to the local clock, at which the sensor is probed. This means that the unavailability has a lower bound,

$$T_i - \mathcal{T}(g_i) \geq 0.$$

The write to the local copy of the shared variable cannot take place until physical time matches or exceeds the timestamp of the write because the timestamp of the write represents the physical time at which the sensor was probed.

Let v_i represent the **floor on unavailability** at node i . In the previous example, $v_i = 0$. If there is no floor, $v_i = -\infty$. We now need to modify (2.8) to become

$$T_i - \mathcal{T}(g_i) = \max(v_i, (T'_i - \mathcal{T}(g'_i)) + L_{ii} + E_{ii} + X_{ii} - (\mathcal{T}(g_i) - \mathcal{T}(g'_i))). \quad (2.10)$$

We also need to modify (2.7) to become

$$\bar{A}_i \geq \max(v_i, \max_{j \in \{1 \dots N\}} (\mathcal{L}_{ij})). \quad (2.11)$$

If we further assume that (2.11) accounts for all constraints on unavailability, then we get an equality,

$$\bar{A}_i = \max(v_i, \max_{j \in \{1 \dots N\}} (\mathcal{L}_{ij})). \quad (2.12)$$

Sidebar: Real-Time Processes

Example 2.8 illustrates an instance of a **real-time constraint**. Real-time constraints are either lower or upper bounds on the physical time at which an event can occur. Example 2.8 illustrates a lower bound. Acquiring sensor data with timestamp t cannot occur earlier than when physical time T matches the timestamp. In other words, $T \geq t$. Such a constraint has been described as **logical time chases physical time** (Lee et al., 2023a). When such a constraint is imposed on an entire system, then no event with timestamp $t = \mathcal{T}(g)$ is processed until physical time T matches or exceeds t .

Note that, in general, there could be further constraints that we have not accounted for. For example, the processes may share computing resources, and contention could result in further degradation of availability. Hence, (2.12) is a *best case* unavailability. No implementation can do better than this, but some implementations may do worse.

Combining (2.12) with Definition 2.5 and (2.5), we get

$$\bar{A}_i = \max(v_i, \max_{j \in \{1 \dots N\}} \left[\max \left((T_j - \mathcal{T}(g_j)) + L_{ij} + E_{ij} + X_{ij} - (\mathcal{T}(g_i) - \mathcal{T}(g_j)) \right) \right]). \quad (2.13)$$

where the third maximization is over all communications from j to i , including self-send events.

Now it is convenient to define the following:

Definition 2.8. *The worst-case dependency from node j to i over a trace is*

$$\Gamma_{ij} = \max \left(L_{ij} + E_{ij} + X_{ij} - (\mathcal{T}(g_i) - \mathcal{T}(g_j)) \right), \quad (2.14)$$

where the maximization is over all communications from j to i in the trace. If there is no such communication, then we define $\Gamma_{ij} = -\infty$.

In the (common) situation where $\mathcal{T}(g_i) - \mathcal{T}(g_j)$ is a constant $\bar{C}_{ij}$ across all communications from j to i , this becomes

$$\Gamma_{ij} = \max(L_{ij} + E_{ij} + X_{ij}) - \bar{C}_{ij}, \quad (2.15)$$

where $\bar{C}_{ij}$ is the **inconsistency**. In words, worst-case dependency is the worst-case network penalty minus the inconsistency. Tolerating higher inconsistency mitigates the network penalty, an intuitive observation.

We next introduce a **constant network approximation**, where we assume that *all* communications from j to i experience this worst-case penalty (see sidebar on page 84). With this approximation, (2.13) becomes

$$\bar{A}_i = \max \left(v_i, \max_{j \in \{1 \dots N\}} \left[\max(T_j - \mathcal{T}(g_j)) + \Gamma_{ij} \right] \right) = \max \left(v_i, \max_{j \in \{1 \dots N\}} \left[\bar{A}_j + \Gamma_{ij} \right] \right),$$

for all nodes i and j , where the third maximization is over all writes and receives in process j . This gives us now the simplest form of the **CAL theorem** (or $\bar{C}\bar{A}\bar{L}$ theorem), which we now state formally:

Theorem 2.1. (CAL theorem) Given a *trace* as defined in Section 2.2.3, if we assume that all communications from node j to node i experience a *worst-case dependency* Γ_{ij} , then the unavailability in Definition 2.5 at process i is,

$$\bar{A}_i = \max \left(v_i, \max_{j \in \{1 \dots N\}} [\bar{A}_j + \Gamma_{ij}] \right), \quad (2.16)$$

where v_i is the *floor on unavailability* for process i (or $-\infty$ if there is no floor), and N is the number of nodes.

This is a system of equations where the unknowns are the unavailabilities and the Γ_{ij} are the assumptions about network penalty and tolerance for inconsistency. This system of equations tells us that as tolerance for inconsistency increases, unavailability will tend to decrease (because Γ_{ij} will decrease). On the other hand, as network penalty increases, Γ_{ij} increases, and hence unavailability will increase. This is reassuringly intuitive.

2.3.5 Self Updates

One consequence of (2.16) is,

$$\bar{A}_i \geq \bar{A}_i + \Gamma_{ii}. \quad (2.17)$$

Here, Γ_{ii} is the worst-case dependency when i sends itself messages. What does it mean for i to send itself messages? Here, we assume that whenever a process is updating its local copy of x based on previous values of x and/or external inputs, then it performs some computation, with execution time at most X_{ii} , and assigns the new value of x a tag greater than its previous tag by at least $\bar{C}_{ii}$. This local reassignment of x corresponds to i sending itself a message.

Adapting Definition 2.8 to this special case, let g_i denote the previous tag on process i and g'_i the new tag after the self-update. Then the *worst-case dependency* is,

$$\Gamma_{ii} = \max \left(L_{ii} + E_{ii} + X_{ii} - (\mathcal{T}(g'_i) - \mathcal{T}(g_i)) \right),$$

where the maximization is over all self messages. When i is sending itself messages, we can assume that $E_{ii} = L_{ii} = 0$, so this becomes,

$$\Gamma_{ii} = \max \left(X_{ii} - (\mathcal{T}(g'_i) - \mathcal{T}(g_i)) \right).$$

Sidebar: Constant Network Approximation

The first form of the [CAL theorem](#) (Theorem 2.1) assumes that all network communications from node j to node i experience the [worst-case dependency](#) given in Definition 2.8. In theory, the dynamics could be much more complicated, and the worst-case unavailability may not occur coincidentally with the worst-case dependency.

Equation (2.13) can be abstracted as a maximization of a sum of three terms:

$$\begin{aligned} X &= (T_j - \mathcal{T}(g_j)), && \text{the unavailability at } j, \\ LEX &= L_{ij} + E_{ij} + X_{ij}, && \text{the network penalty, and} \\ Z &= -(\mathcal{T}(g_i) - \mathcal{T}(g_j)), && \text{the negative of the inconsistency.} \end{aligned}$$

Each of these terms can vary for each communication from j to i . The worst-case dependency is the maximum over the trace of $LEX + Z$. But a communication experiencing this worst-case penalty may not yield the worst-case unavailability. In general, $\max(X + LEX + Z) \neq \max(X) + \max(LEX + Z)$. In particular, it might be possible for an application to use an increased inconsistency (lower Z) whenever the unavailability at j is high, and a lower inconsistency (higher Z) otherwise. Or the application could use a faster network (lower LEX) for communication whenever X is large.

The [constant network approximation](#), therefore, is an approximation that assumes applications that do not play such tricks. Only in situations where

$$\max(X + LEX + Z) = \max(X) + \max(LEX + Z)$$

will this equality actually capture the worst case.

Fortunately, this situation is common. The inconsistency Z is typically a constant (and very commonly, $Z = 0$) because it is a semantic property of an application. Applications often just send messages when they can and experience whatever latencies and overhead that the network imposes. For such applications, the constant network approximation will give a good sense of the worst-case unavailability.




Figure 2.3: Two node system where a source sends a message to a destination periodically.

If this number is positive, then (2.17) tells us that $\bar{A}_i = \infty$. Therefore, we have the following theorem:

Theorem 2.2. *Under the [constant network approximation](#), if the [worst-case dependency](#) Γ_{ii} of process i to itself is positive, then the [unavailability](#) is unbounded.*

This surprising result actually makes sense. If in process i the time X_{ii} that it takes to update a local copy of x is bigger than the timestamp increment, then over repeated such updates, the [lag](#) will increase without bound. Such a system is failing to keep up with real time.

Of course, in practice, the execution time will not be constant, the timestamp increment may also not be constant, and the trace could be finite, in which case the unavailability as in Definition 2.5 may be finite. But this result gives us a conservative guideline nonetheless.

Conversely, if the [worst-case dependency](#) is less than or equal to zero, then (2.17) is trivially satisfied by any $\bar{A}_i$. Hence, the other inequalities in (2.16), where $i \neq j$, may dominate, or the [floor on unavailability](#) v_i may dominate. If $v_i = -\infty$, then it is possible for (2.16) to yield $\bar{A}_i = -\infty$, which we interpret to mean that there is no lower bound on the physical time at which the shared variable can be made available. It has no [real-time constraints](#) nor dependencies on any other events with real-time constraints, and hence the write to the local copy of the shared variable can be performed arbitrarily far ahead of time, ahead of its timestamp.

Example 2.9: Figure 2.3 shows perhaps the simplest possible distributed system where a source sends a message to a destination periodically. Here, assume that every 100 ms, the Source component reads a sensor and performs a computation that takes constant execution time X_1 , writing the

result with a write event with timestamps that are increasing by 100 ms each iteration. The time value X_1 includes the **setup time** for the sensor (the time it takes for a sensor reading to stabilize), plus any overhead involved in reading the sensor value, such as execution time of an interrupt service routine, plus any computation the Source node locally performs on the sensor data.

The Source component then sends this timestamped update to the Sink component. The Sink component has execution time X_2 to process the received message. Assume the communication suffers a constant network latency of L and clock synchronization error of E .

Assume here that there is no manipulation of the tag in transit, so the tags at the sending end equal the tags at the receiving end. If the timestamps are measured in milliseconds, then the **timestamp** for the k^{th} message is $\mathcal{T}(g_k) = k \times 100$ milliseconds, assuming the first message is at time 0.

The graphical notation here is that of **Lingua Franca**, a coordination language that will be discussed in Chapter 8. The clock symbol represents a periodic trigger and the dark gray chevrons represent computations. In Lingua Franca terminology, these computations are called **reactions**. Here, we assume each reaction performs a write or receive event. Hence, the **availability** we are interested in is the time at which these reactions can begin execution. We will identify reactions by the names X_i , where X_i also represents the **worst-case execution time (WCET)** of the reaction.

Assume the Source and Sink nodes maintain some local state that may be read by reactions X_1 and X_2 , and that we require consistency in these two local states. There is no communication back from the Sink to the Source, so we assume this shared state is only modified at the Source and that the modification is communicated to the Sink. Hence, the time at which the local state is available in the Source and the Sink is the time at which the reactions X_1 and X_2 can begin.

To ensure integrity of the state, we impose a constraint that each reaction cannot be restarted until the previous invocation of the same reaction is complete. Let the Source be the index 1 node and the Sink be the index 2 node. Hence,

$$\Gamma_{11} = X_1 - 100 \text{ ms},$$

where the minus 100 ms is the time difference between the timestamp of the k^{th} message and the timestamp of the $(k + 1)^{\text{th}}$ message. Similarly, at the Sink,

$$\Gamma_{22} = X_2 - 100 \text{ ms}.$$

Given the assumptions, the worst-case dependency is, from (2.15),

$$\Gamma_{21} = L + E + X_1$$

There is no communication from the Sink to the Source, so $\Gamma_{12} = -\infty$.

Because the Source is reading sensor data, we can assume a **real-time constraint**. Logical time must chase physical time (see sidebar on page 81). Then the **floor on unavailability** is $v_1 = 0$. Reaction X_1 cannot begin at tag g until physical time T catches up with logical time $\mathcal{T}(g)$. Hence, from the **CAL theorem** (Theorem 2.1), the unavailability at the Source component is

$$\begin{aligned} \bar{A}_1 &= \max(v_1, \max(\bar{A}_1 + \Gamma_{11}, \bar{A}_2 + \Gamma_{12})) \\ &= \max(0, \bar{A}_1 + X_1 - 100 \text{ ms}, \bar{A}_2 - \infty) \\ &= \max(0, \bar{A}_1 + X_1 - 100 \text{ ms}). \end{aligned}$$

If $X_1 > 100 \text{ ms}$, then this diverges, yielding $\bar{A}_1 = \infty$. Otherwise, it yields $\bar{A}_1 = 0$. This means that if the Source can keep up with real time, the unavailability is zero. The reaction X_1 can begin as soon as physical time matches logical time. This is intuitive.

Assume that the destination node has no real-time constraints of its own, so $v_2 = -\infty$. Then the unavailability of the Sink component is given by

$$\begin{aligned} \bar{A}_2 &= \max(v_2, \max(\bar{A}_1 + \Gamma_{21}, \bar{A}_2 + \Gamma_{22})) \\ &= \max(-\infty, \bar{A}_1 + L + E + X_1, \bar{A}_2 + X_2 - 100 \text{ ms}) \\ &= \max(\bar{A}_1 + L + E + X_1, \bar{A}_2 + X_2 - 100 \text{ ms}). \end{aligned}$$

As with the Source, if $X_2 > 100 \text{ ms}$, this diverges, and $\bar{A}_2 = \infty$. The program fails to keep up with real time. Otherwise,

$$\bar{A}_2 = \bar{A}_1 + L + E + X_1.$$

If the Source is keeping up with real time, this becomes

$$\bar{A}_2 = L + E + X_1.$$

This tells us that execution X_2 can begin when physical time matches logical time plus LEX. This result is again intuitive.

We can use these methods on much more complicated examples, but to do so, we need better notation.

2.3.6 Max-Plus Formulation

Summarizing where we are, the CAL theorem gives an algebraic relation between inconsistency, unavailability, and latency. This relation shows that if communication latency becomes unbounded (e.g., the network becomes partitioned), then one of inconsistency and unavailability must also become unbounded, and hence the CAP theorem is a special case of the CAL theorem.

The form of the CAL theorem given in (2.16) can be made much more elegant by observing that the mathematical operations of maximization and addition on numbers form an algebra called a max-plus algebra with very useful properties (Baccelli et al., 1992). In this algebra, addition (written $\oplus$) is maximization,

$$a \oplus b = \max(a, b),$$

and multiplication (written $\otimes$) is addition,

$$a \otimes b = a + b.$$

In max-plus, the additive identity is $-\infty$, and the multiplicative identity is zero. As with ordinary arithmetic, $\oplus$ and $\otimes$ are associative and commutative, and $\otimes$ distributes over $\oplus$, meaning that $a \otimes (b \oplus c) = (a \otimes b) \oplus (a \otimes c)$. Hence, many properties of ordinary algebra carry over to the max-plus algebra. Using max-plus algebra, we can get a much simpler form of the CAL theorem.

Definition 2.9. Let N be the number of nodes. The **dependency matrix** Γ is an $N \times N$ matrix such that its i, j element (row i , column j) Γ_{ij} is the **worst-case dependency** of Definition 2.8.




Figure 2.4: Communication graph for Example 2.9 and Figure 2.3.

In other words, from (2.14), the i, j -th entry in the Γ matrix is an assumed bound on **LEX** adjusted downwards by the specified inconsistency for the communication from j to i .

Theorem 2.3. (CAL Theorem, max-plus formulation) Let $\mathbf{A}$ be a column vector with elements equal to the unavailabilities $\bar{A}_i$, and $\mathbf{V}$ be a column vector with elements equal to the **floor on unavailability** v_i . Then (2.16) of Theorem 2.1 can be written

$$\mathbf{A} = \mathbf{V} \oplus \Gamma \mathbf{A}, \quad (2.18)$$

where the matrix multiplication is in the max-plus algebra.

This is a system of equations in the max-plus algebra.

A graphical interpretation of the dependency matrix Γ is useful.

Definition 2.10. Given an $N \times N$ dependency matrix Γ , the **communication graph** is a directed graph with N nodes where there is an arc from node j to node i if $\Gamma_{ij} > -\infty$. The weight on the arc is Γ_{ij} .

A **cycle** in this graph is a path from node j back to node j , and the **weight of the cycle** is the sum of the weights along the path.

Example 2.10: The dependency matrix for Example 2.9 and Figure 2.3 is

$$\Gamma = \begin{bmatrix} X_1 - 100 \text{ ms} & -\infty \\ L + E + X_1 & X_2 - 100 \text{ ms} \end{bmatrix}.$$

The communication graph is shown in Figure 2.4. It has two cycles, one with weight $X_1 - 100$ ms and the other with weight $X_2 - 100$ ms. Both cycles have length one in this example.

The i, j element of the dependency matrix Γ gives the weight of the path of length one from j to i . To solve (2.18), we will need to consider longer paths.

Definition 2.11. Γ^* is the *transitive closure matrix* given by

$$\Gamma^* = \mathbf{I} \oplus \Gamma \oplus \Gamma^2 \oplus \dots, \quad (2.19)$$

where $\mathbf{I}$ is the *max-plus identity matrix*, which has zeros along the diagonals and $-\infty$ everywhere else.

For any $k \geq 1$, the i, j element of Γ^k is the maximum weight of a path from j to i of length k , or $-\infty$ if there is no such path. If any cycle in the communication graph has positive weight, then at least one element of Γ^k will diverge to ∞ as $k \rightarrow \infty$, and hence Γ^* will have infinite elements. On the other hand, if no cycle has positive weights, then every element of Γ^* will remain bounded from above. If all cycles have *negative* weights, then all elements of Γ^k will tend towards $-\infty$ as $k \rightarrow \infty$. This leads to the following theorem.

Theorem 2.4. If the communication graph has no cycle with positive weight, then

$$\Gamma^* = \mathbf{I} \oplus \Gamma \oplus \dots \oplus \Gamma^{N-1}, \quad (2.20)$$

where N is the number of nodes.

Proof. (From Baccelli et al. (1992), Theorem 3.20) For any $M \geq N$, the i, j element of Γ^M is either $-\infty$ or the weight of a path from j to i of length M . Since there are only N nodes, such a path must include a cycle. Because no cycle has positive weight, that element is no larger than the weight of a shorter path from j to i that does not include the cycle. Therefore, for all $M \geq N$,

$$\Gamma^M \leq \mathbf{I} \oplus \Gamma \oplus \dots \oplus \Gamma^{N-1}.$$

Hence, for any $M \geq N$, adding Γ^M to (2.20) with $\oplus$ does not change the matrix. $\square$

Note that it is easy to check whether the graph has cycles with positive weight if you are given the Γ matrix. Any cycle will have length k , $1 \leq k \leq N$. The i^{th} diagonal of Γ^k is the heaviest weight of all cycles of length k that start and end at

node i . Hence, if any cycle has positive weight, it will show up on the diagonal as a positive number on a diagonal element of Γ^k .

We can use the transitive closure matrix to solve (2.18).

Theorem 2.5. *If no cycle of the communication graph for the dependency matrix Γ has positive weight, then (2.18) has a solution bounded from above given by*

$$A = \Gamma^* V. \quad (2.21)$$

Moreover, if all cycles have negative weight, this solution is the unique finite solution.

Proof. (From Baccelli et al. (1992), Theorem 3.17) Plugging (2.21) into the right side of (2.18), we get

$$V \oplus \Gamma(\Gamma^* V) = (I \oplus \Gamma\Gamma^*)V = \Gamma^* V,$$

which shows that (2.21) is a solution for A in (2.18). From Theorem 2.4, this solution will be bounded from above if no cycle has positive weight.

If every cycle has negative weights, then we can show that this solution is unique. Suppose B is a finite solution. I.e.,

$$B = V \oplus \Gamma B.$$

Then, plugging B into the right side, we get

$$B = V \oplus \Gamma(V \oplus \Gamma B) = V \oplus \Gamma V \oplus \Gamma^2 B.$$

Repeating, for any $k \geq 2$, we get

$$B = V \oplus \Gamma V \oplus \cdots \oplus \Gamma^{k-1} V \oplus \Gamma^k B.$$

Because every cycle has negative weight, Γ^k tends towards a matrix with all elements being $-\infty$ as $k \rightarrow \infty$. Hence, in the limit,

$$B = V \oplus \Gamma V \oplus \Gamma^2 V \oplus \cdots = \Gamma^* V.$$

□




Figure 2.5: Three node system where a source sends a message through a computation with a logical execution time to a destination sporadically.

Example 2.11: For the dependency matrix of Example 2.10, the transitive closure is (assuming non-negative cycle weights)

$$\Gamma^* = \begin{bmatrix} 0 & -\infty \\ -\infty & 0 \end{bmatrix} \oplus \begin{bmatrix} X_1 - 100 \text{ ms} & -\infty \\ L + E + X_1 & X_2 - 100 \text{ ms} \end{bmatrix},$$

where the $\oplus$ addition is pointwise. Since we assume non-negative cycle weights, $X_1 \leq 100 \text{ ms}$ and $X_2 \leq 100 \text{ ms}$, so this becomes

$$\Gamma^* = \begin{bmatrix} 0 & -\infty \\ L + E + X_1 & 0 \end{bmatrix}.$$

Applying Theorem 2.5, we get

$$\mathbf{A} = \begin{bmatrix} \bar{A}_1 \\ \bar{A}_2 \end{bmatrix} = \Gamma^* \mathbf{V} = \begin{bmatrix} 0 & -\infty \\ L + E + X_1 & 0 \end{bmatrix} \begin{bmatrix} 0 \\ -\infty \end{bmatrix} = \begin{bmatrix} 0 \\ L + E + X_1 \end{bmatrix}.$$

This is the same result we got in Example 2.9.

Example 2.12: Figure 2.5 shows an example similar to that in Figure 2.3, but with three nodes and two other significant differences. First, instead of a periodic trigger, this example has a **sporadic** trigger, which is an irregularly timed trigger that has a minimum spacing between events. In this case, the minimum spacing is 100 ms. The notation of the figure is that of *Lingua Franca*, which will be covered in Chapter 8. The triangle with a “P” represents a “physical action” in LF, which represents externally triggered asynchronous events that get assigned a timestamp equal to the local physical time when the event occurs.

The second significant difference is that the Computation node has a **logical execution time (LET)**, which means that its output has a timestamp greater than that of the triggering input. In this case, the LET is 100 ms and is shown using the Lingua Franca notation for logical delays.

The communication from node 1 (Source) to node 2 (Computation) has latency plus clock synchronization error $L_1 + E_1$, whereas the next communication has $L_2 + E_2$. We can immediately construct the [dependency matrix](#):

$$\Gamma = \begin{bmatrix} X_1 - 100 \text{ ms} & -\infty & -\infty \\ L_1 + E_1 + X_1 & X_2 - 100 \text{ ms} & -\infty \\ -\infty & L_2 + E_2 + X_2 - 100 \text{ ms} & X_3 - 100 \text{ ms} \end{bmatrix}.$$

Note how the LET affects the matrix. It introduces a non-zero inconsistency $\bar{C}_{32} = 100 \text{ ms}$.

To simplify notation, let $S = 100 \text{ ms}$ represent the minimum spacing, $LET = 100 \text{ ms}$ represent the logical execution time, and $LEX_i = L_i + E_i + X_i$. This becomes

$$\Gamma = \begin{bmatrix} X_1 - S & -\infty & -\infty \\ LEX_1 & X_2 - S & -\infty \\ -\infty & LEX_2 - LET & X_3 - S \end{bmatrix}.$$

If we assume non-negative cycle weights, the transitive closure becomes

$$\begin{aligned} \Gamma^* &= \begin{bmatrix} 0 & -\infty & -\infty \\ -\infty & 0 & -\infty \\ -\infty & -\infty & 0 \end{bmatrix} \oplus \begin{bmatrix} X_1 - S & -\infty & -\infty \\ LEX_1 & X_2 - S & -\infty \\ -\infty & LEX_2 - LET & X_3 - S \end{bmatrix} \\ &\oplus \begin{bmatrix} X_1 - S & -\infty & -\infty \\ LEX_1 + \max(X_1, X_2) - S & X_2 - S & -\infty \\ LEX_1 + LEX_2 - LET & LEX_2 + \max(X_2, X_3) - LET - S & X_3 - S \end{bmatrix}. \end{aligned}$$

Because we assume non-negative cycle weights,

$$\max(X_1, X_2) < S \text{ and } \max(X_2, X_3) < S,$$




Figure 2.6: Pattern where a Mediator is consulted before processing a request from a Source.

so this becomes

$$\Gamma^* = \begin{bmatrix} 0 & -\infty & -\infty \\ LEX_1 & 0 & -\infty \\ LEX_1 + LEX_2 - LET & LEX_2 - LET & 0 \end{bmatrix}.$$

As before, assume that $v_1 = 0, v_2 = -\infty$, and $v_3 = -\infty$. Applying Theorem 2.5, we get

$$\begin{aligned} A &= \Gamma^* V \\ \begin{bmatrix} \bar{A}_1 \\ \bar{A}_2 \\ \bar{A}_3 \end{bmatrix} &= \begin{bmatrix} 0 & -\infty & -\infty \\ LEX_1 & 0 & -\infty \\ LEX_1 + LEX_2 - LET & LEX_2 - LET & 0 \end{bmatrix} \begin{bmatrix} 0 \\ -\infty \\ -\infty \end{bmatrix} \\ &= \begin{bmatrix} 0 \\ LEX_1 \\ LEX_1 + LEX_2 - LET \end{bmatrix}. \end{aligned}$$

This tells us that the computation in the Source can execute as soon as physical time matches logical time. The execution in the Computation node can begin by the time physical time matches logical time plus $L_1 + E_1 + X_1$, as with the previous example. More interestingly, the computation in the Sink can begin earlier because of the LET. If $LET \geq LEX_1 + LEX_2$ (logical execution time is at least as big as physical execution time), then the computation in the Sink can begin as soon as physical time matches logical time. This is a nice example illustrating where sacrificing consistency (with the logical delay implied by LET) improves availability.

Example 2.13: Figure 2.6 shows a common pattern where a Source broadcasts a request to a Mediator and a Sink. The expectation is that the Sink should wait for the result from the Mediator before processing the request from the Source. If we number the Source, Mediator, and Sink 1, 2, and 3 respectively, then the dependency matrix becomes

$$\Gamma = \begin{bmatrix} X_1 - S & -\infty & -\infty \\ LEX_{2,1} & X_2 - S & -\infty \\ LEX_{3,1} & LEX_{3,2} & X_3 - S \end{bmatrix},$$

where $LEX_{i,j}$ is $L + E + X$ for the communication from j to i . Assuming non-positive cycle weights, we get

$$\Gamma^* = \begin{bmatrix} 0 & -\infty & -\infty \\ LEX_{2,1} & 0 & -\infty \\ \max(LEX_{3,1}, LEX_{2,1} + LEX_{3,2}) & LEX_{3,2} & 0 \end{bmatrix},$$

Applying Theorem 2.5, we get

$$\begin{aligned} \begin{bmatrix} \bar{A}_1 \\ \bar{A}_2 \\ \bar{A}_3 \end{bmatrix} &= \begin{bmatrix} 0 & -\infty & -\infty \\ LEX_{2,1} & 0 & -\infty \\ \max(LEX_{3,1}, LEX_{2,1} + LEX_{3,2}) & LEX_{3,2} & 0 \end{bmatrix} \begin{bmatrix} 0 \\ -\infty \\ -\infty \end{bmatrix} \\ &= \begin{bmatrix} 0 \\ LEX_{2,1} \\ \max(LEX_{3,1}, LEX_{2,1} + LEX_{3,2}) \end{bmatrix}. \end{aligned}$$

This is again exactly what we would expect. Reaction X_1 can execute immediately upon becoming enabled. Reaction X_2 waits, in the worst case, $LEX_{2,1}$. Reaction X_3 waits the larger of the dependency from 1 to 3 or the sum of the 1 to 2 and 2 to 3.

Example 2.14: The final example, shown in Figure 2.7, has the structure of a common pattern, where some number of Replica nodes maintain copies of some shared state that is required to be consistent. Some number of Client nodes offer updates to this shared state and receive updates from the replicas. Lee et al. (2023a) use a similar structure for autonomous vehicles




Figure 2.7: Six node system where a pair of Clients send messages to two nodes that maintain state replicas and feed back the merged results to the Clients.

competing for use of an intersection. [Paladino et al. \(2026\)](#) give a similar structure for a distributed banking application.

In Figure 2.7, the X_1 and X_4 reactions are assumed to be **sporadic** with a minimum spacing of some time interval S . Assume that the **LEX** penalty for a reaction i sending to reaction j is given by $LEX_{i,j}$. Then we can write the **dependency matrix** as follows:

$$\Gamma = \begin{bmatrix} X_1 - S & -\infty & 2X_3 - S & -\infty & -\infty & -\infty \\ LEX_{2,1} & 2X_2 - S & -\infty & LEX_{2,4} & -\infty & -\infty \\ X_1 & LEX_{3,2} & 2X_3 - S & -\infty & -\infty & -\infty \\ -\infty & -\infty & -\infty & X_4 - S & -\infty & 2X_6 - S \\ LEX_{3,1} & -\infty & -\infty & LEX_{5,4} & 2X_5 - S & -\infty \\ -\infty & -\infty & -\infty & X_4 & LEX_{6,5} & 2X_6 - S \end{bmatrix} \quad (2.22)$$

First, let's understand the diagonal elements. Reactions X_1 and X_4 are triggered directly by the sporadic sources with minimum spacing S , and hence their self loops, similar to those in Figure 2.4, have weights equal to their execution time minus the minimum timestamp increment between triggers. All other reactions are potentially triggered twice within any time

interval of length S , so their self loops have weights of the form $2X_i - S$. In other words, to be assured of keeping up with real time, S must be large enough to accommodate two executions of these reactions.

The entries of the form $LEX_{i,j}$ do not need any explanation.

The off-diagonal elements $\Gamma_{3,1} = X_1$ and $\Gamma_{6,4} = X_4$ are a consequence of a rule in Lingua Franca that two reactions belonging to the same reactor cannot execute concurrently and must execute in a pre-defined sequence if they are simultaneously enabled. In this case, we assume that reaction X_1 must execute before X_3 at any tag where both are enabled, and similarly for X_4 and X_6 . Hence, there is a dependence from X_1 to X_3 and X_4 to X_6 with weights X_1 and X_4 respectively.

The remaining off-diagonal elements are $\Gamma_{1,3} = 2X_3 - S$ and $\Gamma_{4,6} = 2X_6 - S$. These are the result of a similar Lingua Franca rule that all executions of X_3 (X_6) with tags earlier than g must complete before an execution of reaction X_1 (X_4) with tag g can start. Since there may be two invocations of X_3 and X_6 in each time interval S , there is again a penalty of two execution times offset by the minimum timestamp increment S for X_1 and X_4 .

This matrix is considerably more complex than that of the previous examples, but it is easy to perform the computations if we have specific numbers. To keep the numbers simple, let's assume:

$$S = 6$$

$$X_i = 1, \text{ for all } i$$

$$LEX_{i,j} = 2, \text{ for all } i, j$$

With these numbers,

$$\Gamma = \begin{bmatrix} -5 & -\infty & -4 & -\infty & -\infty & -\infty \\ 2 & -4 & -\infty & 2 & -\infty & -\infty \\ 1 & 2 & -4 & -\infty & -\infty & -\infty \\ -\infty & -\infty & -\infty & -5 & -\infty & -4 \\ 2 & -\infty & -\infty & 2 & -4 & -\infty \\ -\infty & -\infty & -\infty & 1 & 2 & -4 \end{bmatrix}$$

and

$$\Gamma^* = \begin{bmatrix} 0 & -2 & -4 & 0 & -2 & -4 \\ 2 & 0 & -2 & 2 & 0 & -2 \\ 4 & 2 & 0 & 4 & 2 & 0 \\ 0 & -2 & -4 & 0 & -2 & -4 \\ 2 & 0 & -2 & 2 & 0 & -2 \\ 4 & 2 & 0 & 4 & 2 & 0 \end{bmatrix}$$

as calculated by (2.20). Because no diagonal element is positive, we know there are no cycles with positive weights, so (2.20) provides the right value for Γ^* . We can now calculate the solution to (2.18), given by (2.21),

$$A = \Gamma^* V = \begin{bmatrix} 0 & -2 & -4 & 0 & -2 & -4 \\ 2 & 0 & -2 & 2 & 0 & -2 \\ 4 & 2 & 0 & 4 & 2 & 0 \\ 0 & -2 & -4 & 0 & -2 & -4 \\ 2 & 0 & -2 & 2 & 0 & -2 \\ 4 & 2 & 0 & 4 & 2 & 0 \end{bmatrix} \begin{bmatrix} 0 \\ -\infty \\ -\infty \\ 0 \\ -\infty \\ -\infty \end{bmatrix} = \begin{bmatrix} 0 \\ 2 \\ 4 \\ 0 \\ 2 \\ 4 \end{bmatrix}.$$

Again, these results are intuitive and would be easy to infer without all this formalism. Reactions X_1 and X_4 can execute immediately upon being triggered as long as the local physical clock matches or exceeds their timestamps. Reactions X_2 and X_5 must wait for X_1 and X_4 to complete, for the communication to complete, and for any clock synchronization error. Since $LEX_{i,j} = 2$, their unavailability is therefore also 2. Finally, X_3 and X_5 suffer an additional $LEX_{i,j} = 2$, making their unavailability 4.

Less intuitive, however, is that any value for S less than 6 causes the unavailability to diverge to infinity. If we try $S = 5$, there will be cycles with positive weight. Specifically, with other numbers held the same, we will find that

$$\Gamma^3 = \begin{bmatrix} 1 & -4 & -5 & 1 & \infty & \infty \\ 0 & 1 & -4 & 0 & 1 & -4 \\ 1 & 0 & 1 & 1 & \infty & 1 \\ 1 & \infty & \infty & 1 & -4 & -5 \\ 0 & 1 & -4 & 0 & 1 & -4 \\ 1 & \infty & 1 & 1 & 0 & 1 \end{bmatrix}$$

The diagonals are all positive, meaning that there is a cycle from each node back to itself of length 3 that has positive weight. In other words, if events occur at both clients with intervals of five time units, the overall program will be unable to keep up with real time. This is much harder to see. If the times at which these events occur are not identical, however, then it is possible to trace through the executions and see that a total of six time units will be required at each client before the next input event can be processed. Hence, any spacing less than six will result in failure to keep up with real time.

2.4 Conclusion

Concurrency, parallelism, distribution, and determinism are distinct concepts, and conflating them obscures more than it clarifies. A program is concurrent if it specifies operations with no ordering constraints; it is deterministic if every correct execution yields the same observable outcomes, given the same inputs. Whether such a program runs on a single core, on several cores sharing memory, or on a network of machines is a separate question, one that affects timing but, in principle, not the program's defined behavior.

In a distributed setting, however, timing cannot be wished away. Communication takes time, clocks are imperfect, and there is no shared instant at which physically separated components can agree on anything. We therefore have to replace the naive picture of consistency-as-agreement-at-an-instant with a picture organized around the *order* of events. Following Schwarz and Mattern (1994), an execution is causally consistent if no observer sees an effect before its cause. Following Lamport (1984), a consistent total order can be obtained by attaching to each event a *tag* drawn from a totally ordered set G , with a monotone map $\mathcal{T} : G \rightarrow \mathbb{T}$ relating tags to physical time. Together with eventual consistency, this gives us the *CET* model (causally and eventually consistent with timestamps) that the rest of the book builds on.

CET lets us turn qualitative claims about consistency, availability, and partition tolerance into quantitative ones. For a given *trace*, we define three time intervals: the *inconsistency* $\bar{C}_{ij}$, the smallest gap between corresponding receive and send tags; the *unavailability* $\bar{A}_i$, the worst-case *lag* of write and receive events on process i ; and the *apparent latency* $\mathcal{L}_{ij}$, the worst-case lag of receives on i that originate on

j . A single one-way communication decomposes into the sender's unavailability, the network penalty **LEX** (communication latency, clock synchronization error, and execution time overhead), and a credit for tolerated inconsistency. Adopting the **constant network approximation** and collecting the constraints on $\bar{A}_i$, we arrived at the **CAL theorem**,

$$\bar{A}_i = \max\left(v_i, \max_{j \in \{1 \dots N\}} [\bar{A}_j + \Gamma_{ij}]\right),$$

where Γ_{ij} is the **worst-case dependency** from j to i and v_i is the **floor on unavailability** at i .

Recasting this in the max-plus algebra exposes its structure. The dependency matrix Γ induces a weighted **communication graph**; the system $\mathbf{A} = \mathbf{V} \oplus \Gamma \mathbf{A}$ has a bounded solution $\mathbf{A} = \Gamma^* \mathbf{V}$ precisely when no cycle of that graph has positive weight. A positive-weight cycle is the formal statement that a process (or group of processes) cannot keep up with real time: each round of execution adds more lag than the timestamp budget can absorb, and unavailability diverges. Brewer's **CAP** theorem falls out as the limiting case: when latency is unbounded, at least one of inconsistency and unavailability must be unbounded as well.

The CAL theorem is descriptive, not prescriptive. It tells us what the trade space looks like; it does not, by itself, build a system that occupies a chosen point in that space. The remaining chapters supply mechanisms that do. Process networks (Chapter 4), dataflow (Chapter 5), and the synchronous-reactive model (Chapter 6) give deterministic concurrency that emphasizes consistency. Reactors (Chapter 8) and their distributed coordination strategies extend these guarantees across networks while letting the programmer trade off inconsistency, unavailability, and latency against one another in a quantitatively meaningful way.

Nondeterministic Concurrency

3.1	Threads	101
3.2	Actors	101
3.3	Publish and Subscribe	101
3.4	Remote Procedure Calls	101
3.5	Addressing the Challenges	101

3.1 Threads

3.2 Actors

3.3 Publish and Subscribe

3.4 Remote Procedure Calls

3.5 Addressing the Challenges

Part II

Untimed Models

Process Networks

4.1	Kahn Process Networks	106
4.2	Semantics of Process Networks	110
	<i>Historical Notes: Process Networks</i>	111
4.2.1	Least Fixed Point Semantics	112
4.2.2	Monotonic and Continuous Functions	117
4.3	Execution of Process Networks	122
4.3.1	Turing Completeness of Process Networks	127
4.3.2	Effective Execution	129
4.4	Convergence to the Denotational Semantics	134
	<i>Sidebar: Limit of a Sequence of Real Numbers</i>	136
	<i>Sidebar: Scott Topologies</i>	137
	<i>Sidebar: Three Views of Convergence</i>	138
4.5	Beyond Kahn-MacQueen	140
4.5.1	Sequential Functions	140
4.5.2	Stable Functions	142
4.6	Dynamic Kahn Networks	144
4.7	Nondeterministic Extensions of Kahn Networks	146
4.7.1	Nondeterministic Merge	146
4.7.2	The Brock-Ackerman Anomaly [†]	147
4.8	Rendezvous	150
4.9	Conclusion	150

4.1 Kahn Process Networks

A **Kahn process network (KPN)** is a concurrent composition of sequential processes that communicate using a particular form of message passing. A key property of these networks is that, despite the [concurrency](#), they are assured of being [deterministic](#), in the sense that a KPN defines a unique sequence of messages on each communication channel between processes. That is, the messages that are communicated do not depend on the scheduling of the sequential processes. This chapter develops the theory behind such networks, explaining why they are deterministic, and discusses constraints that an execution engine must satisfy to execute them correctly.

A **process network (PN)** is a collection of components called **processes**, each representing a sequential, step-by-step procedure. A, B, and C in Figure 4.1 depict processes as rectangles. A process may have private state (variables that are invisible to other processes). The steps in its procedure may manipulate its state, send messages to processes (including possibly itself), or receive messages from processes (again including possibly itself).

For modularity, so that processes do not need to have references to each other, here we assume that a process sends and receives messages via named input or output **ports**.¹ In Figure 4.1, p1, p2, and p3 are input ports, and p4, p5, and p6 are

¹The process networks implementation in Ptolemy II uses ports in this way ([Ptolemaeus, 2014](#), Chapter 4).




Figure 4.1: Example of a process network.

output ports. An output port may be connected to an input port via a **channel** that carries a sequence of messages using a **first in first out (FIFO)** policy. Each message is called a **token** and may or may not carry a payload. A sequence of tokens is called a **stream**. In Figure 4.1, s_1 , s_2 , and s_3 represent streams. A stream may consist of a finite or infinite sequence of tokens.

A key assumption in process networks is that tokens are delivered to the destination input port reliably and in order. TCP sockets, for example, realize such semantics with reasonable confidence. The channel will **buffer** tokens for later delivery to the receiving process, and, in our model, will never lose tokens. A process can always send a token. It does not need to wait for the recipient to be ready to receive.

We begin with a special case of process networks that is easier to understand than the general case. This special case was introduced by Kahn and MacQueen (1977), and has been the subject of considerable study ever since. A **Kahn-MacQueen process** is a process with **blocking reads**. This means that the only mechanism that the process has to receive a token is to attempt to read a token from an input port; any such attempt blocks execution of the process if there is no available input token. The process will remain blocked until a token becomes available. In particular, a process cannot determine a priori whether a token is available. The only operation it has available is a blocking read.

A Kahn-MacQueen process is a program in an **imperative** language augmented with a blocking read statement and **nonblocking write** statement that reference the input and output ports of the process. Here, we describe Kahn-MacQueen processes using a C-like pseudo language that has structured control statements like **while** and **if-then-else**, plus statements like

```
t = read(p);
```

which performs a blocking read on input port p and returns a token t , and

```
write(p, t);
```

which writes the token t to port p . The read procedure call does not return until there is a token available at port p . The write statement returns immediately. Our language is a pseudo language in that we do not fully define it, and we omit important details like data types.

Example 4.1: Suppose that A in Figure 4.1 executes the following Kahn-MacQueen procedure,

```

1 write(p5, 0);
2 while(true) {
3   t = read(p1);
4   write(p5, t);
5 }
```

This procedure first produces an output token with payload value 0 and then enters an infinite loop where it reads an input token and sends it to the output port. If the input is a sequence tokens (1, 2, 3), for example, then the output will be (0, 1, 2, 3).

A process with this behavior is called a **unit delay** because the input tokens appear at the output delayed by one step in the sequence. In practice, the value of the initial output (0 in this case) would typically be a **parameter** of the process rather than a built-in constant.

Let $a.b$ refer to the **concatenation** of sequences a and b . That is, $a.b$ is a sequence with **prefix** a followed by b . If a is infinite, then $a.b = a$. Using this notation, the process of Example 4.1 defines its output sequence as a function of its input sequence.

Example 4.2: If the input to the process of Example 4.1 is a sequence denoted by s_1 , then the output is $s_2 = A(s_1) = (0).s_1$, where (0) is a **sequence** of length 1.

We can use the definition of process A from Example 4.1 to build a complete process network in Figure 4.1.

Example 4.3: Suppose that B executes the following procedure,

```

1 while(true) {
2   t2 = read(p2);
3   t3 = read(p3);
4   write(p4, t2+t3);
5 }
```

Suppose further that C executes the following procedure,

```

1 while(true) {
2   write(p6, 1);
3 }
```

An execution of the process network in Figure 4.1 yields the following sequences:

$$\begin{aligned}
 s_1 &= (1, 2, 3, 4, \dots) \\
 s_2 &= (0, 1, 2, 3, \dots) \\
 s_3 &= (1, 1, 1, 1, \dots)
 \end{aligned}$$

In the previous example, an execution of the process network yields infinite sequences. This is not always the case.

Example 4.4: Suppose that instead of the unit delay given in Example 4.1, A in Figure 4.1 executes the following Kahn-MacQueen procedure,

```

1 while(true) {
2   t = read(p1);
3   write(p5, t);
4 }
```

Such a process is called an **identity process**, because the output sequence is the same as the input sequence. An execution of the resulting process network yields

$$\begin{aligned}
 s_1 &= \perp \\
 s_2 &= \perp \\
 s_3 &= (1, 1, 1, 1, \dots)
 \end{aligned}$$

where $\perp$ is the empty sequence. Processes A and B both block immediately attempting to read a token provided by the other. Process C, on the other hand, is able to execute and produce an infinite sequence of outputs. Since the semantics of process networks requires that tokens on channels not be lost, the tokens produced by C must be stored until they are consumed by

B. In this case, however, they will never be consumed by B, and eventually, any execution platform will run out of memory to store the tokens.

The previous example illustrates two phenomena that can occur with process networks. The first is **deadlock**, where the processes on a directed cycle of the network are blocked waiting for tokens for each other. In this case, we have a **local deadlock**, because only part of the process network is deadlocked. In particular, C is not blocked.

The second phenomenon illustrated by this example is **unbounded memory**. This process network cannot correctly execute without an unbounded amount of memory for storing unconsumed tokens.

We will see below that, in general, whether a Kahn-MacQueen process network deadlocks is **undecidable**, even if we constrain the processes to a few very simple primitive ones. It is also undecidable whether such a process network can be executed with bounded memory. These two limitations are actually a consequence of the rich expressiveness of the model of computation. We will see that a very few primitive processes are sufficient to make the model of computation **Turing complete**, which means that it can describe every **effectively computable** function.

4.2 Semantics of Process Networks

Example 4.2 suggests that a Kahn-MacQueen process can be defined as a function that maps input sequences to output sequences. Consider a process with a single input port and single output port. The **data type** of a port p is a set T_p of token values that the port may consume or produce. The set T_p^{**} is the set of finite and infinite **sequences** of tokens of type T_p (see Appendix A). A process is therefore a function defined on such sequences.

Example 4.5: Suppose that in Example 4.2, the input and output data types are $T = \mathbb{N}$, the natural numbers. Then the **unit delay** is a function of form

$$A: \mathbb{N}^{**} \rightarrow \mathbb{N}^{**}$$

where for all $s_1 \in \mathbb{N}^{**}$,

$$A(s_1) = (0).s_1.$$

Historical Notes: Process Networks

The notion of concurrent processes interacting by sending messages is rooted in Conway's **coroutines**. Conway (1963) described software modules that interacted with one another as if they were performing I/O operations. In Conway's words, "When coroutines *A* and *B* are connected so that *A* sends items to *B*, *B* runs for a while until it encounters a read command, which means it needs something from *A*. The control is then transferred to *A* until it wants to write



Gilles Kahn (1946 – 2006), French computer scientist who developed Kahn process networks.

whereupon control is returned to *B* at the point where it left off." The key idea is that both *A* and *B* maintain as part of their state their progress through their procedures, and transfer of control occurs to satisfy demands for data. The least fixed-point semantics is due to Kahn (1974), who developed the model of processes as continuous functions on a CPO. Kahn and MacQueen (1977) defined process interactions using nonblocking writes and blocking reads as a special case of continuous functions and developed a language for defining interacting processes. Their language included recursive constructs, an optional functional notation, and dynamic instantiation of processes. They gave a demand-driven execution semantics, related to the lazy evaluators of Lisp (Friedman and Wise, 1976; Morris and Henderson, 1976). Berry (1976) generalized these processes with stable functions.

Sequences of tokens communicated by process networks are **unbounded lists**. The notion of unbounded lists as data structures first appeared in Landin (1965). This underlies the communication mechanism between processes in a process network. The UNIX operating system, due originally to Ritchie and Thompson (1974), includes the notion of **pipes**, which implement a limited form of process networks.

Kahn (1974) stated but did not prove that a maximal and fair execution of process network yields the least fixed point. This was later proved by Faustini (1982) and Stark (1995). It is known as the **Kahn principle**.

As we pointed out before, normally the initial output 0 would be a parameter of the process rather than a built-in constant.

In general, the **semantics** of a program is its meaning. In the case of process networks, the semantics is the **streams** defined by the network. Each stream is a finite or infinite sequence of **tokens**. In this section, we will show that if processes define functions that satisfy a particular constraint (they are **continuous**), then the semantics of the network is unique (i.e., the network defines exactly one sequence of tokens for each stream), and we can give a constructive procedure for building the streams defined by the network (i.e., a mechanism for executing the network). For this section, you will want to have mastered the material on the **prefix order** and **complete partial orders**, in Chapter A.

4.2.1 Least Fixed Point Semantics

Recall that a **fixed point** of a function $F: X \rightarrow X$ is an element $x \in X$ such that $F(x) = x$. Execution of every process network can be reduced to finding a fixed point of a function. The function is a composition of the functions defined by the individual processes.

Example 4.6: Consider the process network in Figure 4.1. This network is redrawn in Figure 4.2(a). In Figure 4.2(b), we reorganize the drawing and draw a box around the three processes. This box can itself be considered a function with three input sequences and three output sequences, as illustrated in Figure 4.2(c). Figure 4.2(d) further abstracts this by aggregating the three streams into one.

Suppose that the three streams in Figure 4.2 have **data type** T . Then each stream s_i , $i = 1, 2, 3$, is a member of the set T^{**} of sequences of tokens of type T . In Figure 4.2(d), the function F therefore has the form

$$F: (T^{**})^3 \rightarrow (T^{**})^3.$$


Using the process definitions from Examples 4.1 and 4.3, we see that if the input is

$$s = ((a_1, a_2, \dots), (b_1, b_2, \dots), (c_1, c_2, \dots)),$$


a three-tuple of sequences, then the output is

$$F(s) = ((b_1 + c_1, b_2 + c_2, \dots), (0, a_1, a_2, \dots), (1, 1, \dots)),$$


which is also a three-tuple of sequences.




(a)



(b)



(c)



(d)

Figure 4.2: Execution of every process network can be reduced to finding a fixed point of a function, as illustrated by this sequence of transformations of a network.

Because of the feedback loop, we seek a sequence s that satisfies $F(s) = s$, so in this case, we require that

$$\begin{aligned} a_1 &= b_1 + c_1 \\ a_2 &= b_2 + c_2 \\ b_1 &= 0 \\ b_2 &= a_1 \\ c_1 &= 1 \\ c_2 &= 1 \\ &\dots \end{aligned}$$

which implies that

$$\begin{aligned} a_1 &= 1 \\ a_2 &= 2 \\ b_1 &= 0 \\ b_2 &= 1 \\ c_1 &= 1 \\ c_2 &= 1 \\ &\dots \end{aligned}$$

which is the same result we obtained before in Example 4.3.

Execution of a process network with n streams of type T can be reduced to finding a fixed point of a function of the form

$$F: (T^{**})^n \rightarrow (T^{**})^n.$$

If a process network has streams with different types, then the domain and codomain of the function will be a [cartesian product](#) of the sets of sequences of those types. The notation gets more complex, but not the concept. It may be surprising that *all* process networks can be reduced in the same way to finding a fixed point of a function.

Example 4.7: Consider the network in Figure 4.3(a). It may seem that this would be difficult to reduce to a fixed point problem because there is no feedback in the network. However, this can be redrawn as shown in Figure 4.3(b), and abstracted as shown in Figure 4.3(c). In this case, the function F is a rather trivial function. Regardless of the input stream, it always produces the same output stream. It is a **constant function**, where the output is independent of the input.

To be concrete, suppose that the data type is $T = \mathbb{N}$, the natural numbers, and that process A produces the sequence $(0, 1, 2, 3, \dots)$. Then $F: \mathbb{N}^{**} \rightarrow \mathbb{N}^{**}$ is a function such that for all $s \in \mathbb{N}^{**}$,

$$F(s) = (0, 1, 2, 3, \dots).$$

Hence, $s = (0, 1, 2, 3, \dots)$ is a fixed point of F (the only fixed point, in fact).

For the above examples, the function F has exactly one fixed point. This is not always the case.




Figure 4.3: Even a process network with no feedback loops can be reduced to finding a fixed point of a function, as illustrated by this sequence of transformations of a network.




Figure 4.4: A process network with multiple fixed points, where I is the identity process.

Example 4.8: Consider an **identity process** I , described in Example 4.4, where for all $s \in T^{**}$,

$$I(s) = s.$$

Suppose that this process is used in the process network shown in Figure 4.4. This process network has multiple fixed points. In fact, any stream $s \in T^{**}$ is a fixed point.

A **prefix order** is a partial-order relation $\sqsubseteq$ over sets of sequences T^{**} (see Appendix A). For any two $a, b \in T^{**}$, $a \sqsubseteq b$ if the sequence a forms the leading part of sequence b . If a is infinite, then $a \sqsubseteq b$ if and only if $a = b$.

When a process network has multiple fixed points, we will choose the **least fixed point**, where “least” is in the prefix order. The least fixed-point semantics is a **denotational semantics**; it gives the program a meaning without necessarily offering a procedure for finding that meaning. An **operational semantics**, in contrast, gives the meaning of a program in terms of a procedure for finding that meaning. Kahn-MacQueen processes, with their blocking reads, offer an operational semantics.

Example 4.9: For the process network in Figure 4.4, the least fixed point is $s = \perp$, the empty sequence. This choice is indeed the least fixed point because for all other fixed points s , $\perp \sqsubseteq s$. Moreover, it is exactly the fixed point that results from executing the Kahn-MacQueen process. The process deadlocks immediately, blocked on the first read. Hence, the result of execution is the empty sequence.

A match between an operational semantics and a denotational semantics is sometimes called **full abstraction**. The previous example suggests that with Kahn-MacQueen and the least fixed point semantics, we have full abstraction. However, when the semantics involves infinite structures, like infinite streams, what we mean by “a match” gets tricky. An operational semantics can never produce an infinite stream until it has run forever.

We will see next that with Kahn-MacQueen process networks, the least fixed point is always uniquely defined, and execution of the network always produces a *prefix* of that least fixed point. We will see that whether the operational semantics converges in the limit to the least fixed point depends on what we mean by convergence and on details of the execution policies that we have not yet discussed, such as management of buffer sizes. We will also show that some useful processes that cannot be defined as Kahn-MacQueen processes also yield a uniquely defined least fixed point, so the Kahn-MacQueen operational semantics is, in a sense, smaller than the least fixed-point denotational semantics.

4.2.2 Monotonic and Continuous Functions

A function $A: T^{**} \rightarrow T^{**}$ is **monotonic** in the prefix order if $a \sqsubseteq b$ implies that $A(a) \sqsubseteq A(b)$. A Kahn-MacQueen process defines a **function** from input sequences to output sequences as long as the imperative language used to define the process is deterministic. This means that given particular input streams at the input ports, the output streams are fully defined by the process. More interestingly, Kahn-MacQueen processes are monotonic.

Example 4.10: The **unit delay** function considered in Example 4.5 is monotonic. Consider two possible input streams $a, b \in \mathbb{N}^{**}$. The corresponding output sequences are $A(a) = (0).a$ and $A(b) = (0).b$. Clearly, if $a \sqsubseteq b$, then $A(a) \sqsubseteq A(b)$.

The fact that Kahn-MacQueen processes are monotonic is easy to see by considering their execution. Suppose $a, b \in \mathbb{N}^{**}$ are two possible input streams where $a \sqsubseteq b$. A Kahn-MacQueen process that is presented with input b will first read a sequence of tokens equal to a . Since the process is deterministic, once it has read this prefix of b , its output will be exactly the output sequence it would have produced had it been presented with input a . Upon continuing to read tokens

from b (which now differ from a), it will only extend the output stream. It has no mechanism for retracting previously produced tokens, and hence the output produced in response to b is an extension of the output produced in response to a .

Notice that the blocking reads of a Kahn-MacQueen process play an important role in ensuring monotonicity.

Example 4.11: Consider a process A with one input port and one output port that produces on its output the sequence (0) if the input is the empty sequence, and otherwise produces the output (1) . This process clearly defines a function, in that the output is fully defined by the input. However, this function is not monotonic. Suppose that $a = \perp$, the empty sequence, and $b = (0)$. Then $a \sqsubseteq b$, but $A(a) = (0)$ is not a prefix of $A(b) = (1)$.

Notice that this process cannot be implemented as a Kahn-MacQueen process. In order to produce the output (0) , the process needs to know that the input is the empty sequence. But the only operation it has available on the input port is to perform a blocking read. If the input is indeed empty, then it will remain blocked forever, and hence will not be able to produce the output (0) .

The poset $(T^{**}, \sqsubseteq)$ is a **complete partial order** (CPO), as explained in Example A.16 of Appendix A. Moreover, by Proposition A.3, the poset $((T^{**})^n, \sqsubseteq)$ is also a CPO, where $\sqsubseteq$ is the **pointwise prefix order**. A function $F: (T^{**})^n \rightarrow (T^{**})^n$ is **continuous** if for all chains $C \subseteq (T^{**})^n$,

$$F(\bigvee C) = \bigvee \hat{F}(C),$$

where $\hat{F}$ is the **lifted** version of F . By Proposition A.5, every continuous function is monotonic.

In practice, every **Kahn-MacQueen process** is not only monotonic, but also continuous. Why? Intuitively, continuity means that a function does not “wait forever” before producing output. Suppose that the **chain** $C = \{s_0, s_1, s_2, \dots\}$ represents a sequence of partially constructed inputs to a Kahn-MacQueen process. Then $\bigvee C = s$ represents the eventual, complete input. The set

$$\hat{F}(C) = \{F(s_0), F(s_1), F(s_2), \dots\}$$

represents the partially constructed outputs, given partially constructed inputs. By the definition of Kahn-MacQueen processes, these are exactly the outputs

that the process can produce. Then continuity requires that $\bigvee \hat{F}(C)$ be equal to the eventually complete output $F(s)$, and by the definition of Kahn-MacQueen processes, that is exactly the output.

Example 4.12: Suppose that A is a process with one input port and one output port. Suppose that its input is eventually going to be the infinite sequence

$$s = (0, 1, 2, 3, \dots)$$

Suppose that $C = \{s_0, s_1, s_2, \dots\}$, where

$$\begin{aligned} s_0 &= (0) \\ s_1 &= (0, 1) \\ s_2 &= (0, 1, 2) \\ &\dots \end{aligned}$$

C is clearly a chain. Moreover,

$$\bigvee C = s.$$

Suppose further that A is the function

$$A(s) = \begin{cases} \perp & \text{if } s \text{ is finite} \\ (1) & \text{otherwise} \end{cases}$$

This function is clearly monotonic, but it is not continuous. In particular, when the function is applied to each partially constructed input in C , the output is $\perp$. Hence,

$$\hat{F}(C) = \{F(s_0), F(s_1), F(s_2), \dots\} = \{\perp\}.$$

Hence,

$$\bigvee \hat{F}(C) = \perp.$$

However,

$$F(\bigvee C) = F(s) = (1).$$

These two are not equal. Intuitively, the function F has to wait forever to determine whether the input is finite or not.

Notice that this function cannot be implemented by a Kahn-MacQueen process. We might try to specify it as follows,

```
1 while(true) {
2   read(p1);
3 }
4 write(p2, 1);
```

That is, after reading an infinite number of inputs, we output a 1. If the input is finite, this process will block forever on one of the reads. In practice, it will never reach the point of producing the 1, so this procedure does not really implement the function.

The [Kleene fixed-point theorem](#) (Proposition A.6) then provides what we need. It states that if F is a continuous function, then it has a unique least fixed point. Moreover, it asserts that the least fixed point is the result of applying the function first to $\perp$, and then recursively to the result, etc., as follows:

$$\begin{aligned} s_0 &= \perp \\ s_1 &= F(\perp) \\ s_2 &= F(F(\perp)) \\ &\dots \\ s_m &= F^m(\perp) \\ &\dots \end{aligned}$$

Thus, instead of solving a system of equations, as we did in Example 4.6, we can construct the solution by just repeatedly applying the function F .

Example 4.13: Consider the same process network of Example 4.6. Recall that the function $F: (T^{**})^3 \rightarrow (T^{**})^3$ is given by

$$F(s) = ((b_1 + c_1, b_2 + c_2, \dots), (0, a_1, a_2, \dots), (1, 1, \dots)),$$

where the input is

$$s = ((a_1, a_2, \dots), (b_1, b_2, \dots), (c_1, c_2, \dots)),$$

This function is easily shown to be continuous.

In this case, the bottom of the CPO is the three-tuple of empty sequences, so the constructive procedure given by the Kleene fixed-point theorem proceeds as follows,

$$\begin{aligned} F((\perp, \perp, \perp)) &= (\perp, (0), (1, 1, \dots)) \\ F(F((\perp, \perp, \perp))) &= ((1), (0, 1), (1, 1, \dots)) \\ F(F(F((\perp, \perp, \perp)))) &= ((1, 2), (0, 1, 2), (1, 1, \dots)) \\ &\dots \end{aligned}$$

This eventually converges to the same sequences found in Example 4.3.

Notice that this procedure immediately creates a practical problem. The very first invocation of the function F yields as the third element of the output tuple an infinite sequence. If we were to literally implement this procedure in a computer, then we would never get past the first step because we would exhaust available memory in any attempt to generate this infinite sequence. The following section will consider practical execution policies.

Example 4.14: For the variant given by Example 4.4, which has a local deadlock, the Kleene procedure immediately converges to the final answer,

$$F((\perp, \perp, \perp)) = (\perp, \perp, (1, 1, \dots)).$$

Example 4.15: For the process network in Figure 4.4, recall that the least fixed point is $s = \perp$. The Kleene procedure immediately converges to this solution.

Example 4.16: Recall that the [unit delay](#) process of Example 4.1 is a function of form

$$A: \mathbb{N}^{**} \rightarrow \mathbb{N}^{**}$$

where for all $b \in \mathbb{N}^{**}$,

$$A(b) = (0).b.$$

This function is continuous. Suppose we put it in a feedback loop, as we did with the identity function in Figure 4.4. Then the Kleene procedure would

yield

$$\begin{aligned}
 A(\perp) &= (0) \\
 A(A(\perp)) &= (0, 0) \\
 A(A(A(\perp))) &= (0, 0, 0) \\
 &\dots
 \end{aligned}$$

The least upper bound of this chain is the infinite sequence $(0, 0, 0, \dots)$.

4.3 Execution of Process Networks

The least fixed point of a process network often includes infinite sequences. In practice, no computer program can construct an infinite sequence. Our interpretation is that process networks whose semantics include infinite sequences do not terminate. They are able to continue to execute for as long as we would like them to execute. At any point during the execution, the sequences that they have produced are mere approximations of the [denotational semantics](#) of the network. When there are multiple infinite sequences defined by a network, then there are many possible approximations. In this section, we consider execution policies that provide useful approximations.

We need some definitions.

Definition 4.1. *A **correct execution** of a process network is one which, at any time during execution, has constructed a prefix of each stream defined by the [denotational semantics](#) of the network.*

These prefixes will always be finite, even if the semantics of the network includes infinite sequences.

Definition 4.2. *A **maximal execution** is a correct execution that either does not halt, or if it halts, has produced exactly every sequence defined by the [denotational semantics](#) of the network.*

What we mean by “halt” in this case is that the execution no longer appends tokens to streams. A maximal execution will only halt if the denotational semantics of




Figure 4.5: Example of a process network where unfair execution is desirable because fair execution will exhaust available memory.

the network defines only finite streams. An **infinite execution** is an execution that does not halt. All infinite executions are maximal.

Definition 4.3. A *fair execution* is one that ensures that if any process is able to produce an output token or read an input token, then it will eventually be allowed to do so.

The following proposition, known as the **Kahn principle**, stated by Kahn and MacQueen (1977) and later proved by Faustini (1982) and Stark (1995), provides key guidance for execution policies.

Proposition 4.1. Any two fair and maximal executions of a process network produce the same sequences of tokens, equalling the least fixed point that is the *denotational semantics*.

The Kahn principle seems to suggest that practical executions of process networks should always be fair and maximal. This is not in fact obvious. The following example shows that fair and maximal executions are not always desirable.

Example 4.17: Consider the process network in Figure 4.5. Suppose that processes A and B produce infinite sequences, and C is able to consume an infinite sequence. Suppose that the process labeled true produces an output of type $\mathbb{T} = \{true, false\}$ that is a constant infinite sequence

$$s_3 = \{true, true, true, \dots\}.$$

The remaining process in the center, which has no label, is a **Select** process. It accepts a sequence of control tokens on the bottom port, and uses these to merge the sequences at its other two input ports. The merged sequence is the output. Specifically, the Select process has three input ports, T, F, and control, where the control port is shown on the bottom of its icon. It has a single output port p. The process is defined by the following Kahn-MacQueen procedure,

```

1 while(true) {
2   c = read(control);
3   if (c) {
4     t = read(T);
5   } else {
6     t = read(F);
7   }
8   write(p, t);
9 }

```

Any fair execution of this process network has to permit B to produce its infinite sequence of tokens, even though those tokens will never be consumed. This will require unbounded memory to store unconsumed tokens.

There is, however, a maximal execution of this network that runs in bounded memory. Such an execution would only permit process B to produce a finite number of tokens. Such an execution is maximal and correct, but not fair.

Definition 4.4. A *bounded execution* is an execution of a process network where there is a natural number $M \in \mathbb{N}$ such that at all times during the execution there are no more than M unconsumed tokens.

A fair and maximal execution of the network in Figure 4.5 is not bounded.

Definition 4.5. An *effective* process network is one where every token that can be produced by a process will eventually be read by its destination processes (Geilen and Basten, 2003).

The process network in the previous example is not effective. For networks that are not effective, we do not necessarily want fair and maximal execution, because such execution will eventually fail due to running out of memory.

Even for networks that are effective, executing them with bounded memory is not always easy. Suppose that we modify the network in Figure 4.5 so that the process labeled *true* produces a random sequence of *true* and *false* values. How can we ensure that A and B do not overflow available memory on their outputs? Even for much simpler networks there is risk of memory overflow.

Example 4.18: Suppose that in the network in Figure 4.3(a), process A produces tokens faster than process B. Then eventually, we will run out of memory.

We seek, therefore, an execution policy that will deliver maximal executions for all process networks, bounded executions for networks for which bounded execution is possible, and fair executions for effective process networks. We call these **effective executions**. Effective executions are not required to be fair for networks that are not effective. They are not required to be bounded for networks that have no bounded execution. They are required to be maximal for all networks.

There is a long history of failed attempts to achieve this goal. One popular proposed technique is **demand-driven execution**, related to the lazy evaluators of Lisp (Friedman and Wise, 1976; Morris and Henderson, 1976), where no token is produced unless there is a downstream process that is ready to consume it. Unfortunately, this does not solve the problem, as illustrated by the following example.

Example 4.19: A variant of Figure 4.5 is shown in Figure 4.6. The only difference is that a new process D has been added that consumes tokens from process B. Each token produced by B needs to be routed to both the *Select* process and to D. This can be accomplished by a **fork** process, indicated in the figure by a small black diamond. The fork can be implemented, for example, by the Kahn-MacQueen procedure

```

1 while (true) {
2     t = read(in);
3     write(out1, t);
4     write(out2, t);
5 }

```




Figure 4.6: Example of a process network where demand-driven execution fails to deliver a bounded execution.

Thus, there will be two buffers, one storing tokens destined for the Select, and one storing tokens destined for D.

Assume that D is able to read an infinite input sequence. As a consequence, it will demand tokens. In demand-driven execution, the fork will respond by demanding inputs. Once the fork receives an input, the above Kahn-MacQueen implementation will send the token to both outputs, and hence the buffer at the F input of the Select will eventually overflow. If the fork is implemented differently so that it only sends outputs in response to demands, then it will be obligated to store the token locally in case the token is later demanded. This too will overflow available memory.

In the previous example, demand-driven execution can deliver bounded execution only if we somehow regulate the demands issued by the two sink processes, C and D. But how would we do that? Moreover, suppose that network contains cycles, or more interestingly, a multiplicity of cycles. How would demands be generated?

Data-driven execution is the complement of demand-driven execution. Instead of **sink** processes driving the execution (processes with no output ports), **source** processes drive the execution (processes with no input ports). A process is permitted to execute only when it has available input data, except source processes, which are always permitted to execute. This obviously does not achieve the

goal, since even a network as simple as that in Figure 4.3(a) cannot be assured of remaining bounded.

The problem of achieving *effective executions* is harder than it first appears. Fair execution, demand-driven execution, and data-driven execution all fail. It is not surprising that the problem is hard, however. It turns out that whether a bounded execution exists is undecidable, as is whether an infinite execution exists. We defend this claim in the next subsection. After that (Section 4.3.2), we show how to solve the undecidable problem by giving a policy that delivers effective execution.

4.3.1 Turing Completeness of Process Networks

In his PhD thesis, [Buck \(1993\)](#) proved that Kahn process networks could realize a universal Turing machine using only a small set of very simple processes. This means that the KPN model of computation is as expressive as any other programming language capable of realizing all effectively computable functions. It also means that some questions about process networks become *undecidable*, meaning that there is no algorithm that can answer the question for all process networks. One such undecidable question is whether a given process network can be executed with bounded memory. Another undecidable question is whether a process network will halt (through deadlock or starvation).

One way to show the Turing completeness of KPN is to construct a universal Turing machine using only a small set of primitive processes that operate on only a binary data type. This can be quite tedious, but a major step in that direction is to show how to construct an unbounded stack. Two such stacks can be used to represent two sides of the tape in a Turing machine.

One possible implementation of a stack as a KPN is given in Figure 4.7. This example is one of the demos included in the [Ptolemy II distribution](#) ([Ptolemaeus, 2014](#)). The top-level network uses a pair of [Select](#) and [Switch](#) processes in a feedback loop with a [UnitDelay](#). This loop stores the contents of the stack in the KPN buffers between the processes. The [StackCounter](#) controls these to either accept a new input from the [Source](#) when a data value is being pushed onto the stack or pop a value off the stack and route it to the [Display](#) process. The [TestSequence](#) process produces a [true](#) to push a new value onto the stack and a [false](#) to pop a value off the stack.




Figure 4.7: A KPN network implementing an unbounded stack onto which tokens can be pushed and popped.

The **Select** process is the same as that used in Example 4.17. It chooses one of two input ports on which to perform a blocking read, depending on a boolean control input. The **Switch** process does the converse. It routes input tokens to one of two output ports depending on a boolean control input. The **UnitDelay** process is a parameterized version of the **unit delay** that we saw in Example 4.1. It outputs a constant and then just forwards input tokens to the output.

The hard work is done by the StackCounter subsystem. In addition to Switch, Select, and UnitDelay, this subsystem uses several instances of CondConst, the implementation of which is shown at the bottom. This subsystem discards false-valued inputs and, upon receiving a true-valued input, sends a constant value (given by a parameter) to the output port. The StackCounter also uses processes that perform logic operations on boolean-valued inputs and a process that reads and discards inputs (which has an icon that looks like an electrical ground).

How this network implements a stack is complicated. But the key point is that the memory of the stack is provided by buffers that store tokens that have been produced but not yet consumed. None of the individual processes provide any memory that could function as the Turing machine tape. The amount of memory used by these buffers is not bounded, like a Turing machine tape. It depends on the test sequence that is provided as input.

Using two of these stacks, it is possible to emulate a tape, where moving the tape to the left or right pops a token off one stack and pushes it onto the other. With some additional logic to control these pushes and pops, it is possible to build a universal Turing machine.

4.3.2 Effective Execution

Because of the Turing completeness of KPN, whether a process network has a **bounded execution** is undecidable. This means that there is no terminating algorithm that can, for all process networks, determine whether it has a bounded execution. Nevertheless, we seek a coordination mechanism that will deliver a bounded execution for every process network that has a bounded execution. In other words, we want to solve an undecidable problem.

Fortunately, we do not need to deliver an answer in bounded time. Only a process network that runs forever can possibly require unbounded memory. It is OK for

our coordination mechanism to take forever to give a definitive answer to the undecidable problem as long as it delivers a bounded execution when that is possible and delivers an **effective execution** when it is not.

A partial solution was given by **Parks (1995)**. Parks' solution delivers bounded and maximal execution for every process network that has a bounded execution. But it does not guarantee a **fair execution** and hence falls short of delivering an effective execution.

Parks' algorithm is simple. An execution starts with bounded buffers on all connections between processes. It does not matter what the bound is. When a process attempts to write to a buffer that is full, the write blocks. That is, the process stalls, waiting for room to become available. Now, one of three things can happen. First, the network might execute forever, continuing to produce and consume tokens, even though some processes are blocked on a write. This achieves an infinite (hence maximal) bounded execution. Second, the network might deadlock where all processes have either terminated or are blocked on a read. In this case, the process network specifies a finite execution, and the finite execution has been completed. The execution is again maximal and bounded. Third, the network might deadlock where at least one process is blocked on a write. In this third case, Parks' strategy finds the smallest buffer on which a process is write blocked and increases its capacity so as to unblock the write-blocked process. Execution continues until one of these three scenarios occurs again.

For many process networks, Parks' strategy achieves an **effective execution**. But not for all, as illustrated by the following example.

Example 4.20: Consider the process network in Figure 4.8. Suppose that all four processes can produce and consume infinite streams. Moreover, suppose that C is defined by the following Kahn-MacQueen procedure,

```
1 while (true) {  
2   t = read(in);  
3   write(out1, t);  
4   write(out1, t);  
5   write(out2, t);  
6 }
```

Suppose further that D is defined by the following Kahn-MacQueen procedure,

```

1 write(out, 0);
2 while (true) {
3   t = read(in2);
4   t = read(in1);
5   t = read(in1);
6   write(out, t);
7 }

```

It is easy to see that the least fixed point yields the infinite stream $s_2 = (0, 0, 0, \dots)$.

Now suppose that we execute this network using Parks' algorithm, starting with initial buffer sizes of 1. D successfully produces its first output, enabling C to complete the read on line 2. C will then perform the write on line 3, but the write on line 4 will block because the buffer capacity has been reached. As a consequence, D will block on the read on line 3. A **local deadlock** occurs.

Under Parks' algorithm, no corrective action is taken because processes A and B are able to continue running indefinitely. As a consequence, this is not a **fair execution**. To be fair, every process that is able to produce outputs must be allowed to do so. C is able to produce outputs, but we have blocked




Figure 4.8: Example of a process network where Parks' algorithm does not yield an effective execution.

it artificially as part of our execution policy (for this reason, the state of the execution is called a **local artificial deadlock**).

Geilen and Basten (2003) identified this flaw in Parks' algorithm, and proposed an alternative execution strategy. Their strategy is similar to that of Parks in that execution starts with arbitrarily bounded buffers. But while Parks' algorithm looks only for a **global artificial deadlock** (all processes are blocked, at least one on a write), Geilen and Basten's strategy looks for local artificial deadlock.

Specifically, consider a cycle of processes $(A_1, A_2, \dots, A_n)$, $n \geq 1$, where A_i is connected to A_{i+1} for $i < n$, and A_n is connected to A_1 . When we say " A_i is connected to A_{i+1} ," we mean that either A_i sends tokens to A_{i+1} or vice versa. So $(A_1, A_2, \dots, A_n)$ is an **undirected cycle**. If at any time during the execution there is an undirected cycle where all of these processes are blocked, at least one on a write, then we have a local artificial deadlock. When such an artificial deadlock is detected, the Geilen and Basten strategy will increase the size of at least one buffer on which a process in the cycle is write blocked.

Geilen and Basten (2003) prove that this strategy delivers an effective execution for **effective** process networks. But they make no guarantee about networks that are not effective.

One of the challenges in implementing both Parks' algorithm and Geilen and Basten's is performing a distributed deadlock detection. Interesting implementations are described by Olson and Evans (2005), Allen et al. (2007), and Dulloo and Marquet (2004).

For all of the examples we have considered so far, it is relatively easy to determine whether a **bounded execution** exists. This is not always the case, as illustrated by the following example.

Example 4.21: Consider for example the network shown in Figure 4.9. The output is an ordered sequence of integers of the form $2^n 3^m 5^k$, where n, m and k are non-negative integers. These are known as the **Hamming numbers**, and this program for computing them was studied by Dijkstra (1976) and Kahn and MacQueen (1977).

To understand how this network works, we need to understand what each of the processes does. The ones labeled ScaleN simply read an integer-




Figure 4.9: A process network for which it is difficult to determine appropriate buffer sizes.

valued input token, multiply it by N , and send the product to the output port. They repeat this sequence forever, as long as there are input tokens. The `UnitDelay` processes are parameterized versions of the `unit delay` that we saw in Example 4.1. They first produce an output token with value N , then repeat forever the following sequence: read the input token and send it to the output. As before, the black diamonds are `fork` processes. They direct a single token sequence to multiple destinations (over separate, unbounded buffers).

The processes labeled **OrderedMerge M** are the only nontrivial processes in this program. Given a sequence of numerically increasing tokens on each input port, they produce a numerically increasing merge of the two input sequences on their output ports. A Kahn-MacQueen implementation could look like this:

```

1 t1 = read(in1);
2 t2 = read(in2);
3 while(true) {
4   if (t1 < t2) {
5     write(out, t1);
6     t1 = read(in1);
7   } else {
8     write(out, t2);

```

```
9      t2 = read(in2);  
10  }  
11 }
```

You can now understand how this program generates the Hamming numbers. The loop at the upper left in the figure produces the sequence

$$(5, 25, 125, \dots, 5^n, \dots).$$

That sequence is merged into the second loop, which as a result produces the sequence

$$(3, 5, 9, 15, 25, \dots, 5^n 3^m, \dots).$$

That sequence is merged into the final loop, which produces the Hamming numbers, in numerically increasing order, without repetitions.

Although it is far from easy to see at a glance, it is possible to prove that any finite bound on buffer sizes will cause the network to fail to produce all the Hamming numbers. This network intrinsically requires unbounded buffers. It is *effective*, and both Parks' and Geilen and Basten's strategies deliver an *effective execution*.

4.4 Convergence to the Denotational Semantics

The *Kahn principle* states that maximal and fair executions produce sequences that match the least fixed point. But what do we mean by “match?” Every execution is finite at all times, and can only have produced a prefix of the least fixed point if that least fixed point includes infinite sequences. Intuitively, we would like such executions to converge in the limit to the least fixed point. But to make this intuition precise, we have to define a notion of convergence for sequences.

We use the notion of a topology.

Definition 4.6. Let X be any set. A collection τ of subsets of X is called a **topology** if three conditions are satisfied:

1. X and $\emptyset$ are members of τ .
2. The intersection of any two members of τ is in τ .
3. The union of any collection of members of τ is in τ .

For any topology τ , the members of τ are called the **open sets** of the topology. See the sidebar on page 136 for an explanation of how this relates to the usual notion of open sets of real numbers. In our case, we are not concerned with real numbers, but rather with sequences of tokens of arbitrary type.

There is an important but subtle distinction between conditions (2) and (3) above. A “collection of members of τ ” may be infinite. Thus, in a topology, the *union* of an infinite number of open sets is an open set. But the *intersection* of an infinite number of open sets may not be an open set. The intersection of any *finite* number of open sets is required to be an open set, but not of an *infinite* number of open sets.

Consider a set T and the set T^{**} of finite and infinite sequences of elements of T .

Definition 4.7. Given a finite sequence $s \in T^{**}$, an **open neighborhood** N_s around s is the set

$$N_s = \{s' \in T^{**} \mid s \sqsubseteq s'\}.$$

N_s is the set of sequences with prefix s .

N_s is an **upper set**,² meaning that N_s is **upward complete**. That is, it contains every element $a \in T^{**}$ that is greater than any $s' \in N_s$, $s' \sqsubseteq a$.

Let τ be the collection of all sets that are arbitrary unions of such open neighborhoods. Then τ is a topology. This topology is a particular kind of **Scott topology** (see sidebar on page 137).

²For a poset $(X, \leq)$ and subset $A \subseteq X$, an **upper set** $\uparrow A$ of A is defined to be

$$\uparrow A = \{x \in X \mid a \leq x \text{ for some } a \in A\}.$$

That is, an upper set of A is the set of all upper bounds of elements of A .

Notice that an open neighborhood is defined by a *finite* prefix. Were we to allow an infinite prefix, then a singleton set consisting of a single infinite sequence would be an open neighborhood. So infinite prefixes are not allowed.

Note further that an arbitrary intersection of open sets is not necessarily an open set. Consider any infinite chain $C \subset T^{**}$ containing only finite sequences. Each $c_i \in C$ defines an open neighborhood (the set of all sequences with c_i as a prefix). The intersection of all of these open neighborhoods can only contain sequences that have every $c_i \in C$ as a prefix. There is only one such sequence, $\bigvee C$, so this

Sidebar: Limit of a Sequence of Real Numbers

An **open neighborhood** around a in $\mathbb{R}$ is the set

$$N_{a,\varepsilon} = \{x \in \mathbb{R} \mid a - \varepsilon < x < a + \varepsilon\}$$

for some positive real number ε . An **open set** A in the reals is a subset of $\mathbb{R}$ such that for all $a \in A$, there is an open neighborhood around a that is a subset of A . The collection of such open sets in the reals is called the **standard topology** for the reals. It is a **topology** because it includes both $\mathbb{R}$ itself and $\emptyset$, it includes the intersection of any two such open sets, and it includes an arbitrary union of such open sets. Notice that it *does not* include any arbitrary intersection of open sets. Consider for example the intersection of the sets

$$N_{0,\varepsilon}, \quad \text{where } \varepsilon = 1/n, \quad n \in \{1, 2, 3, \dots\}.$$

The intersection of these sets is the singleton set $\{0\}$, which is not an open set.

Consider a sequence of real numbers

$$s: \mathbb{N} \rightarrow \mathbb{R}.$$

This sequence is said to converge to a real number a if for all open sets A containing a there exists an integer $n \in \mathbb{N}$ such that for all integers $m > n$ the following holds:

$$s(m) \in A.$$

This notion of convergence captures the intuition that the sequence eventually gets arbitrarily close to a .

Sidebar: Scott Topologies

Let X be a poset. A subset $O \subseteq X$ is **Scott open** if it is an **upper set** and all **directed sets** $D \subseteq X$ where $\bigvee D \in O$ have a non-empty intersection with O ; i.e., $D \cap O \neq \emptyset$.

The **Scott topology** on X is the set of Scott open subsets of X . The Scott topology is named after Dana Scott (1932–present), a renowned computer scientist and mathematician. It has a number of interesting properties.

First, for the poset $(T^{**}, \sqsubseteq)$, every open set of the Scott topology is an **upper set**. To see this, observe that an **open neighborhood** in this topology, by definition, is an upper set. Next, observe that the union of any family of upper sets is an upper set. Finally, since the open sets in the topology are unions of open neighborhoods, every open set must be an upper set.

A **Hausdorff space** is a topology in which distinct points have disjoint neighborhoods. Specifically, if X is a set and τ is a topology over X , then for any two $x_1, x_2 \in X$ where $x_1 \neq x_2$, there are two open sets $X_1, X_2 \in \tau$ where $x_1 \in X_1$, $x_2 \in X_2$, and $X_1 \cap X_2 = \emptyset$. As a mnemonic, in a Hausdorff space, two distinct points can be “housed off” by open sets.

The Scott topology over T^{**} is *not* a Hausdorff space. This property of the Scott topology accounts for much of the difficulty in reasoning about this topology because the standard topology over the reals *is* Hausdorff.

To see that the Scott topology is not Hausdorff, let $x_1, x_2 \in T^{**}$ be two sequences where $x_1 \neq x_2$ and $x_1 \sqsubseteq x_2$. These two distinct points cannot be “housed off” by open sets. Any open set that contains x_1 also contains x_2 because an open set is an upper set.

A **Kolmogorov space** is a topology in which, given two distinct points, at least one of them is in an open neighborhood that does not contain the other. The Scott topology over T^{**} is a Kolmogorov space. This means that distinct sequences are topologically distinguishable. Specifically, consider any two $x_1, x_2 \in T^{**}$ where $x_1 \neq x_2$. If neither is a prefix of the other, then each has an open neighborhood containing it that does not contain the other. Just find finite prefixes x_3 of x_1 and x_4 of x_2 such that $x_3 \neq x_4$; then the open neighborhood N_{x_3} contains x_1 but not x_2 and the open neighborhood N_{x_4} contains x_2 but not x_1 . If, on the other hand, $x_1 \sqsubseteq x_2$, then x_1 is necessarily finite. Let x_3 be the prefix of x_2 with length one greater than x_1 . Then open neighborhood N_{x_3} contains x_2 but not x_1 .

Sidebar: Three Views of Convergence

There are several ways to think of limits of sequences of sequences. Here, we compare a least upper bound as a limit, a metric space limit, and a topological limit. Consider set $D = \{0, 1\}$ and the CPO $A = D^{**}$, with $\sqsubseteq$, the prefix order. Consider the following sequence of sequences:

$$C = \{c_0, c_1, c_2, \dots\} = \{(0, 1), (0, 0, 1), (0, 0, 0, 1), \dots\}. \quad (4.1)$$

Note that no two elements of C are joinable, so the set itself does not have an upper bound, much less a least upper bound. However, in both a metric-space formulation and a topological formulation, this sequence converges to $c = (0, 0, 0, \dots)$, an infinite sequence of zeros.

For a metric-space formulation, we can use the [Baire distance](#) d of Definition B.5. The sequence (4.1) converges in the metric space (A, d) to c . That is, for any real number $\epsilon > 0$, there is a $K \in \mathbb{N}$ such that for all $i > K$, $d(c_i, c) < \epsilon$. For any $i \in \mathbb{N}$,

$$d(c_i, c) = 1/(i + 2).$$

Moreover, for all $j > i$, $d(c_j, c) < 1/(i + 2)$. For any $\epsilon > 0$, there is an i such that $1/(i + 2) < \epsilon$. Hence, the sequence converges to c in this metric space, even though it has no [LUB](#) in the prefix order.

The sequence (4.1) also converges in the [Scott topology](#) on A to $c = (0, 0, 0, \dots)$. An [open neighborhood](#) anchored by a finite sequence $b \in A$ is the set of all sequences having prefix b . In a Scott topology, any union of such open neighborhoods is called an open set. The sequence $(c_0, c_1, c_2, \dots)$ converges topologically to the infinite sequence $c = (0, 0, 0, \dots)$ if for all open sets N containing c , there is a $K \in \mathbb{N}$ such that for all $i > K$, $c_i \in N$.

Every open set containing c includes at least one open neighborhood anchored by a finite prefix of c . Every finite prefix of c is a finite sequence of zeros. Specifically, let $b_0 = \perp, b_1 = (0), b_2 = (0, 0)$, etc. Then an open neighborhood containing c is N_{b_j} for some $j \in \mathbb{N}$. For any such open neighborhood, let $K = j$, and note that $c_K \in N_{b_j}$ (because c_K has K leading zeros). Moreover, for all $i > K$, $c_i \in N_{b_j}$. Hence, the sequence converges topologically to c , even though it has no least upper bound.

intersection of open sets is a singleton set containing exactly one infinite sequence. This singleton set is not an open set.

Using our Scott topology, we can now define a notion of convergence of a sequence of sequences. Consider a sequence of sequences

$$S: \mathbb{N} \rightarrow T^{**}.$$

This sequence is said to converge to a sequence $a \in T^{**}$ if for all open sets A containing a , there exists an integer $n \in \mathbb{N}$ such that for all integers $m > n$, the following holds:

$$S(m) \in A.$$

Intuitively, this means that for any finite prefix $p \sqsubseteq a$, the sequences in S eventually all have prefix p .

In summary, in the Scott topology, for a sequence of sequences to converge to an infinite sequence, it is necessary for the sequences to grow without bound. They need to eventually produce every finite prefix of the infinite sequence.

This notion of convergence creates an interesting conundrum.

Definition 4.8. *A **convergent execution** of a process network is one that eventually produces every finite prefix of the least fixed point.*

A [maximal execution](#), by contrast, is only required to produce some prefix of the least fixed point. A maximal execution does not necessarily converge to the least fixed-point.

Example 4.22: For the example in Figure 4.5, a maximal execution is permitted to avoid executing process B. In particular, the output s_3 of process true is an infinite sequence of boolean true tokens, so the output s_2 of process B is never needed and will never be consumed by the Select. But if B is capable of producing an infinite sequence, then the least fixed point solution for s_2 is an infinite sequence. Any execution that converges to this least fixed point solution will exhaust available memory because tokens produced by B will not be consumed.

Maximal and fair executions, however, do converge. For [effective](#) process networks, Geilen and Basten's execution strategy, described in Section 4.3.2, will

deliver an [effective execution](#), and that effective execution will converge to the least fixed point. For networks that are not effective, however, the strategy is less satisfactory because it may exhaust available memory. For such networks, [Parks' algorithm](#) may be preferable, even though it does not converge to the least fixed point.

4.5 Beyond Kahn-MacQueen

So far, we have defined [Kahn-MacQueen processes](#) in terms of their implementation using an imperative language with [blocking read](#) and [nonblocking write](#) operations. All functions defined this way are continuous, but there are continuous functions that cannot be defined this way. In this section, we focus on the properties of the *functions* realized by a component rather than on their implementation as a Kahn-MacQueen process. This opens up the possibility of a richer set of implementations, including those that internally are concurrent.

4.5.1 Sequential Functions

Prior to Kahn and MacQueen's work, [Vuillemin \(1973\)](#) formally defined a class of functions that exactly match those implementable as a Kahn-MacQueen process. He called a function in this class a **sequential function**.

Let X^m be the set of m -tuples of sequences in $X = T^{**}$, for some data type T . For an m -tuple $x \in X^m$, let $x_i, i \in \{1, \dots, m\}$, be the i^{th} sequence in the tuple. Consider a function $F: X^m \rightarrow Y^n$.

Definition 4.9. $F: X^m \rightarrow Y^n$ is **sequential** if it is [continuous](#) and for any $x \in X^m$, there exists an $i \in \{1, \dots, m\}$, such that for all $x' \in X^m$ where $x \sqsubseteq x'$,

$$x_i = x'_i \implies F(x) = F(x'). \quad (4.2)$$

Intuitively, at all times during an execution, there is an input channel (the i^{th} one) that blocks further output. This is just like the Kahn-MacQueen blocking read, but it does not require the process to be implemented as a single imperative procedure.

Many useful functions, and even some trivial ones, however, are not sequential. An identity function with two or more input streams, for example, is not sequential.

Example 4.23: Consider a two-input, two-output identity function, $I_2: X^2 \rightarrow X^2$, where for all $x \in X^2$, $I_2(x) = x$. This is a rather trivial function. It is continuous, but not sequential. This identity function is not sequential because there is no i satisfying (4.2). Extending either input will extend the output, by the definition of the function.

Consider the example shown in Figure 4.10, where A is an arbitrary source producing the sequence x_2 . In the figure, $(x_1, x_2) = (y_1, y_2)$. Using this example, it is easy to see that there is no Kahn-MacQueen implementation of this identity function. Any blocking read on the upper port of I_2 will prevent the process from handling tokens on the lower port. And any implementation that does no reads at all of the upper port will not implement the identity function if the input x_1 is not empty.

This identity function, however, is continuous (and, hence, monotonic), so the program in Figure 4.10 is deterministic. The least fixed point for Figure 4.10 yields an empty sequence for x_1 and y_1 and $x_2 = y_2$ is whatever sequence A produces.

The sequential property is not compositional. The one-input, one-output identity function I_1 is sequential (trivially). However, a parallel composition of two such functions, shown in Figure 4.11 realizes the two-input, two-output identity function I_2 , which is not sequential. On the other hand, any such composition of continuous functions is continuous, so continuity *is* a compositional property.




Figure 4.10: An example illustrating that the two-input, two-output identity function is not sequential.




Figure 4.11: An example illustrating that sequential functions are not compositional.

4.5.2 Stable Functions

Berry (1976) defined a larger class of functions than sequential functions that he called **stable functions**.

Definition 4.10. A function $F: X \rightarrow Y$ on a *complete lower semilattice* is **stable** if it is *continuous* and for all *joinable* sets $C \subseteq X$, $\hat{F}(C)$ is joinable and

$$\bigwedge \hat{F}(C) = F\left(\bigwedge C\right). \quad (4.3)$$

Intuitively, if two possible inputs do not contain contradictory information (they can later be joined, i.e., they have a least upper bound), then neither will the two corresponding outputs. Sequential functions are stable, but not all stable functions are sequential. The two-input, two-output identity function on streams is (trivially) stable. But so are more interesting functions.

Example 4.24: Let $X = \{\perp, 0, 1\}$ with the *flat partial order*, where $\perp < 0$ and $\perp < 1$. The **Gustave function** (Berry, 1976) is the least function $G: X^3 \rightarrow X$

s.t.

$$G(0, 1, \perp) = 0$$

$$G(1, \perp, 0) = 0$$

$$G(\perp, 0, 1) = 0$$

This example is easily extended to a function that operates on sequences in $\{0, 1\}^*$ and hence can become a component in a Kahn network. Informally, such a function takes three sequences in. It can output a length-one sequence (0) if the head of the first two input sequences is 0 and 1, respectively, or the head of the first and third sequence is 1 and 0, respectively, or the head of the second and third sequence is 0 and 1, respectively. For all other inputs, the output is an empty sequence. In each of these three cases, whether input tokens are available on the third stream is irrelevant. Hence, a blocking read will not work because it is impossible to tell which input to block on.

Note that although equation (4.3) looks quite like the definition of [continuous](#), it uses the meet rather than the join. Not all functions that satisfy (4.3) are continuous. Consider a function $F : X \rightarrow Y$, where $X = Y = \mathbb{N} \cup \{\infty\}$, with the usual numerical ordering, given by

$$F(x) = \begin{cases} \perp_X = 0 & \text{if } x \text{ is finite} \\ x & \text{otherwise} \end{cases}$$

We can show that this function is not continuous. Consider the chain $C = \mathbb{N}$. This chain is joinable, where $\bigvee C = \infty$. In this case, $\hat{F}(C) = \{\perp_X\} = \{0\}$, which is obviously joinable, where $\bigvee \hat{F}(C) = \perp_X = 0$. But $F(\bigvee C) = F(\infty) = \infty$, which is not equal to $\perp_X$, so the function is not continuous.

However, we can show that for all joinable sets $C \subseteq X$, $\hat{F}(C)$ is joinable, and the function F does satisfy (4.3). Given a joinable set $C \subseteq X$, $\hat{F}(C)$ can only take on three possible values, $\{0\}$, $\{0, \infty\}$, or $\{\infty\}$, all three of which are joinable, so $\hat{F}(C)$ is joinable. The second constraint, that $\bigwedge \hat{F}(C) = F(\bigwedge C)$, is easy to check. Note that $\bigwedge C$ will be finite unless $C = \{\infty\}$. Hence,

$$F(\bigwedge C) = \begin{cases} 0 & \text{if } C \neq \{\infty\} \\ \infty & \text{otherwise} \end{cases}$$

For the right hand side, $\bigwedge \hat{F}(C)$, it is equal to the same expression,

$$\bigwedge \hat{F}(C) = \begin{cases} 0 & \text{if } C \neq \{\infty\} \\ \infty & \text{otherwise} \end{cases}$$

Stable functions like these represent a useful class of Kahn processes that react to patterns across their inputs in well-defined ways. They can be more difficult to implement, however, because an imperative process with blocking reads will not provide an implementation.

4.6 Dynamic Kahn Networks

An interesting feature of Kahn networks is that it is possible to grow them dynamically. [Kahn and MacQueen \(1977\)](#) give an example implementing the **sieve of Eratosthenes** in a manner similar to that shown in Figure 4.12. The program starts with just the Generator process, which, on its first line, spawns a new Filter process and connects its output to the input of that process. The new Filter process accepts integer inputs and discards any that are multiples of its prime parameter. The Generator process then starts feeding the Filter process increasing integers starting with 3.

When the Filter first encounters an integer that is not a multiple of its prime parameter, it has discovered a new prime number, so it spawns another downstream Filter to filter out multiples of that prime number. After that first encounter, it just forwards numbers that are not multiples of its prime. The chain of processes grows indefinitely as new prime numbers are discovered.

Semantically, this process network can be modeled as an infinite chain of processes. Each process realizes a continuous function, so the [least fixed-point semantics](#) still works. Using the strategy illustrated in Figure 4.3, we can construct a composite function with an infinite number of input and output streams. The composite function is continuous, so the program has a single, well-defined, infinite behavior given by the least fixed point of the composite function.




Figure 4.12: Sieve of Eratosthenes as a dynamically growing process network.

4.7 Nondeterministic Extensions of Kahn Networks

Determinism is the most notable feature of Kahn networks. But some applications benefit from nondeterministic mechanisms, for example, situations where components of a system appear and disappear for reasons unrelated to the computation, such as faults. It is possible to extend Kahn networks with nondeterministic mechanisms. The result is systems that are deterministic by default, with nondeterministic mechanisms judiciously introduced only exactly where needed. This contrasts with the techniques of Chapter 3, which are nondeterministic by default and leave it to the programmer to figure out how to achieve determinism when needed. It also contrasts with some stream processing formalisms that embrace nondeterminism from the start, as described for example by [Stephens \(1997\)](#) and [Broy and Stefanescu \(2001\)](#). A basic building block to judiciously introduce nondeterminism is the nondeterministic merge.

4.7.1 Nondeterministic Merge

A common extension to Kahn networks is a **nondeterministic merge**, a component with some number of input streams that nondeterministically merges tokens from those streams onto a single output stream. A Hewitt-Agha [actor](#), as described in Section 3.2, can be modeled as a single-input Kahn process preceded by a nondeterministic merge.

A simple realization of a nondeterministic merge is given in Ptolemy II ([Ptolemaeus, 2014](#), Chapter 4). There, a nondeterministic merge with n input ports spawns a [thread](#) for each port with an infinite loop that performs a blocking read on that port. When an input token arrives, the thread consumes the token, acquires a [mutex](#) lock, and writes the token to the output stream. When inputs arrive nearly simultaneously, the threads compete to acquire the mutex, and the resulting race determines the order in which output tokens are produced.

A nondeterministic merge is useful when upstream components may or may not be present. For example, their network connectivity may be sporadic, in which case, a deterministic Kahn node that reads their token would be blocked until connectivity is restored. The nondeterministic merge only blocks one of its threads, leaving open the possibility of forwarding tokens from other input ports.

The introduction of such nondeterminism, however, undermines fundamental features of the semantics of the network. The nondeterministic merge can no longer be modeled as a function from a tuple of input streams to an output stream. [Plotkin \(1976\)](#) suggests instead a **powerdomain** construction, in which the nondeterministic merge is modeled as a function from tuples of input streams into a set of allowable output streams:

$$F: X^n \rightarrow 2^X, \quad (4.4)$$

where n is the number of input streams, $X = T^{**}$ is the set of streams of tokens of type T , and 2^X is the set of all subsets of X . For a particular n -tuple $x \in X^n$ of input streams, $F(x)$ will be the set of all possible interleavings of the tokens in those streams.

Unfortunately, the powerdomain construction fails to distinguish some components that are different. Two components may be modeled by exactly the same function of the form of (4.4) and yet can exhibit significantly different behavior, as we will see next.

4.7.2 The Brock-Ackerman Anomaly[†]

[Brock and Ackerman \(1981\)](#) give an example now known as the **Brock-Ackerman anomaly** of two nondeterministic stream operators that realize identical powerdomain functions of the form of (4.4) and yet can exhibit significantly different behavior.

Consider the system shown in Figure 4.13. This system has two inputs at the left that can accept streams of integer-valued tokens and one output at the right that produces a stream of integers. The subsystems labeled “Repeat” take each input integer and repeat it twice on their outputs. For example, if the input to the top Repeat is the sequence (1, 2), then the output will be the sequence (1, 1, 2, 2). The subsystem labeled “Merge” is a nondeterministic merge. For example, given the two input sequences (1, 2) and (3, 4), it can produce any of (1, 2, 3, 4), (1, 3, 2, 4), (1, 3, 4, 2), (3, 4, 1, 2), (3, 1, 4, 2), or (3, 1, 2, 4). If Merge receives nothing on one of its inputs, then it will simply produce whatever it receives on the other input. The

[†] This section is adapted from [Lee \(2022\)](#).




Figure 4.13: A system with two distinguishable implementations that realize the same powerdomain function.

subsystem labeled “FirstTwo” simply outputs the first two inputs it receives. For example, given $(1, 2, 3, 4)$, it will produce $(1, 2)$.

[Brock and Ackerman \(1981\)](#) then gave two subtly different realizations of the FirstTwo subsystem:

1. The first realization produces outputs as it receives inputs. That is, as soon as it sees a 1 on its input, it will produce 1 on its output.
2. The second realization waits until there are two inputs available before producing any output. That is, it will not produce any output until both 1 and 2 are available, at which point it will produce the sequence $(1, 2)$.

If you model this “BrockAckerman” as a function of the form (4.4), you will get exactly the same function for these two different realizations. For example, if the system is presented with inputs (5) and (6) , i.e., two sequences of length one, the possible outputs for either realization are $(5, 5)$, $(5, 6)$, $(6, 5)$, and $(6, 6)$. The choice of realization has no effect on these possibilities.

Nevertheless, the two realizations yield different behaviors in some circumstances. Consider the system in Figure 4.14. The subsystem labeled “Increment” simply adds one to each input. For example, given the input sequence $(1, 2)$, it will produce $(2, 3)$. In this usage, it makes a difference which of the two realizations of FirstTwo is used. Suppose that the subsystem labeled “Source” provides on its output the length-one sequence (5) . Under realization (1) of FirstTwo, there are two possible outputs from the BrockAckerman subsystem, $(5, 5)$ and $(5, 6)$. But




Figure 4.14: A use of the system in Figure 4.13 where the two variants of the FirstTwo subsystem yield different behaviors.

under realization (2), there is only one possible output, (5, 5). The powerdomain construction fails to model this difference.

One way to understand the system in Figure 4.14 is to think of the context of the “BrockAckerman” component as **interacting** with the component. It provides inputs that are based on observations of previous outputs via a **feedback** loop. This example illustrates that such interaction is fundamentally different from passive **observation**, where the inputs provided to a system do not depend on observations (see Wegner (1998) for an in-depth discussion of the power of interaction). Hence, with this example, Brock and Ackerman (1981) proved that two systems that are indistinguishable by a passive observer cannot be substituted one for the other without possibly changing the behavior. They are not equivalent.

Note that the Brock-Ackerman system is not a Turing-Church computation because of the nondeterministic merge component. The Turing-Church theory admits no such nondeterminism. In Chapter 12 of Lee (2020), I show that a passive observer, one that can only see the inputs and outputs of a system, cannot tell the difference between such a nondeterministic system and a deterministic one (a deterministic one *would* be a Turing-Church computation). Only through interaction with the system is it possible to tell the difference. The feedback shown in Figure 4.14 is one form of such interaction. That argument depends on another celebrated result in computer science, Milner’s concept of bisimulation, but this is beyond the scope of this book.

4.8 Rendezvous

In Section 4.3, we considered the use of blocking writes to keep buffers bounded. An extreme form of such blocking writes uses no buffer at all, requiring instead that the sending process **rendezvous** with the receiving process. A write to a stream blocks until a corresponding read from the stream is performed, and vice versa, a read from a stream blocks until a corresponding write is performed. Such rendezvous semantics underlie Hoare's **communicating sequential processes (CSP)** (Hoare, 1978) and Milner's **calculus of communicating systems (CCS)** (Milner, 1980).

By itself, a rendezvous-based framework for distributed computation is not Turing complete. Its memory use is bounded (at zero) by construction, and hence it cannot possibly implement an arbitrary Turing machine. Such frameworks rely instead on the programs that constitute the communicating processes for expressiveness.

It is possible to restrict rendezvous-based frameworks to be deterministic. To do so, each component in the system can be constrained to be a sequential imperative program that can rendezvous (read or write) with only one other process at a time. However, this restriction makes it quite difficult to construct useful programs. As a consequence, most such frameworks are nondeterministic. They include primitives where a process can declare a willingness to rendezvous with any of several other processes at one step. With this, the order in which multiple rendezvous occur is nondeterministic. Most processes, therefore, cannot be modeled as functions except by using the powerdomain construction. But that construction fails to distinguish processes that may behave differently.

4.9 Conclusion

This chapter has developed a rigorous foundation for understanding process networks as a model of deterministic concurrency. Beginning with Kahn-MacQueen process networks, we show how a concurrent system composed of sequential processes communicating through unbounded FIFO channels can nevertheless define a unique, deterministic semantics. The key insight is that such networks correspond to continuous functions on complete partial orders of streams, and

their meaning is given by the least fixed point of a global functional representation of the network.

By framing process networks in terms of monotonicity, continuity, and least fixed points, we separate concerns of what a network means from how it is executed. This separation allows executions to be analyzed as approximations of a semantic object that may itself be infinite. The Kahn principle provides a powerful guarantee: any fair and maximal execution converges to the least fixed point, regardless of scheduling. This result explains why determinism can be preserved even in highly concurrent settings.

At the same time, the chapter emphasizes that execution is not merely a technical detail. Practical execution strategies must confront fundamental limitations arising from Turing completeness, including the undecidability of deadlock, termination, and bounded memory execution. Through a series of examples, we see that fairness, maximality, and boundedness can be mutually incompatible, and that naive demand-driven or data-driven policies are insufficient. More sophisticated strategies, such as those of Parks and of Geilen and Basten, demonstrate how effective execution can be achieved in many cases, while also clarifying the inherent tradeoffs.

We also broaden the model beyond imperative Kahn-MacQueen processes by characterizing components directly in terms of the functions they define. Sequential and stable functions generalize the original process model while preserving determinacy, revealing continuity as the essential compositional property. This functional perspective accommodates richer internal implementations, including concurrent ones, without sacrificing the semantic guarantees of process networks.

Finally, by introducing a notion of convergence based on the Scott topology, we make precise what it means for an execution to approach its semantic ideal. This notion explains why some maximal executions fail to converge and why convergence itself may be undesirable for networks that intrinsically require unbounded resources.

Taken together, these results establish process networks as a mathematically robust and practically relevant model of deterministic concurrency. They also expose the deep connections between concurrency, computation, and domain theory that will recur throughout the remainder of this book. Subsequent chapters build

on these ideas to explore alternative models, richer forms of interaction, and different tradeoffs between expressiveness, analyzability, and implementability.

Dataflow

5.1	Firing and Firing Rules	155
	<i>Sidebar: Dataflow Languages and Frameworks</i>	155
5.1.1	Dataflow Processes	156
5.1.2	Feedback and Stateful actors	163
5.1.3	Fixed Point Semantics with Initial Tokens	165
5.1.4	Commutative Firings	167
5.1.5	Scheduling	171
5.2	Synchronous Dataflow	172
5.2.1	Balance Equations	174
5.2.2	Solving the Balance Equations	177
5.2.3	Properties of the Connectivity Matrix	179
5.2.4	Dynamics of Execution	181
5.2.5	Parallel Scheduling	185
	<i>Sidebar: SDF Schedulers</i>	186
5.3	Boolean Dataflow	189
	<i>Sidebar: StreamIt</i>	190
	<i>Sidebar: Multidimensional Synchronous Dataflow</i>	190
5.4	Modal Dataflow	196
5.4.1	Cyclo-Static Dataflow	196
	<i>Sidebar: Petri Nets</i>	197
5.4.2	Heterochronous Dataflow	199
5.4.3	SDF Scenarios	202
5.4.4	Parameterized SDF	202

5.5 Conclusion 202

Broadly, the term **dataflow** (or **data flow**) is used in the literature for any model where each chunk of computation is (at least conceptually) performed when its input data are available. This contrasts with **imperative** models of computation, where the sequence of computational steps is explicitly specified. The **Kahn-MacQueen process** networks of the previous chapter combine imperative and dataflow models, in that each process is assumed to be specified imperatively, but processes stall to await availability of data.

In this chapter, we consider a purer form of dataflow originally attributed to **Dennis (1974)**, who envisioned it as an alternative to the standard Von Neumann computer architecture. In Dennis dataflow, primitive operations, such as addition or multiplication, are enabled when their input data are available rather than when a program counter reaches an instruction specifying the operation. A program is given as a graph or network connecting such operations. Dennis called the primitive operations **actors**, but they are quite different from Hewitt-Agha **actors**, so we will use the more verbose term “dataflow actor”.

Dennis envisioned dataflow actors to be primitive operations performed in hardware, like the instructions in a Von Neumann **instruction set architecture (ISA)**. His work led to considerable experimentation with dataflow computer architectures (**Arvind et al., 1991; Srini, 1986**). Although such machines were capable, in principle, of exploiting parallel hardware more effectively, none of them were commercially successful. This may have been partly due to the lack of competitive programming models. The programming languages of the day were not easily compiled for execution on a dataflow machine, and languages developed specifically for the machines (**Johnston et al., 2004**) were slow to catch on.

The dataflow concept need not be restricted to the microarchitecture level. Dataflow actors can equally well be complex operations, even those specified using imperative code in conventional languages. As a consequence, the connection with the Kahn networks of the previous chapter is quite strong. Dennis dataflow can be viewed as a special case of Kahn process networks (**Lee and Parks, 1995**).

The concepts behind dataflow have had practical realizations in stream languages (**Stephens, 1997**) and operating systems (such as Unix pipes). Interest in these models of computation was rekindled by the re-emergence of parallel computing

in the form of multicore architectures (Creeger, 2005). Dataflow models of computation have been explored for programming such parallel machines (Gordon et al., 2002), distributed systems (Lázaro Cuadrado et al., 2007; Olson and Evans, 2005; Parks and Roberts, 2003), and embedded systems (Lin et al., 2006; Jantsch and Sander, 2005). Considerable effort has gone into improved execution policies (Thies et al., 2005; Geilen and Basten, 2003; Turjan et al., 2003; Lee and Parks, 1995) and standardization (Object Management Group (OMG), 2007; Hsu et al., 2004).

In this chapter, we consider programs given as networks of dataflow actors and use the theory developed in the previous chapter to find conditions under which such a network is **deterministic**. We then consider special cases that yield to useful analysis and scheduling optimizations.

5.1 Firing and Firing Rules[†]

Let $S = T^{**}$ be the set of finite or infinite sequences of **tokens** with some data type T .

Definition 5.1. A *dataflow actor* with m inputs and n outputs is a pair (R, f) , where

1. $R \subseteq S^m$ is a non-empty set of m -tuples of finite sequences;
2. $f : S^m \rightarrow S^n$ is a (possibly partial) function defined at least on R ;
3. The sequences in the tuple $f(r)$ are finite for every $r \in R$;
4. for all $r, r' \in R$, if $r \neq r'$, then $\{r, r'\}$ does not have an upper bound in $(S^m, \sqsubseteq)$.

We call each $r \in R$ a **firing rule**, and f the **firing function** of the dataflow actor.

The last condition is equivalent to the following statement: for any given m -tuple $s \in S^m$, there is at most one firing rule r in R such that r is a prefix of s , $r \sqsubseteq s$. If

[†]This section is based on Lee and Matsikoudis (2009).

there were two such firing rules in R , then s would be an upper bound for these two, contradicting condition 4.

If $m = 0$, then R is the singleton set $\{\emptyset\}$, so condition 4 is trivially satisfied. If $n = 0$, then condition 3 is trivially satisfied.

5.1.1 Dataflow Processes

Let (R, f) be a dataflow actor with m inputs and n outputs. We can define a Kahn process $F : \mathcal{S}^m \rightarrow \mathcal{S}^n$ based on this actor, and a reasonable condition to impose is that for any $s \in \mathcal{S}^m$,

$$F(s) = \begin{cases} f(r).F(s') & \text{if there exists } r \in R \text{ such that } s = r.s'; \\ \lambda^n & \text{otherwise.} \end{cases} \quad (5.1)$$

Here, λ^n is an n -tuple of empty sequences. This is not a definition of F because F appears on both sides of the equality. The following, however, is a definition:

Definition 5.2. A *dataflow process* corresponding to a *dataflow actor* (R, f) is the least function $F : \mathcal{S}^m \rightarrow \mathcal{S}^n$ that satisfies (5.1).

This requires some elaboration. It is not obvious that such an F exists, nor, if it does exist, that it is unique, nor that any such F is *continuous*, and therefore defines a Kahn process. Although not obvious, all three of these statements are true.

First, for there to be a “least function,” we need an order relation on *functions* from tuples of sequences to tuples of sequences rather than merely on tuples of sequences. Let $\mathcal{S}^m \rightarrow \mathcal{S}^n$ denote the set of all functions from $\mathcal{S}^m$ into $\mathcal{S}^n$. We can define a **prefix order** on this set of functions that is analogous to the *pointwise order* on tuples of sequences.

Definition 5.3. The function $F : \mathcal{S}^m \rightarrow \mathcal{S}^n$ is a *prefix* of the function $G : \mathcal{S}^m \rightarrow \mathcal{S}^n$, written $F \sqsubseteq G$, if and only if $F(s) \sqsubseteq G(s)$ for any m -tuple s .

With this definition $(\mathcal{S}^m \rightarrow \mathcal{S}^n, \sqsubseteq)$ is a *poset*. Moreover, it is a *pointed* poset. Its bottom element is the **empty function**

$$\Lambda_m^n : \mathcal{S}^m \rightarrow \mathcal{S}^n \quad \text{such that for all } s \in \mathcal{S}^m, \Lambda_m^n(s) = \lambda^n. \quad (5.2)$$

The bottom element is simply a function that yields an n -tuple of empty sequences for all arguments.

Proposition 5.1. *The ordered set $(\mathcal{S}^m \rightarrow \mathcal{S}^n, \sqsubseteq)$ is a CPO.*

Proof. This follows immediately from the fact that $(\mathcal{S}^n, \sqsubseteq)$ is a CPO. $\square$

Now consider a function that operates on functions,

$$\phi : (\mathcal{S}^m \rightarrow \mathcal{S}^n) \rightarrow (\mathcal{S}^m \rightarrow \mathcal{S}^n).$$

Functions that operate on functions are sometimes called **functionals**.

Definition 5.4. *For a dataflow actor (R, f) with m inputs and n outputs, we can define a **dataflow actor functional** ϕ such that for any $F \in \mathcal{S}^m \rightarrow \mathcal{S}^n$ and any m -tuple s ,*

$$\phi(F)(s) = \begin{cases} f(\mathbf{r}).F(s') & \text{if there exists } \mathbf{r} \in R \text{ such that } s = \mathbf{r}.s'; \\ \lambda^n & \text{otherwise.} \end{cases} \quad (5.3)$$

Unlike (5.1), this *is* a definition because ϕ does not appear to the right of the equality.

Theorem 5.1. *ϕ is monotonic.*

Proof. Let F_1 and F_2 be arbitrary functions of type $\mathcal{S}^m \rightarrow \mathcal{S}^n$, and suppose that $F_1 \sqsubseteq F_2$. We need to show that $\phi(F_1) \sqsubseteq \phi(F_2)$.

For any $s \in \mathcal{S}^m$, if there is a firing rule $\mathbf{r} \in R$ such that $\mathbf{r} \sqsubseteq s$, then by condition 4 in definition 5.1, $\mathbf{r}$ is the only firing rule that is a prefix of s , and hence $\phi(F_1)(s) = f(\mathbf{r}).F_1(s')$ and $\phi(F_2)(s) = f(\mathbf{r}).F_2(s')$, where $s = \mathbf{r}.s'$. However, by assumption, $F_1(s') \sqsubseteq F_2(s')$ for any m -tuple s' , and hence $\phi(F_1)(s) \sqsubseteq \phi(F_2)(s)$.

If there is no firing rule that is a prefix of s , then $\phi(F_1)(s) = \lambda^n = \phi(F_2)(s)$.

In either case, $\phi(F_1)(s) \sqsubseteq \phi(F_2)(s)$, and hence ϕ is monotonic. $\square$

Theorem 5.2. ϕ is *continuous*.

Proof. Let $D \subseteq \mathcal{S}^m \rightarrow \mathcal{S}^n$ be a *directed set* in $(\mathcal{S}^m \rightarrow \mathcal{S}^n, \sqsubseteq)$, and $s \in \mathcal{S}^m$ an arbitrary m -tuple. We need to show that $\bigvee \hat{\phi}(D) = \phi(\bigvee D)$, where $\hat{\phi}$ is the *lifted* version of ϕ .

If there is a firing rule $r \in R$ such that $r \sqsubseteq s$, then by condition 4 in definition 5.1, r is unique, and hence for every $F \in \mathcal{S}^m \rightarrow \mathcal{S}^n$, $\phi(F)(s) = f(r).F(s')$, where $s = r.s'$. Thus,

$$\begin{aligned} \bigvee \{\phi(F)(s) \mid F \in D\} &= \bigvee \{f(r).F(s') \mid F \in D\} \\ &= f(r). \bigvee \{F(s') \mid F \in D\} \\ &= f(r).(\bigvee D)(s') \\ &= \phi(\bigvee D)(s). \end{aligned}$$

Notice that since D is directed in $(\mathcal{S}^m \rightarrow \mathcal{S}^n, \sqsubseteq)$, it has a least upper bound therein, $(\mathcal{S}^m \rightarrow \mathcal{S}^n, \sqsubseteq)$ being a CPO.

If there is no firing rule $r \in R$ such that $r \sqsubseteq s$, for every $F \in \mathcal{S}^m \rightarrow \mathcal{S}^n$, $\phi(F)(s) = \lambda^n$, and hence

$$\bigvee \{\phi(F)(s) \mid F \in D\} = \lambda^n = \phi(\bigvee D)(s).$$

In both cases,

$$\bigvee \{\phi(F)(s) \mid F \in D\} = \phi(\bigvee D)(s),$$

and hence ϕ is continuous. □

Since ϕ is a continuous function over the CPO $(\mathcal{S}^m \rightarrow \mathcal{S}^n, \sqsubseteq)$, by the *Kleene fixed-point theorem* (proposition A.6), it has a least fixed point F in $(\mathcal{S}^m \rightarrow \mathcal{S}^n, \sqsubseteq)$ which must satisfy (5.1).

Moreover, we can give a constructive procedure for finding that least fixed point. We can start with the least element in $(\mathcal{S}^m \rightarrow \mathcal{S}^n, \sqsubseteq)$, the bottom function Λ given

in (5.2), and iterate ϕ to obtain the following sequence of functions:

$$\begin{aligned}
 F_0 &= \Lambda_m^n, \\
 F_1 &= \phi(F_0), \\
 F_2 &= \phi(F_1), \\
 &\dots
 \end{aligned}
 \tag{5.4}$$

Since ϕ is monotonic, and Λ_m^n is a prefix of every other function in $\mathcal{S}^m \rightarrow \mathcal{S}^n$, the set $\{F_0, F_1, F_2, \dots\}$ is a chain and hence directed in $(\mathcal{S}^m \rightarrow \mathcal{S}^n, \sqsubseteq)$. Thus, it has a least upper bound therein, and that least upper bound is the least fixed point of ϕ , by the [Kleene fixed-point theorem](#).

Let us examine this chain more closely for some fixed m -tuple s . Suppose that there is some sequence of firing rules $(r_1, r_2, \dots)$ such that $s = r_1.r_2.\dots$. Then, for this particular m -tuple, we can rewrite (5.4) in the following form:

$$\begin{aligned}
 s_0 &= F_0(s) = \lambda^n \\
 s_1 &= F_1(s) = f(r_1) \\
 s_2 &= F_2(s) = f(r_1).f(r_2) \\
 &\dots
 \end{aligned}
 \tag{5.5}$$

This is an exact description of the [operational semantics](#) in Dennis dataflow for a single actor. Start with the dataflow actor producing only the empty sequence. Then find the prefix of the input that matches a firing rule, and invoke the firing function on that prefix, producing a partial output. Notice here that because of condition 4 in definition 5.1, no more than one firing rule can match a prefix of the input at any time. Then find the prefix of the remaining input that matches another firing rule, invoke the firing function on that prefix, and concatenate the result with the output.

In general, even when s is infinite, it is possible that there is only a finite sequence of firing rules $(r_0, \dots, r_p)$ such that $s = r_0.\dots.r_p.s'$, with s' having no prefix in R . In both the [operational semantics](#) of Dennis dataflow and a [least fixed-point semantics](#), the firings simply stop, and the output is finite.

When $m = 0$ (there are no inputs), then $\mathcal{S}^m$ is a [singleton set](#) $\{\emptyset\}$. In this case, the least fixed point of ϕ is a source process, and if $\emptyset \in R$, then it produces the sequence $f(\emptyset).f(\emptyset).\dots$. If $f(\emptyset)$ is non-empty, then the dataflow process produces a tuple of output sequences that is infinite and periodic. This might seem limiting

for dataflow processes that act as sources, but in fact it is not; a source with a more complicated output sequence can be constructed using a feedback composition, as we will see.

When $n = 0$ (there are no outputs), then the least fixed point of ϕ is a sink process, producing the sequence $\emptyset.\emptyset.\dots$. We can define the concatenation of empty sets to be the empty set, $\emptyset.\emptyset.\dots = \emptyset$. Note that this is not the same as an empty sequence nor a tuple of empty sequences. It is no sequences at all.

In view of this perfect coincidence with the operational semantics, we are tempted to define a Kahn process based on the dataflow actor (R, f) as this least fixed point of ϕ . But in order to do this, we still need to prove that, in the general case, this least fixed point of ϕ is itself a continuous function, and thus a Kahn process. It suffices to prove the following theorem:

Theorem 5.3. *For any $F : \mathcal{S}^m \rightarrow \mathcal{S}^n$, if F is *continuous*, then $\phi(F)$ is also continuous for any *dataflow actor functional* ϕ .*

Proof. Let $F : \mathcal{S}^m \rightarrow \mathcal{S}^n$ be a continuous function, and $D \subseteq \mathcal{S}^m$ a *directed set* in $(\mathcal{S}^m, \sqsubseteq)$.

First, we show that there is at most one $r \in R$ that is a prefix of any tuple in D . Suppose, toward contradiction, that there are $r_1, r_2 \in R$ and $s_1, s_2 \in D$ such that $r_1 \neq r_2$, but $r_1 \sqsubseteq s_1$ and $r_2 \sqsubseteq s_2$. Then since D is directed in $(\mathcal{S}^m, \sqsubseteq)$, $\{s_1, s_2\}$ has an upper bound in D , which is also an upper bound of $\{r_1, r_2\}$, in contradiction to condition 4 in definition 5.1.

If there is such an $r \in R$, then

$$\begin{aligned} \bigvee \{\phi(F)(s) \mid s \in D\} &= \bigvee \{f(r).F(s') \mid r.s' \in D\} \\ &= f(r). \bigvee \{F(s') \mid r.s' \in D\} \\ &= f(r).F(\bigvee \{s' \mid r.s' \in D\}) \\ &= \phi(F)(\bigvee D). \end{aligned}$$

The first equality follows because for any $s \in D$ where $r \not\sqsubseteq s$, $\phi(F)(s) = \lambda^n$, and hence does not contribute to the least upper bound.

The last equality follows because $\bigvee D$ must have prefix r , where $\bigvee D = r.s''$ and $s'' = \bigvee \{s' \mid r.s' \in D\}$. To see this, let $s_0 \in D$ be such that $r \sqsubseteq s_0$. For any other

$s \in D$, because D is directed, $s_0 \vee s \in D$ and has prefix r . From this, it follows that $\bigvee D = r.s''$ and $s'' = \bigvee \{s' \mid r.s' \in D\}$.

If there is no $r \in R$ that is a prefix of some tuple in D , then there is no $r \in R$ that is a prefix of $\bigvee D$. This follows because condition 1 in definition 5.1 requires that r be finite, and hence, if any $r \in R$ is a prefix of $\bigvee D$, then it must also be a prefix of some tuple in $s \in D$. Hence,

$$\phi(F)(\bigvee D) = \lambda^n.$$

Moreover,

$$\bigvee \{\phi(F)(s) \mid s \in D\} = \lambda^n.$$

Hence, whether is an $r \in R$ that is a prefix of some tuple in D or not,

$$\bigvee \{\phi(F)(s) \mid s \in D\} = \phi(F)(\bigvee D),$$

and hence $\phi(F)$ is continuous. □

Since Λ_m^n is trivially continuous, an easy induction suffices to see that the least fixed point of ϕ is a continuous function.

Note here that the firing function f need not be continuous. In fact, it does not even need to be monotonic. The continuity of the least fixed point of ϕ is guaranteed if (R, f) is a valid dataflow actor satisfying the conditions in definition 5.1.

Example 5.1: Consider a system where the set of token values (the data type) is the binary digits, $\mathbb{B} = \{0, 1\}$. Let us examine some possible sets $R \subset \mathcal{S}$ of firing rules for unary firing functions $f : \mathcal{S} \rightarrow \mathcal{S}$.

The following sets of firing rules all satisfy condition 4 in definition 5.1:

$$\begin{aligned} R &= \{\lambda\}; \\ R &= \{(0)\}; \\ R &= \{(0), (1)\}; \\ R &= \{(0, 0), (0, 1), (1, 0), (1, 1)\}. \end{aligned} \tag{5.6}$$

The first of these corresponds to a firing function that can fire infinitely regardless of the length of the input sequence and consumes no tokens. The second consumes only the leading zeros from the input sequence and then stops firing. The third consumes one token from the input on every firing, regardless of its value. The fourth consumes two tokens on the input on every firing, again regardless of the values.

The following example shows the firing rules that would be needed for two-input [nondeterministic merge](#).

Example 5.2: Consider a two-input, one output actor operating on streams of type $\mathbb{B}$. A set of firing rules that does not satisfy condition 4 is:

$$R = \{(\lambda, (0)), (\lambda, (1)), ((0), \lambda), ((1), \lambda)\}. \quad (5.7)$$

Such firing rules would allow an actor to nondeterministically merge two input sequences.

Writing firing rules like those in the previous example is tedious and, for larger datatypes, impractical. We therefore rewrite using a **wildcard** $*$ as follows:

$$R = \{(\lambda, (*)), ((*), \lambda)\}. \quad (5.8)$$

The wildcard matches any token in the datatype of the actor inputs.

The firing rules in example 5.1 can all be realized in a [Kahn-MacQueen process](#) with [blocking reads](#). But dataflow actors are more general and need not be implementable using blocking reads.

Example 5.3: Consider the following firing rules:

$$R = \{((1), (0), \lambda), ((0), \lambda, (1)), (\lambda, (0), (1))\}. \quad (5.9)$$

These firing rules satisfy condition 4 in definition 5.1 and correspond to the [Gustave function](#), a function defining a process which is [stable](#) but not [sequential](#).

Given any dataflow actor conforming with definition 5.1, we now have a continuous function describing its input-output behavior over all firings. Any network

of such actors can be rearranged as shown in Figure 4.2 yielding a composite function of form $F: \mathcal{S}^N \rightarrow \mathcal{S}^N$, where N is the total number of streams in the network. This function is a parallel composition of continuous functions and is therefore itself continuous. Consequently, F has a unique least fixed point, as explained in the previous chapter. We take that fixed point to be the semantics of the network. We have therefore demonstrated the following proposition:

Proposition 5.2. *Every network of **dataflow actors** is deterministic in the sense that it has a unique **denotational semantics** given by the **least fixed point** of a **monotonic function** on streams.*

We do not yet quite have an **operational semantics** because we have only talked about the execution of a single actor using the procedure (5.5). We have a bit more work to do.

5.1.2 Feedback and Stateful actors

The definition of a **dataflow actor**, so far, looks extremely limiting. The **firing function** f and **firing rules** are the same for each of a sequence of firings. This means that a dataflow actor's behavior cannot change over time. Many useful actors seem to be excluded by this model.

Example 5.4: A **unit delay** actor with type $\mathbb{B}$ produces an initial output and then subsequently forwards each input token to the output. This seems to require firing rules that do not satisfy condition 4 in definition 5.1:

$$R = \{\lambda, (*)\}. \quad (5.10)$$

The first firing would use the first rule, consuming no input and producing the initial output. Subsequent firings would use the second rule, consuming and forwarding an input token.

The unit delay actor in the previous example is deterministic. But such firing rules would also correspond to an actor that could nondeterministically consume or not consume an input token upon firing. Hence, allowing the firing rules of (5.10) opens the door to nondeterministic networks.




Figure 5.1: Initial token on a feedback loop enables stateful behavior.

Fortunately, a small addition to the model vastly improves its expressiveness. All we need is to permit the connections between dataflow actors to have **initial tokens**.

Example 5.5: The **unit delay** can be implemented as a dataflow actor conforming to definition 5.1 using a feedback loop with an initial token, as shown in Figure 5.1. The type of the feedback path, in this case, can be boolean, $\mathbb{B}$. The firing rules are

$$R = \{((1), \lambda), ((0), (*))\}.$$

The firing function, which has form $f: \mathcal{S}^2 \rightarrow \mathcal{S}^2$, is given by

$$f(x, y) = \begin{cases} ((0), (i)) & \text{if } x_0 = 1 \\ ((0), (y_0)) & \text{otherwise,} \end{cases}$$

where x_0 and y_0 denote the first token in the upper and lower input streams, respectively and i denotes the initial output of the unit delay.

This example is easily generalized so that any dataflow actor can have state. The initial state is the value of the initial token on a feedback connection.

Example 5.6: For the unit delay actor, the feedback connection in Figure 5.1 represents a state variable that keeps track of whether a firing is the first firing. The initial token has value $i = 1$, and all subsequent values on this feedback path are 0.

Using such a feedback path with an initial token, the firing function can be completely different for each value of the feedback token as long as the conditions in definition 5.1 are satisfied. The feedback path can have an arbitrary data type

and therefore can represent the state of the actor, with the initial token giving the initial state.

5.1.3 Fixed Point Semantics with Initial Tokens

The initial tokens that we have added are not taken into account in the demonstration of determinism in proposition 5.2. In this section, we adapt the fixed point semantics to allow initial tokens.

Consider an arbitrary network of [dataflow actors](#) with zero or more initial tokens on any connection from the output of one actor to the input of another. Rearrange the actors as shown in Figure 5.2 to form a composite function of form $F: \mathcal{S}^N \rightarrow \mathcal{S}^N$, where N is the total number of streams in the network. This is nearly identical




Figure 5.2: Fixed point semantics adapted to allow initial tokens.

to what we did with Kahn networks in Figure 4.2, with the exception that there is now an initial token on the connection from B to A . As before, the F function is a parallel composition of continuous functions and hence is itself continuous. But now, assume that the N feedback paths are initialized with an N -tuple $\mathbf{d}$ of finite sequences, each containing zero or more initial tokens.

Let the semantics of the network be defined to be the N -tuple $\mathbf{s}$ of sequences coming out of F . Unlike Figure 4.2, the input to F is no longer the same as its output $\mathbf{s}$ because of the initial tokens. The network, therefore, defines $\mathbf{s}$ as an N -tuple satisfying

$$\mathbf{s} = F(\mathbf{d}.\mathbf{s}),$$

where $\mathbf{d}.\mathbf{s}$ is the [concatenation](#) of the sequences in $\mathbf{d}$ and $\mathbf{s}$. With this model, we are no longer looking for a [fixed point](#) of F .

Define a new function $F_{\mathbf{d}}: \mathcal{S}^N \rightarrow \mathcal{S}^N$ as follows:

$$F_{\mathbf{d}}(\mathbf{s}) = F(\mathbf{d}.\mathbf{s}). \quad (5.11)$$

With this definition, we seek a fixed point $\mathbf{s}$ of $F_{\mathbf{d}}$, where

$$\mathbf{s} = F_{\mathbf{d}}(\mathbf{s}).$$

Proposition 5.3. *If F is [continuous](#), then so is $F_{\mathbf{d}}$ for any N -tuple of finite sequences $\mathbf{d}$.*

Proof. Let C be a chain in $\mathcal{S}^N$. Note that

$$\bigvee \hat{F}_{\mathbf{d}}(C) = \bigvee \hat{F}(\mathbf{C}_{\mathbf{d}}),$$

where $\hat{F}_{\mathbf{d}}$ is the [lifted](#) version of $F_{\mathbf{d}}$, and

$$\mathbf{C}_{\mathbf{d}} = \{\mathbf{d}.c \mid c \in C\}.$$

Because C is a chain, so is $\mathbf{C}_{\mathbf{d}}$. Because F is continuous,

$$\bigvee \hat{F}(\mathbf{C}_{\mathbf{d}}) = F\left(\bigvee \mathbf{C}_{\mathbf{d}}\right).$$

But now,

$$\bigvee \mathbf{C}_{\mathbf{d}} = \mathbf{d}.\bigvee C,$$

so

$$\bigvee \hat{F}_{\mathbf{d}}(C) = F(\mathbf{d}.\bigvee C) = F_{\mathbf{d}}\left(\bigvee C\right),$$

which means that $F_{\mathbf{d}}$ is continuous. □




Figure 5.3: A two-input two-output identity process described as an aggregation of two one-input one-output identity processes.




Figure 5.4: If F is an identity process, the appropriate firing rules are (5.7).

Because F_d is continuous, it has a unique least fixed point, and hence we have established the following:

Proposition 5.4. *A network of **dataflow actors** with initial tokens is deterministic in the same sense as in proposition 5.2.*

If we assume there is some sequence of firing rules $(r_1, r_2, \dots)$ such that $d.s = r_1.r_2.\dots$, then exactly the same constructive procedure as in (5.5) will work.

5.1.4 Commutative Firings

While dataflow actors conforming to definition 5.1 yield continuous Kahn processes, condition 4 is more restrictive than what is really necessary. The firing rules in (5.7), for example, are not only the firing rules for the dangerous non-determinate merge, but also the firing rules for a perfectly harmless two-input two-output identity process.

Example 5.7: Let $I: \mathcal{S} \rightarrow \mathcal{S}$ be a simple identity function. We can define a **dataflow actor** with firing rule $R = \{(*)\}$ and firing function $f: \mathcal{S} \rightarrow \mathcal{S}$ given

by,

$$f(s) = (s_0),$$

where s_0 is the first token in s . Suppose now that we want to define a dataflow actor that is the parallel composition of two such identity actors, as shown in Figure 5.3. At first glance, it might seem that this sort of identity process could be implemented using the following firing rules:

$$R = \{((*), (*))\}.$$

With these firing rules, however, the two-input, two-output identity function would only produce outputs when it can consume one token from each input. The two examples in Figure 5.4 show why this will not work. In the first example, the first (top) input and output should be the empty sequence under the [least fixed-point semantics](#), so there will never be a token to trigger any firing rule. In the second example, the second (bottom) input and output have the same problem. The firing rules of (5.7), however, have no difficulty with these cases.

To overcome this limitation, we can replace condition 4 with a more general rule.

Definition 5.5. A *commutative dataflow actor* is a [dataflow actor](#) according to definition 5.1, but with rule 4 replaced by the following more elaborate condition:

4 for all $r, r' \in R$, if $r \neq r'$ and $\{r, r'\}$ has an upper bound in $(S^m, \sqsubseteq)$, then $f(r).f(r') = f(r').f(r)$ and $r \wedge r' = \lambda^m$.

This revised condition states that if any two firing rules are consistent, namely they have a common upper bound, and therefore can possibly be enabled at the same time, then it makes no difference in what order we use these firing rules; the values of the firing function at these consistent rules commute with respect to the concatenation operator. Furthermore, any two consistent firing rules have no common prefix other than the m -tuple of empty sequences.

Definition 5.5 makes our dataflow model more **compositional** (Talcott, 1996), in that an aggregation of dataflow actors is itself a dataflow actor. The two-input, two-output identity is perhaps the simplest example illustrating the compositionality problem. One-input one-output identity processes are trivially described as

dataflow actors that satisfy definition 5.1, but a two-input two-output identity process cannot be so described.

It is easy to see that when condition 4 in definition 5.5 is satisfied,

$$r \vee r' = r.r' = r'.r; \quad (5.12)$$

that is, the least common extension (least upper bound) of any two consistent firing rules is their concatenation, in either order.

To prove determinism with this generalization, we need to reconstruct the functional that we used to define the Kahn process. Let $P_R(s)$ denote the set $\{r \in R \mid r \sqsubseteq s\}$. This is a possibly empty finite set. The functional ϕ' is defined such that for any function $F : S^m \rightarrow S^n$ and any m -tuple s ,

$$\phi'(F)(s) = \begin{cases} f(r_1) \cdots f(r_p).F(s') & \text{if } P_R(s) \neq \emptyset \text{ and } \{r_1, \dots, r_p\} = P_R(s); \\ \lambda^n & \text{otherwise.} \end{cases}$$

Here, we assume, as before, that s' is defined so that $s = r_1 \dots r_p.s'$. Notice that because of (5.12), for any permutation π on $\{1, \dots, p\}$,

$$r_1 \dots r_p = r_{\pi(1)} \dots r_{\pi(p)},$$

and similarly, because of condition 4 in definition 5.5,

$$f(r_1) \cdots f(r_p) = f(r_{\pi(1)}) \cdots f(r_{\pi(p)}).$$

Therefore, it makes no difference in what order we invoke the enabled firing rules.

As before, we define the Kahn process F for the commutative dataflow actor (R, f) to be the least fixed point of the functional ϕ' . It is straightforward to extend the results on ϕ to conclude that both the functional ϕ' and its least fixed point F are continuous; the proofs are practically identical, albeit notationally more complicated.

Example 5.8: For the two-input, two-output identity function of Figure 5.3, we see that we can use the firing rules of (5.7), and a firing function $f : S^2 \rightarrow S^2$ such that for any firing rule $r \in R$, $f(r) = r$. This gives us a **commutative dataflow actor** for the two-input two-output identity.




Figure 5.5: A composition that does not form a dataflow actor but forms a commutative dataflow actor.

Example 5.9: More interestingly, we can use the same firing rules (5.7) to implement a commutative dataflow actor with two inputs and one output of type $\mathbb{B}$ and firing function $f : \mathcal{S}^2 \rightarrow \mathcal{S}$ such that for each firing rule $r \in R$,

$$f(r) = \begin{cases} (1) & \text{if } r = ((1), \lambda) \text{ or } r = (\lambda, (1)); \\ \lambda & \text{otherwise.} \end{cases}$$

This actor outputs a token with value 1 for each input token with value 1 regardless of which port receives the input. Such a process is a version of what is called a **parallel or** because it realizes a logical disjunction that is tolerant of unknown inputs. As soon as either input is known to be *true*, it outputs *true*. Contrast this with the logical or operator in many programming languages, often written $x \mid\mid y$, which evaluates x and, only if x is *false*, evaluates y . No output is produced until x can be evaluated.

The resulting Kahn process is interesting because it is neither **sequential** nor **stable**, and yet it is continuous.

Example 5.10: As a final example, consider the composition of Figure 5.5. The top process is an identity process, and the bottom one a source of the infinite sequence $(0, 0, \dots)$. How can we define the firing rules R and firing function f for the composition?

A first, naive attempt would be to let $R = \{(*)\}$. However, with the feedback arc in Figure 5.5, this results in no firing rule ever becoming enabled. Instead, we need $R = \{(*), \lambda\}$, which violates condition 4 in definition 5.1. However, if we define the firing function such that

$$\begin{aligned} f((0)) &= ((0), \lambda), \\ f((1)) &= ((1), \lambda), \text{ and} \\ f(\lambda) &= (\lambda, (0)), \end{aligned}$$

then condition 4 of definition 5.5 is satisfied and the composition forms a **commutative dataflow actor**.

The previous example is easily generalized to show that for any parallel composition of **dataflow actors**, like that shown in Figure 5.2(b), we can define firing rules and a firing function, based on those of the original dataflow actors, that satisfy definition 5.5.

Any network of dataflow actors can be restructured into a parallel composition with feedback, as illustrated in Figure 5.2. Any such network can be modeled as a single **commutative dataflow actor** with feedback, possibly containing initial tokens. We now have (finally!) a candidate **operational semantics**. The function F in Figure 5.2(d) can be realized by a simple firing function that invokes any of the firing functions of the component actors that is enabled. This firing function is (trivially) commutative. Repeated execution of this firing function provides a constructive procedure for finding the least fixed point.

In plain language, a network of dataflow actors can be executed using a very simple policy. Invoke the firing function of any **enabled actor** (one that has sufficient inputs). Repeat. Stop only if and when there are no more enabled actors. This procedure will always yield a prefix of the **denotational semantics**, but whether it converges in any sense to the denotational semantics depends on the choices made when more than one actor is enabled. This is the scheduling problem.

5.1.5 Scheduling

The constructive procedure given by (5.5) ensures that repeated firings always yield a prefix of the least fixed point semantics. Whether the procedure *converges*

to the least fixed point semantics depends on the same notions of convergence considered in Section 4.4 of the previous chapter and on scheduling decisions.

In the previous chapter, we assumed that a Kahn process had its own thread of execution that would forge ahead until blocked on a read (or write, if [blocking writes](#) are used). For dataflow actors, however, the execution is divided into firings. Rather than devoting a thread to each actor, we could have some number of **worker threads** invoke firing functions for various actors according to some scheduling policy. Finding reasonable scheduling policies turns out to be a sophisticated problem.

As we saw in the previous chapter, seemingly reasonable techniques like [data-driven execution](#) and [demand-driven execution](#) can result in undesirable executions. It may even be preferable to choose schedules that fail to converge to the least fixed point semantics or executions that fail to be fair.

The same [effective execution](#) methods considered in Section 4.3.2 of the previous chapter work for dataflow actors. Instead of [blocking writes](#), however, a scheduler just chooses to not invoke the firing function of an actor whose output buffers are full. However, for dataflow actors, it is often possible to provide much more sophisticated schedulers. The rest of this chapter is devoted to considering special cases of dataflow networks that yield to such sophisticated scheduling techniques.

5.2 Synchronous Dataflow

Synchronous dataflow (SDF), also called **static dataflow**,² was introduced by [Lee and Messerschmitt \(1987b\)](#) and is implemented, among other places, in Ptolemy II ([Lee et al., 2014b](#)). SDF is a simple special case of dataflow where each actor firing consumes and produces a fixed number of tokens. The questions of whether

²The term “synchronous dataflow” can cause confusion because it is not synchronous in the sense of [SR](#), considered in Chapter 6. There is no global clock in SDF models, and actors are fired asynchronously. For this reason, some authors prefer the term “static dataflow.” This does not avoid all confusion, however, because [Dennis \(1974\)](#) had previously coined the term “static dataflow” to refer to dataflow graphs where buffers could hold at most one token. Since there is no way to avoid a collision of terminology, we stick with the original “synchronous dataflow” terminology used in the literature. The term SDF arose from a signal processing concept, where two signals with sample rates that are related by a rational multiple are deemed to be synchronous.

buffers are bounded and whether executions are finite become decidable, and optimized static scheduling becomes possible, albeit sometimes intractable.

Example 5.11: A simple SDF actor consumes N input tokens with numerical values and outputs one token containing their average. A scheduler simply needs to ensure that the upstream actor that provides the tokens executes enough times that the required N input tokens are available when this averaging actor fires.

A further special case is the **homogeneous SDF** model, where a dataflow actor consumes exactly one token on each of its input ports and produces exactly one token on each output port. In this case, the scheduler simply has to ensure that each actor fires after the actors that supply it with data. An iteration of the model consists of exactly one firing of each actor in some data-precedence-constrained order.

As with more general dataflow models, SDF allows **initial tokens** on connections between actors.

Definition 5.6. *A complete iteration of a network of SDF actors, if it exists, is a finite non-zero number of firings that begins and ends with each connection having the same number of tokens as its initial number of tokens.*

When an iteration exists, the model can be executed infinitely with bounded buffers and the scheduling problem becomes a bounded **job-shop scheduling** problem, which is well studied in the literature. To find an iteration, we write and solve balance equations, as we explore next.




Figure 5.6: SDF actor A produces M tokens when it fires, and actor B consumes N tokens when it fires.

5.2.1 Balance Equations

Consider a single connection between two actors, A and B , as shown in Figure 5.6. The notation here means that when A fires, it produces M tokens on its output port, and when B fires, it consumes N tokens on its input port. M and N are non-negative integers. Suppose that A fires q_A times and B fires q_B times. All tokens that A produces are consumed by B if and only if the following **balance equation** is satisfied,

$$q_A M = q_B N. \quad (5.13)$$

Given values q_A and q_B satisfying (5.13), the system remains in balance; A produces exactly as many tokens as B consumes.

Suppose we wish to process an arbitrarily large number of tokens, a situation that is typical of streaming applications. A naive strategy is to fire actor A an arbitrarily large number q_A times, and then fire actor B q_B times, where q_A and q_B satisfy (5.13). This strategy is naive, however, because it requires storing an arbitrarily large number of unconsumed tokens in a buffer. A better strategy is to find the smallest non-negative q_A and q_B that satisfy (5.13) and are not both zero. Then we can construct a schedule that fires actor A q_A times and actor B q_B times, and we can repeat this schedule as many times as we like without requiring any more memory to store unconsumed tokens. That is, we can achieve an **unbounded execution** (an execution processes an arbitrarily large number of tokens) with bounded buffers (buffers with a bound on the number of unconsumed tokens). In each **complete iteration**, actor B consumes exactly as many tokens as actor A produces.

Example 5.12: Suppose that in Figure 5.6, $M = 2$ and $N = 3$. There are many possible solutions to the corresponding balance equation, one of which is $q_A = 3$ and $q_B = 2$. With these values, the following schedule can be repeated forever:

$A, A, A, B, B.$

An alternative schedule could also be used:

$A, A, B, A, B.$

The latter schedule has an advantage in that it requires less memory for storing intermediate tokens; B fires as soon as there are enough tokens, rather than waiting for A to complete its entire cycle.

Another solution to (5.13) is $q_A = 6$ and $q_B = 4$. This solution includes more firings in the schedule than are strictly needed to keep the system in balance.

The equation is also satisfied by $q_A = 0$ and $q_B = 0$, but if the number of firings of actors is zero, then no useful work is done. Clearly, this is not a solution we want. Negative solutions are also not meaningful.

A reasonable strategy finds the least integer solution to the balance equations that is not all zero and constructs a schedule that fires the actors in the model the requisite number of times, given by this solution.

In a more complicated SDF model, every connection between actors results in a balance equation. Hence, the model defines a system of equations, and finding the least integer solution is not entirely trivial.

Example 5.13: Figure 5.7 shows a network with three SDF actors. The connections result in the following system of balance equations:

$$\begin{aligned} q_A &= q_B \\ 2q_B &= q_C \\ 2q_A &= q_C. \end{aligned}$$

The least positive integer solution to these equations is $q_A = q_B = 1$, and $q_C = 2$, so the following schedule can be repeated forever to get an unbounded execution with bounded buffers,

$$A, B, C, C.$$

The balance equations do not always have a non-trivial solution, as illustrated in the following example.

Example 5.14: Figure 5.8 shows a network with three SDF actors where the only finite solution to the balance equations is the trivial one, $q_A = q_B = q_C = 0$. A consequence is that there is no unbounded execution with bounded buffers for this model. It cannot be kept in balance.

Definition 5.7. An SDF model that has a non-zero finite solution to the balance equations is said to be **consistent**. If the only solution is zero, then it is **inconsistent**.

An inconsistent model has no unbounded execution with bounded buffers.

We can compactly represent the balance equations in matrix form. Define a **connectivity matrix** Γ that has one row for each connection and one column for each actor. The entry $\Gamma_{i,j}$ at row i and column j has a non-negative number p if actor j produces p tokens on connection i and $-c$ if actor j consumes c tokens from connection i . Then define the vector q to represent the number of firings of each actor and the vector o to be a zero vector, with one zero for each connection. The balance equations can then be written in compact form,

$$\Gamma q = o. \quad (5.14)$$

Example 5.15: For the dataflow network in Figure 5.7, the balance equations can be written

$$\Gamma q = \begin{bmatrix} 1 & -1 & 0 \\ 0 & 2 & -1 \\ 2 & 0 & -1 \end{bmatrix} \begin{bmatrix} q_A \\ q_B \\ q_C \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \\ 0 \end{bmatrix} = o \quad (5.15)$$




Figure 5.7: A consistent SDF model.




Figure 5.8: An inconsistent SDF model.

A solution to the balance equations is

$$\begin{bmatrix} q_A \\ q_B \\ q_C \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix}$$

Example 5.16: For the dataflow network in Figure 5.8, the balance equations differ from the previous equations in only one number, $\Gamma_{2,2}$:

$$\Gamma \mathbf{q} = \begin{bmatrix} 1 & -1 & 0 \\ 0 & 1 & -1 \\ 2 & 0 & -1 \end{bmatrix} \begin{bmatrix} q_A \\ q_B \\ q_C \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \\ 0 \end{bmatrix} = \mathbf{o} \quad (5.16)$$

This equation has no nontrivial solution.

Note that in general the Γ matrix need not be square because the number of connections need not equal the number of actors. Moreover, any permutation of the rows of Γ yields an identical system of equations because the numbering of the connections is arbitrary. You can also permute the columns, but then you must also permute the elements of the $\mathbf{q}$ vector.

5.2.2 Solving the Balance Equations

Each balance equation has the form

$$q_A p = q_B c,$$

where p and c are non-negative integers. For a connected graph, a straightforward way to solve these equations is to pick one actor, say A , and set $q_A = 1$. Then

$$q_B = q_A p / c = p / c.$$

This will be a positive rational number. Then, for all other equations involving B , solve for the other similarly. Continue until either you have a rational solution for all q variables or you reach a contradiction, where the other side already has an incompatible solution. If you reach a contradiction, then there is no nontrivial solution and the dataflow graph is **inconsistent**. If you do not reach a contradiction, and the graph is connected, then you have a rational solution for all q variables. Now find the least common multiple (LCM) of the denominators of the rational solutions and multiply all q variables by this LCM. You will now have the smallest positive integer solution to the balance equations. If the graph is disconnected, then repeat the procedure for each connected component. Note that if any actor produces or consumes zero tokens on a connection, this is equivalent to not having the connection in the graph.

Example 5.17: Consider the dataflow graph shown in Figure 5.9 (we will consider the initial tokens later). The production and consumption parameters are shown next to each port. The balance equations can be written,

$$\begin{bmatrix} -1 & 0 & 3 \\ 0 & 0 & 0 \\ 2 & -3 & 0 \\ 0 & 1 & -2 \end{bmatrix} \begin{bmatrix} q_A \\ q_B \\ q_C \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \\ 0 \\ 0 \end{bmatrix}$$




Figure 5.9: A cyclic SDF model.

Notice that the second row corresponds to the self loop on actor A , a connection in which A produces one token and consumes one token, for a net production or consumption of zero tokens.

We can solve the balance equations using the procedure outlined above. Start with $q_A = 1$. The connection from A to B requires the following:

$$q_A = 1 \Rightarrow q_B = 2/3$$

The connection from B to C requires the following:

$$q_B = 2/3 \Rightarrow q_C = 2/6 = 1/3$$

The connection from C to A requires the following:

$$q_C = 1/3 \Rightarrow q_A = 1,$$

so the graph is consistent. The LCM of the denominators is 3, so multiplying through by 3 we get

$$q_A = 3, \quad q_B = 2, \quad q_C = 1.$$

5.2.3 Properties of the Connectivity Matrix

In linear algebra, a collection of vectors is **linearly independent** if no vector in the set is a linear combination of the other vectors in the set. Equivalently, vectors $v_1, v_2, \dots, v_N$ are **linearly dependent** if there exist non-zero numbers $q_1, q_2, \dots, q_N$ such that

$$q_1 v_1 + q_2 v_2 + \dots + q_N v_N = \mathbf{o}, \quad (5.17)$$

where $\mathbf{o}$ is the zero vector.

Example 5.18: The following set X of vectors is linearly independent, whereas Y is linearly dependent:

$$X = \left\{ \begin{bmatrix} 1 \\ 0 \\ 2 \end{bmatrix}, \begin{bmatrix} -1 \\ 1 \\ 0 \end{bmatrix}, \begin{bmatrix} 0 \\ -1 \\ -1 \end{bmatrix} \right\} \quad Y = \left\{ \begin{bmatrix} 1 \\ 0 \\ 1 \end{bmatrix}, \begin{bmatrix} -1 \\ 1 \\ 0 \end{bmatrix}, \begin{bmatrix} 0 \\ -1 \\ -1 \end{bmatrix} \right\}.$$

To see that Y is linearly dependent, note that the negative of the sum of the second two vectors is equal to the first. To see that X is linearly independent, we can use the above procedure to attempt to find a non-zero solution to (5.17) for these vectors. We will reach a contradiction.

The balance equations are identical to equation (5.17); the vectors $v_1, v_2, \dots, v_N$ are the columns of Γ . Hence, a non-trivial solution to the balance equations exists if and only if the columns of the **connectivity matrix** Γ are linearly dependent.

In linear algebra, the number of linearly independent columns of a matrix is called the **rank** of the matrix. This is also equal to the number of linearly independent rows, so the rank of the matrix cannot exceed the lesser of the number of columns and the number of rows. Hence, the rank of the connectivity matrix Γ is less than or equal to the smaller of the number of actors a and the number of connections c . If the graph is **connected**, then $c \geq a - 1$.

Proposition 5.5. *The rank of the **connectivity matrix** Γ of a connected dataflow graph is either $a - 1$ or a , where a is the number of actors.*

Proof. For $a = 2$, the Γ_a matrix will have two columns, so the rank cannot be bigger than 2. If there is only one connection, it will have one row, be **consistent**, and the rank will be $a - 1 = 1$. If there is more than one connection, then the rank will still be one if the graph is consistent and two if it is not.

For arbitrary $a > 2$, we can use induction. Assume the rank of Γ_a is a or $a - 1$. If we add one more actor and a single connection to it from any preexisting actor, Γ_{a+1} will have one more column, so the rank cannot get larger than $a + 1$. The new connectivity matrix will have the form,

$$\Gamma_{a+1} = \left[\begin{array}{c|c} \Gamma_a & \mathbf{o} \\ \hline \boldsymbol{\rho}^T & \end{array} \right],$$

where $\mathbf{o}$ is a zero column vector and $\boldsymbol{\rho}^T$ is a row vector representing the added connection. This row vector will have a non-zero entry in the last column, and hence this new row vector must be linearly independent of all rows above it, all of which have a zero in the last column. Hence, the addition of the new actor and single connection increases the rank of the matrix by at least one. Addition of more connections to the new actor may further increase the rank, but not to more than $a + 1$. Hence, the rank of Γ_{a+1} must be either a or $a + 1$.

Note that for this induction to work for any connected dataflow graph, we need the fact that every connected graph with $a + 1$ actors has a spanning-tree leaf—an actor whose deletion, together with its incident connections, leaves a connected graph with a actors. Applying the induction hypothesis to that graph yields the block form above.

Note further that self-loops connections add all-zero rows, which affect neither the rank nor the connectivity. □

Proposition 5.6. *A connected dataflow graph with a actors is **consistent** if and only if the rank of its **connectivity matrix** is equal to $a - 1$.*

Proof. If the rank of Γ is a , then all a columns are **linearly independent** and hence there cannot be any nontrivial solution to the balance equations. If it is not a , then by proposition 5.5 it must be $a - 1$. □

Example 5.19: The X vectors in Example 5.18 form the columns of a connectivity matrix for an inconsistent dataflow graph with three actors. The Y vectors form the columns of a connectivity matrix for a consistent dataflow graph. We leave it as an exercise to construct these graphs.

5.2.4 Dynamics of Execution

Given a dataflow graph with a actors and c connections, suppose we wish to construct a sequential schedule. I.e., we want to fire one actor at a time, possibly forever. Suppose v is a **one-hot** column vector of dimension a that is all zero except for a one in position n , where n is the index of an actor. Then a sequential schedule can be represented as a list of such vectors,

$$v_0, v_1, \dots$$

Define another column vector b of dimension c that contains nonnegative integers giving the number of tokens on the buffers for each connection. Let b_0 represent the

number of initial tokens on each connection. Then the execution of the sequential schedule has the following effect on the buffers:

$$b_{n+1} = b_n + \Gamma v_n \quad (5.18)$$

for any $n \geq 0$. The iteration in (5.18) can be thought of as a **symbolic execution** of the dataflow program. It doesn't actually invoke any firings, but rather emulates the effect of such firings on the buffers.

Definition 5.8. A *periodic admissible sequential schedule (PASS)* of length $K > 0$ is a sequence

$$v_0, v_1, \dots, v_{K-1}$$

such that $v_n \geq \mathbf{o}$ for each $n \in \{0, \dots, K-1\}$, and

$$b_K = b_0 + \Gamma(v_0 + \dots + v_{K-1}) = b_0.$$

A PASS is a finite schedule that brings buffers back to their initial sizes, and hence it can be repeated indefinitely with bounded memory. Let $q = v_0 + \dots + v_{K-1}$, and note that a PASS requires that $\Gamma q = \mathbf{o}$. Hence, a necessary condition for the existence of a PASS is that the balance equations have a non-zero solution. Hence, a PASS can only exist for a **consistent** SDF graph.

It is easy now to see how to systematically construct a PASS. First, solve the balance equations to find q . Then begin a series of steps where you select an **enabled actor** and append it to the schedule. An enabled actor i is one that has not yet been scheduled q_i times and has sufficient tokens on its input buffers to satisfy its **firing rules**. If there is no enabled actor, stop. Otherwise, for the selected actor, construct the **one-hot** vector and update the buffer sizes using (5.18). The algorithm terminates when every actor i has been appended to the schedule q_i times, in which case a PASS has been found, or no actor is enabled, in which case the program deadlocks.

Example 5.20: Figure 5.10 shows an SDF model with initial tokens on a feedback loop. The balance equations are

$$\begin{aligned} 3q_A &= 2q_B \\ 2q_B &= 3q_A. \end{aligned}$$

The least positive integer solution exists and is $q_A = 2$, and $q_B = 3$, so the model is consistent. With four initial tokens on the feedback connection, as shown, the following schedule is a [PASS](#) that can be repeated forever,

$$A, B, A, B, B.$$

This schedule starts with actor *A* because at the start of execution, only actor *A* can fire. Actor *B* does not have sufficient tokens. When *A* fires, it consumes three tokens from the four initial tokens, leaving one behind. It sends three tokens to *B*. At this point, only *B* can fire, consuming two of the three tokens sent by *A*, and producing two more tokens on its output. At this point, actor *A* can fire again because there are exactly three tokens on its input. It will consume all of these and produce three tokens. At this point, *B* has four tokens on its input, enabling two firings. After those two firings, both actors have been fired the requisite number of times, and the buffer on the feedback arc again has four tokens. The schedule has therefore returned the dataflow graph to its initial buffer sizes.

Were there any fewer than four initial tokens on the feedback path, the model would deadlock. If there were only three tokens, for example, then *A* could fire, followed by *B*, but neither would have enough input tokens to fire again.

At each step of this scheduling procedure, unlike the previous rather trivial example, there is usually more than one enabled actor. This offers opportunities for optimization. For example, you might pick actors in an attempt to minimize buffer memory or to give the most compact schedule (see the sidebar on page [186](#)). These optimization problems can be quite complex.




Figure 5.10: An SDF model with initial tokens on a feedback loop.

Example 5.21: Compact disks (CDs) and many digital audio encodings sample audio signals at 44.1 kilosamples per second. In 1987, a new standard was introduced for digital audio tapes (DATs) that deliberately used a different sample rate, 48 kilosamples per second, in a desperate and ultimately futile attempt to discourage unauthorized copying of digital audio data. Digitally converting a CD to DAT requires sophisticated signal processing that is easy today, but was expensive at the time. The conversion can be realized by a series of four digital filters, each of which converts its input to a new sample rate.

A realization of such a series is shown in Figure 5.11. The first filter, labeled *A*, doubles the sample rate. The second downsamples its input to $2/3$ the sample rate. The third upsamples by $8/7$. And the final one downsamples by $5/7$. If the input sample rate is 44,100 samples per second, the output sample rate is 48,000 samples per second.

Solving the balance equations for this SDF graph yields a smallest integer solution,

$$\mathbf{q}^T = [147, 147, 98, 28, 32, 160].$$

A schedule that minimizes buffer memory is this one (Bhattacharyya et al., 1996b):

```
ABABCABABCABABCDEAFFFFFFBABCABABCABABCDEAFFFFFFBABCABABCABABC
DEAFFFFFFBABCABABCABABCDEAFFFFFFBABCABABCABABCDEAFFFFFFBABCAB
CABABCABABCDEAFFFFFFBABCABABCABABCDEAFFFFFFBABCABABCABABCDEAFFFF
BCABABCABABCDEAFFFFFFBABCABABCABABCDEAFFFFFFBABCABABCABABCDEAFFFF
FEBCAFFFFFFBABCABABCABABCDEAFFFFFFBABCABABCABABCDEAFFFFFFBABCABABC
ABCDEAFFFFFFBABCABABCABABCDEAFFFFFFBABCABABCABABCDEAFFFFFFBABCAB
```




Figure 5.11: An SDF model that converts a CD sample rate signal to a DAT sample rate signal.

```
CABCDEAFFFFBABCABCABABCABCDEAFFFFFEBAFFFFFBCABCABABCDEAFFFFF
BCABABCABCABABCDEAFFFFFBCABABCABCDEAFFFFFBCABCABABCABCDEAF
FFFFBABCABCABABCDEAFFFFFBCABABCABCABABCDEAFFFFFBCABABCABCDEF
FFFFEFFFFF
```

This schedule is complex to find and quite chaotic.

For SDF graphs, as a consequence of the methods above, it is **decidable** whether the program can be executed with bounded memory, and it is decidable whether the program halts (deadlocks). Thus, SDF, as a model of computation, is less expressive than dataflow in general and less expressive than [Kahn process networks](#). It is not [Turing complete](#).

5.2.5 Parallel Scheduling

The procedure outlined in Section 5.2.4 can be adapted to formulate a parallel scheduling problem, where dataflow actors are mapped onto parallel computing resources. One strategy is to use the [symbolic execution](#) of (5.2.4) to construct an **acyclic precedence graph (APG)** rather than directly constructing a schedule. You can then use standard [job-shop scheduling](#) techniques on that graph.

To construct an APG, when performing the [symbolic execution](#) of (5.18), for each token recorded in the $\mathbf{b}$ vector, keep track of whether it is an initial token and, if not, which firing produced the token. When encountering an [enabled actor](#), instead of adding it to a list schedule, add it to a graph and add directed arc from each firing that produced at least one of the tokens it consumes. When the procedure terminates, you will have a graph that represents a finite number of firing and the dependencies between them.

Example 5.22: Consider the dataflow graph shown in Figure 5.9. As before, we solve the balance equations to obtain

$$q_A = 3, \quad q_B = 2, \quad q_C = 1.$$


Initially, the only enabled actor is C and the tokens it consumes are all initial tokens. So we add C to the APG with no incoming arcs. Upon firing C, A

Sidebar: SDF Schedulers

An advantage of using SDF is that there may be many possible schedules for a given model, including some that execute actors in parallel. In this case, actors in the dataflow graph can be mapped onto different processors in a multicore or distributed architecture for improved performance. [Lee and Messerschmitt \(1987a\)](#) adapt classical job-shop scheduling algorithms ([Coffman, 1976](#)), particularly those introduced by [Hu \(1961\)](#), to SDF by converting the SDF graph into an acyclic precedence graph (APG). [Lee and Ha \(1989\)](#) classify scheduling strategies into **fully dynamic scheduling** (all scheduling decisions are made at run time), **static assignment scheduling** (all decisions except the assignment to processors are made at run time), **self-timed scheduling** (only the timing of an actor firing is determined at run time), and **fully-static scheduling** (every aspect of the schedule is determined before run time). [Sih and Lee \(1993a\)](#) extend the job-shop scheduling techniques to account for interprocessor communication costs (see also [Sih and Lee \(1993b\)](#)). [Pino et al. \(1994\)](#) show how to construct schedules for heterogeneous multiprocessors. [Falk et al. \(2008\)](#) give a parallel scheduling strategy based on clustering and demonstrate significant performance gains for multimedia applications.

In addition to parallel scheduling, other scheduling optimizations are useful (see [Bhattacharyya et al. \(1996b\)](#) for a collection of these). [Ha and Lee \(1991\)](#) relax the constraints of SDF to allow data-dependent iterative firing of actors (a technique called **quasi-static scheduling**). [Bhattacharyya and Lee \(1993\)](#) develop optimized schedules for iterated invocations of actors (see also [Bhattacharyya et al. \(1996a\)](#)). [Bhattacharyya et al. \(1993\)](#) optimize schedules to minimize memory usage and later apply these optimizations to code generation for embedded processors ([Bhattacharyya et al., 1995](#)). [Murthy and Bhattacharyya \(2006\)](#) collect algorithms that minimize the use of memory through scheduling and buffer sharing. [Geilen et al. \(2005\)](#) show that model checking techniques can be used to optimize memory. [Stuijk et al. \(2008\)](#) explore the tradeoff between throughput and buffering (see also [Moreira et al. \(2010\)](#)). [Sriram and Bhattacharyya \(2009\)](#) develop scheduling optimizations that minimize the number of synchronization operations in parallel SDF.

becomes enabled and will consume tokens that were produced by the firing of C . Hence, we add A to the graph with an arc from C to A . Continuing in this manner, we obtain the following APG:



The previous example has an interesting subtlety. Two of the firings of A depend on the previous firing of A because of the self loop on A . Such a self loop represents [state](#). B has no such self loop, and hence the second firing of B does not depend on the first firing. These two firings can, in principle, execute in parallel. This would not be safe if actor B were to have any internal state.

As a consequence, even though self loops representing state have no effect on the balance equations, it is important to include them in a model to represent dependencies introduced by state variables.

Once an APG has been constructed, building a parallel schedule becomes a [job-shop scheduling](#) problem. This is a well studied problem with a vast literature giving many algorithms for approximating optimal solutions depending on what you wish to optimize and what you know about firings and communication overhead (see sidebar on page [186](#)).

A common optimization criterion that is used is to minimize the **makespan**, defined to be the time at which all executions in the APG have completed. Constructing such an optimal schedule requires knowing the execution times of the firings. With bounded processing resources, this is a well-known NP-complete problem, meaning that the best known optimal algorithms have complexity that is an exponential function of the number of nodes in the graph. For small and medium sized graphs, brute-force solutions that find the optimal schedule may be tractable, but it is relatively rare to know precise execution times, so the value of such schedules may be questionable.

Example 5.23: For the APG constructed in the previous example, assume that the execution times are all one time unit. We can easily construct the minimum makespan schedule. A two-processor schedule with the first processor executing C, A_1, A_2, A_3, B_2 and the second executing B_1 (starting at time 2) achieves the minimum makespan of five.

This example is simple enough that we can consider all possible schedules. It also becomes clear that there is no value in having more than two processors and that the minimum achievable makespan is five (it is the total execution time of the longest path in the graph). But most problems are not so forgiving, so heuristics may be necessary.

One classic heuristic for job-shop scheduling is the **Hu level scheduling** strategy (Hu, 1961). In this strategy, each node in the APG is assigned a **level**, which is the length of the longest upstream chain. The schedule then consists of executing all nodes with level 0 first in any order. When these are complete, then execute all nodes with level 1. Etc. This strategy ensures that, regardless of execution times, all precedences will be satisfied. The strategy rarely yields the minimum makespan, although in the previous example, with everything having unit execution time, it will.

Example 5.24: In the APG of the previous example, C_1 has level 0, A_1 has level 1, A_2 and B_1 have level 2, A_3 has level 3, and B_2 has level four. The same schedule as before results from the Hu level scheduling algorithm.

Many more sophisticated algorithms have been developed over time. For example, Sih and Lee (1993b) give a heuristic that takes into account the time that communication between processors takes and contention for limited communication resources.


One limitation of the scheduling methods described so far is that they consider only a single **complete iteration** of an SDF graph. Overlapping of these iterations could yield overall faster execution. One way to do this is to consider bigger complete iterations. In other words, instead of choosing the smallest positive integer solution q to the balance equations, choose a larger solution Kq , where $K > 1$ is an integer **vectorization factor**.

Example 5.25: Define the **throughput** to be the number of minimal **complete iterations** per unit time. With the minimum makespan schedule for Figure 5.9 that we constructed in the previous examples, the throughput is one complete iteration for every five time units. Can we do better?

An alternative solution to the balance equations is

$$q_A = 6, \quad q_B = 4, \quad q_C = 2.$$

An APG for this vectorization factor of two is given below:



The minimum makespan for this APG is 8, yielding a throughput of $1/4$, which is better than the previous schedule, which yielded a throughput of $1/5$.

In principle, for the example in Figure 5.9, the limit on throughput is imposed by the self loop on A . In the limit, a throughput of $1/3$ should be achievable. However, no finite **vectorization factor** will yield this throughput.

5.3 Boolean Dataflow

The limited expressiveness of SDF can be quite limiting at times. Harmless programs that are easy to schedule often do not conform with the SDF model.

Example 5.26: Consider a decision-making program, shown in Figure 5.12, that uses a stream of boolean conditions c generated by actor C to choose whether to forward the output of actors A or B to D . This is not an SDF program, though it is not hard to see how to schedule it.

Sidebar: StreamIt

Thies et al. (2002) give a textual programming language, **StreamIt**, based on SDF and intended for use with streaming data applications such as multimedia. Software components (called **filters** rather than actors) produce and consume fixed amounts of data. The language provides compact structured constructs for common patterns of actor composition, such as chains of filters, parallel chains of filters, or feedback loops.

A key innovation in StreamIt is the notion of a **teleport message** (Thies et al., 2005). Teleport messages improve the expressiveness of SDF by allowing one actor to sporadically send a message to another; that is, rather than sending a message on every firing, only some firings send messages. The teleport message mechanism nonetheless ensures determinism by ensuring that the message is received by the receiving actor in exactly the same firing that it would have if the sending actor had sent messages on every firing. But it avoids the overhead of sending messages on every firing. This approach models a communication channel where tokens are sometimes, but not always, produced and consumed. But it preserves the determinism of SDF models, where the results of execution are the same for any valid schedule.

Sidebar: Multidimensional Synchronous Dataflow

Murthy and Lee (2002) describe **multidimensional SDF (MDSDF)**. Whereas a channel in SDF carries a sequence of tokens, a channel in MDSDF carries a multidimensional array of tokens. That is, the history of tokens can grow along multiple dimensions. This model is effective for expressing certain kinds of signal processing applications, particularly image processing, video processing, radar and sonar.




Figure 5.12: A decision-making program that is easily handled with BDF.

This example uses a **Select** actor similar to the **Select** process used in the previous chapter. Its firing rules are

$$R = \{((*), \lambda, (true)), (\lambda, (*), (false))\}.$$

If the third input (the bottom one in the figure) has a *true*-valued token, then it consumes one token from the first input (the top one, labeled 'T') and nothing from the second input. If the third input has a *false*, then it consumes from the 'F' input.

I have previously shown (Lee, 1991) that the balance equations can be generalized to handle actors like Select. Let c represent the boolean token consumed by the Select. Let p_c represent the probability that this token is *true*. Then the system will be balanced on average if the following balance equations have a non-trivial solution:

$$\Gamma \mathbf{q} = \begin{bmatrix} 1 & 0 & 0 & -p_c & 0 \\ 0 & 1 & 0 & p_c - 1 & 0 \\ 0 & 0 & 1 & 1 & 0 \\ 0 & 0 & 0 & 1 & -1 \end{bmatrix} \begin{bmatrix} q_A \\ q_B \\ q_C \\ q_S \\ q_D \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \\ 0 \\ 0 \end{bmatrix} = \mathbf{o} \quad (5.19)$$

Here, the consumption of the Select on its T and F ports is represented symbolically by $-p_c$ and $-(1 - p_c)$. It is easy to see that these balance equations are satisfied by

$$\mathbf{q}^T = (p_c, 1 - p_c, 1, 1, 1). \quad (5.20)$$

This is a positive solution, but not an integer solution, in general.

Dataflow programs where production and consumption numbers are given in terms of probabilities for boolean variables are called **boolean dataflow (BDF)** programs. Buck (1994) has further generalized this idea to allow for the production and consumption numbers to be given in terms of probability mass functions for integer-valued variables. He called the resulting graphs **integer dataflow (IDF)** programs.

How to construct a schedule for BDF programs is addressed in the PhD theses of Ha (1992) and Buck (1993). For this example, if we use the solution (5.20) to the balance equations, then constraining p_c to be either zero or one gives an integer repetition vector q . From (5.18), it is easy to see that a schedule constructed using either zero or one for p_c will yield a **complete iteration**. Thus, we only need to construct two schedules, one for the outcome where c is *true* ($p_c = 1$) and one for the outcome where c is *false* ($p_c = 0$):

$$S_0 = (C, B, S, D), \quad S_1 = (C, A, S, D).$$

For situations where exactly one of each of N controlling booleans is produced in each iteration, and these booleans are not stored on buffers for use in future iterations, we would need, in the worst case, 2^N schedules.

In many cases, the schedules can be compactly represented as a **quasi-static schedule**, where some firings are annotated with a boolean expression indicating whether they will occur in any particular iteration. For this example, the following will work:

$$S = (C, c?A, !c?B, S, D),$$

where $c?A$ means that A fires if c is *true* and $!c?B$ means that B fires if c is *false*. However, not all programs yield to such a simple treatment.

Example 5.27: Consider the more problematic program shown in Figure 5.13. The balance equations can be written:

$$\Gamma q = \begin{bmatrix} 1 & 0 & 0 & -p_c \\ 0 & 1 & 0 & p_c - 1 \\ 0 & 0 & 1 & 1 \\ -1 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} q_A \\ q_B \\ q_C \\ q_S \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \\ 0 \\ 0 \end{bmatrix} = \mathbf{o} \quad (5.21)$$




Figure 5.13: A problematic decision-making program.

This Γ matrix has full rank (rank four) for all values of p_c except $p_c = 1$. Hence, the system is in balance only if C only produces *true*-valued tokens. But in that case, the program immediately deadlocks. In fact, it is assured of deadlocking unless the very first output of C is *false*.

Much more interesting programs sometimes yield to such a quasi-static scheduling strategy.

Example 5.28: Consider the program of Example 4.21 and Figure 4.9, which produces the [Hamming numbers](#), an ordered sequence of integers of the form $2^n 3^m 5^k$, where n , m and k are non-negative integers. These numbers can be produced by a BDF program.

The only nontrivial actor in Figure 4.9 is the [OrderedMerge](#), which merges two streams of increasing numbers into one stream of increasing numbers. This can be implemented as shown in Figure 5.14. Actors A and B provide increasing sequences of numbers and C consumes the merged increasing sequence. The **MinMax** and **Xor** actors here are ordinary SDF actors. MinMax consumes one integer-valued token from each input port, outputs the larger of the two on the **max** output port and the smaller on the **min** output port. It does not matter what it does if they are equal. On the **swapped** output port, it produces a *true* if the lower input went to the **max** port and a *false* otherwise. The **Xor** actor takes one boolean-valued token from each input and outputs the exclusive or of the two. The initial token at the output of




Figure 5.14: A BDF implementation of an ordered merge towards a program generating the Hamming numbers.

Xor has value *false*. The initial token on the upper feedback path has value *min*, the smallest integer value allowed, or $-\infty$ if that is allowed by the data type.

This program can be kept in balance by a quasi static schedule. Let x represent the value of the boolean token at the output of Xor at the beginning of each iteration. Then the following schedule provides a [complete iteration](#):

$$S = (x?A, !x?B, \text{Select}, \text{MinMax}, C, \text{Xor}).$$

However, the full program in Figure 4.9 that generates the Hamming numbers does not have the same structure as Figure 5.14. The output of the ordered merge is fed back to *B*, thereby giving it structure more like Figure 5.13. The program in Figure 4.9 has no bounded quasi-static schedule.

Analyzing programs like that in Figure 5.14 can get arbitrarily complicated. Here, the boolean token that is used in the schedule starts with an initial token. In general, the notion of a [complete iteration](#) may need to be augmented to not only restore buffers to their initial sizes, but also to restore boolean tokens stored across iterations to their initial values, though this is not needed in this example. Logic operations on the boolean tokens, such as the Xor in the figure, can be taken into account using symbolic algebra. When Switch and Select are used to conditionally route streams of boolean tokens, they function as logic operations themselves that




Figure 5.15: A BDF program that has no finite complete cycle.

can be analyzed. See [Buck and Lee \(1992, 1993\)](#) for more illustrations of such analysis.

Unlike SDF, for BDF, having a non-trivial solution to the balance equations does not provide assurance that, given enough initial tokens, there exists a finite [complete iteration](#). Consider the example in Figure 5.15. Except for the [Switch](#) and [Select](#), all are SDF actors that, unless otherwise indicated, produce or consume a single token on each port. But actors *B* and *C* each consume and produce two tokens per firing. The symbolic balance equations have a solution where all actors fire twice while *B* fires p times and *C* $(1 - p)$ times, where p is the probability of a *true* token from *D*. Suppose that the boolean sequence produced by *D* is *(true, false, false, false, ...)*. There is no finite number of firings that form a [complete iteration](#), and hence there is no quasi-static schedule.

Whether a BDF program can be kept in balance without deadlocking is, in general, [undecidable](#) ([Buck, 1993](#)). This is easy to see by recognizing that the stack program given in Figure 4.7 of the previous chapter is, in fact, a BDF program. The fact that it implements a stack means that its buffers cannot be bounded. In fact, there are many programs that will not yield to quasi-static scheduling and will not remain balanced.

An alternative to constructing a quasi-static schedule is to use a run-time scheduler, which makes all scheduling decisions at runtime based on buffer status and firing rules. An example implementation is given in Ptolemy II ([Lee et al., 2014b](#)), where it is called **dynamic dataflow (DDF)**. As with [Kahn process networks](#), we can apply [Parks' algorithm](#) to ensure that such a scheduler delivers a bounded

memory schedule when that is possible. Such a strategy would work well for the examples in Figures 5.14 and 5.15. All the same convergence questions apply to dynamic dataflow in a manner very similar to process networks, as discussed in Section 4.4.

In summary, the symbolic balance equations can be used to analyze, optimize, and schedule some programs, but not all. We next consider another way to improve the expressiveness of SDF without admitting undecidable problems.

5.4 Modal Dataflow

A disadvantage of SDF is that every actor must produce and consume a fixed amount of data; the production and consumption rates cannot depend on the data. BDF relaxes this constraint, but its *quasi-static schedules* can be difficult to find. DDF further relaxes the constraint, abandons static scheduling and static analysis for deadlock or boundedness. A number of variants of dataflow, however, are more expressive than SDF but still amenable to some forms of static analysis.


In this section, we outline several models, all of which conform to a simple pattern. Actors in these models are **modal**, meaning that they have more than one mode of operation. In each mode, the behavior of an actor is SDF; it produces and consumes a fixed number of tokens. But these numbers can differ across modes. The differences between the modal dataflow models we present here concern primarily how the switching between modes is governed.

5.4.1 Cyclo-Static Dataflow

One of the first interesting modal dataflow models is **Cyclo-static dataflow** (CSDF), introduced by Bilsen et al. (1996). Whereas an SDF actor is given by single pair (R, f) of *firing rules* and *firing function*, a CSDF actor provides a sequence $((R_1, f_1) \cdots (R_P, f_P))$ for some period $P \geq 1$. The first firing is governed by (R_1, f_1) , the second by (R_2, f_2) , until we reach (R_P, f_P) . Then we begin again with (R_1, f_1) . Each pair (R_i, f_i) is required to correspond with an SDF actor (the number of tokens consumed and produced must be fixed).

Sidebar: Petri Nets

Petri nets, named after Carl Adam Petri, are a popular modeling formalism related to dataflow ([Murata, 1989](#)). They have two types of elements, **places** and **transitions**, depicted as white circles and rectangles, respectively. A place can contain any number of tokens, depicted as black circles. A transition is **enabled** if all places connected to it as inputs contain at least one token.



An enabled transition can **fire**, consuming one token from each input place and depositing one token on each output place. The state of a network, called its **marking**, is the number of tokens on each place in the network. The figure above shows the marking of a simple network before and after a firing. If a place provides inputs to more than one transition, then a token on that place may trigger a firing of either destination transition (chosen nondeterministically).

Petri net transitions are like dataflow actors; they fire when sufficient inputs are available. In basic Petri nets, tokens have no value. A firing just moves tokens from one place to another. Also, places do not preserve token ordering, unlike dataflow buffers. Like [homogeneous SDF](#), transitions are enabled by a single token on each input place. Unlike SDF, a place may feed more than one transition, yielding nondeterminism.

There are many variants of Petri nets, at least one of which is equivalent to SDF. In particular, **colored tokens** have values, where the color represents the value. Transitions can manipulate colors (analogous to the firing function of a dataflow actor). Arcs connecting places to transitions can be weighted to indicate that more than one token is required to fire a transition, or that a transition produces more than one token, like SDF production and consumption rates. And finally, the Petri net can be constrained so that for each place, there is exactly one source transition and exactly one destination transition. With order-preserving places, such Petri nets are SDF graphs.




Figure 5.16: Two simple examples of CSDF actors.




Figure 5.17: Two simple examples of CSDF actors.

Example 5.29: A CSDF version of the [Switch](#) actor is a **Commutator**, illustrated in Figure 5.16, which splits an input stream into two output streams. It consumes one input token in each firing and routes the token to the top output in odd firings and to the bottom output in even firings. The **Multiplexor**, also illustrated in the figure, merges two input streams by alternating their tokens. In the figure, the number of tokens produced and consumed for each port is shown as a sequence of length $P = 2$. The corresponding firing functions just read and write the appropriate ports.

It is straightforward to generalize the balance equations to handle CSDF. Simply replace the production and consumption numbers with their average. When finding an integer solution, find the least integer solution so that the number of firings of each actor is an integer multiple nP of its period P .

Example 5.30: Consider the program in Figure 5.17, which multiplexes outputs from A , B , and C in a pattern that looks like $(a_1, c_1, b_1, c_2, a_2, c_3, b_2, c_4, \dots)$. We can write down the balance equations in matrix form as follows:

$$\Gamma \mathbf{q} = \begin{bmatrix} 1 & 0 & 0 & 0 & -1/2 & 0 \\ 0 & 1 & 0 & 0 & -1/2 & 0 \\ 0 & 0 & 1 & 0 & 0 & -1/2 \\ 0 & 0 & 0 & 0 & 1 & -1/2 \\ 0 & 0 & 0 & -1 & 0 & 1 \end{bmatrix} \begin{bmatrix} q_A \\ q_B \\ q_C \\ q_D \\ q_{M_1} \\ q_{M_2} \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \\ 0 \\ 0 \\ 0 \end{bmatrix} = \mathbf{o}$$

The smallest integer solution is

$$\begin{bmatrix} q_A \\ q_B \\ q_C \\ q_D \\ q_{M_1} \\ q_{M_2} \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 2 \\ 4 \\ 2 \\ 4 \end{bmatrix} \quad (5.22)$$

In this solution, q_{M_1} and q_{M_2} are multiples of the periods ($P = 2$) of their CSDF specifications, so this solution can be used directly to construct a schedule that will form a [complete iteration](#).

Once we have a suitable solution to the balance equations, all of the same scheduling techniques can be used as for SDF. For example, we could perform [symbolic execution](#) and construct an [acyclic precedence graph](#) and then construct a schedule using [job-shop scheduling](#) techniques.

5.4.2 Heterochronous Dataflow

A related technique called **heterochronous dataflow (HDF)** is described by [Girault et al. \(1999\)](#). In their model, each dataflow actor has associated with it a **finite state machine (FSM)**, which is a finite set of states and a transition function that chooses a next state as a function of the current state and the inputs. Each state s is associated with an SDF actor (R_s, f_s) . And each state machine has an initial state.

With each actor in its initial state, we have an SDF graph which, if it is **consistent**, will have a uniquely defined minimal **complete iteration** for which we can construct a static schedule. To execute this graph, we first execute this static schedule for the initial states. We then evaluate the transition functions, which possibly causes each of the actors to change state. We now have a new SDF graph which, again if it is consistent, will have a uniquely defined minimal complete iteration. We execute a schedule for this second complete iteration. Upon completion of this schedule, we again evaluate the transition functions for each actor that executed at least once, ending up with a new SDF graph. We can repeat this process indefinitely with bounded memory as long as each of the resulting SDF graphs is consistent. In the worst case, the total number of SDF graphs that must be analyzed and scheduled is equal to the product of the number of states of all the SDF graphs. This number is finite, so, in principle, this is doable.

Example 5.31: The program in Figure 5.17 can be realized as an HDF graph. The **Multiplexor** actors have two states. The initial state consumes one token from the top input port and zero from the bottom port and produces one output token. The other state consumes one token from the *bottom* input port and zero from the top port and produces one output token. The transition function is such that whenever it is evaluated, it returns the state that is not the current state.

With each of the two multiplexors in their initial states, we have an SDF graph where the least integer solution is

$$\begin{bmatrix} q_A \\ q_B \\ q_C \\ q_D \\ q_{M_1} \\ q_{M_2} \end{bmatrix} = \begin{bmatrix} 1 \\ 0 \\ 0 \\ 1 \\ 1 \\ 1 \end{bmatrix}.$$

Notice the zeros in this solution for actors *B* and *C*. The corresponding schedule is (A, M_1, M_2, D) .

In the new state, we again have an SDF graph. This time, the least solution to the balance equations is

$$\begin{bmatrix} q_A \\ q_B \\ q_C \\ q_D \\ q_{M_1} \\ q_{M_2} \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \\ 1 \\ 1 \\ 0 \\ 1 \end{bmatrix}. \quad (5.23)$$

The corresponding schedule is (C, M_2, D) .

We now evaluate the transition function for M_2 but not M_1 because M_1 did not execute in this schedule. Solving the balance equations for the new SDF graph, we get

$$\begin{bmatrix} q_A \\ q_B \\ q_C \\ q_D \\ q_{M_1} \\ q_{M_2} \end{bmatrix} = \begin{bmatrix} 0 \\ 1 \\ 0 \\ 1 \\ 1 \\ 1 \end{bmatrix}$$

and a schedule (B, M_1, M_2, D) .

The next iteration gives us (5.23) again. Evaluating the transition functions now takes us back to the initial state, so we can start all over.

Notice that if we sum the solutions to the balance equations over the four iterations before returning to the initial state, we get the same result as the CSDF solution (5.22),

$$\begin{bmatrix} q_A \\ q_B \\ q_C \\ q_D \\ q_{M_1} \\ q_{M_2} \end{bmatrix} = \begin{bmatrix} 1 \\ 0 \\ 0 \\ 1 \\ 1 \\ 1 \end{bmatrix} + \begin{bmatrix} 0 \\ 0 \\ 1 \\ 1 \\ 0 \\ 1 \end{bmatrix} + \begin{bmatrix} 0 \\ 1 \\ 0 \\ 1 \\ 1 \\ 1 \end{bmatrix} + \begin{bmatrix} 0 \\ 0 \\ 1 \\ 1 \\ 0 \\ 1 \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 2 \\ 4 \\ 2 \\ 4 \end{bmatrix}.$$

Notice that HDF requires special care in handling consumption or production rates of zero. Any subgraph that is connected on the far side of such a production or consumption should be given zero repetitions. It also requires care when choosing

when to evaluate the transition functions. They should be evaluated only after a complete iteration has completed, but only for actors with one or more firings in that iteration. Because of these subtleties, HDF can be considerably more difficult to use than CSDF.

5.4.3 SDF Scenarios

Geilen and Stuijk (2010) introduce **synchronous dataflow scenarios**, which are like a flattened version of HDF. A program has some set of **scenarios**, each of which is an SDF graph. An SDF graph is executed through a **complete iteration** before the next scenario takes over. The structure of the SDF graph can be completely different in each scenario. A practical realization must provide a mechanism for tokens to be carried from one scenario to the next.

5.4.4 Parameterized SDF

A similar variant of SDF is called **parameterized SDF (PSDF)**, introduced by Bhattacharya and Bhattacharyya (2000). PSDF allows the production and consumption rates of ports to be given by a parameter rather than being a constant, but the overall structure is held constant. The value of the parameter is permitted to change, but only between **complete iterations**. When the value of such a parameter changes, a new schedule must be used for the next complete iteration.

5.5 Conclusion

This chapter has examined dataflow models of computation as a foundation for deterministic concurrent systems. Starting from a precise definition of **dataflow actors** in terms of firing rules and firing functions, we show how individual actors induce **continuous** Kahn processes and how networks of such actors admit a unique **denotational semantics** given by the least fixed point of a continuous functional. This establishes determinism at the semantic level, independent of execution order or scheduling decisions.

A key theme is the separation between semantics and execution. Although the denotational semantics of a dataflow network is well defined and deterministic,

operational behavior depends critically on scheduling policies. By introducing [commutative firings](#), we broaden the class of actors that can be composed while preserving determinism, enabling networks to be executed by repeatedly firing any enabled actor. This yields a simple and flexible operational model, but one whose convergence properties depend on scheduling choices and resource constraints.

The chapter then focuses on [synchronous dataflow](#) (SDF) as an important and practically useful specialization. By restricting actors to fixed token production and consumption rates, SDF admits static analysis techniques that are not available in general dataflow networks. [Balance equations](#) and the associated connectivity matrix characterize whether a model is consistent and can execute indefinitely with bounded memory. For [consistent](#) graphs, [periodic admissible sequential schedules](#) can be constructed, and both deadlock freedom and boundedness become decidable properties. These results make SDF particularly attractive for signal processing and embedded systems, where predictability, resource bounds, and compile-time scheduling are essential.

SDF also admits parallel scheduling through well-studied job-shop scheduling techniques. The chapter shows how to construct an [acyclic precedence graph](#) for an SDF graph. This APG can then be turned over to a variety of parallel scheduling methods.

More expressive variants, including [boolean](#), [cyclo-static](#), and [modal dataflow](#), relax the rigidity of SDF while retaining many of its analyzable properties. These models occupy a useful middle ground between the full generality of Kahn process networks and the strict regularity of SDF, enabling systems whose structure or rates vary over time but remain amenable to disciplined execution strategies.

Overall, dataflow models illustrate how strong semantic guarantees—specifically determinism—can coexist with concurrency and parallelism. By carefully constraining the interaction between computation and communication, dataflow provides a framework in which nondeterminism is avoided not by limiting parallelism, but by structuring it. The next chapter will explore a contrasting approach, synchronous reactive models, which impose global coordination rather than local firing rules to achieve deterministic behavior.

Part III

Timed Models

Synchronous Reactive Models

6.1	The Global Clock Model	208
	<i>Sidebar: About Synchrony</i>	209
	<i>Sidebar: Synchronous-Reactive Languages</i>	210
6.1.1	An SR Component Model	212
6.1.2	Components with State	214
6.1.3	Fixed-Point Semantics	215
6.1.4	Constructive Semantics	221
	<i>Sidebar: Information Order</i>	223
6.1.5	Concurrent State Machines	227
	<i>Sidebar: Intuitionistic Logic</i>	228
	<i>Sidebar: Synchronous State Machines</i>	229
	<i>Sidebar: Nondeterministic State Machines</i>	231
6.1.6	Scheduling Strategies	235
6.2	Alternative Clock Models	239
6.2.1	Clock Domains	239
6.2.2	Sparse Synchronous Models	242
6.2.3	Clock Calculus	243
6.3	Conclusion	252

The **synchronous-reactive (SR)** models of computation offer deterministic concurrency with potential massive parallelism. We use the plural, “models,” because there are at least two distinct models that have been extensively studied. Note

that the term “synchronous” has a number of inconsistent usages (see sidebar on page 209), and the SR models we consider cover only a subset of the possible interpretations of this term.

The first and simpler model has a conceptual **global clock** defining a sequence of **ticks**. Each component **fires** at some subset of these ticks. This model is concretely realized in the Esterel language (Berry and Gonthier, 1992) and various languages and frameworks inspired by Esterel. In the second model, a clock is a property of communication between two components. This model is realized in Lustre (Halbwachs et al., 1991) and Signal (Benveniste and Le Guernic, 1990). The second model is more subtle, requires a clock calculus, and, in the case of Signal, admits nondeterminism. In this chapter, we describe both models, beginning with the global clock model. In both cases, we focus not on the low-level programming languages like Esterel, Lustre, and Signal, but rather treat the SR model as a coordination semantics that governs the interaction of individual SR components. The components themselves could be language primitives (such as arithmetic operators like addition), circuit elements (such as gates) or software components realized using ordinary programming languages such as C, Rust, or Java. The examples we give all use very simple components, but the coordination principles apply even when the components have much more complicated logic.

6.1 The Global Clock Model

In this section, a global clock “ticks” repeatedly, potentially indefinitely. At each tick, each SR component reacts to any inputs that may be present and potentially produces outputs that are logically simultaneous with the inputs. All components conceptually react simultaneously and instantaneously. The idea that all components react simultaneously may trouble readers unfamiliar with this abstraction. As we will see, however, only direct hardware implementations in the form of synchronous circuits are actually implemented this way. Executing every component at every tick of a global clock could be absurdly inefficient, but the sparse synchronous model (Edwards and Hui, 2020) has very efficient implementations, as we will see. Component reactions can be scheduled in sophisticated ways so that they react only when there is work to be done. The *instantaneity* of the reactions may also be troubling. We will show that this abstraction leads to an elegant deterministic fixed-point semantics not unlike the ones given in Chapters 4 and 5 for process networks and dataflow.

Sidebar: About Synchrony

The term **synchronous** is used to mean (1) occurring or existing at the same time or (2) moving or operating at the same rate. In engineering and computer science, usages are mostly consistent with these definitions, but oddly inconsistent with one another. In referring to concurrent software using threads or processes, synchronous communication refers to a **rendezvous** style of communication, where the sender of a message must wait for the receiver to be ready to receive, and the receiver must wait for the sender. Conceptually, the communication occurs at the same time from the perspective of each of the two threads, consistent with definition (1). In Java, however, the keyword `synchronized` defines blocks of code that are *not* permitted to execute simultaneously, which is inconsistent with both definitions.

The **synchronous languages** (see box on page 210) have two key properties. First, the outputs of components in a program are (conceptually) simultaneous with their inputs (this is called the **synchrony hypothesis**). Second, components in a program execute (conceptually) **simultaneously and instantaneously**. Even though this cannot occur in reality, a correct execution must behave as though it did. This interpretation is consistent with *both* definitions (1) and (2) above because executions of components occur at the same time and at the same rate.

In circuit design, “synchronous” refers to a style where a clock signal that is distributed throughout a circuit causes circuit components called “latches” to record their inputs on the edges of the clock. The time between clock edges needs to be sufficient for circuit gates between latches to settle. Conceptually, this model is very similar to the model in synchronous languages. Assuming that the gates between latches have zero delay is equivalent to the synchrony hypothesis, and global clock distribution gives simultaneous and instantaneous execution of those gates. Hence, SR is often useful for modeling digital circuits.

In power systems engineering, “synchronous” means that electrical waveforms have the same frequency and phase. In signal processing, “synchronous” means that the sample rates of signals are the same or fixed multiples of one another. The term **synchronous dataflow**, described in Chapter 5, is based on this latter meaning of the word synchronous. This usage is consistent with definition (2).

Sidebar: Synchronous-Reactive Languages

The synchronous-reactive model of computation dates back to at least the mid-1980s, when a number of programming languages were developed. The term “reactive” comes from a distinction in computational systems between **transformational systems**, which accept input data, perform a computation, and produce output data, and **reactive systems**, which engage in an ongoing dialog with their environment (Harel and Pnueli, 1985). Manna and Pnueli (1992) state, “the role of a reactive program ... is not to produce a final result but to maintain some ongoing interaction with its environment.”

The distinctions between transformational and reactive systems led to the development of a number of innovative programming languages. The **synchronous languages** (Benveniste and Berry, 1991) take a particular approach to the design of reactive systems, in which pieces of the program react **simultaneously and instantaneously** at each tick of a global clock. Primary among these languages are Lustre (Halbwachs et al., 1991), Esterel (Berry and Gonthier, 1992), and Signal (Le Guernic et al., 1991). Statecharts (Harel, 1987) and its implementation in Statemate (Harel et al., 1990) also have a strongly synchronous flavor.

SCADE (Berry, 2003) (Safety Critical Application Development Environment), a product created by Esterel Technologies, later acquired by Ansys and then by Synopsys, builds on Lustre, borrows concepts from Esterel, and provides a graphical syntax in which state machines are drawn and components are composed synchronously. One of the main attractions of synchronous languages is their strong formal properties that facilitate formal analysis and verification techniques. For this reason, SCADE models are used in the design of safety-critical flight control software systems for commercial aircraft made by Airbus.

In Ptolemy II, SR is a form of coordination language rather than a programming language, (see also ForSyDe (Sander and Jantsch, 2004), which also uses synchrony in a coordination language). This allows for “primitives” in a system to be complex components rather than built-in language primitives. This, in turn, enables heterogeneous combinations of MoCs, since the complex components may themselves include components developed under another model of computation. The reactor model, covered in Chapter 8, can be viewed as a generalization of SR.

SR systems are good at orchestrating concurrent actions, managing shared resources, and detecting and adapting to faults in a system. Whereas [dataflow](#) models are good for managing streams of data, SR systems are good at managing sporadic data, where events may be present or absent, and where the absence of events has meaning (more than just transport delay). For example, detecting the absence of an event may be an essential part of a **fault management** system.

The SR model given here has been influenced by the family of so-called **synchronous languages** (see sidebar on page 210) and in particular Esterel ([Berry and Gonthier, 1992](#)). The SR model here most closely resembles the model of **synchronous block diagrams** as described by [Edwards and Lee \(2003b\)](#). The model of computation is closely related to synchronous digital circuits. In fact, this chapter will illustrate some of the ideas using circuit analogies.

SR can be viewed as describing **logically timed** systems, but there is no measurement of time involved. The ticks of the global clock are simply a sequence with no meaning to any time gap between ticks. The [discrete-event system](#) model considered in Chapter 8 generalizes this concept to make the time gap between ticks part of the semantics.

The similarities and differences with [dataflow](#) models are:

1. Like [homogeneous SDF](#), an iteration of an SR model consists of one iteration of each component in the model. Each iteration of the model corresponds to one tick of the logical clock.
2. Unlike dataflow and [process networks](#), there is no buffering on the communication between components. In SR, an output produced by one component is observed by the destination components in the same tick. Unlike [rendezvous](#), which also does not have buffered communication, SR is **deterministic**.
3. Unlike dataflow, an input or output may be **absent** (denoted ϵ) at a tick. In dataflow, the absence of an input means simply that the input hasn't arrived yet. In SR, however, an absent signal has more meaning. Its absence is not a consequence of accidents of scheduling or of the time that computation or communication may take. Instead, the absence of a signal at a tick is *defined* deterministically by the model. As a consequence, in SR, components may react to the *absence* of a signal. This is quite different from dataflow, where actors react only to the *presence* of a signal.




Figure 6.1: Model for a synchronous component under the global clock semantics.

6.1.1 An SR Component Model

An SR component under the global clock semantics can be modeled, as shown in Figure 6.1, as a sequence of functions $f(n)$, where $n \in \mathbb{N}$ is the index of the clock tick. Let the input and output **data types** be the sets T_1 and T_2 , respectively.

Definition 6.1. For a data type T , let the **extended data type** be $\tilde{T} = T \cup \{\perp, \epsilon\}$, where $\perp$ represents “unknown” and ϵ represents “absent.”

The function $f(n)$ has the form,

$$f(n): \tilde{T}_1 \rightarrow \tilde{T}_2.$$

Then at each tick $n \in \mathbb{N}$,

$$s_2(n) = (f(n))(s_1(n)).$$

In this model, at a tick, the input may be unknown or known. If it is known, then it is either absent or present with a value from the set T_1 . The resulting output may be unknown, absent, or present with a value from the set T_2 .

Do not confuse “unknown” and “absent.” Here, “unknown” will be a transient state, a temporary placeholder while we try to find the value of a signal $s(n)$ at tick n , whereas “absent” will be a final value indicating that the signal has no value, no event at tick n .

This model is easily generalized to tuples of inputs and outputs. For example, Figure 6.2 shows a component with two inputs and one output. Here, the function has the form

$$f(n): \tilde{T}_1 \times \tilde{T}_2 \rightarrow \tilde{T}_3.$$




Figure 6.2: Model for a synchronous component with a two-tuple of inputs and one output.



Figure 6.3: Models for synchronous components with no inputs or no outputs.

This model even extends to components with no inputs or no outputs, as shown in Figure 6.3. These are represented, respectively, by the functions,

$$f(n): T^0 \rightarrow \tilde{T}_1 \quad \text{and} \quad f(n): \tilde{T}_1 \rightarrow T^0,$$

where T^0 is a [singleton set](#).

Example 6.1: Consider a **parallel OR** gate with two inputs and one output. The data type of the inputs and output is $T = \{true, false\}$. The function $f(n)$ is given by the following truth table, showing $((f(n))(a, b))$,

$a \backslash b$	$\perp$	ϵ	<i>true</i>	<i>false</i>
$\perp$	$\perp$	$\perp$	<i>true</i>	$\perp$
ϵ	$\perp$	ϵ	<i>true</i>	ϵ
<i>true</i>	<i>true</i>	<i>true</i>	<i>true</i>	<i>true</i>
<i>false</i>	$\perp$	ϵ	<i>true</i>	<i>false</i>

Like an ordinary OR gate, the output is *true* if either input is *true*, and the output is *false* if both inputs are *false*. However, what makes this a **parallel** OR is that the output is *true* even if the other input is unknown.

In this design, when one input is absent (ϵ), then the output equals *true* if the other input is *true*, and ϵ if the other input is *false* or ϵ . An alternative design would define the output to be ϵ if either input is ϵ .

The parallel OR in the previous example is an instance of a **non-strict** component. Non-strict components are components that can produce outputs even when some of the inputs are unknown. These become important when there are cyclic dependencies between components.

Example 6.2: A **parallel AND** gate is similar, with the following truth table,

$a \backslash b$	$\perp$	ϵ	<i>true</i>	<i>false</i>
$\perp$	$\perp$	$\perp$	$\perp$	<i>false</i>
ϵ	$\perp$	ϵ	ϵ	<i>false</i>
<i>true</i>	$\perp$	ϵ	<i>true</i>	<i>false</i>
<i>false</i>	<i>false</i>	<i>false</i>	<i>false</i>	<i>false</i>

Like an ordinary AND gate, the output is *false* if either input is *false*, and the output is *true* if both inputs are *true*. However, what makes this a **parallel AND** is that the output is *false* even if the other input is unknown.

In this design, when one input is absent (ϵ), then the output is present and *false* if the other input is *false*, and ϵ if the other input is *true* or ϵ . An alternative design would define the output to be ϵ if either input is ϵ .

6.1.2 Components with State

The function $f(n)$ can be different at each tick. How does it change from tick to tick? Typically, a component maintains local **state** that changes from tick to tick. Let V be the set of possible states. The state is a function $v(n): \mathbb{N} \rightarrow V$. An initial state $v(0)$ is given. At each tick, the state changes according to a **state transition function** g that maps the current state and the input to the new state. For example, if there is a single input $s_1(n)$, then

$$v(n+1) = g(n, v(n), s_1(n)).$$

The input is required to be known at the time of the transition, so that the function g above has the form

$$g: \mathbb{N} \times V \times (T \cup \{\epsilon\}) \rightarrow V.$$

The function $f(n)$ is then updated using the new state,

$$f(n+1) = F(v(n+1)),$$

where F is a **functional** with the form (for a component with a single input and a single output),

$$F: V \rightarrow (\tilde{T}_1 \rightarrow \tilde{T}_2)$$

where $\tilde{T}_1$ and $\tilde{T}_2$ are the extended data types of the input and output, respectively. This concept generalizes to components with multiple inputs and outputs (or no inputs or no outputs).

Example 6.3: Consider a **counter** with a single output $s_1(n)$ and data type $T = \mathbb{N}$. The initial state is $v(0) = 0$. The state transition function is given by

$$g(n, v(n), s_1(n)) = v(n) + 1.$$

The function F is given by

$$F(v(n)) = f(n), \quad \text{where} \quad (f(n))(*) = v(n).$$

Because there are no inputs, the function $f(n)$ maps a singleton set (the member of which is indicated by $*$ above) to the output, which in this case is the natural number $v(n)$.

6.1.3 Fixed-Point Semantics

Consider a model with three components with the structure shown in Figure 6.4(a). As illustrated in Figures 6.4(b) through (d), any such model may be rearranged so that it becomes a single component with function $f(n)$ at tick n . The domain of this function is the cross product of the extended data types of the inputs, $\tilde{T} = \tilde{T}_1 \times \tilde{T}_2 \times \tilde{T}_3$. So is the codomain. Therefore, at tick n , the values of the inputs and outputs are given by the tuple $s(n) = (s_1(n), s_2(n), s_3(n))$ such that

$$s(n) = (f(n))(s(n)).$$




Figure 6.4: An SR model is reducible to a fixed point problem at each tick of the logical clock.

At each tick of the logical clock, an execution of the program must find a **fixed point** $s(n)$ of the function $f(n)$. The subtleties around SR models concern whether such a fixed point exists and whether it is unique. As we will see, in a well-constructed SR model, there is a unique least fixed point that can be found in a finite number of steps.

Even an SR model without feedback is reducible to a fixed point problem at each tick of the logical clock.

Example 6.4: A simple case of SR is a model without feedback, as shown in Figure 6.5. Even such a model is reducible to a fixed-point problem, but in this case it becomes a rather simple problem. The function $f_1(n)$ at tick n




Figure 6.5: Even an SR model without feedback is reducible to a fixed point problem at each tick of the logical clock.

only needs to be evaluated once at each tick, and it immediately finds the fixed point.

Example 6.5: Consider the circuit example from [Malik \(1994\)](#) shown in Figure 6.6(a). Assume the AND and OR gates are the parallel AND and OR gates described in Examples 6.2 and 6.1, respectively. Suppose that the input $s_1(n)$ is *true* at tick n . Does a fixed point exist? If so, is it unique? These questions can be answered in this case by considering all possible values for $s_2(n)$ and $s_3(n)$. If we first assume that $s_1(n) = \text{true}$, by trying all possible values for $s_2(n)$ and $s_3(n)$, we can determine that the only values that work (satisfy the logic of the gates) are $s_2(n) = \text{true}$ and $s_3(n) = \text{true}$. Therefore, the unique fixed point is $s(n) = (\text{true}, \text{true}, \text{true})$.

Consider now the case where the input $s_1(n)$ is *false* at tick n . In this case, it must be that $s_2(n) = \text{false}$, whereas $s_3(n)$ can be either *true* or *false*. Thus, there are two possible fixed points: $s(n) = (\text{false}, \text{false}, \text{true})$ and $s(n) = (\text{false}, \text{false}, \text{false})$.





Figure 6.6: Two circuit examples from Malik (1994) used to illustrate existence and uniqueness of fixed points.

Example 6.6: Consider the circuit shown in Figure 6.6(b), also from Malik (1994). Assume the AND and OR gates are similar to the parallel AND and OR gates described in Examples 6.2 and 6.1, respectively, but with the inputs and outputs inverted where the small circles appear. That is, a *true* on the second input of an AND gate is treated as a *false* and vice versa, and a *true* output from the OR gate becomes a *false*.

Now suppose that the input $s_1(n)$ is *true* at tick n . Does a fixed point exist? If so, is it unique? Again, we can consider all possible values for $s_2(n)$ and $s_3(n)$. If we first assume that $s_1(n) = \text{true}$, the only possible values are $s_2(n) = \text{true}$ and $s_3(n) = \text{false}$. Therefore, the unique fixed point is $s(n) = (\text{true}, \text{true}, \text{false})$.

Consider now the case where the input $s_1(n)$ is *false* at tick n . In this case, it must be that $s_2(n) = \text{false}$, but neither *true* nor *false* work for $s_3(n)$. Thus, there is no fixed point.

The above examples illustrate that whether a fixed point exists and is unique is not always straightforward to answer. Moreover, the method used in these examples, searching over all possible values for the inputs and outputs, is not practical for most models. In the next section, we will develop a constructive procedure for finding fixed points that is practical and scalable, but we will see that this method rejects some models that have unique fixed points. To include



A cyclic token-ring system composed of three blocks. This system arbitrates fairly among requests for exclusive access to a shared resource by marching a token around a ring. At each "tick" of the clock, the arbiter grants access to the first requestor downstream of the block with the token.

From (Edwards and Lee, *The Semantics and Execution of a Synchronous Block-Diagram Language*, Science of Computer Programming, 2003)


Figure 6.7: A token ring implementation in Ptolemy II from Edwards and Lee (2003a).

such models, the only known solution is to use a non-constructive method, such as exhaustive search.

Example 6.7: A practical example that includes the subtleties of the toy examples above is a **token ring** implementation in Ptolemy II from Edwards and Lee (2003a). The token-ring system is composed of three blocks, each of which is an arbiter, arranged in a ring. This system arbitrates fairly among requests for exclusive access to a shared resource by marching a token around a ring. At each tick of the global clock, the arbiter grants access to the first requestor downstream of the block that holds the token.

An implementation of the arbiter is shown in Figure 6.8. This uses **non-strict** logic to implement the parallel OR and AND gates. If an arbiter owns the token and a request is made to that arbiter, then it will grant access. If it owns the token and no request is made to that arbiter, then it will pass out permission to use the token to the downstream arbiter. If an arbiter does not own the token, but the permission to use the token is passed in to that

● initiallyOwnsToken: false



If this owns the token and a request is made, then grant access.
 If this owns the token and no request is made, then pass out permission to use the token. If this does not own the token, but the permission to use the token is passed in, then if a request is made, grant access.
 In all cases, pass the permission to use the token to the next arbiter.

Figure 6.8: A token arbiter implementation in Ptolemy II from [Edwards and Lee \(2003a\)](#).

arbiter, then if a request is made to that arbiter, it will grant access. In all cases, an arbiter will pass the permission to use the token to the next arbiter if it does not grant access.

Key to this model is that permissions are passed around in a cycle within a single tick of the global clock. The starting point of the cycle is the arbiter that owns the token at the beginning of the tick, and this changes on each tick, so there is no fixed position at which the causal chain can be broken. This logic is not implementable without some form of the logic in parallel OR and AND gates.




Figure 6.9: A flat partial order for the data type $T = \{true, false\}$.

6.1.4 Constructive Semantics

For a model described as the feedback system of Figure 6.4(d), we can construct a CPO and use the Kleene fixed-point theorem (Proposition A.6) to find a fixed point via a constructive procedure. First, we construct a flat partial order for each extended data type in the model. In this partial order, the bottom element is the unknown value $\perp$, and the other elements are the possible values for the data type plus the absent symbol ϵ . An example of a flat partial order is shown in Figure 6.9 for the data type $T = \{true, false\}$.

The flat partial orders for each signal in the model are then combined into a pointwise order by taking the Cartesian product of the partial orders for each signal. An example of a pointwise order is shown in Figure 6.10 for the data type $T_1 \times T_2 = \{true, false\}^2$.

To apply the Kleene fixed-point theorem on models like that of Figure 6.4(d), we need for the function $f(n)$ to be monotonic. The function $f(n)$ will be monotonic automatically if each of the component functions $f_i(n)$ is monotonic.

In the pointwise order, it is natural for a component to implement a monotonic function. Intuitively, a component that implements a monotonic function does not “change its mind” about outputs when inputs become known. In other words, if it asserts a non-unknown output when an input is unknown, then it will assert the same non-unknown output when the input becomes known. A function that does this is always monotonic in the pointwise order.

Example 6.8: Consider the parallel OR function from Example 6.1. This function is monotonic in the pointwise order. This is easy to check (albeit




Figure 6.10: A pointwise order for the data types $T_1 \times T_2 = \{true, false\}^2$, where t and f abbreviate *true* and *false*, respectively.

somewhat tedious) using the truth table given in Example 6.1. Intuitively, however, the parallel OR only asserts a known output when it has enough information about the inputs to be sure of the output value. Given more information about the inputs, it does not “change its mind.”

The [parallel AND](#) function from Example 6.2 is also monotonic in the pointwise order.

Note that the pointwise order for any finite tuple constructed from a flat partial order always has only finite [chains](#). Consequently, every monotonic function is also [continuous](#). As a result, the [Kleene fixed-point theorem](#) provides a constructive procedure for finding a fixed point of a monotonic function. We just need to construct the [Kleene chain](#) of the function and find its least upper bound.

A procedure, therefore, for determining the behavior of any network of monotonic SR components is to:

1. Redraw the network as shown in Figure 6.4 to obtain a single [endofunction](#), the $f(n)$ in Figure 6.4(d), representing the behavior of the network.
2. Construct the [Kleene chain](#) of the function and find its least upper bound:
 - (a) Start with the bottom element of the pointwise order, $(\perp, \perp, \dots, \perp)$.

- (b) Apply the function $f(n)$ iteratively to get the next element.
 - (c) Repeat until a fixed point is reached (no progress is made).
3. The resulting least fixed point is the behavior of the network.

Sidebar: Information Order

When formulating the constructive fixed-point semantics, an alternative to the [pointwise order](#) is an **information order**. In an information order, each node in the [Hasse diagram](#) is a set of possible values for a signal or tuple of signals. The bottom element, replacing $(\perp, \perp)$ in Figure 6.10, is the set of all possible tuples, specifically, for that data type,

$$\{(\epsilon, \epsilon), (\epsilon, t), (\epsilon, f), (t, \epsilon), (t, t), (t, f), (f, \epsilon), (f, t), (f, f)\}.$$

This set is simply another way of writing that both elements of the tuple are unknown. The elements in the second layer of the Hasse diagram will have one element unknown and the other known. For example, corresponding to $(\perp, t)$ in Figure 6.10, the set of possible values is

$$\{(\epsilon, t), (t, t), (f, t)\}.$$

This set is simply another way of writing that the first element of the tuple is unknown, and the second element is known. The elements in the third layer of the Hasse diagram will have both elements known. For example, corresponding to (t, t) in Figure 6.10, the set of possible values is a singleton set,

$$\{(t, t)\}.$$

If we add a top element, the empty set, then the information order becomes a [lattice](#).

Notice that the order relation is reverse set inclusion. That is, an element x is less than or equal to an element y if and only if x is a superset of y . A monotonic function, therefore, is one where if you are given more information about the inputs, you can only get more information about the outputs. Here, more information means that the set of possible values gets smaller.




Figure 6.11: The circuit from Figure 6.6(a) redrawn as suggested in Figure 6.4.

This procedure terminates in a finite number of steps if the height of the pointwise order is finite. Because of the [Kleene fixed-point theorem](#) (Proposition A.6), the procedure is guaranteed to find the least fixed point, which is guaranteed to be unique.

There remains a potential problem, however. The least fixed point may include unknowns. It is unsatisfying to have a model that yields unknown behavior. This motivates the following definition.

Definition 6.2. *A synchronous-reactive system consisting of monotonic components is said to be **constructive** if its least fixed point includes no unknowns (i.e., it is maximal in the pointwise product of flat CPOs).*

Example 6.9: The circuit example in Figure 6.6(a) can be redrawn to set up the fixed-point iteration as shown in Figure 6.11(b). For notational compactness, let the data type of the circuit be $T = \{0, 1\}$ rather than $\{\text{true}, \text{false}\}$.

The first problem we must address is that there is an externally provided input signal $s_1(n)$, and Figure 6.4 does not provide for such inputs. To model this, we can add a source box that produces the value of $s_1(n)$ on each tick n , as shown in Figure 6.11(b). The box is labeled “0 or 1” to indicate that it produces either 0 or 1 on each tick n , depending on the value of $s_1(n)$.

The feedback structure in Figure 6.11(b) is chosen so that the grey box represents a function of the form $f(n): \tilde{T}^3 \rightarrow \tilde{T}^3$. Notice that the output $s_2(n)$ is fed back but ignored by the function. Making it an (ignored) input enables us to model the whole system with a single **endofunction**, so that we can define the behavior as a fixed point of that endofunction.

To find the fixed point, we can construct the **Kleene chain** of the function $f(n)$ and find its least upper bound. The Kleene chain is constructed by starting with the bottom element of the pointwise order, which is $(x_0, y_0, z_0) = (\perp, \perp, \perp)$. The second element of the chain is $(x_1, y_1, z_1) = f(0)(x_0, y_0, z_0) = (1, \perp, \perp)$, assuming that $s_1(0) = 1$. Then, we apply the function $f(0)$ iteratively to get the next element, $(x_2, y_2, z_2) = f(0)(x_1, y_1, z_1) = (1, \perp, 1)$, where $z_2 = 1$ arises because the parallel OR can resolve its output even though one input is unknown. We apply the function $f(0)$ one more time to get $(x_3, y_3, z_3) = f(0)(x_2, y_2, z_2) = (1, 1, 1)$. One more application gets $(x_4, y_4, z_4) = (1, 1, 1)$ which is the same (x_3, y_3, z_3) , so we have reached a fixed point. We now declare the behavior of the circuit to be the fixed point $(s_1(0), s_2(0), s_3(0)) = (1, 1, 1)$.

Now suppose that $s_1(1) = 0$. Using the same procedure, we find the following sequence:

$$\begin{aligned}(x_0, y_0, z_0) &= (\perp, \perp, \perp) \\ (x_1, y_1, z_1) &= (0, \perp, \perp) \\ (x_2, y_2, z_2) &= (0, 0, \perp) \\ (x_3, y_3, z_3) &= (0, 0, \perp)\end{aligned}$$

We have reached a fixed point, so $(s_1(1), s_2(1), s_3(1)) = (0, 0, \perp)$. Because the fixed point includes an unknown, the behavior of the circuit is not constructive.

In the early development of synchronous languages, many researchers assumed that the only requirement for a model to be correct was that it had a unique fixed point. [Berry \(1999\)](#), after several attempts to create an Esterel compiler that was




Figure 6.12: Non-constructive circuit with a single fixed point from [Malik \(1994\)](#).

assured of finding any unique fixed point, finally gave up and declared any non-constructive program to be defective. Later, [Mendler et al. \(2012\)](#) showed that all non-constructive circuits have behaviors that are either nondeterministic or that oscillate, giving a solid justification for Berry's decision. The circuit considered in the previous example is nondeterministic because either solution $s_3(1) = 0$ or $s_3(1) = 1$ is valid. This will generally be true for circuits with more than one fixed point.

What about circuits with no fixed points? In Example 6.6 (Figure 6.6(b)), the circuit has no fixed points when the input $s_1(n)$ is *false*. If there is a time delay in the OR gate, the circuit will oscillate between *true* and *false*. This will generally be true for circuits with no fixed points. What about circuits that have a unique fixed point but are not constructive?

Example 6.10: Consider the circuit shown in Figure 6.12. This circuit has a single fixed point for all input values if we consider only the data type $T = \{0, 1\}$ and not the extended data type $\tilde{T} = \{\perp, \epsilon, 0, 1\}$. But the circuit is not constructive. When the input $s_1(n)$ is 0, the unique fixed point is $s(n) = (0, 1, 0)$. When the input $s_1(n)$ is 1, the unique fixed point is $s(n) = (1, 0, 0)$.

We can use our intuition to find these fixed points by noting that the output of the AND gate has to always be zero. Its two inputs are opposite of each other because of the small circle on the lower input, which negates the value, so one of the inputs to the AND gate must be zero and the other must be one. Hence, the output of the AND gate must be zero.

However, if we try to evaluate the circuit using the Kleene constructive procedure, we will find the least fixed point $s(n) = (0, \perp, \perp)$ when $s_1(n) = 0$. The output of the AND gate cannot be determined using only its truth table when the two inputs are unknown, so $s_3(n)$ and hence also $s_2(n)$ remain unknown.

Conforming with the result of [Mendler et al. \(2012\)](#), the circuit can oscillate if one of the inputs to the AND gate incurs a delay that is not exactly equal to the delay of the other input. Specifically, if the input $s_1(n)$ is 0, then if we assume all circuits initially have low voltages (corresponding to 0), then after some gate and wire delays, $s_2(n)$ will settle to 1. But if the two inputs to the AND gate have different delays, then they may briefly result in an output of 1, which will then force $s_2(n)$ to 0. This will force $s_3(n)$ to 0, which will then force $s_2(n)$ to 1 again, and the cycle will repeat. This oscillation will continue indefinitely.

The problems identified by [Mendler et al. \(2012\)](#) are not limited to Boolean circuits. They apply to any synchronous-reactive system consisting of monotonic components. Hence, we will reject any such system that is not constructive. This motivates the following proposition.

Proposition 6.1. *A synchronous-reactive system consisting of monotonic components that is not [constructive](#) should be rejected by a compiler.*

6.1.5 Concurrent State Machines

The synchronous-reactive model of computation with a global clock is a good fit for circuit design, but it is much more general than that. Components in the model can be arbitrarily complex, including components with state, as long as they realize monotonic functions in the pointwise extension of the flat partial order. Intuitively, any software component can realize such a monotonic function as long as the state of the component is not changed until all its inputs are known. To illustrate this general idea, we will borrow notation from [Lee and Seshia \(2017\)](#) for concurrent state machines.

Example 6.11: Consider the state machine shown in Figure 6.13. This machine has two states, s_1 and s_2 , one input, x , and one output, y . We have dropped the tick index n for notational simplicity, but please keep in mind that these values change with each tick. The machine transitions between states based on the value of the input x and produces an output y that depends on which transition it takes.

The input x and output y are declared to be **pure**, meaning that they have no meaningful value. They are simply present or absent. We can model this as a singleton data type $T = \{present\}$, giving the extended data type $\tilde{T} = \{\perp, \epsilon, present\}$.

The initial state is s_1 . There are two possible transitions from state s_1 . There is a transition to state s_2 that will be taken if x is ϵ and a self-loop transition back to state s_1 that will be taken if x is *present*. The notation on the transitions is of the form *guard/action*, where the guard is a boolean expression and the action is an assertion about the output y . The notation $x/$, with nothing after the slash, means that the guard is x (x is present) and the action is that y is ϵ , absent. The notation $\neg x/$, with nothing after the

Sidebar: Intuitionistic Logic

In classical (binary) logic, given a predicate p , the predicate $p \vee \neg p$ is always *true*. This is called the **law of the excluded middle**.

In **intuitionistic logic**, $p \vee \neg p$ is not necessarily *true*. It may be unknown. In intuitionistic logic, if $\neg\neg p$ is *true*, this does not imply that p is *true*. It, likewise, may be unknown. For a predicate to be *true* in intuitionistic logic, you must give direct evidence of its truth by a constructive procedure. The Kleene constructive procedure provides such evidence. Hence, Proposition 6.1 feels like applying intuitionistic logic to the problem of finding a fixed point.

Suppose, for example, that we interpret the circuit in Figure 6.12 as a statement in classical logic. In classical logic, by the law of the excluded middle, the value of $s_2(n)$ must be either 0 or 1, and hence, the value of $s_3(n)$ must be 0. But if we interpret the circuit in Figure 6.12 as a statement in intuitionistic logic, then we cannot assume that $s_2(n)$ is either 0 or 1, and hence, the value of $s_3(n)$ remains unknown.

slash, means that the guard is $\neg x$ (x is ϵ) and the action is that y is ϵ . The notation x/y , with y after the slash, means that the transition can be taken if x is *present* and that y will then be asserted to be *present*. Etc.

In the diagram, a feedback path labeled s has been added. This is an assertion that at every tick, $x = y = s$. The input and output are required to be identical. Hence, the value of the output y must be a fixed point of the function f that maps input values to output values. The function f depends on the state of the machine. In state $s1$, the function f is given by

Sidebar: Synchronous State Machines

Finite state machines (FSMs) have a long history in the theory of computation (Hopcroft and Ullman, 1979). An early model for concurrent FSMs is **Statecharts**, developed by Harel (1987). A modern version of Statecharts is **SCCharts**, developed by Hanxleden et al. (2014a). SCCharts, like our model here, restricts systems to constructive compositions. Another realization of constructive concurrent FSMs is found in Ptolemy II (Feng et al., 2014a), which enables hierarchical compositions of FSMs with other models of computation (Feng et al., 2014b), including **dataflow**, **discrete-event**, and continuous-time models. Another recent version is found in the modal models of Lingua Franca (Schulz-Rosengarten et al., 2023). That version, like Ptolemy II, is hierarchical, and within any state, an input-output behavior can be specified using an imperative programming language.

The **Esterel** synchronous language (Berry and Gonthier, 1992) also has the semantics of hierarchical state machines, although it is given a textual syntax rather than a graphical one. Esterel has a rigorous SR semantics for concurrent composition of state machines (Berry, 1999). **SyncCharts**, which came later, provides a visual syntax (André, 1996).

PRET-C (Andalam et al., 2010), **Reactive C (RC)** (Boussinot, 1991), and **Synchronous C (SC)** (von Hanxleden, 2009; Hanxleden et al., 2014b) are C-based languages inspired by Esterel that support concurrent state machines. In both RC and PRET-C, state refinements (which are called “threads”) are executed sequentially in a fixed, static order.




Figure 6.13: A simple well-formed feedback model.

the following table:

x	y
$\perp$	ϵ
<i>present</i>	ϵ
ϵ	ϵ

Notice that the output y is ϵ for all input values, including the unknown value $\perp$. Hence, we can apply the Kleene constructive procedure to find the least fixed point of the function f , which is ϵ .

In state $s2$, the function f is given by the following table:

x	y
$\perp$	<i>present</i>
<i>present</i>	<i>present</i>
ϵ	<i>present</i>

Notice that the output y is *present* for all input values, including the unknown value $\perp$. Hence, we can again apply the Kleene constructive procedure to find the least fixed point of the function f , which is *present*.

Hence, the least fixed point of the function f is ϵ in state $s1$ and *present* in state $s2$.

In the previous example, the state machine is deterministic (see sidebar on page 231), and it is embedded in a feedback model that is constructive. The state

machine together with the feedback connection is equivalent to the state machine shown in Figure 6.14. It simply alternates between the two states, and the feedback signal s alternates between ϵ and *present*.

In the previous example, it is important to note that the input x and output y have the *same value* at every tick. This is what is meant by the feedback connection. Any connection from an output port to an input port means that the value at the input port is the same as the value at the output port at all ticks.

Also, it is important to note that when the input is unknown, which state transition will be taken is usually not determined. It happens that, in this case, the output is known despite not knowing which transition will be taken.

It is easy to see how to generalize this so that the signal s can have any type. Signal s can even be an aggregation of signals, as in Figure 6.4(d).

Sidebar: Nondeterministic State Machines

A state machine is deterministic if in any state, at most one outgoing transition can become enabled for any input. All the state machines in Section 6.1.5 are deterministic. If a state machine is nondeterministic, the Kleene constructive procedure will have to be modified using a **powerdomain construction** (Plotkin, 1976).

Specifically, the extended data type $\tilde{T}$ is replaced by the powerdomain $P = 2^{(T \cup \{\epsilon\})}$, where 2^X is the set of all subsets of X . In P , the unknown value $\perp$ is replaced by the empty set $\emptyset$. The function f is replaced by a function $F: P \rightarrow P$ that maps sets of possible input values to sets of possible output values (including ϵ). The partial order on P is defined by set inclusion. Hence, a monotonic function F is one that given more possible input values, does not remove output values that were possible before it was given those input values.

A constructive composition of nondeterministic state machines is one that yields a non-empty set of possible values for every input and output. The values in the least fixed point define the possible behaviors of the state machine composition. If the least fixed point provides a singleton set for each input and output, then the composition is deterministic even if the state machines themselves are nondeterministic.




Figure 6.14: The semantics of the model in Figure 6.13.




Figure 6.15: An ill-formed feedback model that has no fixed point in state s_2 .

As with the circuit example, there are two potential problems that may occur when seeking a fixed point. First, there may be no fixed point. Second, there may be more than one fixed point. If either case occurs in a reachable state, we call the system an **ill-formed composition**. Otherwise, it is a **well-formed composition**.

Example 6.12: Consider machine B shown in Figure 6.15. In state s_1 , we get the unique fixed point $s(n) = \epsilon$. In state s_2 , however, there is no fixed point. If we attempt to choose $s(n) = \text{present}$, then the machine will transition to s_1 and its output will be ϵ . But the output has to be the same as the input, and the input is present , so we get a contradiction. A similar contradiction occurs if we attempt to choose $s(n) = \epsilon$.




Figure 6.16: An ill-formed feedback model that has more than one fixed point in state s_1 .

Since state s_2 is reachable, this feedback model is ill formed. It is also not constructive. If we attempt to apply the Kleene constructive procedure, we will find in state s_2 that the output is unknown when the input is unknown, so the least fixed point is $\perp$.

Example 6.13: Consider machine C shown in Figure 6.16. In state s_1 , both $s(n) = \epsilon$ and $s(n) = \text{present}$ are fixed points. Either choice is valid. Since state s_1 is reachable, this feedback model is ill formed. Again, it is also not constructive.

If in a reachable state there is more than one fixed point, we declare the machine to be ill formed. An alternative semantics would not reject such a model, but rather would declare it to be nondeterministic. This would be a valid semantics, but it would have the disadvantage that a composition of deterministic state machines is not assured of being deterministic. In fact, C in Figure 6.16 is deterministic, and under this alternative semantics, the feedback composition in the figure would not be deterministic. Determinism would not be a compositional property. Hence, we prefer to reject such models. This provides further support for rejecting non-constructive models, Proposition 6.1.

Notice that calling these models “compositions” is a bit odd because there is only one state machine in the composition. However, following the pattern of




Figure 6.17: A well-formed feedback model that is not constructive.

Figure 6.4, we can take any set of interconnected state machines and model them as a single state machine with a single function. The state space of the single state machine is the cross product of the state spaces of the individual state machines. The function f is the composition of the functions of the individual state machines.

As we saw with circuits before, rejecting non-constructive models may reject models that have a unique fixed point. As illustrated by the following example, the same is true for state machines.

Example 6.14: Consider machine D and its feedback composition shown in Figure 6.17. This composition is well-formed, but it is not constructive. In state $s1$, if the input is unknown, we cannot immediately tell what the output will be, so the output remains unknown, so the least fixed point is $\perp$. However, there is only one fixed point, specifically $s(n) = \epsilon$, if we consider all known possible values.

Determining the outputs of a state machine when (some) inputs are unknown follows a procedure known as a **must-may analysis**. Examining the machine, we can determine what *must* be true of the outputs and what *may* be true of the outputs.

Example 6.15: In Figure 6.17 in state s_1 , the output *may* be 1 or may be absent. If we do not know the input, we cannot determine which of these outputs will prevail. In state s_2 , in contrast, the output *must* be absent.

In Figure 6.13 in state s_1 , we can immediately determine that the output *must* be ϵ . Of course, once we have determined that the output is ϵ , we now know that the input is ϵ , and hence the procedure concludes.

In state s_2 , we can immediately determine that the output *must* be *present*.

The fixed-point iteration procedure can be generalized to an arbitrary model structure. In Figure 6.4(a), there is no real need to convert it to the form of Figure 6.4(d). Instead, we can just begin by labeling all signals unknown, and then in arbitrary order, examine each component to determine whatever can be determined about the outputs, given the initial state of each machine. We repeat this until no further progress can be made, at which point either all signals become known, or we reject the model as non-constructive. Once we know all signals, then all components can make state transitions, and we repeat the procedure in the new state for the next reaction.

Evaluating components in arbitrary order is rather inefficient. We consider in the next section scheduling strategies that evaluate components in a more efficient order.

6.1.6 Scheduling Strategies

It is easy to show that given N components that define M outputs arranged in a feedback loop as shown in Figure 6.4(b), the least fixed point can be found by evaluating the function given by each component no more than M times, so the total cost of a brute-force evaluation is $O(MN)$. If we simply evaluate the functions in an arbitrary order, if no new output moves from $\perp$ to something known, then we can stop. We have found a fixed point. Otherwise, at least one output must have moved from $\perp$ to something known. This procedure need not be repeated more than M times because, at this point, all outputs are known or a fixed point has been reached.

In practice, we can do much better than this. First, by examining the connections between components and determining which outputs depend on which inputs,




Figure 6.18: A refactorization can convert the problem of finding a fixed point into a succession of smaller fixed-point problems.

we can use these dependencies to choose a more efficient order for evaluating the functions. For example, in Figure 6.4(b), the third component does not depend on any inputs at all, so it can be evaluated first with assurance that it will produce a known output. Once the output $s_3(n)$ is known, we have a smaller fixed-point problem to solve involving only the first two components. This smaller problem has two components and two outputs, so at most four function evaluations are needed to find a fixed point. Adding the first evaluation of the third component, we see that at most five function evaluations are needed to find a fixed point, vs. the worst case of nine function evaluations for the arbitrary order evaluation.

The conversion of the problem into smaller fixed-point problems requires a mechanism for finding the least fixed point of a function that is not an [endofunction](#). This is illustrated in Figure 6.18, where the boxes from Figure 6.4 are refactored so




Figure 6.19: A parallel OR with a feedback loop.

that only two of the three boxes are combined into one, whose function now has the form $F: \tilde{T}^3 \rightarrow \tilde{T}^2$. Once the input is determined, this can be used as a constant part of the definition of the function F , and, for that particular input, we can find a fixed point using the Kleene constructive procedure.

An alternative that sometimes works is to analytically find a function $G: \tilde{T} \rightarrow \tilde{T}^2$ that realizes the box labeled G in Figure 6.18(d). In other words, we seek a function $G: \tilde{T} \rightarrow \tilde{T}^2$ such that

$$G(s_3) = (s_1, s_2),$$

where we have omitted the tick index n for brevity. This function has to satisfy,

$$F(G(s_3), s_3) = G(s_3).$$

This is a new kind of fixed-point problem, where now we seek a function G that is a fixed point of a **functional** $\phi: [\tilde{T} \rightarrow \tilde{T}^2] \rightarrow [\tilde{T} \rightarrow \tilde{T}^2]$ that is defined by,

$$(\phi(G))(s_3) = F(G(s_3), s_3).$$

As shown in Section A.5.3, the set of continuous total functions $[\tilde{T} \rightarrow \tilde{T}^2]$ is a CPO under the pointwise order. Therefore, we can use the Kleene constructive procedure (Proposition A.6) to find the least fixed point of the functional ϕ .

Example 6.16: Consider a specific simpler example where the function $F: \tilde{T}^2 \rightarrow \tilde{T}$ is a **parallel OR** with a feedback loop, as shown in Figure 6.19.

The functional $\phi: [\tilde{T} \rightarrow \tilde{T}] \rightarrow [\tilde{T} \rightarrow \tilde{T}]$ is now given by,

$$(\phi(G))(x) = \begin{cases} 1 & \text{if } x = 1 \\ \perp & \text{if } x = \perp \text{ and } G(x) = 0 \\ G(x) & \text{otherwise} \end{cases}$$

Applying the Kleene constructive procedure, we evaluate

$$\phi(\perp_{[\tilde{T} \rightarrow \tilde{T}]}), \phi(\phi(\perp_{[\tilde{T} \rightarrow \tilde{T}]}), \dots$$

until fixed point, getting (immediately),

$$G(x) = \begin{cases} 1 & \text{if } x = 1 \\ \perp & \text{otherwise} \end{cases}$$

Because of the possibility of a $\perp$ result, this circuit is not constructive.

In general, applying this procedure to find a closed-form representation of the input-output behavior of the system can be quite difficult. It has been shown by [Shiple et al. \(1996\)](#) to be effective for small circuits, where the functions of the components are known and simple and the data types are small. Nevertheless, refactoring the model to convert it into a succession of smaller fixed-point problems is a general technique that can be applied to any model, although it may be hard to know exactly which of many possible refactorings to choose. The problem of how to do this best has been studied extensively in the synchronous language community. See [Edwards and Lee \(2003a\)](#) for a detailed examination of this problem.

One common approach is to restrict the semantics of the model by requiring that every directed cycle have at least one component that can produce a known output given only unknown inputs. In this case, a compiler can statically determine a schedule that will produce known outputs for each component in the cycle. This approach is taken in Lustre ([Halbwachs et al., 1991](#)), for example. A more refined approach breaks each component into chunks (called [reactions](#)) that provide a partial evaluation of the component function, and restrict dependencies between reactions to be acyclic. This approach is used in Lingua Franca, which will be discussed in the next chapter. In both of these approaches, there is some loss of expressiveness. In particular, cyclic dependencies like that in the token-ring

example (Example 6.7) cannot be handled directly, although there are relatively simple work-arounds.

6.2 Alternative Clock Models

So far, we have assumed a global clock where every component reacts at every tick. In practice, this is not always the most useful model. We will now consider some alternative clock models that allow for more flexibility in the scheduling of components.

6.2.1 Clock Domains

In many applications, different components have different rates of operation. The single clock model that we have been using can be easily extended by dividing an application into **clock domains**, where each clock domain has its own clock. The most straightforward way to do this is using **subclocks**, where clock domains are nested hierarchically, and a clock ticks in a clock domain only when some predicate evaluates to *true* in the parent clock domain. This strategy has been implemented in Esterel (Berry and Sentovich, 2001) and in Ptolemy II (Edwards et al., 2014), for example. In Ptolemy II, subclocks are created using the **EnabledComposite** component, as illustrated in the next example.

Example 6.17: Figure 6.20 shows a **guarded count** example with a clock domain hierarchy realized in Ptolemy II. The **EnabledComposite** component is used to create a subclock within which the starting number is provided to the **CountDown** component. The **CountDown** component is then used to count down from the starting number. The **Display** components are used to display the count down and the enable signal for the **EnabledComposite** component. This example is often used to illustrate features of synchronous-reactive languages.

The first component to understand is the **NonStrictDelay**. This component provides a one-tick delay. On each firing, it produces on the output port whatever value it read on the input port in the previous tick of the clock. If the input was absent on the previous tick of the clock, then the output will




Figure 6.20: A Guarded Count example with a clock domain hierarchy realized in Ptolemy II.

be absent. On the first tick, the output is the *initialValue* parameter (shown in the icon), if it is given, and absent otherwise.

On the first tick, the `NonStrictDelay` inside the `CountDown` component produces a zero-valued output, because the *initialValue* parameter is 0. It is *non-strict* in that it does not require the input to be known in order to produce an output. The downstream `Expression` component will produce a *true* output, because the input is ≤ 0 . This results in the *ready* port of the `CountDown` component becoming known and *true*, which causes the `EnabledComposite` component to become enabled in this first tick.

When it is enabled, the `EnabledComposite` component will produce the numerical sequence specified by the `Sequence` component, which in the first tick will be 5, followed by 3 in the next (local) tick, followed by 2 in the next (local) tick.

Hence, in the first tick, the `CountDown` component will receive a start value of 5. In the **Default** component, when the *preferred* input has a token, then the output is equal to that token. If the *preferred* input is absent, then the output is equal to the *alternate* input (whether it is absent or not). This component is non-strict, in that it can produce an output even if the *alternate* input is unknown. As a consequence, the `Default` component will select the 5 rather than the -1 provided by the feedback from the `NonStrictDelay` because the 5 appears on its *preferred* port.

In the **When** component, the output is equal to the input when the *control* port has the value *true*. Thus, the count output port will produce a 5 in the first tick.

In the second tick, the `NonStrictDelay` inside the `CountDown` component will produce a 5 output, because the input was 5 in the previous tick. As a consequence, the *ready* output will be *false* in this second tick. Hence, the `EnabledComposite` component will not be enabled in this second tick. The 5 will be fed back, decremented by 1 by the `Expression2` component, and sent to the `Default` component, which will select the resulting 4 because its *preferred* port will be absent at this tick. The 4 will be produced on the count output port of the `CountDown` component in this second tick.

In the third tick, the `NonStrictDelay` inside the `CountDown` component will produce a 4 output, because the input was 4 in the previous tick, and the previous process will repeat, producing a 3 on the `count` output port of the `CountDown`. This continues until the count reaches 0, at which point the `ready` output will become *true* in the next tick, which causes the `EnabledComposite` component to become enabled to provide the next starting number in the next tick. Hence, the count will be displayed as 5, 3, 2, 1, 0, 3, 2, 1, 0, 2, 1, 0, $\dots$.

In summary, the `EnabledComposite` component creates a subclock that will tick only on the first tick and the tick after the count down from the previously provided starting number reaches 0. The `NonStrictDelay`, `Default`, and `When` components correspond to primitives found in many synchronous-reactive languages, including Esterel and Lustre. Nearly all synchronous-reactive languages provide a mechanism for creating subclocks, and something analogous to the `EnabledComposite` component is a standard primitive in many of them.

6.2.2 Sparse Synchronous Models

The global clock model considered so far has a property that can make it inefficient for some applications. In particular, many components will never produce non-absent outputs unless at least one of their inputs is present. So far, we have made no assumption about the components, and hence the component function needs to be evaluated at every tick of the clock, even if the inputs are all absent. A more intelligent scheduler, therefore, could invoke component functions only when necessary, and skip the evaluation of components that will never produce non-absent outputs.

Components without inputs or components that produce outputs even when the inputs are all absent may be able to predict which future ticks of the clock will result in non-absent outputs. Such components could inform the scheduler of which future ticks of the clock will result in non-absent outputs, and the scheduler could then skip the evaluation of the component on other ticks.

This approach is called the **sparse synchronous model** by [Edwards and Hui \(2020\)](#). In the next chapter, we will see that the [reactor model](#) leverages this idea






Figure 6.21: Example used to explain the clock calculus.



(a)



(b)

Figure 6.22: Plot output under the global clock model.

and goes further to define a vast number of ticks, where the vast majority of them will be skipped.

6.2.3 Clock Calculus

Some synchronous-reactive languages, particularly Lustre ([Halbwachs et al., 1991](#)) and Signal ([Benveniste and Le Guernic, 1990](#)), have a different interpretation of clocks that is much more subtle than the global clock model. In these languages, a clock is associated with each signal in the system rather than with the whole system or a subsystem with a subclock. Components then impose constraints on the clocks of the signals they use, and the scheduler and/or compiler must respect these constraints. To understand the difference, consider an example.

Example 6.18: Consider the example in Figure 6.21. The global clock model will produce a plot like that shown in Figure 6.22(a). The clock calculus model will produce a plot like that shown in Figure 6.22(b). In both cases, two **Ramp** components are used to generate a counting sequence, and these counting sequences are added together.

Under the global clock model, the two Ramp components are evaluated at every tick of the global clock. The **When** component is the same as that considered in guarded count, Example 6.17. Here it is controlled by a **Bernoulli** component, which produces a pseudo-random sequence of *true* or *false* with a probability of 0.5 at each tick. When the control input to the When component is *false*, the output is absent. The Add component, therefore, adds one reliable counting sequence to another, unreliable counting sequence, which is absent at approximately half the ticks. This produces a plot like that shown in Figure 6.22(a).

Under the clock calculus model, the Add component imposes a constraint on the clocks of its two incoming signals. Specifically, it requires that they both be present whenever it reacts. The When component, however, is much more subtle. The simplest version is where, like the Add, it requires its two inputs to be present at any tick where it reacts. Its output will be absent at approximately half the ticks, whenever the input from the Bernoulli component is *false*. Because of the constraints imposed by the Add component, these absent outputs do not align with any output from Ramp2. The result is a plot like that shown in Figure 6.22(b). Here, the effect is that approximately half of the outputs from the Ramp are discarded by the When component, and the remaining outputs are added to the outputs from Ramp2.

The clock calculus model is quite subtle. In the previous example, the relationship between the ticks of Ramp and Ramp2 is peculiar. When Ramp produces an output, Ramp2 may or may not produce a simultaneous output, depending on whether the output of the Bernoulli component is *true* or *false*. There is still a global clock, but when components react is not determined by this clock (or a subclock) but rather by the inferred clocks of the signals they interact with. At least one of these signals must be present for a reaction to occur. In the case the Ramp components, they react when their *output* is present, and when this occurs depends on how those outputs are used!

The situation gets worse. In Lustre, it is not actually required for the two inputs of the *When* to have the same clock. They can have different clocks, as long as the clock of the control input is a subclock of the clock of the other input. If the input is present and the control input is absent, then the output is absent. If there are no other constraints on those clocks, then Lustre assumes they have the same clock.

To deal with these subtleties, researchers have developed a calculus for analyzing the clocks associated with signals (Caspi, 1992; Colaço and Pouzet, 2003; Le Guernic et al., 1991). Here, we give one version of such a calculus.

First, we define a **master clock** which is the union of all the clocks in the system. This is really the same as the *global clock* considered previously, but it does not drive reactions. Let $n \in \mathbb{N}$ be the index of a tick of the master clock. Associated with each signal $s_i(n)$ is a clock $c_i(n)$ that indicates whether this signal is present at tick n . In addition, the clock $c_i(n)$ encodes the value of boolean signals, like s_2 in Figure 6.21, which controls the *When* component and therefore affects clocks.

Formally, a signal s_i is a function $s_i: \mathbb{N} \rightarrow T_i \cup \{\epsilon\}$, where T_i is the data type for the signal. The clock $c_i(n)$ is a function $c_i: \mathbb{N} \rightarrow \{-1, 0, 1\}$ defined as follows when $T_i = \{true, false\}$:

$$c_i(n) = \begin{cases} 0 & \text{if } s_i(n) = \epsilon, \\ 1 & \text{if } s_i(n) = true, \\ -1 & \text{if } s_i(n) = false. \end{cases} \quad (6.1)$$

If T_i is some other data type, $c_i(n)$ is either 1 or -1 when the signal is present. It will not matter which one it is.

The set of clock values $C = \{-1, 0, 1\}$ is treated as a **Galois field** with the commutative and associative operation of addition defined as follows:

$$\begin{aligned} 0 + x &= x, \\ 1 + 1 &= -1, \\ -1 + -1 &= 1, \\ -1 + 1 &= 0. \end{aligned} \quad (6.2)$$

Multiplication is also commutative and associative and is defined as follows:

$$\begin{aligned} 0 \cdot x &= 0, \\ 1 \cdot x &= x, \\ -1 \cdot x &= -x. \end{aligned} \tag{6.3}$$

The additive identity is 0, and the multiplicative identity is 1.

Using this algebraic structure, we can define a **clock calculus** for the system. To start, we can express constraints that signals associated with a component must either all be present or all be absent at any tick.

Example 6.19: Consider the Add component in Figure 6.21. The Add component imposes a constraint on the clocks of its two incoming signals. Specifically, it requires that they both be present whenever it reacts. This can be expressed as the following constraint:

$$c_3^2(n) = c_4^2(n) = c_5^2(n). \tag{6.4}$$

Similarly, the When component may impose a constraint on the clocks of its two incoming signals as follows:

$$c_1^2(n) = c_2^2(n). \tag{6.5}$$

More interestingly, the clock of signal s_3 depends on the boolean signal s_2 . If s_2 is *true*, then s_3 must be present. If s_2 is *false*, then s_3 must be absent. This can be expressed as the following constraint:

$$c_3(n) = c_1(n)(-c_2(n) - c_2^2(n)). \tag{6.6}$$

This states that if s_1 is absent, then so is s_3 . If s_1 is present, then s_3 is present if and only if s_2 is present and *true*. These three equations relate the clocks of all five signals in the system.

If we wish to allow a more general **When** component, one that permits the clock of the control signal to be a subclock of the clock of the other input, then we can replace (6.5) with the following constraint:

$$c_1^2(n) = c_2^2(n) + c_1^2(n)(1 - c_2^2(n)). \tag{6.7}$$




Figure 6.23: Inconsistent program revealed by the clock calculus.

This states that if s_1 is absent, then so is s_2 , but it allows s_1 to be present when s_2 is absent. In this latter case, (6.6) already ensures that s_3 will be absent when s_2 is absent.

Because these programs still have a master clock, we can use the same Kleene fixed-point iteration procedure to find the least fixed point of the system at each tick of the master clock. But the clock calculus enables sophisticated compiler optimizations that can significantly improve the performance of the system. More interestingly, the clock calculus can reveal inconsistencies in a design that are not apparent in the global clock model.

Example 6.20: Consider the program in Figure 6.23. Under the global clock model, there is no difficulty with this program, but it might be an error if the intent of the Add component is to always add two numbers. If we constrain the Add component to always add two numbers, then the clock calculus will reveal an inconsistency. Specifically, the constraint of the Add and When components are,

$$c_1^2(n) = c_3^2(n) = c_4^2(n) \quad (6.8)$$

and

$$c_3(n) = c_1(n)(-c_2(n) - c_2^2(n)). \quad (6.9)$$

Combining these, we get,

$$c_1(n) = c_1(n)(-c_2(n) - c_2^2(n)). \quad (6.10)$$

This states that s_1 can be present only when s_2 is present and *true*. Constraining s_2 to never be *false* may not be possible in a program, so this result reveals an inconsistency.

Clock calculus constraints are similar in spirit to the [balance equations](#) of the dataflow model studied in Chapter 5. A significant difference, however, is that inconsistencies in a dataflow model reveal either deadlocks or unbounded memory. Synchronous-reactive programs, by construction, have bounded memory, so an inconsistency in a clock calculus model reveals a likely bug in the program.

Notice that the *When* constraint of (6.6) preserves the boolean value of s_1 by passing it through to s_3 , if this signal happens to have a boolean data type. If, for example, s_3 happens to feed the control input of another *When* component, then the boolean value of s_1 will be preserved through to that component.

The clock calculus can similarly handle boolean operators like AND, OR, and NOT.

Example 6.21: A logical NOT (an **inverter**) with input $s_1(n)$ and output $s_2(n)$ has the following simple clock calculus constraints:

$$c_2(n) = -c_1(n). \quad (6.11)$$

This states that s_2 is present if and only if s_1 is present, and that if s_1 is *true*, then s_2 is *false* and vice versa.

Example 6.22: The [parallel OR](#) of Example 6.1 is a bit more complicated. Assume the inputs are $s_1(n)$ and $s_2(n)$, and the output is $s_3(n)$. The clock calculus constraints for this component are,

$$c_3 = -c_1 - c_2 - c_1^2 - c_2^2 + c_1c_2 + c_1^2c_2^2, \quad (6.12)$$

where we have omitted the tick index n for brevity. This captures the logic of an OR gate exactly as given in the truth table in Example 6.1, ignoring unknown ($\perp$) inputs, which are relevant to a fixed-point iteration but not to the clock calculus.

If both inputs are absent, $c_1 = c_2 = 0$, and all terms are 0, so $c_3 = 0$. If one input is absent and the other is *true*, say, without loss of generality, $c_1 = 1, c_2 = 0$, then

$$c_3 = -1 - 0 - 1 - 0 + 0 + 0 = 1,$$

as desired. If one input is absent and the other is *false*, say $c_1 = -1, c_2 = 0$, then

$$c_3 = 1 - 0 - 1 - 0 + 0 + 0 = 0,$$

so the output is absent. If both inputs are *true*, we get

$$c_3 = -1 - 1 - 1 - 1 + 1 + 1 = 1.$$

If both inputs are *false*,

$$c_3 = 1 + 1 - 1 - 1 + 1 + 1 = -1.$$

Finally, if one input is *false* and the other *true*, say $c_1 = 1, c_2 = -1$, then

$$c_3 = -1 + 1 - 1 - 1 - 1 + 1 = 1.$$

If we wish to capture the alternative semantics where the output is absent if either input is absent, then we can use the following constraint:

$$c_3 = (c_1 c_2)^2 ((c_1 - 1)(c_2 - 1) + 1). \quad (6.13)$$

The first term, $(c_1 c_2)^2$, ensures that the output is absent if either input is absent. The second term, $(c_1 - 1)(c_2 - 1) + 1$, yields 1 if either input is *true* and -1 if both inputs are *false*.

Note that neither of these constraints contradict the *parallel* nature of the parallel OR, which enables it to produce an output when an input is *unknown*.

Example 6.23: The *parallel AND* of Example 6.2 is similarly complicated. Again assume the inputs are $s_1(n)$ and $s_2(n)$, and the output is $s_3(n)$. To get




Figure 6.24: An inconsistent program using Switch and Select.

the truth table given in the example, the following constraint is used:

$$c_3 = c_1^2 + c_2^2 - c_1 - c_2 - c_1c_2 - c_1^2c_2^2. \quad (6.14)$$

If we wish to capture the alternative semantics where the output is absent if either input is absent, then we can use the following constraint:

$$c_3 = (c_1c_2)^2(-(c_1 + 1)(c_2 + 1) - 1). \quad (6.15)$$

The [Switch](#) and [Select](#) components that were used in Chapters 4 and 5 have analogies in synchronous-reactive languages. Under the global clock model, their behavior is straightforward. Under the clock calculus model, their behavior is more subtle. For example, the [Select](#) component requires that the *true* input be present when the component reacts with a *true* control input. This can lead to some interesting clock calculus constraints.

Example 6.24: Consider the program in Figure 6.24. The clock calculus constraints for the Switch component are,

$$\begin{aligned} c_1^2 &= c_2^2, \\ c_3 &= -c_2(c_2 + 1)c_1, \\ c_4 &= c_2(1 - c_2)c_1. \end{aligned} \quad (6.16)$$

The first of these asserts that the control input (the boolean signal s_2) must be present when the input s_1 is present and vice versa. The second asserts that if the control signal (s_2) is *true*, then $s_3 = s_1$, and if s_2 is *false*, then s_3 is




Figure 6.25: A non-constructive and inconsistent program using Select.

absent. The third asserts that if (s_2) is *false*, then $s_4 = s_1$, and if s_2 is *true*, then s_4 is absent.

The clock calculus constraints for the Select component are a bit more complicated. Using the signal names from Figure 6.24, the constraints are:

$$\begin{aligned} c_6 &= c_5(c_4(1 - c_5) - c_3(1 + c_5)), \\ -(c_5 + 1)c_5 &= c_3^2, \\ -(c_5 - 1)c_5 &= c_4^2. \end{aligned} \tag{6.17}$$

The last equation, for example, asserts that if s_4 is present, then s_5 must be present and *false*.

Putting these together, and using the fact that, because of the NOT, $c_5 = -c_2$, we can derive that $c_6 = 0$. The output is always absent. Moreover, there is no value for c_4 and c_5 that satisfies these constraints, so the program is inconsistent.

Example 6.25: Consider the program in Figure 6.25. The clock calculus in this case reveals that $c_1 = 0$ and that either $c_3 = 0$ or $1 + c_3 = 0$. This means that s_1 is absent and s_3 is either absent or *false*. Note that if s_3 is *false*, this model is nonconstructive.

The clock calculus can get much more elaborate. For example, if there are delays, using for example the [NonStrictDelay](#) component, then the boolean values of

signals that pass through these delays can become state variables in a dynamical system. In addition, in the Signal language (Le Guernic et al., 1991), clocks can be nondeterministically interleaved, eliminating the notion of a master clock. These elaborations are beyond the scope of this book.

6.3 Conclusion

The synchronous-reactive (SR) model of computation provides a rigorous framework for deterministic concurrency, offering a powerful alternative to the nondeterminism often found in threaded software. By adopting the [synchrony hypothesis](#)—the abstraction that components react simultaneously and instantaneously—SR models enable the design of complex systems that remain predictable and amenable to formal verification.

Key takeaways from this chapter include:

- **The Global Clock vs. Clock Calculus:** We explore two primary SR styles: the simpler global clock model used in languages like Esterel, where every component conceptually fires at each tick, and the more nuanced clock calculus model used in Lustre and Signal, where clocks are properties of individual signals.
- **Fixed-Point Semantics:** Central to SR is the requirement that at every tick, the system must find a fixed point for all signal values. While feedback loops can introduce complexities, the use of monotonic functions and flat partial orders allows for the application of the Kleene fixed-point theorem to find a unique least fixed point.
- **Constructive vs. Non-constructive Models:** We distinguish between [constructive](#) models, which can be solved through iterative procedures, and non-constructive models that may be logically sound but can lead to physical instabilities like oscillation. Following the principles of intuitionistic logic, we prioritize constructive models to ensure reliability.
- **Practical Utility:** SR is uniquely suited for orchestrating concurrent actions, managing sporadic data where the absence of an event is as meaningful as its presence, and modeling safety-critical systems like digital circuits and flight control software.

While the SR model is highly effective for logically timed systems, it does not inherently account for the physical passage of time between ticks. In the next chapter, we will see how the discrete-event (DE) and reactor models generalize these concepts by making the time gap between events a core part of the semantics.

Discrete-Event Systems

7.1	Discrete-Event Signals and Components	256
7.1.1	Tags and Logical Time	256
7.1.2	Signals and Events	257
7.1.3	Discreteness and Zeno Signals	258
7.1.4	Discrete-Event Components	259
	<i>Sidebar: Discrete-Event Languages and Simulators</i>	260
7.2	Metric Space Semantics	262
7.2.1	Composition of Discrete-Event Components	262
7.2.2	Causality and Contracting Functions	263
	<i>Sidebar: Background on Metric Space Semantics</i>	263
7.2.3	Fixed-Point Semantics	267
7.2.4	Eventually Delta Contracting Systems	269
7.2.5	Zeno Behaviors	276
7.3	Generalized Ultrametric Semantics	278
7.3.1	A Generalized Distance Function	278
	<i>Sidebar: Difficulty of Determining Zeno Behaviors</i>	279
	<i>Sidebar: Lipschitz Continuous Functions in Metric Spaces</i>	280
7.3.2	Prefix Order and the Generalized Ultrametric Space	282
7.3.3	Completeness of the Generalized Ultrametric Space	285
7.3.4	Fixed Points in the Generalized Ultrametric Space	286
7.3.5	Constructive Procedure to Find the Fixed Point	287
7.4	Conclusion	293

In Chapters 4 and 5, the communication between components in a concurrent system was via *streams* of *tokens*, where a stream is a *sequence*, and a token carries any payload data value. Given two streams, there is no intrinsic relationship between the position of a value in one stream and the position of a value in another. That is, the n^{th} value in stream s_1 may or may not align with the n^{th} value in stream s_2 .

Chapter 6 changes this situation. In the *global clock* model, the communication between components may again be thought of as streams of values, but now there is a strong alignment property. The n^{th} value in stream s_1 is *logically simultaneous* with the n^{th} value in stream s_2 . Moreover, the notion of a value is augmented with a special value ϵ , representing the absence of a value. Hence, the streams become alternations of values (or tokens) and ϵ so that, across an entire *clock domain*, the n^{th} values in all streams are logically simultaneous.

In this chapter, we go further and model the communication between components as signals that are functions of a *logical time*. Logical time here will be all about the alignment of values across signals. Any two values in distinct signals are either logically simultaneous or definitively ordered so that one precedes the other. These order relations are *logical*, not physical. Only when there are causal relationships between signals will we find physical ordering constraints. Hence, the use of logical time does not inherently reduce the concurrency in a system or the ability to exploit that concurrency with parallel execution. These ideas benefit from some formalization, which is what we will do in this chapter.

7.1 Discrete-Event Signals and Components

We begin by defining the communication signals between components in a discrete-event system.

7.1.1 Tags and Logical Time

This chapter assumes a logical time model like that described in Chapter 2, Section 2.2.6, where logical time is represented with a totally ordered tag set G .

Example 7.1: Examples of tag sets are:

- **real-valued time:** $\mathbb{G} = \mathbb{R}$, with ordinary numeric order.
- **real-valued time with an origin:** $\mathbb{G} = \mathbb{R}_+$, the non-negative reals.
- **superdense time:** $\mathbb{G} = \mathbb{R}_+ \times \mathbb{N}$, with **lexicographic order**.
- **discrete superdense time:** $\mathbb{G} = \mathbb{N} \times \mathbb{N}$, with lexicographic order.

As in Section 2.2.6, a tag may be associated with a physical time value. If $\mathbb{T}$ represents the set of all possible physical time values, and if there is a nondecreasing function $\mathcal{T} : \mathbb{G} \rightarrow \mathbb{T}$ that gives a **physical time interpretation** to a tag, then we call $\mathcal{T}(g)$ the **timestamp** of the tag g .

In this chapter, we will further require the function $\mathcal{T}$ to be injective (one-to-one). I.e., for any two distinct tags $g \neq g'$, their timestamps are distinct, $\mathcal{T}(g) \neq \mathcal{T}(g')$. This constraint will be relaxed in Section 7.3, where timestamps no longer play any role in the fixed-point semantics.

7.1.2 Signals and Events

Definition 7.1. Let $\mathbb{G}$ be a totally ordered **tag** set. Let V represent a **data type** (a set of values). A **signal** is a function $s : \mathbb{G} \rightarrow V \cup \{\epsilon\}$, where ϵ represents **absent**, for no value. The **empty signal** λ is the signal that maps every tag to ϵ . The **set of signals** over $\mathbb{G}$ and V is denoted by $\mathcal{S}$, where

$$\mathcal{S} = (\mathbb{G} \rightarrow V \cup \{\epsilon\}).$$

Definition 7.2. An **event** is a pair $(g, v) \in \mathbb{G} \times V$, where $g \in \mathbb{G}$ is a tag, and $v \in V$ is a value. The **set of events** in a signal $s \in \mathcal{S}$ is

$$\mathcal{E}(s) = \{(g, v) \in \mathbb{G} \times V \mid v = s(g)\}.$$

A signal is fully defined by its event set. The **tag set** of a signal s is

$$\mathcal{G}(s) = \{g \in \mathbb{G} \mid s(g) \neq \epsilon\}.$$

Note that $\mathcal{E}(\lambda) = \mathcal{G}(\lambda) = \emptyset$. Note further that because a signal is a function, no two events in $\mathcal{E}(s)$ have the same tag.

7.1.3 Discreteness and Zeno Signals

In this chapter, we are interested in discrete-event signals.

Definition 7.3. A *discrete-event signal (DE signal)* is a signal s where the tag set $\mathcal{G}(s)$ is a *discrete set*.

A discrete-event signal is one whose events can be enumerated in order and has an earliest event.

Example 7.2: For the tag set $\mathbb{G} = \mathbb{R}_+$, signals with the following tag sets are discrete-event signals:

- $\mathcal{G}(s_1) = \{0, 1, 2, \dots\}$.
- $\mathcal{G}(s_2) = \emptyset$.
- $\mathcal{G}(s_3) = \{1 - 1/2^n \mid n \in \mathbb{N}, n > 0\} = \{0, 0.5, 0.75, \dots\}$.

Here, s_1 might be called a **clock signal**, and $s_2 = \lambda$ is the empty signal. The signal s_3 has an infinite number of events in a finite amount of time, but those events can nevertheless be enumerated in order.

The following signals are not discrete-event signals:

- $\mathcal{G}(s_4) = \mathbb{R}_+$
- $\mathcal{G}(s_5) = \{1 - 1/2^n \mid n \in \mathbb{N}, n > 0\} \cup \{1\}$.

Here, s_4 is a continuous-time signal, and s_5 is s_3 with an additional event at time 1. The tag set of s_5 is not discrete because there is no way to put the events in an enumerated list in order while including the event with tag 1. There are an infinite number of events prior to tag 1.

The signals s_3 and s_5 are particularly problematic. Signal s_5 has an infinite number of events that occur prior to at least one other event in the signal, and hence it is not a discrete-event signal. Signal s_3 is a discrete-event signal, but it is nevertheless problematic in computational systems because it has an infinite number of events in a finite amount of time, which is not possible in practice. We call such signals **Zeno signals**.

Definition 7.4. A **Zeno signal** is a signal s where the tag set $\mathcal{G}(s)$ is a countable set with an infinite number of events prior to some finite tag $g \in \mathbb{G}$.

Both s_3 and s_5 are Zeno signals, where s_3 is a discrete-event signal, and s_5 is not.

We can form a **tuple of signals** $s = (s_1, s_2, \dots, s_n) \in \mathcal{S}^n$, for some $n \in \mathbb{N}$. The tag set for the tuple is the union of the tag sets of the individual signals,

$$\mathcal{G}(s) = \mathcal{G}(s_1) \cup \mathcal{G}(s_2) \cup \dots \cup \mathcal{G}(s_n).$$

Definition 7.5. A **discrete-event tuple** is a tuple $s \in \mathcal{S}^n$, for some $n \in \mathbb{N}$, where the tag set $\mathcal{G}(s)$ is a **discrete set**. A **non-Zeno tuple** is a discrete-event tuple where the tag set $\mathcal{G}(s)$ has a finite number of events prior to any finite tag $g \in \mathbb{G}$.

The above definition deliberately allows n to be zero, representing a tuple with no signals. By convention, we define

$$\mathcal{S}^0 = \{\lambda\}.$$

This is a singleton set containing the empty signal. All members of this set are non-Zeno tuples.

Note that even if the individual signals in a tuple are discrete-event signals, the tuple may not be a discrete-event tuple.

Example 7.3: Let $s = (s_1, s_3)$, where s_1 and s_3 are the signals defined in Example 7.2. Even though s_1 and s_3 are discrete-event signals, the tuple s is not a discrete-event tuple because the tag set $\mathcal{G}(s)$ is not a discrete set.

However, if all elements of the tuple are non-Zeno signals, then the tuple is a non-Zeno tuple.

7.1.4 Discrete-Event Components

A discrete-event system is a composition of components that operate on discrete-event signals.

Sidebar: Discrete-Event Languages and Simulators

The **discrete-event** model of computation is one of the oldest and most widely used approaches to modeling concurrent and reactive systems. Rather than aligning values on a global clock, as in the [synchronous](#) model of Chapter 6, discrete-event systems advance [logical time](#) in jumps from one event to the next. Events may occur at irregular intervals, and components react when events arrive on their inputs or when a scheduled internal event becomes due. The model is well suited to systems where activity is sporadic or periodic or a mix.

Discrete-event ideas appear across many application domains. In **hardware description languages** such as VHDL (IEEE Std 1076–2000) and SystemC (IEEE Std 1666–2011), processes suspend until an event occurs and then resume, often using **delta cycles** or other microstep mechanisms to order reactions that are logically simultaneous. Graphical modeling environments such as Simulink (from The MathWorks) use block diagrams in which blocks execute in response to timed events. Network simulators such as ns-2, ns-3 ([NS-3 Consortium, 2015](#)), and OPNET organize protocol models as collections of modules that communicate through timestamped events. General-purpose formalisms such as DEVS (Discrete Event System Specification) ([Zeigler, 1976](#); [Zeigler et al., 2000](#)) provide a framework for hierarchical, modular discrete-event models, and have influenced many specialized tools. SysML ([Object Management Group \(OMG\), 2010](#)) extends UML with notations for modeling discrete events in systems engineering workflows. SimPy ([Matloff, 2008](#)) is a lightweight Python library that implements a classic discrete-event simulation engine with a priority queue of future events.

These languages and tools differ considerably in their operational semantics. Some treat logically simultaneous events as a single simultaneous reaction; others order them using delta time or other microsteps; still others admit non-deterministic orderings. Some require every component to introduce delay; others allow zero-delay reactions. These differences matter when composing models or when seeking deterministic behavior. The formal model developed in this chapter abstracts from any particular language, but it is informed by experience with Ptolemy II ([Lee et al., 2014a](#)) and Lingua Franca ([Lohstroh et al., 2021](#)), where discrete-event models are used as a coordination language rather than as a programming language.

Definition 7.6. For some set of signals $\mathcal{S}$, a **discrete-event component** is a function

$$F_i: \mathcal{S}^m \rightarrow \mathcal{S}^n, \quad m, n \in \mathbb{N},$$

where, if $s \in \mathcal{S}^m$ is a **discrete-event tuple**, then $F_i(s)$ is a discrete-event tuple.

Definition 7.7. For some set of signals $\mathcal{S}$, a **non-Zeno component** is a function

$$F_i: \mathcal{S}^m \rightarrow \mathcal{S}^n, \quad m, n \in \mathbb{N},$$

where, if $s \in \mathcal{S}^m$ is a **non-Zeno tuple**, then $F_i(s)$ is a non-Zeno tuple.

The above definitions deliberately allow m or n to be zero, representing a component with no inputs or no outputs, respectively. A function of the form

$$F: \mathcal{S}^0 \rightarrow \mathcal{S}^n$$

maps a singleton set to the set of n -tuples of signals. Because the domain is a singleton set, the function is constant. Similarly, a function of the form

$$F: \mathcal{S}^m \rightarrow \mathcal{S}^0$$

maps the set of m -tuples of signals to a singleton set. Because the codomain is a singleton set, the function is again constant.

It can be quite subtle to determine whether a function is a discrete-event component, as illustrated by the following example.

Example 7.4: Consider the tag set $\mathbb{G} = \mathbb{R}_+$ and value set $V = \{0\}$, a singleton set, where for the n^{th} event $(t_n, 0)$ of a discrete-event input signal s , there is an output event $(t_n + \tau_n, 0)$ (i.e., delayed by τ_n time units), where $0 < \tau_n < 1$. Specifically, let

$$\tau_n = 1/(n+1), \quad n \in \mathbb{N}.$$

Then the function F is not a **discrete-event component**. Suppose the input signal s has tag set

$$\mathcal{G}(s) = \{1 - 1/(n+1)^2\} = \{0, 0.75, 0.89, 0.94, \dots\}.$$

This event set is discrete, so this is a **discrete-event signal**. But the output will have tag set

$$\mathcal{G}(F(s)) = \{1 - 1/(n+1)^2 + 1/(n+1)\} = \{1, 1.25, 1.22, 1.19, \dots\},$$

which is not discrete. The output has a first event at time 1, but no second event.

Note that it might seem that this component should be discrete because for each input event, there will be at most one output event. But there is no assurance of being able to construct an **order embedding** $f: \mathcal{G}(F(s)) \rightarrow \mathbb{N}$ for the output, as the above example demonstrates.

7.2 Metric Space Semantics

We now consider the composition of discrete-event components to form a **discrete-event system** and develop a formal semantics based on metric spaces.

7.2.1 Composition of Discrete-Event Components

Consider the composition of discrete-event components $F_1, F_2, \dots, F_k$ to form a discrete-event system by connecting output signals to input signals. Exactly as done for synchronous systems in Chapter 6, Figure 6.4, and for process networks in Chapter 4, Figures 4.2 and 4.3, the determination of the overall behavior of the system reduces to finding a fixed point of a function.

Suppose each component F_i has m_i inputs and n_i outputs, for $i = 1, 2, \dots, k$. The composition forms a **closed system** if every input signal of a component is connected to an output signal of a component in the system. Note that even if we wish to model a system that accepts input from the environment, we can still model it as a closed system by adding a component with zero inputs where the outputs are the signals from the environment. Let

$$N = \sum_{i=1}^k n_i$$

be the total number of outputs of the components, and let $s = (s_1, s_2, \dots, s_N)$ be the tuple of signals that are the outputs of the components. Then, following the pattern of Figures 4.2 and 6.4, there is a function $F: \mathcal{S}^N \rightarrow \mathcal{S}^N$ that is the parallel composition of the functions $F_1, F_2, \dots, F_k$. The behavior of the system is given by the fixed point of F ,

$$F(s) = s, \quad (7.1)$$

where s is the tuple of signals that are the outputs of the components. Using the mathematics of metric spaces, we can find a fixed point of this function F . If that fixed point exists, it is the behavior of the system. If the fixed point is unique, then the system is deterministic.

7.2.2 Causality and Contracting Functions

Suppose that for the tag set $\mathbb{G}$ of a discrete-event signal or tuple of signals, there is an order-preserving one-to-one mapping $\mathcal{T}: \mathbb{G} \rightarrow \mathbb{R}_+$ from tags to non-negative real numbers representing time. We call $\mathcal{T}(g)$ the **timestamp** corresponding to the tag g . In such a model, we can use the **Cantor metric** (see Appendix B) to define

Sidebar: Background on Metric Space Semantics

Early work using metric spaces in the context of concurrent systems is due to [Reed and Roscoe \(1988a\)](#), who applied the concept to a timed version of communicating sequential processes (CSP) ([Hoare, 1978](#)). This work was extended by [Yates \(1993\)](#), who related the metric space approach to **Kahn process networks** ([Kahn, 1974](#)). Their work was later extended by [Lee \(1999\)](#), who used it to define a fixed-point semantics for discrete-event systems. [Naundorf \(2000\)](#) showed that **strictly contracting** functions have a unique fixed point, thereby extending the results beyond **delta contracting** functions. However, he did not provide a constructive procedure to find the fixed point. [Cataldo et al. \(2006\)](#) extended the results to use generalized ultrametric spaces from [Priess-Crampe and Ribenboim \(1996\)](#) to handle what they call **eventually delta contracting** functions. [Liu et al. \(2006a\)](#) and [Matsikoudis and Lee \(2013, 2015\)](#) further extend this work to encompass continuous-time systems and Zeno systems. Most of these results are reviewed in this chapter.

a metric space on the set of discrete-event signals. Then we can apply the [Banach fixed-point theorem](#) to find a fixed point of a function, at least for some systems.

We can define a metric space on the set of discrete-event signals $\mathcal{S} = (\mathbb{G} \rightarrow V \cup \{\epsilon\})$, where the Cantor metric is given by

$$d(s_1, s_2) = 2^{-\tau},$$

where τ is the timestamp of the earliest event in s_1 not in s_2 or vice versa. In other words, τ is the time at which the two signals first differ. This is the Cantor metric of Section B.2.1, but with a somewhat simpler definition because of the discrete nature of the tag set of the signals. This metric extends easily to tuples of signals and gives us a useful characterization of components in a discrete-event system.

Definition 7.8. A *causal component* is a [discrete-event component](#)

$$F_i: \mathcal{S}^m \rightarrow \mathcal{S}^n$$

such that for all discrete-event tuples $s_1, s_2 \in \mathcal{S}^m$,

$$d(F_i(s_1), F_i(s_2)) \leq d(s_1, s_2),$$

where d is the Cantor metric.

The component is **strictly causal** if for all discrete-event tuples $s_1, s_2 \in \mathcal{S}^m$ where $s_1 \neq s_2$,

$$d(F_i(s_1), F_i(s_2)) < d(s_1, s_2).$$

The component is **delta causal** if there exists a non-negative real number $\delta < 1$ such that for all discrete-event tuples $s_1, s_2 \in \mathcal{S}^m$,

$$d(F_i(s_1), F_i(s_2)) \leq \delta d(s_1, s_2).$$

There is a powerful intuition behind these definitions. A causal component is one whose output at any time cannot depend on future inputs. Consider two possible m -tuple inputs s and s' to a component F_i . The statement

$$d(F_i(s), F_i(s')) \leq d(s, s')$$

declares that the time at which the outputs $F_i(s)$ and $F_i(s')$ for these two inputs differ cannot be earlier than the time at which the inputs s and s' differ. This is the essence of causality. An output cannot depend on the future.

A strictly causal component is one whose output at any time cannot depend on the current or future inputs. A delta causal component is one whose output at any time can only depend on inputs that are at least a non-zero amount of time old.

Example 7.5: The function F given in Example 7.4 is **strictly causal** under the **Cantor metric** because every output occurs strictly later than the input that causes it.

Note that these definitions work also for components with either no inputs or no outputs. If $m = 0$ (there are no inputs), then the domain of the function F_i is a **singleton set**, so all inputs are identical. Hence,

$$d(F_i(s), F_i(s')) = d(s, s') = 0.$$

Such a function is causal, strictly causal, and delta causal with $\delta = 0$.

If $n = 0$ (there are no outputs), then the range of the function F is a singleton set, so all outputs are identical. Hence,

$$d(F(s_1), F(s_2)) = 0 \leq d(s_1, s_2).$$

Such a function is again causal, strictly causal, and delta causal with $\delta = 0$.

Notice that causality here is defined only for functions on tuples of **discrete-event signals**, and only for those with non-negative real-valued timestamps that are distinct for distinct tags. For more general signals, such as continuous-time signals or signals with no earliest event, the situation is more subtle, and it is possible to define components that are strictly causal, for example, but do not meet the above definition. Consider the following example, taken from **Matsikoudis and Lee (2015)**.

Example 7.6: Consider the tag set $G = \mathbb{Z}$, value set $V = \mathbb{Z}$, and the function $F: \mathcal{S} \rightarrow \mathcal{S}$ given by

$$F(s)(g) = \begin{cases} s(g-1) + 1 & \text{if } s(g-1) \neq \epsilon \\ 0 & \text{otherwise} \end{cases}$$

This function is intuitively causal, strictly causal, and delta causal because there is a delay of one time unit from any input event to the corresponding output event. But it does not satisfy the above definition for strictly causal.

First, it is not clear how to define timestamps for this tag set, since the timestamps must be given by an order-preserving mapping from tags to non-negative real numbers. We can try using the identity function $\mathcal{T}(g) = g$ and mechanically applying the Cantor metric, but this leads to difficulties.

Consider two possible inputs s_1 and s_2 where $s_1(g) = 1$ and $s_2(g) = 2$ for all $g \in \mathbb{G}$. These are not discrete-event signals because they do not have an earliest event. Attempting to apply the Cantor metric to these signals can only yield $d(s_1, s_2) = \infty$, which is not a valid metric. These signals differ over all tags, so the earliest time at which the two signals differ must be $\tau = -\infty$, and $2^{-\tau} = 2^\infty = \infty$. Similarly, $d(F(s_1), F(s_2)) = \infty$ for the same reason. Hence, it is not true that $d(F(s_1), F(s_2)) < d(s_1, s_2)$ for all $s_1, s_2 \in \mathcal{S}$, so this function does not satisfy the definition of strictly causal.

Another problematic example is given by [Liu et al. \(2006a\)](#).

Example 7.7: Consider the tag set $\mathbb{G} = \mathbb{R}_+$, value set $V = \mathbb{R}$, and the function $F: \mathcal{S} \rightarrow \mathcal{S}$ given by

$$F(s)(g) = \begin{cases} \lim_{h \rightarrow g^+} s(h) & \text{if the right limit exists} \\ 0 & \text{otherwise} \end{cases}$$

where $\lim_{h \rightarrow g^+} s(h)$ is the limit of $s(h)$ as h approaches g from the right. Assume the timestamps are given by $\mathcal{T}(g) = g$.

This function satisfies the property that $d(F(s), F(s')) \leq d(s, s')$ for all $s, s' \in \mathcal{S}$ under the Cantor metric, but it is not intuitively causal. To see this, consider two possible inputs s and s' where for all $g \in \mathbb{G}$, $s(g) = 0$ and

$$s'(g) = \begin{cases} 0 & \text{if } g \leq 1 \\ 1 & \text{if } g > 1 \end{cases}$$

By the [Cantor metric](#), $d(s, s') = 2^{-\tau} = 1/2$, where τ is the greatest lower bound of the set of all g such that $s(g) \neq s'(g)$. Then $F(s)(g) = 0$ for all $g \in \mathbb{G}$ and $F(s')$ is given by

$$F(s')(g) = \begin{cases} 0 & \text{if } g < 1 \\ 1 & \text{if } g \geq 1 \end{cases}$$

By the Cantor metric, $d(F(s), F(s')) = 2^{-\tau} = 1/2$, the same as the distance between the inputs. Nevertheless, while the inputs are identical for all g in the interval $[0, 1]$, the outputs differ at $g = 1$. Hence, the output depends on future inputs, so this function is not intuitively causal, despite seeming to satisfy the definition of causality. However, these signals are not discrete-event signals because their tag sets are not [discrete sets](#).

In Definition 7.8, causality is defined for discrete-event components, which map m -tuples of discrete-event signals to n -tuples of discrete-event signals, for some $m, n \in \mathbb{N}$. It bears a strong resemblance to the definition of [contracting functions](#) (see Section B.2.3 in the metric spaces appendix), but those are defined only on [endofunctions](#), where the domain and codomain are the same set. When analyzing compositions of components, however, we ultimately want to consider an endofunction, as explained above in Section 7.2.1.

A [closed system](#) F composed of causal components $F_1, F_2, \dots, F_k$ is itself causal. Similarly, a closed system composed of strictly causal components is strictly causal, and a closed system composed of delta causal components is delta causal. A closed system F is an [endofunction](#) on $\mathcal{S}^N$,

$$F: \mathcal{S}^N \rightarrow \mathcal{S}^N. \quad (7.2)$$

For such an endofunction, causality corresponds to the definition of contracting functions in Section B.2.3 of the metric spaces appendix. A causal system is a [contraction](#). A strictly causal system is a [strictly contracting](#) function. A delta causal system is a [delta contracting](#) function.

7.2.3 Fixed-Point Semantics

Given a closed system F modeled as an endofunction (7.2), if it is delta causal, we can find a unique fixed point of F by applying the Banach fixed-point theorem

(Theorem B.2 in the metric space appendix). This fixed point is the behavior of the system. However, we must be careful, because the Banach fixed-point theorem only applies to delta causal functions on **complete metric spaces**.

Proposition 7.1. *The set of **discrete-event signals** (or tuples) with the Cantor metric forms a complete metric space.*

Proof. Let $\{s_1, s_2, \dots\}$ be a **Cauchy** sequence of discrete-event signals in $\mathcal{S}^N$. This means that for all $\epsilon > 0$, there exists an integer k such that for all $m, n \geq k$,

$$d(s_m, s_n) < \epsilon,$$

where d is the Cantor metric. We must show that this sequence converges to a signal $s \in \mathcal{S}^N$. This means that for all $\epsilon' > 0$, there exists an integer k' such that for all $m \geq k'$,

$$d(s_m, s) < \epsilon'.$$

Choose some $\epsilon' > 0$ and let

$$\tau = \log_2(1/\epsilon').$$

Note that by choosing a small enough ϵ' , we can make τ as large as we want. Because the sequence is Cauchy, we can choose $\epsilon = \epsilon'$, and then there is an integer k such that all signals $s_m, s_n, m, n \geq k$, are identical up to but not necessarily including time τ .

Let s be a signal where, for any choice of $\epsilon' > 0$, and hence any arbitrarily large τ , s is identical to s_k up to but not necessarily including time τ . Then, for all $m \geq k$,

$$d(s_m, s) < \epsilon'.$$

Hence, the sequence s_m converges to s .

It remains to show that the limit s is a discrete-event signal. Suppose, toward a contradiction, that s is not a discrete-event signal. Then there exists some tag $g \in \mathcal{G}(s)$ that has infinitely many predecessors. Since $\mathcal{T}: \mathbb{G} \rightarrow \mathbb{R}_+$ is order-preserving, the predecessors of g all have timestamps less than $\mathcal{T}(g)$. Choose any $\tau > \mathcal{T}(g)$. By construction of s , there is an s_k that agrees with s up to but not including time τ . But such an s_k would not be a discrete-event signal, a contradiction. Hence, s must be a discrete-event signal. $\square$

The following theorem now follows immediately from the [Banach fixed-point theorem](#) (Theorem B.2).

Theorem 7.1. *A delta-causal discrete-event system F has a unique fixed point that is the limit of the sequence,*

$$\begin{aligned} s_1 &= F(s_0) \\ s_2 &= F(s_1) \\ &\vdots \\ s_{n+1} &= F(s_n), \quad n \geq 1 \\ &\vdots \end{aligned}$$

for any starting point $s_0 \in \mathcal{S}^N$.

7.2.4 Eventually Delta Contracting Systems

To use the Banach fixed-point theorem, we need to constrain the system function F to be [delta contracting](#). This turns out to be a rather severe constraint, one that is not automatically met by many systems. However, most practical systems yield to a simple transformation that gives us a delta contracting system.

First, note that every [delta causal](#) component introduces a time delay from input events to output events that are affected by those input events. Suppose $F_D: \mathcal{S} \rightarrow \mathcal{S}$ is a delta causal component with some positive $\delta < 1$. Suppose that g is the first tag at which two discrete-event signals s_1 and s_2 differ. Then, because F_D is delta causal,

$$d(F_D(s_1), F_D(s_2)) \leq \delta d(s_1, s_2) = \delta 2^{-\mathcal{T}(g)}.$$

Let $D = -\log_2(\delta)$, so $\delta = 2^{-D}$. Note that D is positive because $\delta < 1$. Then,

$$d(F_D(s_1), F_D(s_2)) \leq 2^{-D} 2^{-\mathcal{T}(g)} = 2^{-(\mathcal{T}(g)+D)}.$$

Hence, the earliest tag at which $F_D(s_1)$ and $F_D(s_2)$ differ is at least D time units after the earliest tag at which s_1 and s_2 differ. Consequently, F_D is a **delay component** with delay at least D . Such delay components immediately yield to the Banach fixed-point theorem for finding the fixed point of the system. Consider the following example.




Figure 7.1: A unit delay feedback system.

Example 7.8: Consider the rather trivial feedback system shown in Figure 7.1, where F_D is a **unit delay** component with delay $D = 1$. That is, for each input event with tag g and value v , the output event has the same value v and tag g' such that $\mathcal{T}(g') = \mathcal{T}(g) + 1$. The unique fixed point of this system is the empty signal, λ . To see this, let s_0 be any signal whose earliest event has tag g_0 . Then consider the sequence,

$$\begin{aligned} s_1 &= F_D(s_0) \\ s_2 &= F_D(s_1) \\ &\vdots \\ s_n &= F_D(s_{n-1}) \\ &\vdots \end{aligned}$$

By the **Banach fixed-point theorem** (Theorem B.2), this sequence converges to the unique fixed point of F_D . The earliest event in s_n has tag g_n such that $\mathcal{T}(g_n) \geq \mathcal{T}(g_0) + nD$. Hence,

$$d(s_n, \lambda) = 2^{-\mathcal{T}(g_n)} \leq 2^{-\mathcal{T}(g_0) - nD} = 2^{-\mathcal{T}(g_0)} 2^{-nD}.$$

This distance converges to zero as $n \rightarrow \infty$, so the sequence converges to the empty signal.

If you observe that real-world components take time to process inputs and produce outputs, then every real-world component is a delay component. Hence, it might seem sensible to assume a lower bound on the delay from input to output. This assumption was first used by [Zeigler \(1976\)](#), then by [Yates \(1993\)](#) and [Yates and Gao \(1993\)](#), then by [Müller and Scholz \(1997\)](#), and later even in my own work ([Lee, 1999](#); [Lee and Sangiovanni-Vincentelli, 1998](#); [Liu and Lee, 2003](#)). The




Figure 7.2: A composition that does not immediately lead to a delta causal model.

same idea had also been used in the context of timed [communicating sequential processes](#) by [Reed and Roscoe \(1988b\)](#).

However, models and software systems are often designed without any such assumption. The [synchronous](#) semantics of Chapter 6 is an example, where components react logically instantaneously. And many discrete-event modeling languages, such as Ptolemy II ([Lee et al., 2014a](#)), also include components that react instantaneously. There is no difficulty in realizing these, so it is unfortunate if the formal semantics cannot handle them. Consider the following simple example.

Example 7.9: Consider the composition shown in Figure 7.2(a), where F_D is the unit delay component of Example 7.8 and F_1 is defined as follows:

$$F_1(s)(g) = \begin{cases} 1 & \text{if } g = \perp_G \\ s(g) & \text{otherwise} \end{cases}$$

This component outputs an initial value of $v = 1$, and then subsequently forwards the input value to the output. The component is causal, but neither strictly nor delta causal. As a consequence, when we arrange the system as shown in Figure 7.2(b), which follows the pattern of Figures 4.2 and 6.4, the composition function F is not delta causal. Therefore, we cannot immediately apply the Banach fixed-point theorem to find the fixed point of the system.

Suppose, however, that we cascade two instances of the function F , as shown in Figure 7.2(c). This still represents the same system, but now the composition function F^2 is delta causal. We can now apply the [Banach fixed-point theorem](#) (Theorem B.2) to find the fixed point of the system.

For simplicity, let $s_0 = \lambda$ be the empty signal (any signal will do, but the empty signal is simplest). Also, for simplicity, let $\mathbb{G} = \mathbb{R}_+$, and let $\mathcal{T}$ be the identity function. Then the sequence of iterates is:

$$s_1 = F^2(s_0) \text{ where } \mathcal{E}(s_1) = \{(0, 1), (1, 1)\}$$

$$s_2 = F^2(s_1) \text{ where } \mathcal{E}(s_2) = \{(0, 1), (1, 1), (2, 1), (3, 1)\}$$

$$\vdots$$

$$s_n = F^2(s_{n-1}) \text{ where } \mathcal{E}(s_n) = \{(0, 1), (1, 1), (2, 1), (3, 1), \dots, (2n-2, 1), (2n-1, 1)\}$$

$$\vdots$$

The limit of this sequence is the signal s with event set

$$\mathcal{E}(s) = \{(0, 1), (1, 1), (2, 1), (3, 1), \dots\}.$$

This signal is the unique fixed point of F^2 . It is the [clock signal](#) of Example 7.2, which has an event every one time unit.

This approach can be generalized.

Definition 7.9. ([Cataldo et al., 2006](#)) A discrete-event system endofunction F is *eventually delta contracting* if there exists an $E \in \mathbb{N}$, $E > 0$, such that the composition F^E is delta contracting.




Figure 7.3: Another composition that does not immediately lead to a delta causal model.

Fortunately, there is a very practical rule that ensures that a system is eventually delta contracting. Informally, we need for every cycle formed by a composition of causal components $F_1, F_2, \dots, F_n$ to include at least one delta causal function. Applying this rule, however, requires a bit more work, as illustrated by the following example.

Example 7.10: Consider the composition in Figure 7.3(a), where $F_1: \mathcal{S} \rightarrow \mathcal{S}^2$ is causal but not delta causal. As before, the simple feedback construction in Figure 7.3(b) does not immediately yield to application of the Banach fixed point theorem because F is not a delta contraction.

Suppose, however, that the path from the input to the *first* output of F_1 is delta causal. That is, if the function F_1 is decomposed into two functions

of the form $F_{1,i}: \mathcal{S} \rightarrow \mathcal{S}$ for $i \in \{1,2\}$, as shown in Figure 7.3(c), then $F_{1,1}$ is delta causal, while $F_{1,2}$ is only causal. This model is identical to that of Figure 7.3(b), and it still does not provide a delta contraction.

However, if we now form F^2 , as shown in Figure 7.3(d), it becomes immediately evident that F^2 is a delta contraction. Every path from one of its inputs to one of its outputs passes through a delta causal component.

Any function of the form $F_j: \mathcal{S}^K \rightarrow \mathcal{S}^M$ can be similarly decomposed into M distinct functions $F_{j,i}: \mathcal{S}^K \rightarrow \mathcal{S}$, $i \in \{1, \dots, M\}$, to construct an identical model. So now, the only requirement for constructing a delta contracting system is that every cycle in the decomposed model contain at least one delta causal component.

Proposition 7.2. Consider a *closed system* composed of causal components $F_1, F_2, \dots, F_n$, and let G be the dependency graph constructed using the decomposed functions $F_{j,i}$, $j \in \{1, \dots, n\}$, $i \in \{1, \dots, M_j\}$, where M_j is the number of outputs of F_j . In the graph G , we have one node per subcomponent and an edge from $F_{j,i}$ to $F_{j',i'}$ whenever the output of $F_{j,i}$ is connected to an input of $F_{j',i'}$. If every directed cycle in G contains at least one *delta causal* component, then the system is *eventually delta contracting*. Moreover, F^E is delta contracting for $E = p + 1$, where p is the number of components on the longest path in G that contains no delta causal component.

Proof. Passing a signal difference through a component adds a delay to the earliest time at which two signals differ: a *causal* component adds a delay of at least 0, and a *delta causal* component $F_{j,i}$ with constant $\delta_{j,i}$ adds a delay of at least $D_{j,i} = -\log_2 \delta_{j,i} > 0$ (with $D_{j,i} = \infty$ when $\delta_{j,i} = 0$). Let $D = \min\{D_{j,i} : F_{j,i} \text{ is delta causal}\}$, which is positive. Let G' be the subgraph of G obtained by deleting every delta causal component and its incident edges. Since every cycle of G contains a delta causal component, G' is acyclic. Hence, every directed walk in G' is a **simple path** (a path with no repeated nodes) with at most p nodes, where p is the number of components on the longest path in G' . (If G itself is acyclic, no cycle contains any component and the same bound holds with $p \leq n$.) Now take any two signal tuples $s, s' \in \mathcal{S}^N$ and let τ be the earliest time at which they differ, so $d(s, s') = 2^{-\tau}$. Each application of F advances the dependence of an output by one component, so a signal in $F^E(s)$ depends on s only through directed walks $j_1 \rightarrow j_2 \rightarrow \dots \rightarrow j_E$ of E successive components in G , and the earliest time that signal can differ is at least τ plus the total delay accumulated

along the walk. If such a walk contained no delta causal component, it would be a walk of $E = p + 1$ nodes in G' ; but G' is acyclic, so its walks are simple and have at most p nodes, a contradiction. Hence, every such walk passes through at least one delta causal component and accumulates delay at least D . Therefore, the earliest time at which $F^E(s)$ and $F^E(s')$ differ is at least $\tau + D$, so

$$d(F^E(s), F^E(s')) \leq 2^{-(\tau+D)} = 2^{-D} d(s, s').$$

Since $2^{-D} < 1$, the composition F^E is delta contracting, and the system is eventually delta contracting. □

Example 7.11: In Example 7.9, Figure 7.2(a), when we remove the delay component, we get a graph with just one node. The longest path in this remaining graph has length $p = 1$, so the system is eventually delta contracting with $E = p + 1 = 2$, just as we found in Example 7.9.

Example 7.12: In Example 7.10, Figure 7.3, when we remove the delay subcomponent $F_{1,1}$, we get a graph with two nodes. Note, however, that F_2 , which has no outputs, is also a delta causal component (with $\delta = 0$). So we should remove it as well, leaving us with a graph with just one node. The longest path in this remaining graph has length $p = 1$, so the system is eventually delta contracting with $E = p + 1 = 2$, just as we found in Example 7.10.

Given an eventually delta contracting system, we can use the Banach fixed-point theorem to find the fixed point of F^E . Fortunately, this also gives us a fixed point of the original system F .

Proposition 7.3. *If F is eventually delta contracting with some $E > 0$, and s is the unique fixed point of F^E , then s is a fixed point of F .*

Proof. We only need to consider $E > 1$. Because F^E is delta contracting, by the Banach fixed-point theorem (Theorem B.2), it has a unique fixed point s , where




Figure 7.4: A system that is not eventually delta contracting and also not Zeno.

$F^E(s) = s$. Suppose that $F(s) = s'$. We can now show that $s' = s$. Note that $F^{E-1}(F(s)) = F^{E-1}(s') = s$. Now apply F to both sides to get,

$$F(F^{E-1}(s')) = F(s) = s',$$

or

$$F^E(s') = s'.$$

Since s is the *unique* fixed point of F^E , it follows that $s' = s$. Hence, s is a fixed point of F . □

With this proposition, we now have a constructive procedure to find the fixed point of any eventually delta contracting system. We just form the composition F^E and apply the Banach fixed-point theorem to find the fixed point.

7.2.5 Zeno Behaviors

Note that having an [eventually delta contracting](#) system provides some protection against Zeno behaviors, but not an assurance. A trivial example is a system with one component with no inputs and one output, where the output is a [Zeno signal](#). This system is eventually delta contracting with $E = 1$, but it is also Zeno.

We might conjecture that if a system is not eventually delta contracting, then it must be Zeno. But this is not the case, as shown in the following example.

Example 7.13: Consider the system shown in Figure 7.4, where F_1 is defined identically to the function F_1 in Example 7.9. F_1 is causal, but not strictly




Figure 7.5: A plot of the signal from Example 7.13 from time 0 to time 5.




Figure 7.6: System discussed in sidebar on page 279 that will be Zeno if the Collatz conjecture is false and non-Zeno if it is true.

causal or delta causal. Assume again that $G = \mathbb{R}_+$ and $\mathcal{T}$ is the identity function.

Now define F_A as follows: if the input signal s has the following event set:

$$\mathcal{E}(s) = \{(g_1, v_1), (g_2, v_2), (g_3, v_3), \dots\},$$

then the output signal $F_A(s) = s'$ has the following event set:

$$\mathcal{E}(s') = \{(g_1 + 1/2, v_1), (g_2 + 1/3, v_2), (g_3 + 1/4, v_3), \dots\}.$$

In other words, the n -th input is delayed by $1/(n+1)$ time units.

This system is not eventually delta contracting, but it is also not Zeno. There are a finite number of events in any finite time, even though the number of events per unit time grows without bound. A plot of the signal from time 0 to time 5 is shown in Figure 7.5.

It can be extraordinarily difficult to determine if a system is Zeno, as discussed in the sidebar on page 279.




Figure 7.7: A plot of the signal that will be Zeno if the Collatz conjecture is false.

7.3 Generalized Ultrametric Semantics

Most but not all practical discrete-event systems can be modeled as [eventually delta contracting](#) systems. But there are exceptions. The system in Example 7.13 is not eventually delta contracting, but also not practically useful. Zeno systems are also not eventually delta contracting and usually not practically useful. However, in the next chapter, we will see a programming model that makes it easy to construct useful programs that are not eventually delta contracting. Moreover, that model uses [superdense time](#), a tag system that does not map easily to non-negative real numbers, a requirement for using the [Cantor metric](#) and the [Banach fixed-point theorem](#). Specifically, with superdense time, there is no injective order preserving function $\mathcal{T} : \mathbb{G} \rightarrow \mathbb{R}_+$.

The principles we have shown so far in this chapter can be generalized to handle a larger class of models. Using a different metric, it is possible to handle all systems that are eventually *strictly* contracting, even Zeno systems. A constructive procedure exists for finding the unique fixed point. The methods even extend beyond discrete-event systems to include systems operating on non-discrete signals, such as continuous-time systems, but we will not pursue those here.

7.3.1 A Generalized Distance Function

The Cantor metric maps the amount of time over which two signals are identical into a non-negative real number. This requires first that there be a non-decreasing

Sidebar: Difficulty of Determining Zeno Behaviors

We give here an example, due to Ben Lickly, which illustrates intuitively why it is difficult to determine whether a given DE model is Zeno. It depends on an unsolved mathematical puzzle known as the Collatz conjecture. If this conjecture is true, this system is discrete and non-Zeno, even though it is not [eventually delta contracting](#). If the conjecture is false, then this model defines a [Zeno signal](#) where an infinite number of events occur in a finite time.

The structure of the model is shown in Figure 7.6. The tag set is $G = \mathbb{R}_+$, the value set is $V = \mathbb{N}$, and the timestamp function is $\mathcal{T}(g) = g$. F_C receives a new positive integer once per time unit and begins a process of determining whether this integer is a counterexample to the Collatz conjecture. In the plot in Figure 7.7, the numbers start at 5 at time 0 and increase by 1 each time unit. For each input event with value v , if v is even, then F_C divides it by 2, otherwise it multiplies it by 3 and adds 1. If the result is 1, then F_C stops working with this positive integer; it is not a counterexample to the Collatz conjecture, which is that this process will always eventually reach 1. If the result is not 1, then F_C feeds the result back to itself via a delay component $F_{D'}$ to repeat the calculation.

The $F_{D'}$ component works as follows: Each value received on the input is forwarded to the output at a later logical time. The delay applied to the n -th event is $1/2^n$, where n counts the number of events received since (and including) the most recent event whose logical time was an integer multiple of one time unit. For example, with inputs starting at time 2:

- input at 2.000 s ($n = 1$) is delayed 0.500 s $\rightarrow$ emitted at 2.500 s
- input at 2.500 s ($n = 2$) is delayed 0.250 s $\rightarrow$ emitted at 2.750 s
- input at 2.750 s ($n = 3$) is delayed 0.125 s $\rightarrow$ emitted at 2.875 s $\dots$

The output times therefore form a geometric series that asymptotically approaches the next integer time unit but never reaches it. If a fresh input does arrive exactly on an integer-time-unit mark (typically because some other component scheduled an event there), the counter n restarts at 1.

As shown in Figure 7.7, for the first number 5 at time 0, the calculated numbers are 16, 8, 4, 2, 1 (the last one not shown) at times 0.5, 0.75, 0.875, 0.9375, and 0.96875. If any counterexample is found, then an infinite number of events will be circulated around the feedback loop before the next one-time-unit boundary.

function mapping the tag set G into a time value that is a real number. We can instead work directly with the set of tags over which two signals are identical.

Recall that the set of signals in Definition 7.1 is the set of functions given by

$$\mathcal{S} = (G \rightarrow V \cup \{\epsilon\}),$$

where G is a totally ordered tag set, V is the value set, and ϵ is the absent value.

Sidebar: Lipschitz Continuous Functions in Metric Spaces

The fixed-point theorems used here bear a close resemblance to the fixed-point theorems used in Chapters 4 and 6. As shown in Proposition B.5, a delta causal component realizes a [Lipschitz continuous](#) function. Definition B.7 of continuity is consistent with Definition A.9, where metric spaces replace CPOs. To see this, use the notion of a [Scott topology](#) (see sidebar on page 137) and the topological definition of the limit of a sequence of real numbers (see sidebar on page 136).

Consider a metric space (X, d) . An [open neighborhood](#) around an $x \in X$ is a set of points given by

$$N(x, r) = \{x' \in X \mid d(x, x') < r\},$$

for some real number r (called the **radius**). For the Cantor metric with [discrete-event signals](#), such an open neighborhood is the set of all signals that share a common prefix. Specifically, for some signal s_1 , $N(s_1, r)$ is the set of signals that are identical to s_1 up to and including time $\tau = \log_2(1/r)$.

Unlike the CPOs of Chapter 4, the set of discrete-event signals with the Cantor metric forms a [Hausdorff space](#). This means that given any two distinct signals s_1 and s_2 , there are open neighborhoods U_1 and U_2 such that $s_1 \in U_1$, $s_2 \in U_2$, and $U_1 \cap U_2 = \emptyset$. Let $r = d(s_1, s_2) > 0$ (since $s_1 \neq s_2$). Let $x_1 \in N(s_1, r)$ and $x_2 \in N(s_2, r)$. We can now show that $x_1 \neq x_2$, so $N(s_1, r)$ and $N(s_2, r)$ must be disjoint. To show that $x_1 \neq x_2$, observe that x_1 is identical to s_1 up to time τ , and x_2 is identical to s_2 up to time τ . But since $r = d(s_1, s_2)$, $\tau = \log_2(1/r)$, and s_1 and s_2 differ at time τ , it follows that $x_1 \neq x_2$. Note that this latter statement relies on the assumption that the signals are discrete.

Definition 7.10. The *generalized distance* $D(s_1, s_2)$ between two signals $s_1, s_2 \in \mathcal{S}$ is the largest *downset* over which they are identical:

$$D(s_1, s_2) = \{g \in \mathbb{G} \mid \forall g' \in \mathbb{G}, g' \leq g \implies s_1(g') = s_2(g')\}.$$

A downset is a subset of a poset that contains all elements below each of its elements. Hence, $D(s_1, s_2)$ is the set of tags over which s_1 and s_2 are identical, stopping before the first tag where they differ. This generalized distance does not require any mapping to real-valued time.

Let Δ_D be the set of all downsets of $\mathbb{G}$ ordered by reverse set inclusion. That is, for $\delta_1, \delta_2 \in \Delta_D$,

$$\delta_1 \leq \delta_2 \iff \delta_1 \supseteq \delta_2.$$

Then Δ_D is a *pointed* poset with bottom element $\perp_{\Delta_D} = \mathbb{G}$. The generalized distance function D has the form $D: \mathcal{S} \times \mathcal{S} \rightarrow \Delta_D$.

Proposition 7.4. $(\mathcal{S}, D, \Delta_D)$ is a *generalized ultrametric space*.

Proof. The first two properties of a generalized ultrametric distance function are obvious. To show the third property, let $s_1, s_2, s_3 \in \mathcal{S}$ be three signals. Let $D(s_1, s_2) = \delta_1$ and $D(s_2, s_3) = \delta_2$. The *join* of δ_1 and δ_2 , under the reverse set inclusion order, is the intersection, $\delta_1 \vee \delta_2 = \delta_1 \cap \delta_2$. For every tag in this intersection, s_1 is identical to s_2 and s_2 is identical to s_3 . Hence, s_1 is identical to s_3 for every tag in $\delta_1 \cap \delta_2$. Hence, $D(s_1, s_3) \leq \delta_1 \cap \delta_2$ (they could be identical for more tags outside this intersection). Therefore, $D(s_1, s_3) \leq \delta_1 \vee \delta_2$. $\square$

If the tag set $\mathbb{G}$ happens to be the non-negative reals, $\mathbb{R}_+$, then it is easy to see that if $D(s_1, s_2) \leq D(s'_1, s'_2)$, then $d(s_1, s_2) \leq d(s'_1, s'_2)$, where d is the Cantor metric. But the reverse implication does not hold in general unless the signals are *discrete-event signals*. But even for discrete-event signals, the distance function D is useful because the tag set $\mathbb{G}$ can be any totally ordered set, including *superdense time*. There is no need to map the tag set to real numbers.

7.3.2 Prefix Order and the Generalized Ultrametric Space

The set of signals $\mathcal{S}$ is a **poset** under a prefix order. This poset has interesting connections to the generalized ultrametric space. The following definition will prove useful:

Definition 7.11. The *restriction* of a signal s to a **downset** δ is the signal $s \upharpoonright \delta$ defined by

$$(s \upharpoonright \delta)(g) = \begin{cases} s(g) & \text{if } g \in \delta \\ \epsilon & \text{otherwise} \end{cases}$$

Definition 7.12. The *prefix order on signals* $\mathcal{S}$ is the relation $\sqsubseteq$ defined by

$$s_1 \sqsubseteq s_2 \iff \exists \delta \in \Delta_D, s_1 = s_2 \upharpoonright \delta,$$

where Δ_D is the set of all **downsets** of $\mathbb{G}$, and $s_2 \upharpoonright \delta$ is the restriction of s_2 to the downset δ .

In other words, s_1 is a prefix of s_2 if it matches an initial segment of s_2 . $(\mathcal{S}, \sqsubseteq)$ is a **pointed poset**. The bottom element is the empty signal λ . Moreover, this poset is a **lower semilattice**, meaning that any two signals have a unique greatest lower bound.

Definition 7.13. The longest **common prefix** of two signals s_1 and s_2 is the greatest lower bound, or **meet**, of two signals s_1 and s_2 in the prefix order $(\mathcal{S}, \sqsubseteq)$. This greatest lower bound is written $s_1 \sqcap s_2$ instead of $s_1 \wedge s_2$ to match $\sqsubseteq$ vs. $\leq$.

Once we have a poset, it might be tempting to consider **monotonic** functions and use the **Kleene fixed-point theorem** to find the fixed point. However, many useful discrete-event components are not monotonic. Nevertheless, a theory that is more-or-less equivalent to the one developed previously in this chapter and is based on the Kleene fixed-point theorem has been given by **Liu and Lee (2008)**. As with the theory using the Cantor metric, it deals with any system that can be modeled as an eventually delta contracting system. But here, we are interested in a more general theory that can handle all systems that are eventually strictly contracting.




Figure 7.8: An example used to illustrate the Matsikoudis construction.

Example 7.14: Consider a system where the tag and value sets are the natural numbers, $G = V = \mathbb{N}$, and the function F is given by

$$F(s)(g) = \begin{cases} 1 & \text{if } g \geq 1 \text{ and } s(g-1) = \epsilon, \\ 1 + s(g-1) & \text{if } g \geq 1 \text{ and } s(g-1) \neq \epsilon, \\ \epsilon & \text{otherwise.} \end{cases}$$

An implementation of this function in Ptolemy II (Ptolemaeus, 2014) is shown in Figure 7.8, where it is put into a feedback loop to produce a fixed point.

This function is not monotonic in the prefix order. For example, consider two inputs s_1 and s_2 , where $\mathcal{E}(s_1) = \{(0,0)\}$ and $\mathcal{E}(s_2) = \{(0,0), (1,1)\}$. Then $s_1 \sqsubseteq s_2$. But $F(s_1)$ is not a prefix of $F(s_2)$; they are instead given by the following event sets:

$$\mathcal{E}(F(s_1)) = \{(1,1), (2,1), (3,1), (4,1), \dots\}$$

$$\mathcal{E}(F(s_2)) = \{(1,1), (2,2), (3,1), (4,1), \dots\}.$$

Consequently, we cannot use the Kleene procedure to find the fixed point of F .

Example 7.15: The function F in Example 7.4 is not **monotonic** in the prefix order, despite being **strictly causal**. For a counterexample, consider two

signals s_1, s_2 given by their graphs as follows

$$\begin{aligned}\mathcal{G}(s_1) &= \{(0,0)\} \\ \mathcal{G}(s_2) &= \{(0,0), (0.1,0)\}.\end{aligned}$$

The corresponding outputs are

$$\begin{aligned}\mathcal{G}(F(s_1)) &= \{(1,0)\} \\ \mathcal{G}(F(s_2)) &= \{(1,0), (0.6,0)\}.\end{aligned}$$

Clearly, $s_1 \sqsubseteq s_2$, but it is not the case that $F(s_1) \sqsubseteq F(s_2)$.

The prefix order and the generalized ultrametric work well together.

Proposition 7.5. *For any $s_1, s_2, s_3 \in \mathcal{S}$,*

$$D(s_1, s_2) \leq D(s_1, s_3) \implies s_1 \sqcap s_3 \sqsubseteq s_1 \sqcap s_2.$$

Proof. $D(s_1, s_2)$ is the downset of tags over which s_1 and s_2 are identical. $D(s_1, s_3)$ is the downset of tags over which s_1 and s_3 are identical. $D(s_1, s_2) \leq D(s_1, s_3)$ uses the reverse set inclusion order on downsets. Hence, $D(s_1, s_2)$ is a superset of $D(s_1, s_3)$, which implies that the common prefix $s_1 \sqcap s_3$ is a prefix of the common prefix $s_1 \sqcap s_2$. □

Proposition 7.6. *For any $s_1, s_2, s_3 \in \mathcal{S}$,*

$$D(s_1 \sqcap s_2, s_1 \sqcap s_3) \leq D(s_2, s_3).$$

Proof. Let $g \in D(s_2, s_3)$. Then s_2 and s_3 agree on every $g' \leq g$. For any such g' , s_1 agrees with s_2 up to g' if and only if it agrees with s_3 up to g' , or equivalently,

$$g' \in D(s_1, s_2) \iff g' \in D(s_1, s_3).$$

Hence, at each $g' \leq g$, the two **restrictions** $s_1 \upharpoonright D(s_1, s_2)$ and $s_1 \upharpoonright D(s_1, s_3)$ are either both $s_1(g')$ or both ϵ , so they agree up to and including g . Thus, $g \in D(s_1 \sqcap s_2, s_1 \sqcap s_3)$, so

$$D(s_2, s_3) \subseteq D(s_1 \sqcap s_2, s_1 \sqcap s_3),$$

which under reverse set inclusion is the claim. □

7.3.3 Completeness of the Generalized Ultrametric Space

The completeness of this space is important because we will be considering iterative evaluations of the system function F on this space, and we need to know that the sequence of iterates converges to a fixed point. The notion of convergence that we use here, however, is a little different from convergence in a metric space. It is more like the topological notion (see sidebar on page 138). At each point in the computation, we will have determined a prefix of the signals that ultimately are the behavior of the system. Each prefix has events in a [downset](#) of the tag set $\mathbb{G}$, and the set of all signals with this prefix is a [spherically closed ball](#), which is similar to an [open neighborhood](#) in a metric space, except that it is closed rather than open. Hence, as the prefix grows during the computation, we get a chain of closed balls. To show completeness, we need to show that this chain has a non-empty intersection, at which point we can use Theorem B.3 to conclude that the system has a unique fixed point.

First notice that a [spherically closed ball](#) $B_\delta(s)$ of radius δ is the set of all signals s' that share a common prefix with s over the downset of tags δ . Any one of these signals in the ball can be used as the center of the ball. That is, if $s' \in B_\delta(s)$, then $B_\delta(s') = B_\delta(s)$. A chain of closed balls is a set $\{B_{\delta_i}(s_i) \mid i \in I\}$ of spherically closed balls, for some totally ordered index set I , such that for all $i, j \in I$, $i \leq j$ implies $B_{\delta_i}(s_i) \supseteq B_{\delta_j}(s_j)$. Each ball in the chain contains all balls with larger indices. Larger indices therefore correspond to longer prefixes.

Clearly, $(s \upharpoonright \delta) \in B_\delta(s)$. Also, if $B_{\delta_i}(s_i) \supseteq B_{\delta_j}(s_j)$, we have $(s_j \upharpoonright \delta_j) \in B_{\delta_i}(s_i)$. Hence, because any member of a ball can be used as its center, $B_{\delta_i}(s_i) = B_{\delta_i}(s_j \upharpoonright \delta_j)$. In other words, the center of any ball $B_{\delta_i}(s_i)$ can be any $(s_j \upharpoonright \delta_j)$ corresponding to a later ball $B_{\delta_j}(s_j)$ in the chain. This suggests that a signal s whose events consist of all the events in the restrictions $(s_j \upharpoonright \delta_j)$ —if such a signal exists and is in the intersection of the balls—might be a fixed point.

Theorem 7.2. (S, D, Δ_D) is a [spherically complete generalized ultrametric space](#).

Proof. (From Liu et al. (2006b), lemma 2.) Consider a chain of closed balls $\{B_{\delta_i}(s_i) \mid i \in I\}$, where each s_i is a signal in $\mathcal{S}$, each δ_i is a downset in Δ_D , and I is a totally ordered index set. Let

$$U = \bigcup_{i \in I} \mathcal{E}(s_i \upharpoonright \delta_i).$$

We assert that the signal s such that $\mathcal{E}(s) = U$ is a valid signal and is in each ball $B_{\delta_i}(s_i)$ in the chain. Hence, the intersection of the chain of closed balls is non-empty.

To prove this assertion, we first show that U is a valid event set for a signal. This means that for each $(g, v), (g, v') \in U$, $v = v'$. That is, the event set U is the graph of a function s . We will then have to show that s is in each ball $B_{\delta_i}(s_i)$ in the chain.

For the first part, towards a contradiction, suppose there exist $(g, v), (g, v') \in E$ where $v \neq v'$. Without loss of generality, suppose $i, j \in I$ and $i < j$ such that $s_i(g) = v$ and $s_j(g) = v'$. Because of the restrictions, it must be true that $g \in \delta_i$, in which case, because s_i and s_j differ at g ,

$$D(s_i, s_j) \not\leq \delta_i.$$

Hence, $s_j \notin B_{\delta_i}(s_i)$, which contradicts the assumption that $B_{\delta_i}(s_i) \supseteq B_{\delta_j}(s_j)$. Consequently, the signal s such that $\mathcal{E}(s) = E$ is a valid signal.

We now show that $s \in B_{\delta_i}(s_i)$ for all $i \in I$, or equivalently, that $D(s, s_i) \leq \delta_i$. Toward contradiction, suppose there is an $i \in I$ such that $D(s, s_i) \not\leq \delta_i$. Then there must be a $j \in I$ and $g \in \delta_i \cap \delta_j$ such that $s_i(g) \neq s_j(g)$. It follows that $D(s_i, s_j) \not\leq \delta_i$ and $D(s_j, s_i) \not\leq \delta_j$, in contradiction to the assumption that either $B_{\delta_i}(s_i) \supseteq B_{\delta_j}(s_j)$ or $B_{\delta_j}(s_j) \supseteq B_{\delta_i}(s_i)$. Hence, $D(s, s_i) \leq \delta_i$ for all $i \in I$. Therefore, $s \in \bigcap_{i \in I} B_{\delta_i}(s_i)$. Consequently, the intersection of the chain of closed balls is non-empty, so the space is spherically complete. $\square$

7.3.4 Fixed Points in the Generalized Ultrametric Space

Now that we have a spherically complete generalized ultrametric space $(\mathcal{S}, D, \Delta_D)$, it is trivial to generalize to consider N -tuples of signals.

Proposition 7.7. $(\mathcal{S}^N, D, \Delta_D)$ is a spherically complete generalized ultrametric space, where D is the *generalized distance* function generalized to N -tuples of signals.

All of the above arguments generalize trivially to the case of N -tuples of signals.

We now again apply the usual construction of the system function $F: \mathcal{S}^N \rightarrow \mathcal{S}^N$. Analogous to the Banach fixed-point formulation, we can define contracting and strictly contracting functions.

Definition 7.14. A *contracting function* $F: \mathcal{S}^N \rightarrow \mathcal{S}^N$ is one where for all $s_1, s_2 \in \mathcal{S}^N$,

$$D(F(s_1), F(s_2)) \leq D(s_1, s_2),$$

where D is the *generalized distance* function. A *strictly contracting function* is one where for all $s_1, s_2 \in \mathcal{S}^N$ where $s_1 \neq s_2$,

$$D(F(s_1), F(s_2)) < D(s_1, s_2).$$

By Theorem B.3, if F is strictly contracting, then F has a unique fixed point, which establishes the following theorem:

Theorem 7.3. A closed system modeled by a strictly contracting function $F: \mathcal{S}^N \rightarrow \mathcal{S}^N$ in the generalized ultrametric space $(\mathcal{S}^N, D, \Delta_D)$ has a unique fixed point.

As before, we can consider eventually strictly contracting functions.

Definition 7.15. An *eventually strictly contracting function* $F: \mathcal{S}^N \rightarrow \mathcal{S}^N$ is one where there exists an $E \in \mathbb{N}$, $E > 0$, such that the composition F^E is strictly contracting.

Using the same reasoning as before, we can conclude that a system modeled by an eventually strictly contracting function has a unique fixed point.

7.3.5 Constructive Procedure to Find the Fixed Point

Theorem B.3 does not provide a constructive procedure to find the fixed point. A procedure for finding the fixed point is given by Matsikoudis and Lee (2015). They combine the generalized ultrametric space with a poset based on the prefix order.

Theorem 7.4. (*Matsikoudis and Lee (2015)*) For a closed system given by a strictly contracting function $F: \mathcal{S}^N \rightarrow \mathcal{S}^N$ in the generalized ultrametric space $(\mathcal{S}^N, D, \Delta_D)$, F has a unique fixed point that can be found by the following procedure: first, define the function

$$G(s) = F(s) \sqcap F(F(s)).$$

Then, start with any $s_0 \in \mathcal{S}^N$ such that $s_0 \sqsubseteq G(s_0)$. Iterate G starting with s_0 to get a sequence $s_0, s_1, s_2, \dots$,

$$s_{i+1} = G(s_i) = F(s_i) \sqcap F(F(s_i)), \quad i \in \mathbb{N}.$$

Let s be the least upper bound of the sequence. If s is a fixed point of F , we are done. Otherwise, it will be true that $s \sqsubseteq G(s)$, so we can start over with $s_0 = s$. Repeat this procedure, getting a potentially transfinite sequence of least upper bounds. The least upper bound of that sequence is the unique fixed point of F .

Our proof is adapted from *Matsikoudis and Lee (2015)*. First, we need one definition and some propositions that will be used in the proof.

Definition 7.16. A signal tuple $s \in \mathcal{S}^N$ is a **post-fixed point** of F if $s \sqsubseteq F(s)$.

Throughout, G denotes the function $G(s) = F(s) \sqcap F(F(s))$, and we write $\sqcup C$ for the least upper bound, when it exists, of a set $C \subseteq \mathcal{S}^N$ under the prefix order $\sqsubseteq$. Propositions 7.5 and 7.6 and the generalized ultrametric inequality of Proposition 7.4 were stated for signals in $\mathcal{S}$, but, as noted above, they hold verbatim for signal tuples in $\mathcal{S}^N$, with $\sqcap$ and D taken componentwise. Recall also that a **contracting function** satisfies $D(F(s_1), F(s_2)) \leq D(s_1, s_2)$, and a **strictly contracting function** satisfies $D(F(s_1), F(s_2)) < D(s_1, s_2)$ whenever $s_1 \neq s_2$.

The first proposition records that G always produces a post-fixed point, and that it advances every post-fixed point.

Proposition 7.8. Let $F: \mathcal{S}^N \rightarrow \mathcal{S}^N$ be a contracting function. Then for every $s \in \mathcal{S}^N$:

1. $G(s) = F(s) \sqcap F(F(s))$ is a post-fixed point of F ; and
2. If s is a post-fixed point of F , then $s \sqsubseteq G(s)$; i.e., s is a post-fixed point of G .

Proof. Write $a = F(s)$ and $b = F(F(s))$, so that $b = F(a)$ and $G(s) = a \sqcap b$.

For part 1, since F is contracting, applied to the pair $a \sqcap b$ and a ,

$$D(F(G(s)), b) = D(F(a \sqcap b), F(a)) \leq D(a \sqcap b, a) = D(a \sqcap b, a \sqcap a) \leq D(b, a),$$

where the last inequality is Proposition 7.6 (with $s_1 = a$, $s_2 = b$, $s_3 = a$). Hence $D(b, F(G(s))) \leq D(b, a)$, and Proposition 7.5 (with $s_1 = b$, $s_2 = F(G(s))$, $s_3 = a$) gives

$$G(s) = b \sqcap a \sqsubseteq b \sqcap F(G(s)) \sqsubseteq F(G(s)),$$

so $G(s)$ is a post-fixed point of F .

For part 2, suppose $s \sqsubseteq F(s) = a$, so that $a \sqcap s = s$. Since F is contracting,

$$D(a, b) = D(F(s), F(F(s))) \leq D(s, F(s)) = D(a, s),$$

and Proposition 7.5 (with $s_1 = a$, $s_2 = b$, $s_3 = s$) yields

$$s = a \sqcap s \sqsubseteq a \sqcap b = G(s).$$

□

The next proposition shows that taking least upper bounds keeps us within the post-fixed points, which is what allows the construction to continue past a limit.

Proposition 7.9. *Let $F: \mathcal{S}^N \rightarrow \mathcal{S}^N$ be a contracting function, and let P be a set of post-fixed points of F whose least upper bound $\sqcup P$ exists in $\mathcal{S}^N$. Then $\sqcup P$ is a post-fixed point of F .*

Proof. Write $u = \sqcup P$, and let $s \in P$. Since F is contracting,

$$D(F(s), F(u)) \leq D(s, u). \quad (7.3)$$

By Proposition 7.6 (with $s_1 = u$, $s_2 = F(s)$, $s_3 = F(u)$) and (7.3),

$$D(u \sqcap F(s), u \sqcap F(u)) \leq D(F(s), F(u)) \leq D(s, u). \quad (7.4)$$

Because s is a post-fixed point, $F(s) \sqcap s = s$, so by Proposition 7.6 (with $s_1 = F(s)$, $s_2 = s$, $s_3 = u$),

$$D(s, u \sqcap F(s)) = D(F(s) \sqcap s, F(s) \sqcap u) \leq D(s, u). \quad (7.5)$$

Applying the **generalized ultrametric inequality** to s , $u \sqcap F(s)$, and $u \sqcap F(u)$, and using (7.4) and (7.5),

$$D(s, u \sqcap F(u)) \leq D(s, u \sqcap F(s)) \vee D(u \sqcap F(s), u \sqcap F(u)) \leq D(s, u),$$

because the join (intersection of downsets) of two distances that are each less than or equal to $D(s, u)$ is again less than or equal to $D(s, u)$. A second application of the ultrametric inequality, to u , s , and $u \sqcap F(u)$, gives

$$D(u, u \sqcap F(u)) \leq D(u, s) \vee D(s, u \sqcap F(u)) \leq D(s, u) = D(u, s).$$

Hence, by Proposition 7.5 (with $s_1 = u$, $s_2 = u \sqcap F(u)$, $s_3 = s$), and using $s \sqsubseteq u$,

$$s = u \sqcap s \sqsubseteq u \sqcap (u \sqcap F(u)) = u \sqcap F(u).$$

As this holds for every $s \in P$, the tuple $u \sqcap F(u)$ is an upper bound of P . Since $u = \sqcup P$ is the *least* upper bound, $u \sqsubseteq u \sqcap F(u)$; combined with $u \sqcap F(u) \sqsubseteq u$, this gives $u = u \sqcap F(u)$, that is, $u \sqsubseteq F(u)$. Thus u is a post-fixed point of F . $\square$

When F is strictly contracting, the fixed points of G are exactly the fixed points of F , so it suffices to drive G to a fixed point.

Proposition 7.10. *Let $F: \mathcal{S}^N \rightarrow \mathcal{S}^N$ be strictly contracting. Then s is a fixed point of F if and only if s is a fixed point of G .*

Proof. If $F(s) = s$, then $F(F(s)) = F(s) = s$, so $G(s) = F(s) \sqcap F(F(s)) = s \sqcap s = s$.

Conversely, suppose $G(s) = s$, that is, $s = F(s) \sqcap F(F(s))$. By Proposition 7.6 (with $s_1 = F(s)$, $s_2 = F(F(s))$, $s_3 = F(s)$),

$$D(s, F(s)) = D(F(s) \sqcap F(F(s)), F(s) \sqcap F(s)) \leq D(F(s), F(F(s))). \quad (7.6)$$

Suppose, toward a contradiction, that $s \neq F(s)$. Then, applying strict contraction to the distinct tuples s and $F(s)$,

$$D(F(s), F(F(s))) < D(s, F(s)),$$

which together with (7.6) gives $D(s, F(s)) \leq D(F(s), F(F(s))) < D(s, F(s))$, a contradiction. Hence $F(s) = s$. $\square$

Finally, we record that the limit stages of the construction are always defined, because the prefix order on signal tuples has least upper bounds of chains.

Proposition 7.11. *Every nonempty $\sqsubseteq$ -chain $C \subseteq \mathcal{S}^N$ has a least upper bound $\sqcup C$, a tuple of signals where each has the union of the events of the corresponding signals in C .*

Proof. It suffices to argue componentwise, so take C to be a chain of signals in $\mathcal{S}$. Any two members of C are prefix-comparable, and hence agree over their common prefix: if $s \sqsubseteq s'$ then $s = s' \upharpoonright \delta$ for the downset $\delta = D(s, s')$. Therefore, the union $\sqcup C$ of the event sets $\mathcal{E}(s)$ and $\mathcal{E}(s')$ is single-valued, so it is the event set of a valid signal u . Each $s \in C$ satisfies $s \sqsubseteq u$, and any upper bound of C must contain every event of every $s \in C$, hence contains u as a prefix. Thus $u = \sqcup C$. $\square$

Proof. (of Theorem 7.4) Uniqueness of the fixed point is already given by Theorem B.3; it remains to show that the stated procedure constructs it.

Choose a post-fixed point s_0 of F to start; the empty signal tuple λ is one, since $\lambda \sqsubseteq F(\lambda)$, and more generally $G(\hat{s})$ is a post-fixed point for every $\hat{s} \in \mathcal{S}^N$ by Proposition 7.8.1. Define a transfinite sequence (s_α) by

$$s_{\alpha+1} = G(s_\alpha) = F(s_\alpha) \sqcap F(F(s_\alpha)), \quad s_\lambda = \bigsqcup_{\beta < \lambda} s_\beta \text{ for limit ordinals } \lambda.$$

The iteration over the finite ordinals produces the sequence $s_0, s_1, s_2, \dots$ of the theorem, and s_ω is the least upper bound of that sequence; the limit stages are the “start over” steps.

We show by transfinite induction that, for every ordinal α , s_α is a well-defined post-fixed point of F and that $s_\beta \sqsubseteq s_\alpha$ whenever $\beta \leq \alpha$. The base case holds by the choice of s_0 . For a successor $\alpha + 1$, the induction hypothesis makes s_α a post-fixed point, so $s_\alpha \sqsubseteq G(s_\alpha) = s_{\alpha+1}$ by Proposition 7.8.2, and $s_{\alpha+1} = G(s_\alpha)$ is itself a post-fixed point by Proposition 7.8.1; monotonicity follows from $s_\alpha \sqsubseteq s_{\alpha+1}$ and the induction hypothesis. For a limit ordinal λ , the induction hypothesis makes $\{s_\beta \mid \beta < \lambda\}$ a $\sqsubseteq$ -chain of post-fixed points, so $s_\lambda = \bigsqcup_{\beta < \lambda} s_\beta$ exists by Proposition 7.11 and is a post-fixed point by Proposition 7.9; it is an upper bound of the chain, so $s_\beta \sqsubseteq s_\lambda$ for all $\beta < \lambda$.

The sequence (s_α) is therefore $\sqsubseteq$ -nondecreasing. Because $\mathcal{S}^N$ is a set, this sequence cannot be strictly increasing through all the ordinals: a proper class of ordinals cannot be injected into a set. Hence there are ordinals $\beta < \gamma$ with $s_\beta = s_\gamma$, and by monotonicity $s_\beta = s_{\beta+1}$, that is, $G(s_\beta) = s_\beta$. So s_β is a fixed point of G , and since F is strictly contracting, s_β is a fixed point of F by Proposition 7.10.

If s_β happens to be reached at a finite stage, or as the least upper bound $s = s_\omega$ of the first ω iterates, we are done. Otherwise $s = s_\omega$ satisfies $s \sqsubseteq G(s)$ but $s \neq G(s)$, and the procedure continues from $s_0 = s$, exactly as in the statement of the theorem. In all cases the construction stabilizes at a fixed point of F , which by uniqueness is *the* fixed point of F , and it is the least upper bound of the transfinite sequence of iterates. □

Theorem 7.4 gives a constructive procedure to find the fixed point of a strictly contracting function, even if the fixed point includes **Zeno signals** and if the function is not **continuous** in the prefix order. If F is continuous (and hence **monotonic**), then $G(s) = F(s)$, so the construction reduces to the Kleene procedure (Proposition A.6). But as we will see in the next example, F is often not monotonic (and hence not continuous).

It is straightforward to extend this result to any **eventually strictly contracting function**. Therefore, the theorem is very general, applicable to a wide range of systems. Moreover, although the use of $G(s) = F(s) \sqcap F(F(s))$ may seem a bit mysterious, it is actually quite intuitive. Each step of the construction gives an incremental step towards the fixed point, adding a (typically) finite number of events to the prefix of the final answer, as illustrated by the following example.

Example 7.16: Consider again the function F in Example 7.14 and Figure 7.8, which is not monotonic, so we cannot use the Kleene procedure to find the fixed point of F . Instead, we use the Matsikoudis construction to find the fixed point. The first step of the construction is shown in Figure 7.9. The initial signal is $s_0 = \lambda$, the empty signal. The first iterate is $s_1 = F(s_0) \sqcap F(F(s_0))$, the common prefix of $F(s_0)$ and $F(F(s_0))$, which are

given by:

$$\begin{aligned}\mathcal{E}(F(s_0)) &= \{(1, 1), (2, 1), (3, 1), (4, 1), \dots\} \\ \mathcal{E}(F(F(s_0))) &= \{(1, 1), (2, 2), (3, 2), (4, 2), \dots\}.\end{aligned}$$

The common prefix is given by the event set:

$$\mathcal{E}(s_1) = \{(1, 1)\},$$

as illustrated in Figure 7.9. Continuing, $s_2 = F(s_1) \sqcap F(F(s_1))$, $s_3 = F(s_2) \sqcap F(F(s_2))$, giving:

$$\begin{aligned}\mathcal{E}(s_2) &= \{(1, 1), (2, 2)\}, \\ \mathcal{E}(s_3) &= \{(1, 1), (2, 2), (3, 3)\}, \\ \mathcal{E}(s_4) &= \{(1, 1), (2, 2), (3, 3), (4, 4)\}, \\ &\vdots\end{aligned}$$

The fixed point is the signal given by

$$\mathcal{E}(s) = \{(1, 1), (2, 2), (3, 3), (4, 4), \dots\},$$

as shown in Figure 7.8. For this example, the construction converges in the first infinite iteration. There is no Zeno behavior.

7.4 Conclusion

This chapter has developed a semantic foundation for discrete-event systems, a model of computation in which components communicate through **signals** that are functions of **logical time**. By representing logical time with a totally ordered **tag** set G , we obtain a precise notion of when values in distinct signals are logically simultaneous and when one definitively precedes another. These orderings are logical rather than physical, so imposing them does not by itself reduce the concurrency available in a system.

The central technical question is the same one that recurs throughout this book: a composition of components is a **closed system** whose behavior is the fixed point of a function, and we must establish when that fixed point exists, is unique, and can be constructed. For discrete-event systems, the key structural property is **causality**, the requirement that an output cannot depend on future inputs. We




Figure 7.9: The Matsikoudis construction for the example in Figure 7.8.

made this precise using the **Cantor metric**, under which causal, **strictly causal**, and **delta causal** components lead to system models that are **contractions**, **strictly contracting** functions, and **delta contracting** functions on a **complete metric space**. For a delta causal system, the **Banach fixed-point theorem** then guarantees a unique fixed point that is the limit of a convergent sequence of iterates.

Because delta contracting is a severe constraint that many useful systems do not directly satisfy, we introduced **eventually delta contracting** systems, where some finite composition F^E is delta contracting. This condition is met whenever every cycle in the system contains at least one delta causal component, and it yields a constructive procedure for finding the fixed point of models. This enables inclusion of components with instantaneous reactions, which are not themselves delta causal. Along the way, we saw that eventual delta contraction offers only partial protection against Zeno behaviors, and that determining whether a given

system is Zeno can be extraordinarily difficult, as illustrated by a model whose Zeno-ness hinges on the Collatz conjecture.

Finally, we generalized these results by replacing real-valued distances with a [generalized distance](#) that records the [downset](#) of tags over which two signals agree. This yields a [spherically complete generalized ultrametric space](#) that requires no mapping from tags to real-valued time, and hence accommodates tag sets such as [superdense time](#). In this setting, the Matsikoudis construction ([Matsikoudis and Lee, 2015](#)) builds the unique fixed point of any strictly contracting function, even when that function is not [monotonic](#) in the prefix order and even when the fixed point contains Zeno signals. When the function happens to be continuous, the construction collapses to the familiar Kleene procedure, revealing discrete-event semantics as a strict generalization of the order-theoretic semantics developed earlier in the book.

These foundations set the stage for the next chapter, which introduces a programming model built on superdense time. That model makes it easy to write practical programs that are useful yet not eventually delta contracting, and the generalized ultrametric semantics developed here is precisely what gives such programs a well-defined, deterministic meaning.

Reactors

8.1	TBD	297
-----	---------------	-----

8.1 TBD

Part IV

Appendices



Partially Ordered Sets

A.1	Sets	302
A.2	Relations and Functions	303
A.2.1	Restriction and Projection	306
A.3	Sequences	307
A.4	Partial Orders	307
	<i>Insight: Exponential Notation for Sets of Functions</i>	308
A.4.1	Orders on Tuples	311
A.4.2	Upper and Lower Bounds	311
A.4.3	Complete Partial Orders	313
A.4.4	Flat Partial Orders	316
A.4.5	Lattices	317
A.5	Functions on Posets	318
A.5.1	Monotonic Functions	318
A.5.2	Continuous Functions	319
A.5.3	Posets of Functions	320
A.6	Fixed Points	322
A.7	Cardinality	323
	<i>Sidebar: Fixed Point Theorems</i>	324
A.7.1	Countable and Uncountable Sets	325
A.8	Discrete Sets	328
A.9	Computability	329
A.10	Conclusion	330

This chapter provides mathematical preliminaries used in subsequent chapters to develop the theory of concurrent systems. It reviews basic ideas and notation in logic, with particular emphasis on sets, functions, partial orders, and fixed-point theorems. The applications in subsequent chapters are essential to develop a full understanding of the role that these mathematical models play in concurrent systems.

A.1 Sets

In this section, we review the notation for sets. A **set** is a collection of objects. When object a is in set A , we write $a \in A$. We define the following sets:

- $\mathbb{B} = \{0, 1\}$, the set of **binary digits**.
- $\mathbb{T} = \{\text{false}, \text{true}\}$, the set of **truth values**.
- $\mathbb{N} = \{0, 1, 2, \dots\}$, the set of **natural numbers**.
- $\mathbb{Z} = \{\dots, -1, 0, 1, 2, \dots\}$, the set of **integers**.
- $\mathbb{R}$, the set of **real numbers**.
- $\mathbb{R}_+$, the set of **non-negative real numbers**.

When set A is entirely contained by set B , we say that A is a **subset** of B and write $A \subseteq B$. For example, $\mathbb{B} \subseteq \mathbb{N} \subseteq \mathbb{Z} \subseteq \mathbb{R}$. The sets may be equal, so the statement $\mathbb{N} \subseteq \mathbb{N}$ is true, for example. The **powerset** of a set A is defined to be the set of all subsets. It is written $\wp(A)$ or 2^A (for a justification of the latter notation, see the sidebar on page 308). The **empty set**, written $\emptyset$, is always a member of the powerset, $\emptyset \in \wp(A)$.

We define **set subtraction** as follows,

$$A \setminus B = \{a \in A : a \notin B\}$$

for all sets A and B . This notation is read “the set of elements a from A such that a is not in B .”

A **cartesian product** of sets A and B is a set written $A \times B$ and defined as follows,

$$A \times B = \{(a, b) : a \in A, b \in B\}.$$

A member of this set (a, b) is called a **tuple**. This notation is read “the set of tuples (a, b) such that a is in A and b is in B .” A cartesian product can be formed with three or more sets, in which case the tuples have three or more elements. For example, we might write $(a, b, c) \in A \times B \times C$. A cartesian product of a set A with itself is written $A^2 = A \times A$. A cartesian product of a set A with itself n times, where $n \in \mathbb{N}$ is written A^n . A member of the set A^n is called an **n -tuple**. By convention, A^0 is a **singleton set**, or a set with exactly one element, regardless of the size of A . Specifically, we define $A^0 = \{\emptyset\}$. Note that A^0 is not itself the empty set. It is a singleton set containing the empty set (for insight into the rationale for this definition, see the box on page 308).

A.2 Relations and Functions

A **relation** from set A to set B is a subset of $A \times B$. A **partial function** f from set A to set B is a relation where $(a, b) \in f$ and $(a, b') \in f$ imply that $b = b'$. Such a partial function is written $f: A \rightarrow B$. A **total function** or just **function** f from A to B is a partial function where for all $a \in A$, there is a $b \in B$ such that $(a, b) \in f$. Such a function is written $f: A \rightarrow B$, and the set A is called its **domain** and the set B its **codomain**. Rather than writing $(a, b) \in f$, we can equivalently write $f(a) = b$.

Example A.1: An example of a partial function is $f: \mathbb{R} \rightarrow \mathbb{R}$ defined by $f(x) = \sqrt{x}$ for all $x \in \mathbb{R}_+$. It is undefined for any $x < 0$ in its domain $\mathbb{R}$.

A partial function $f: A \rightarrow B$ may be defined by an **assignment rule**, as done in the above example, where an assignment rule simply explains how to obtain the value of $f(a)$ given $a \in A$. Alternatively, the function may be defined by its **graph**, which is a subset of $A \times B$.

Example A.2: The same partial function from the previous example has the graph $f \subseteq \mathbb{R}^2$ given by

$$f = \{(x, y) \in \mathbb{R}^2 : x \geq 0 \text{ and } y = \sqrt{x}\}.$$

Note that we use the same notation f for the function and its graph when it is clear from context which we are talking about.

The **set of all functions** $f: A \rightarrow B$ is written $(A \rightarrow B)$ or B^A . The former notation is used when the exponential notation proves awkward. For a justification of the notation B^A , see the box on page 308.

The **function composition** of $f: A \rightarrow B$ and $g: B \rightarrow C$ is written $(g \circ f): A \rightarrow C$ and defined by

$$(g \circ f)(a) = g(f(a))$$

for any $a \in A$. Note that in the notation $(g \circ f)$, the function f is applied first.

A function whose domain and codomain are the same set is called an **endofunction**. For an endofunction $f: A \rightarrow A$, the composition with itself can be written $(f \circ f) = f^2$, or more generally

$$\underbrace{(f \circ f \circ \dots \circ f)}_{n \text{ times}} = f^n$$

for any $n \in \mathbb{N}$. In case $n = 1$, $f^1 = f$. For the special case $n = 0$, the function f^0 is by convention the **identity function**, so $f^0(a) = a$ for all $a \in A$. Note that $f^n \in A^A$ for all $n \in \mathbb{N}$.

For every function $f: A \rightarrow B$, there is an associated **image function** $\hat{f}: \wp(A) \rightarrow \wp(B)$ defined on the **powerset** of A as follows,

$$\forall A' \subseteq A, \quad \hat{f}(A') = \{b \in B : \exists a \in A', f(a) = b\}.$$

We call $\hat{f}$ the **lifted** version of f . When there is no ambiguity, we may write the lifted version of f simply as f rather than $\hat{f}$.

For any $A' \subseteq A$, $\hat{f}(A')$ is called the **image** of A' for the function f . The image $\hat{f}(A)$ of the domain A is called the **range** of the function f .

Example A.3: The image $\hat{f}(\mathbb{R})$ of the function $f: \mathbb{R} \rightarrow \mathbb{R}$ defined by $f(x) = x^2$ is $\mathbb{R}_+$.

A function $f: A \rightarrow B$ is **onto** (or **surjective**) if $\hat{f}(A) = B$. A function $f: A \rightarrow B$ is **one-to-one** (or **injective**) if for all $a, a' \in A$,

$$a \neq a' \Rightarrow f(a) \neq f(a'). \quad (\text{A.1})$$

That is, no two distinct values in the domain yield the same value in the codomain. A function that is both one-to-one and onto is called a **bijection**.

Example A.4: The function $f: \mathbb{R} \rightarrow \mathbb{R}$ defined by $f(x) = 2x$ is a bijection. The function $f: \mathbb{Z} \rightarrow \mathbb{Z}$ defined by $f(x) = 2x$ is one-to-one, but not onto. The function $f: \mathbb{R}^2 \rightarrow \mathbb{R}$ defined by $f(x, y) = xy$ is onto but not one-to-one.

The previous example underscores the fact that an essential part of the definition of a function is its domain and codomain.

Proposition A.1. *If $f: A \rightarrow B$ is onto, then there is a one-to-one function $h: B \rightarrow A$.*

Proof. Let h be defined by $h(b) = a$ where a is any element in A such that $f(a) = b$. There must always be at least one such element because f is onto. We can now show that h is one-to-one. To do this, consider any two elements $b, b' \in B$ where $b \neq b'$. We need to show that $h(b) \neq h(b')$. Assume to the contrary that $h(b) = h(b') = a$ for some $a \in A$. But then by the definition of h , $f(a) = b$ and $f(a) = b'$, which implies $b = b'$, a contradiction. $\square$

The converse of this proposition is also easy to prove.

Proposition A.2. *If $B \neq \emptyset$, and $h: B \rightarrow A$ is one-to-one, then there is an onto function $f: A \rightarrow B$.*

Any bijection $f: A \rightarrow B$ has an **inverse** $f^{-1}: B \rightarrow A$ defined as follows,

$$f^{-1}(b) = a \in A \text{ such that } f(a) = b, \quad (\text{A.2})$$

for all $b \in B$. This function is defined for all $b \in B$ because f is onto. And for each $b \in B$ there is a single unique $a \in A$ satisfying (A.2) because f is one-to-one. For any bijection f , its inverse is also a bijection.

A.2.1 Restriction and Projection

Given a function $f: A \rightarrow B$ and a subset $C \subseteq A$, we can define a new function $f \upharpoonright C$ that is the **restriction** of f to C . It is defined so that for all $x \in C$, $(f \upharpoonright C)(x) = f(x)$.

Example A.5: The function $f: \mathbb{R} \rightarrow \mathbb{R}$ defined by $f(x) = x^2$ is not one-to-one. But the function $f \upharpoonright \mathbb{R}_+$ is.

Consider an n -tuple $a = (a_0, a_1, \dots, a_{n-1}) \in A_0 \times A_1 \times \dots \times A_{n-1}$. A **projection** of this n -tuple extracts elements of the tuple to create a new tuple. Specifically, let

$$I = (i_0, i_1, \dots, i_m) \in \{0, 1, \dots, n-1\}^m$$

for some $m \in \mathbb{N} \setminus \{0\}$. That is, I is an m -tuple of indexes. Then we define the projection of a onto I by

$$\pi_I(a) = (a_{i_0}, a_{i_1}, \dots, a_{i_m}) \in A_{i_0} \times A_{i_1} \times \dots \times A_{i_m}.$$

The projection may be used to permute elements of a tuple, to discard elements, or to repeat elements.

Projection of a tuple and restriction of a function are related. An n -tuple $a \in A^n$ where $a = (a_0, a_1, \dots, a_{n-1})$ may be considered a function of the form $a: \{0, 1, \dots, n-1\} \rightarrow A$, in which case $a(0) = a_0$, $a(1) = a_1$, etc. Projection is similar to restriction of this function, differing in that restriction, by itself, does not provide the ability to permute, repeat, or renumber elements. But conceptually, the operations are similar, as illustrated by the following example.

Example A.6: Consider a 3-tuple $a = (a_0, a_1, a_2) \in A^3$. This is represented by the function $a: \{0, 1, 2\} \rightarrow A$. Let $I = \{1, 2\}$. The projection $b = \pi_I(a) = (a_1, a_2)$, which itself can be represented by a function $b: \{0, 1\} \rightarrow A$, where $b(0) = a_1$ and $b(1) = a_2$.

The restriction $a \upharpoonright I$ is not exactly the same function as b , however. The domain of the first function is $\{1, 2\}$, whereas the domain of the second is $\{0, 1\}$. In particular, $(a \upharpoonright I)(1) = b(0) = a_1$ and $(a \upharpoonright I)(2) = b(1) = a_2$.

A projection may be **lifted** just like ordinary functions. Given a set of n -tuples $B \subseteq A_0 \times A_1 \times \dots \times A_{n-1}$ and an m -tuple of indexes $I \in \{0, 1, \dots, n-1\}^m$, the **lifted**

projection is

$$\hat{\pi}_I(B) = \{\pi_I(b) : b \in B\}.$$

A.3 Sequences

A tuple $(a_0, a_1) \in A^2$ can be interpreted as a sequence of length 2. The order of elements in the sequence matters, and is in fact captured by the natural ordering of the natural numbers. The number 0 comes before the number 1. We can generalize this and recognize that a **sequence** of elements from set A of length n is an n -tuple in the set A^n . A^0 represents the set of empty sequences, a [singleton set](#) (there is only one empty sequence). We denote the **empty sequence** λ .

The set of all **finite sequences** of elements from the set A is written A^* , where we interpret $*$ as a wildcard that can take on any value in $\mathbb{N}$. Since $0 \in \mathbb{N}$, $\lambda \in A^*$. A member of this set with length $n > 0$ is an n -tuple. The $*$ operator here is known as the **Kleene star** (or **Kleene operator** or **Kleene closure**), after American mathematician Stephen Cole Kleene (1909-1994).

The set of **infinite sequences** of elements from A is written $A^{\mathbb{N}}$ or A^{ω} . The set of **finite and infinite sequences** is written

$$A^{**} = A^* \cup A^{\omega}.$$

Finite and infinite sequences play an important role in the semantics of concurrent programs. They can be used, for example, to represent streams of messages sent from one part of the program to another. Or they can represent successive assignments of values to a variable. For programs that terminate, finite sequences will be sufficient. For programs that do not terminate, we need infinite sequences.

A.4 Partial Orders

The notion of **order** is central to much of the structure in concurrent systems. Examples of ordering relationships include the usual intuitive ordering of numbers (0 is less than 1, etc.), but also many other concepts. For example, event a causes event b , 3.14159 is a better approximation of π than 3.14, and John is Sarah's father are all elements of ordering relations. For many such ordering relations, not all

Insight: Exponential Notation for Sets of Functions

The exponential notation B^A for the set of functions of form $f: A \rightarrow B$ is worth explaining. Recall that A^2 is the **cartesian product** of set A with itself, and that $\wp(A)$ or 2^A is the **powerset** of A . These two notations are naturally thought of as sets of functions. A construction attributed to John von Neumann defines the natural numbers as follows,

$$\begin{aligned} \mathbf{0} &= \emptyset \\ \mathbf{1} &= \{\mathbf{0}\} = \{\emptyset\} \\ \mathbf{2} &= \{\mathbf{0}, \mathbf{1}\} = \{\emptyset, \{\emptyset\}\} \\ \mathbf{3} &= \{\mathbf{0}, \mathbf{1}, \mathbf{2}\} = \{\emptyset, \{\emptyset\}, \{\emptyset, \{\emptyset\}\}\} \\ &\dots \end{aligned}$$

With this definition, the powerset 2^A is the set of functions mapping the set A into the set $\mathbf{2}$. Consider one such function, $f \in 2^A$. For each $a \in A$, either $f(a) = \mathbf{0}$ or $f(a) = \mathbf{1}$. If we interpret “ $\mathbf{0}$ ” to mean “nonmember” and “ $\mathbf{1}$ ” to mean “member,” then the set of functions 2^A represents the set of all subsets of A .

Similarly, the cartesian product A^2 can be interpreted as the set of functions of form $f: \mathbf{2} \rightarrow A$, or using von Neumann’s numbers, $f: \{\mathbf{0}, \mathbf{1}\} \rightarrow A$. Consider a tuple $a = (a_0, a_1) \in A^2$. The function $a: \{\mathbf{0}, \mathbf{1}\} \rightarrow A$ naturally represents this tuple, where $a(\mathbf{0}) = a_0$ and $a(\mathbf{1}) = a_1$. The argument to the function is the index into the tuple. We can now interpret the set of functions B^A of form $f: A \rightarrow B$ as a set of tuples indexed by the set A instead of by the natural numbers.

Let $\omega = \{\emptyset, \{\emptyset\}, \{\emptyset, \{\emptyset\}\}, \dots\}$ represent the set of **von Neumann numbers**. This set is closely related to the set $\mathbb{N}$. Given a set A , it is now natural to interpret A^ω as the set of all infinite sequences of elements from A , equal to $A^\mathbb{N}$.

The **singleton set** A^0 can now be interpreted as the set of all functions whose domain is the empty set and codomain is A . There is exactly one such function (no two such functions are distinguishable), and that function has an empty **graph**. Before, we defined $A^0 = \{\emptyset\}$. Using von Neumann numbers, $A^0 = \mathbf{1}$, corresponding nicely with the definition of a zero exponent on ordinary numbers. You can think of $A^0 = \{\emptyset\}$ as the set of all functions with an empty graph. It is customary in the literature to omit the bold face font for A^0 , 2^A , and A^2 , writing instead simply A^0 , 2^A , and A^2 .

elements of the sets of interest are ordered. For example, John and Jane may not be related at all, unlike John and Sarah. Such partial orders play a central role in mathematical models of concurrent systems.

Definition A.1. A *partial order* on a set A is a *relation* from A to A satisfying the following properties. For all $a, b, c \in A$, the relation is

1. *reflexive*: $a \leq a$
2. *antisymmetric*: $a \leq b$ and $b \leq a$ implies that $a = b$.
3. *transitive*: $a \leq b$ and $b \leq c$ implies that $a \leq c$.

In this definition, we have written the relation using the symbol $\leq$. Specifically, the relation is $\leq$, where $\leq \subseteq A \times A$, and $(a_0, a_1) \in \leq$ is equivalently written $a_0 \leq a_1$. The latter notation is referred to as **infix notation**.

A **partially ordered set** or **poset** is a set A and a partial order relation $\leq$ on that set. We can write a poset as a tuple $(A, \leq)$, or $(A, \leq_A)$ if we need to make a distinction between a partial order $\leq_A$ on set A and another partial order $\leq_B$ on another set B .

Example A.7: A set A of sets and the subset relation $\subseteq$ form a poset $(A, \subseteq)$. The relation $\subseteq$ is referred to as the **subset order**.

For sets of sequences, there is a natural partial order relation based on the notion of a **prefix**.

Definition A.2. For a set A and the set of finite and infinite *sequences* A^{**} of elements of A , the *prefix order* is a *relation* $\sqsubseteq$ from A^{**} to A^{**} such that for any $s, s' \in A^{**}$, $s \sqsubseteq s'$ if either s is the empty sequence, or for all $n \in \mathbb{N}$ where $s(n)$ is defined, $s'(n)$ is defined and is equal to $s(n)$.

It is easy to show that $(A, \sqsubseteq)$ is a poset by showing that the prefix order conforms with Definition A.1.

Given a poset $(A, \leq)$ and two elements $a, a' \in A$, if either $a \leq a'$ or $a' \leq a$, then a and a' are said to be **comparable**. Otherwise, they are **incomparable**.

Definition A.3. A *chain* $C \subseteq A$ is a subset of a poset $(A, \leq)$ where any two members of the subset are comparable.

A **total order** is a poset $(A, \leq)$ where A itself is a chain.

Example A.8: The posets $(\mathbb{N}, \leq)$, $(\mathbb{Z}, \leq)$, and $(\mathbb{R}, \leq)$, where $\leq$ is the ordinary **numeric order** are total orders.

Example A.9: The **powerset** $2^{\mathbb{N}}$ is the set of all sets of natural numbers. $(2^{\mathbb{N}}, \subseteq)$, where $\subseteq$ is the subset order, is a poset but not a total order.

A finite poset may be represented by a **Hasse diagram**, where elements of the poset are placed above one another with line segments indicating the order relation. If two elements a and b are joined by a line segment, and a is below b in the diagram, then $a < b$.

Example A.10: Figure A.1 gives the Hasse diagram of a poset

$$A = \{\perp, a, b, c, d, e, f, g, h\},$$

where $\perp$ is the **bottom** element (below everything else in the poset), e, f, g , and h are just above $\perp$, and a, b, c , and d are each above two other elements, as shown in the diagram. In this poset, $\{\perp, e, b\}$, for example, is a chain. The elements a and g are incomparable.




Figure A.1: A Hasse diagram of a simple partial order.

equal to every element in A . When the set is understood, the bottom element may be written simply $\perp$, as it is in the previous example.

Given a poset $(A, \leq)$, we can induce another relation $<$ called the **strict partial order** relation defined as follows.

$$\forall a, a' \in A, \quad a < a' \Leftrightarrow a \leq a' \text{ and } a \neq a'.$$

$(A, <)$ is called a **strict poset**.

A.4.1 Orders on Tuples

Given a poset $(A, \leq_A)$, the set A^2 has a **pointwise order** defined by

$$(a_0, a_1) \leq (a'_0, a'_1) \Leftrightarrow a_0 \leq_A a'_0 \text{ and } a_1 \leq_A a'_1,$$

for all $(a_0, a_1), (a'_0, a'_1) \in A^2$. With this order relation, $(A^2, \leq)$ is clearly a poset. An alternative order for A^2 is the **lexicographic order**, defined by

$$(a_0, a_1) \leq (a'_0, a'_1) \Leftrightarrow a_0 <_A a'_0 \text{ or } (a_0 = a'_0 \text{ and } a_1 \leq_A a'_1),$$

for all $(a_0, a_1), (a'_0, a'_1) \in A^2$.

Both pointwise order and lexicographic order generalize trivially to A^m for $m > 2$. They also generalize to A^* , A^ω , and A^{**} .

Example A.11: Let $A = \{a, b, c, \dots, z\}$ be the letters of the alphabet and $\leq$ be the usual **alphabetical order** for those letters. Then A^* is the set of all finite sequences of letters, which includes the set of all words. The lexicographic order for A^* is called the **dictionary order** because it gives the order in which words appear in a dictionary. The dictionary order is a **total order**, whereas the pointwise order for this poset would not be. In particular, under the pointwise order, the sequence (a, b) would be incomparable to the sequence (b, a) , whereas under the dictionary order, the first is less than the second.

A.4.2 Upper and Lower Bounds

Given a poset $(A, \leq)$ and a subset $B \subseteq A$, an **upper bound** of B , if it exists, is an element $a \in A$ such that for all $b \in B$, $b \leq a$. A **least upper bound** or **LUB**, if it

exists, is an upper bound a such that for all other upper bounds a' we have $a \leq a'$. A set may have an upper bound and no LUB.

Example A.12: Consider the poset $(\mathbb{Q}, \leq)$ of **rational numbers**, where $\leq$ is the natural numeric order. Let $B \subset \mathbb{Q}$ be the subset of rational numbers whose value is less than the irrational number π (this relation “less than” is not the $\leq$ relation in this poset, since π is not a member of the poset, but we understand the definition of this subset anyway). In this poset, B has many upper bounds, but no least upper bound. Its LUB would have to be the greatest rational number less than π , and there is no such number.

If a set $B \subseteq A$ has a least upper bound in the poset $(A, \leq)$, then it is said to be **joinable** in $(A, \leq)$, and the LUB is called the **join** of B and written $\bigvee B$.

Example A.13: Consider the poset $(2^{\mathbb{N}}, \subseteq)$ of sets of natural numbers under the **subset order**. Every subset of $B \subseteq 2^{\mathbb{N}}$ is joinable, and the join is the union of the sets in B ,

$$\bigvee B = \bigcup_{b \in B} b.$$

It is easy to see that this is a LUB. Any other bound must contain at least this union. It is not accidental that the notation $\bigvee$ is similar to $\cup$. This similarity can be very helpful in getting used to the notation.

Correspondingly, a subset $B \subseteq A$ may have a **lower bound** in the poset $(A, \leq)$. This will be an element $a \in A$ such that for all $b \in B$, $a \leq b$. The set B may also have a **greatest lower bound** or **GLB**, defined similarly to the LUB. The GLB, if it exists, is called the **meet** of B in $(A, \leq)$ and is written $\bigwedge B$.

Example A.14: Any subset B as in the previous example has a meet given by

$$\bigwedge B = \bigcap_{b \in B} b.$$

Again, it is not accidental that the notation $\bigwedge$ reminds us of $\cap$.

A.4.3 Complete Partial Orders

Recall that a **pointed** poset $(A, \leq)$ is one that has a **bottom element**. The bottom element is the meet of the entire set,

$$\perp_A = \bigwedge A \in A.$$

Definition A.4. A nonempty subset $D \subseteq A$ of poset $(A, \leq)$ is a **directed set** if every pair of elements in D has an upper bound in D . Equivalently, D is directed if every non-empty finite subset of D is joinable in D .

Every **chain** is a directed set. In fact, directed sets can be viewed as generalizations of the idea of chains.

Definition A.5. A **complete partial order** or **CPO** $(A, \leq)$ is a pointed poset where every directed subset is **joinable** in A .

CPOs have extremely useful properties and are widely used in computer science to study semantics. To understand why they are so useful, note first that any *finite* directed subset is trivially joinable, by the definition of a directed set. To be a CPO, this property has to extend to *infinite* directed subsets. This may seem like a small step, but it is not. Directed subsets in a CPO may be thought of as being “directed” towards some goal. That goal is the least upper bound, whose existence is guaranteed by the fact that the set is a CPO. This LUB is in a sense the “**limit**” of the directed set. This notion of a limit turns out to be very powerful and surprisingly widely applicable.

The intuition behind the LUB of a directed set is easiest to understand for a **chain** $C \subseteq A$, which is a particularly simple kind of directed set. In this case, interpreting the LUB as a limit of the chain is completely natural. Being able to rely on the existence of this limit is valuable. Any ordered sequence in a CPO, therefore, has a limit. This is not a trivial property.

Example A.15: None of the sets $\mathbb{N}$, $\mathbb{Z}$, $\mathbb{R}$ with **numeric order** is a CPO. Neither is the set of **von Neumann numbers** ω under the subset order. Each of these sets is itself a chain with no least upper bound, so clearly it cannot be that every chain has a least upper bound. If we augment $\mathbb{N}$ with an infinite element ∞ , bigger than all elements in $\mathbb{N}$, then the resulting set $\mathbb{N} \cup \{\infty\}$

with the (extended) numeric order is a CPO. Similarly, the set $\omega \cup \{\omega\}$ is a CPO under the subset order, because every element of ω is a subset of ω .

The sets considered in the previous example are all **totally ordered sets**. For such sets, every directed subset is a chain. In fact, every subset is a chain. So to determine whether the set is a CPO, we only need to consider whether every chain has a least upper bound. Some sets that are not totally ordered also have the property that every directed set is a chain.

Example A.16: Given any set A , the set of **finite sequences** A^* of elements from A with the **prefix order** $\sqsubseteq$ is not a CPO. It is easy to construct a chain of growing finite sequences that has no upper bound that is itself a finite sequence. However, the set of finite and infinite sequences A^{**} is a CPO under the prefix order. To show this, first note that every directed subset of A^{**} is a chain (given two sequences that are both a prefix of a third, one of the two must be a prefix of the other). Thus, we simply note that given any chain, if the chain is finite, its least upper bound is its maximal element; if the chain is infinite, its least upper bound is the infinite sequence defined by the union of the elements of the chain. The notion of limits of such chains plays a major role in the semantics of concurrent systems.

It turns out that to decide whether any poset is a CPO, it is enough to consider only whether every chain has a least upper bound. It is not necessary to consider directed sets that are not chains. This follows from the following alternative definition of a CPO.

Definition A.6. A pointed poset $(A, \leq)$ is a **complete partial order (CPO)** if every chain in A is joinable in A .

This definition captures the intuition that every chain has a limit, in the sense of a **least upper bound**. Proving that this definition is equivalent to Definition A.5 appears to be quite difficult. It seems to require the machinery of ordinals and the axiom of choice and is beyond the scope of this text. For a discussion, see [Davey and Priestly \(2002\)](#). In this text, we will use both definitions, though most of the time the second one will prove easier to use. In the case of the **prefix order**, every **directed set** is a chain, so the equivalence of the two definitions is trivial.

The following proposition enables us to construct more complicated CPOs from simpler ones.

Proposition A.3. *Given a CPO $(A, \leq_A)$, $(A^n, \leq)$ is a CPO for any $n \in \mathbb{N}$, where $\leq$ is the pointwise order.*

Proof. First note that if $n = 0$, the proposition is trivial since the set has only one element, and every finite pointed poset is a CPO. It is also trivial for $n = 1$.

For $n > 1$, first note that $(A, \leq_A)$ is pointed with bottom element $(\perp_A, \dots, \perp_A)$. Next, we use Definition A.6 and consider only chains in A^n . Denote such a chain by

$$C = \{(a_{1,1}, a_{1,2}, \dots, a_{1,n}), (a_{2,1}, a_{2,2}, \dots, a_{2,n}), \dots (a_{i,1}, a_{i,2}, \dots, a_{i,n}), \dots\}$$

where $i \leq j$ implies that $a_{i,m} \leq_A a_{j,m}$ for all $m \in \{1, \dots, n\}$. Let $A_j = \{a_{1,j}, a_{2,j}, \dots, a_{i,j}, \dots\}$ for $j = 1, 2, \dots$. It is easy to see then that

$$\bigvee C = (\bigvee A_1, \bigvee A_2, \dots, \bigvee A_n).$$

Hence, the chain has a LUB. Since every such chain has a LUB, the poset is a CPO. □

The following proposition follows trivially using the same proof technique, and proves quite useful.

Proposition A.4. *Given a CPO $(A, \leq_A)$ and any set B , $(A^B, \leq)$ is a CPO, where $\leq$ is the pointwise order.*

The following example illustrates how we can pull together several of these ideas to begin addressing questions of program semantics.

Example A.17: Let B be a set of variable names in a computer program and let A be the set of values that those variables can take on. During the (possibly nonterminating) execution of the program, a variable $b \in B$ takes on a (possibly infinite) sequence of values. The sequence of values is a member of the set A^* . We can define the **semantics** of the program to be a

function $f: B \rightarrow A^{**}$ that for every $b \in B$ yields a sequence $f(b) \in A^{**}$. Since A^{**} is a CPO under the prefix order, by Proposition A.4 $(A^{**})^B$ is also a CPO under the **pointwise prefix order**.

Let a **partial execution** of the program be a function $f_i: B \rightarrow A^*$ for $i \in \mathbb{N}$. A partial execution always yields a finite sequence. If the execution of the program can be described as a chain of such partial executions (in the pointwise prefix order), then the limit of this chain in A^{**} can be taken to be the semantics f .

Describing an execution of a program in this way is natural for many programs. Consider a determinate sequential program that updates values for variables as it executes. Each such update appends a new value to the end of the sequence consisting of the previous values of that variable. The sequence of growing sequences of values of variables is clearly a chain in the prefix order.

A.4.4 Flat Partial Orders

Any set can be turned into a CPO by choosing a **flat order relation**. Given an arbitrary set A , augment the set with one additional element that serves as the bottom element. Specifically, let B be the augmented set and the one additional element be $\perp_B$, so $B = A \cup \{\perp_B\}$. Define a poset $(B, \leq)$, where the order relation is such that $\perp_B \leq a$ for all $a \in A$ and $a \leq a' \Rightarrow a = a'$ for all $a, a' \in A$. This is called a **flat partial order**. Any two distinct elements a and a' in A are incomparable. Moreover, every directed subset of B is a chain with either one or two elements. Trivially, each such chain has an upper bound. Hence, $(B, \leq)$ is a CPO. In this CPO, all chains are finite, so the notion of a limit of a chain is rather trivial.

Example A.18: An example of a flat partial order is illustrated by the [Hasse diagram](#) in Figure A.2(a), where $A = \{a_1, a_2\}$ and $B = \{\perp_B, a_1, a_2\}$.

Example A.19: Suppose that $(B, \leq_B)$ is the flat partial order of the previous example. Then $(B^2, \leq)$, where $\leq$ is the pointwise order, is illustrated by the Hasse diagram in Figure A.2(b). By Proposition A.3, $(B^2, \leq)$ is also a CPO.




Figure A.2: (a) Flat partial order for $B = \{\perp_B, a_1, a_2\}$. (b) Pointwise order on B^2 .

A.4.5 Lattices

Some partial orders have more structure than we have so far assumed.

Definition A.7. A **lattice** is a poset $(A, \leq)$ where any two elements $a, a' \in A$ have a unique greatest lower bound $a \wedge a' \in A$ and a unique least upper bound $a \vee a' \in A$.

Example A.20: Given a set of sets A , the poset $(A, \subseteq)$ formed with the **subset order** is a lattice. Given two sets $a, a' \in A$, their GLB is their intersection $a \cap a'$, and their LUB is the union $a \cup a'$. If A is a **powerset**, then this lattice is known as the **powerset lattice**.

Example A.21: The posets in Figure A.2 are not lattices. For example, in Figure A.2(a), the subset $\{a_1, a_2\}$ has no upper bound, much less a least upper bound. Adding an artificial **top** element (commonly denoted $\top$), would change these posets so that they both become lattices.

Some posets have some of the structure of lattices, but not all of it. A **lower semilattice** or **meet semilattice** is a poset $(A, \leq)$ where any two elements $a, a' \in A$ have a unique greatest lower bound $a \wedge a' \in A$.

Example A.22: The posets in Figure A.2 are lower semilattices. You can exhaustively verify that for any pair of elements, there is a single unique GLB.

Example A.23: The poset $(A^*, \sqsubseteq)$ of finite and infinite sequences from the set A is a lower semilattice under the prefix order. Any two sequences in A^* have a unique common prefix (which may be the empty sequence).

An **upper semilattice** or **join semilattice** is dually defined as a poset $(A, \leq)$ where any two elements $a, a' \in A$ have a unique least upper bound $a \vee a' \in A$.

A **complete lattice** is a poset $(A, \leq)$ where every subset (not just every pair of elements) has a LUB and GLB (and similarly for **complete lower semilattice** or **complete upper semilattice**). A complete lattice is also a CPO.

Example A.24: The powerset $2^{\mathbb{R}}$ of the reals with the **subset order**, $(2^{\mathbb{R}}, \subseteq)$ is a complete lattice. Given any set of sets of reals, the GLB is the intersection of the sets and the LUB is the union.

A.5 Functions on Posets

In this section, we define classes of functions on posets that have particularly useful properties.

A.5.1 Monotonic Functions

A function from a poset to a poset may preserve order.

Definition A.8. A function $f: A \rightarrow B$ from poset $(A, \leq_A)$ to poset $(B, \leq_B)$ is **monotonic** or **order preserving** if $a \leq_A a' \implies f(a) \leq_B f(a')$.

Example A.25: Consider a function $f: A^* \rightarrow A^*$ that is monotonic in the **prefix order**. This means that if a sequence s is a prefix of another s' , i.e. that $s \sqsubseteq s'$, then $f(s) \sqsubseteq f(s')$. Functions with this property have many special

qualities. The property implies that it is “safe” to evaluate the function with partial information about its input. Given only a prefix of what the input will eventually be, we can nonetheless evaluate the function using only the prefix, and we will get partial information (a prefix) about what the output will eventually be when the input is complete.

A function $f: A \rightarrow B$ is an **order embedding** if

$$a \leq_A a' \iff f(a) \leq_B f(a'). \quad (\text{A.3})$$

Note that an order embedding is monotonic, but a monotonic function is not necessarily an order embedding. An order embedding is also necessarily **one-to-one**. An order embedding $f: A \rightarrow B$ that is **onto** is called an **order isomorphism** from A to B . Since it is also **one-to-one**, an order isomorphism is necessarily a **bijection**. Two posets are **order isomorphic** if there is an order isomorphism from one to the other. Order isomorphism is a rather strong relationship between posets. The posets are essentially the same, in that one can be obtained from the other by just renaming the elements.

Example A.26: Let $A = \{0, 1, 2, \dots, 255\} \subset \mathbb{N}$ and $B = \mathbb{B}^8$. Here, B is the set of all sequences of 8 binary digits. Let $f: B \rightarrow A$ be a function that for each $b \in B$ yields $f(b) \in A$ that is the number represented by b when b is interpreted as a binary unsigned number with the high-order bit first. Assume $\mathbb{B}$ is ordered so that $0 < 1$ and that B is endowed with a **lexicographic order**. Assume A is endowed with the usual **numeric order**. Then f is an order isomorphism.

A.5.2 Continuous Functions

If we are considering functions on CPOs rather than arbitrary posets, then an even more useful property than monotonicity is continuity.

Definition A.9. A function $f: A \rightarrow B$ from **CPO** $(A, \leq_A)$ to **CPO** $(B, \leq_B)$ is **continuous** if for all chains $C \subseteq A$,

$$f(\bigvee C) = \bigvee f(C),$$

where $\hat{f}$ is the **lifted** version of f .

The **set of all continuous functions** from A to B is denoted $[A \rightarrow B]$, and is obviously a subset of the set B^A of all functions from A to B .

Proposition A.5. *Every continuous function $f: A \rightarrow B$, where $(A, \leq_A)$ and $(B, \leq_B)$ are CPOs, is monotonic.*

Proof. Consider any $a, a' \in A$ where $a \leq a'$. Let $C = \{a, a'\}$, a chain. Note that $\bigvee C = a'$. Since f is continuous, we know that

$$\bigvee f(C) = \bigvee \{f(a), f(a')\} = f(\bigvee C) = f(a').$$

Hence, $f(a) \leq f(a')$, so the function is monotonic. □

Not every monotonic function is continuous however.

Example A.27: Consider the CPO $(A, \leq)$, where $A = \mathbb{N} \cup \{\infty\}$ and $\leq$ is the **numeric order**. Let $f: A \rightarrow A$ be given by

$$f(a) = \begin{cases} 1 & \text{if } a \neq \infty \\ 2 & \text{otherwise} \end{cases}$$

This function is obviously monotonic. But it is not continuous. To see that, let $C = \mathbb{N}$ and note that $\bigvee C = \infty$. Hence, $f(\bigvee C) = 2$. However, $\hat{f}(C) = \{1\}$, because every element of C is finite. Hence, $\bigvee \hat{f}(C) = 1 \neq 2$. So the function is not continuous.

A.5.3 Posets of Functions

So far, we have considered posets of sets. Now we consider posets of functions.

Suppose $(A, \leq)$ and $(C, \leq)$ are CPOs. Consider two functions $f, g \in (A \rightarrow C)$. Define the **pointwise order** on these functions to be

$$f \leq g \iff f(a) \leq g(a) \quad \forall a \in A.$$

Let $[A \rightarrow C]$ be the **set of all continuous functions** from A to C .

Theorem A.1. For CPOs $(A, \leq)$ and $(C, \leq)$, let $X \subset [A \rightarrow C]$ be the set of all continuous total functions from A to C . Then $(X, \leq)$ is a CPO under the pointwise order $\leq$.

Proof.

First we need to show that X has a bottom element. This is easy. The bottom element is a function $g \in X$ where for all $a \in A$, $g(a) = \perp$. This function is obviously continuous and total, and hence is in X .

Second, we need to show that any chain of functions in X has a LUB and that the LUB is continuous. Consider a chain of functions

$$K = \{f_1, f_2, \dots\} \subset X.$$

Since each of these functions is continuous (and hence monotonic), then for any $a \in A$, the following set is also a chain,

$$K'_a = \{f_1(a), f_2(a), \dots\} \subset C.$$

Since C is a CPO, this set has a LUB. Define the function $g: A \rightarrow C$ such that for all $a \in A$,

$$g(a) = \bigvee K'_a.$$

Then in the pointwise order, it must be that

$$g = \bigvee K = \bigvee \{f_1, f_2, \dots\}.$$

It remains to show that g is in X . To show this, we must show that it is continuous. We must show that for all chains $D \subset A$,

$$g(\bigvee D) = \bigvee \hat{g}(D).$$

Writing the elements of $D = \{d_1, d_2, \dots\}$, observe that

$$\begin{aligned} \bigvee \hat{g}(D) &= \bigvee \{g(d_1), g(d_2), \dots\} \\ &= \bigvee \{\bigvee \{f_1(d_1), f_2(d_1), \dots\}, \bigvee \{f_1(d_2), f_2(d_2), \dots\}, \dots\} \\ &= \bigvee \{\bigvee \{f_1(d_1), f_1(d_2), \dots\}, \bigvee \{f_2(d_1), f_2(d_2), \dots\}, \dots\} \\ &= \bigvee \{\bigvee \hat{f}_1(D), \bigvee \hat{f}_2(D), \dots\} \\ &= \bigvee \{f_1(\bigvee D), f_2(\bigvee D), \dots\} \\ &= g(\bigvee D). \end{aligned}$$

Note that the above implicitly uses the axiom of choice, which states that given a set of sets, one can construct a new set by collecting one element from each of the sets in the set of sets. □

The bottom element of $(X, \leq)$ is the function $g \in X$ where for all $a \in A$, $g(a) = \perp$.

A.6 Fixed Points

Given an **endofunction** $f: A \rightarrow A$, if there is value $a \in A$ such that $f(a) = a$, that value is called a **fixed point**. Existence and uniqueness of fixed points play a central role in the semantics of programs. Existence of a fixed point will be interpreted as “the program has a meaning.” Uniqueness will be interpreted as “the program has no more than one meaning.” For functions that have multiple fixed points, we are often interested in the **least fixed point**, which is a fixed point $a \in A$ such that for any other fixed point $a' \in A$, $a \leq a'$. The following proposition assures the existence and uniqueness of such a least fixed point for continuous functions, and moreover gives a constructive procedure to determine that least fixed point.

Proposition A.6. Kleene fixed-point theorem. *For any monotonic function $f: A \rightarrow A$ on a CPO $(A, \leq)$, let*

$$C = \{f^n(\perp) : n \in \mathbb{N}\}.$$

Then if $\bigvee C = f(\bigvee C)$, $\bigvee C$ is the least fixed point of f . Moreover, if f is also continuous, then $\bigvee C = f(\bigvee C)$.

Proof. The first part of this theorem does not require that f be continuous, but only that it be monotonic. Suppose $\bigvee C = f(\bigvee C)$. This is a fixed point. Let a be any other fixed point, i.e. $f(a) = a$. We can show that $\bigvee C \leq a$, and hence $\bigvee C$ is the least fixed point. First, observe that $\perp \leq a$. Since f is monotonic, this implies that $f(\perp) \leq f(a) = a$. Again, since f is monotonic, this implies that $f(f(\perp)) \leq f(f(a)) = f(a) = a$. Continuing in this fashion, for any $n \in \mathbb{N}$,

$$f^n(\perp) \leq f^n(a) = a.$$

Hence, a is an upper bound of C . Since $\bigvee C$ is the least upper bound of C , it follows that $\bigvee C \leq a$, and hence $\bigvee C$ is the least fixed point of f .

For the second part of this theorem, we require that f be continuous, and not just monotonic. First, we observe that C is a chain in the CPO $(A, \leq)$. To see that, note that $\perp \leq f(\perp)$. Since f is monotonic, this implies that $f(\perp) \leq f(f(\perp))$. Continuing, we see that for all $n \in \mathbb{N}$, $f^n(\perp) \leq f^{n+1}(\perp)$, so C is a chain. Since C is a chain, it has a LUB $\bigvee C$.

Next note that $\hat{f}(C) \cup \{\perp\} = C$. Moreover, $\bigvee(\hat{f}(C) \cup \{\perp\}) = \bigvee \hat{f}(C)$ (an additional $\perp$ in a chain will not change its least upper bound). Combining these two facts, we conclude that $\bigvee \hat{f}(C) = \bigvee C$. But since f is continuous, we also know that $\bigvee \hat{f}(C) = f(\bigvee C)$. Hence, $f(\bigvee C) = \bigvee C$, and hence $\bigvee C$ is a fixed point. $\square$

This theorem has profound consequences. It states that for a continuous function f , the least fixed point of this function can be found by first evaluating $f(\perp)$, then $f(f(\perp))$, then $f^3(\perp)$, etc. The result of this sequence of evaluations is a chain in A , and since A is a CPO, this chain has a limit. That limit is the least fixed point of the function. Hence, this theorem offers a **constructive procedure** for finding the least fixed point of a continuous function. The ascending chain $C = \{f^n(\perp) : n \in \mathbb{N}\}$ is known as the **Kleene chain** of the function f . Applications of this theorem are scattered throughout the text.

A.7 Cardinality

In mathematics, the **cardinality** of a set is a measure of the number of elements of the set. For example, set $A = \{1, 2, 3, 4\}$ has four elements, and therefore A has a cardinality of 4. More interestingly, the set $\mathbb{N}$ has an infinite number of elements, and so does the set $\mathbb{R}$. And yet, in a rather strong sense, $\mathbb{N}$ has fewer elements than $\mathbb{R}$. Even more interestingly, although we can show that the set $\mathbb{N}$ and $\mathbb{Z}$ have the same number of elements, the elements of $\mathbb{N}$ can be listed in numeric order, but the elements of $\mathbb{Z}$ cannot.

The study of semantics benefits considerably from a solid understanding of the cardinality of sets. For example, we compare the number of programs expressible in a particular programming language to the number of functions of form $f: \mathbb{N} \rightarrow$


IN and show that the former is distinctly smaller than the latter. It then follows immediately that not all such functions can be expressed by computer programs in that language.

More particular to concurrent systems, it is useful to understand how many behaviors of programs are enabled by concurrency. It is also useful to understand how to distinguish finite behaviors (ones that either terminate or deadlock) from infinite ones.

Sidebar: Fixed Point Theorems

The [Kleene fixed-point theorem](#), Proposition A.6, is named after American mathematician Stephen Cole Kleene (1909–1994). This theorem is often attributed to Alfred Tarski, (1901–1983), a Polish-American logician and mathematician and a professor of mathematics at the University of California, Berkeley. **Tarski’s fixed-point theorem** is similar to the Kleene fixed-point theorem, but in its original statement, it is about monotone functions on [complete lattices](#). Because of this attribution, techniques used in computer science that are based on fixed-point theorems are often called **Tarskian**.

Another well-known theorem in this family is the **Knaster-Tarski fixed-point theorem**, developed earlier by Tarski and Bronislaw Knaster. This theorem is a special case of Tarski’s fixed-point theorem that applies to the [powerset lattice](#).



Stephen Cole Kleene (1909–1994). Photo by Konrad Jacobs, Erlangen, copyright (1978) MFO, Mathematisches Forschungsinstitut Oberwolfach, licensed under the Creative Commons Attribution-Share Alike 2.0 Germany license.

A.7.1 Countable and Uncountable Sets

A set A is **countable** if there exists a **one-to-one** function $f: A \rightarrow \mathbb{N}$. Such a function associates with each element $a \in A$ a unique number $n \in \mathbb{N}$.

If there is a function that is one-to-one and also **onto** (and hence a **bijection**), then the set A is **countably infinite**. To determine whether a set is countable or countably infinite, we simply need to find such functions or show that none exists.

Example A.28: The set of integers $\mathbb{Z}$ is countable. To show this, let $f: \mathbb{Z} \rightarrow \mathbb{N}$ be defined by

$$\forall n \in \mathbb{Z}, \quad f(n) = \begin{cases} 2n - 1 & \text{if } n > 0 \\ -2n & \text{otherwise} \end{cases}$$

This function is easily shown to be a bijection, and hence $\mathbb{Z}$ is countably infinite.

Given two functions $f: A \rightarrow B$ and $g: B \rightarrow \mathbb{N}$ that are known to be one-to-one (or onto), the composition $f \circ g$ is also one-to-one (or onto). Hence, if a set B is known to be countable (countably infinite), then if we find a function $f: A \rightarrow B$ that is one-to-one (or onto), then we can conclude that A is countable (or countably infinite).

Example A.29: The set of finite sequences $\mathbb{B}^*$ of binary digits $\mathbb{B}$ is countable. To show this, first let $\mathbb{N}_+ = \mathbb{N} \setminus \{0\}$ and note that it is easy to show that $\mathbb{N}_+$ is countably infinite. Next, let $g: \mathbb{B}^* \rightarrow \mathbb{N}_+$ be the function that yields the number represented by the sequence of binary digits as a digital encoding as follows

$$\forall b = (b_0, b_1, \dots, b_{n-1}) \in \mathbb{B}^*, \quad f(b) = 2^n + \sum_{i=0}^{n-1} b_i 2^{(n-i-1)}.$$

Specifically,

$$\begin{aligned}
 f(\lambda) &= 2^0 = 1 \\
 f((0)) &= 2^1 + 0 \cdot 2^0 = 2 \\
 f((1)) &= 2^1 + 1 \cdot 2^0 = 3 \\
 f((0,0)) &= 2^2 + 0 \cdot 2^0 + 0 \cdot 2^1 = 4 \\
 &\text{etc.}
 \end{aligned}$$

It is easy to see that this function is one-to-one (any two distinct bit patterns yield a different positive natural number). It is also easy to see that it is onto (any positive natural number has such a binary representation). Hence, this function is a bijection. Since $\mathbb{N}_+$ is countably infinite, so is $\mathbb{B}^*$.

A bijection $f: A \rightarrow \mathbb{N}$ can be viewed as a list of the elements of the set A , since it assigns a unique natural number to each. It is said to **enumerate** the set. The enumeration is a list, possibly infinite, that includes all elements of A . Moreover, if the function is only one-to-one and not onto, it can still enumerate A by considering the **numeric order** of each $f(a)$ for $a \in A$ as giving the position of a in the list. This list may be finite if the set is countable but not countably infinite. Thus we see that to show that a set A is countable, we need simply give a way to enumerate the elements in A .

Proposition A.7. *The set $\mathbb{N}^*$ of finite sequences of natural numbers is countable.*

Proof. To enumerate $\mathbb{N}^*$, let $f: \mathbb{N}^* \rightarrow \mathbb{N}$ be a function defined by

$$f(s) = \sum_{i=0}^{n-1} (s_i + 1)$$

for any sequence $s = (s_0, s_1, \dots, s_{n-1}) \in \mathbb{N}^*$ of length n . Note that for each $j \in \mathbb{N}$, there are a finite number of sequences s in $\mathbb{N}^*$ where $f(s) = j$. For $j = 0$, only the empty sequence $s = \lambda$ yields $f(s) = 0$. For $j = 1$, only the sequence $s = (0)$ results in $f(s) = 1$. For $j = 2$, the sequences (1) and $(0, 0)$ yield $f(s) = 2$. For $j = 3$, we get (2) , $(1, 0)$, $(0, 1)$, and $(0, 0, 0)$. Continuing, it is easy to see that we can enumerate the set $\mathbb{N}^*$ by listing first the sequences s where $f(s) = 0$, then the ones where $f(s) = 1$, then the ones where $f(s) = 2$, etc. Since every sequence s yields a finite $f(s)$, every sequence s is included in this list. Hence, $\mathbb{N}^*$ is countable.

□

If a set A is countable, then any subset $B \subseteq A$ is also countable. More generally, if a set A is countable, then any set B for which there exists a one-to-one function $f: B \rightarrow A$ is also countable. This is evident because the existence of such a function shows that A is at least as large as B , since every $b \in B$ must yield a unique $a = f(b) \in A$. These facts make it easy to prove the following rather useful proposition.

Proposition A.8. *The set B^* of finite sequences of elements from any countable set B is countable.*

Proof. Since B is countable, there is a one-to-one function $g: B \rightarrow \mathbb{N}$. From this function, we can easily construct a one-to-one function $f: B^* \rightarrow \mathbb{N}^*$ as follows,

$$f(x) = (g(b_0), g(b_1), \dots, g(b_m))$$

for all $x = (b_0, b_1, \dots, b_m) \in B^*$. Since this function is one-to-one, and by Proposition A.7 $\mathbb{N}^*$ is countable, it follows that B^* is countable. □

A set that is not countable is **uncountable**. Let $\mathbb{B}^\omega$ be the set of all infinite sequences of bits.

Proposition A.9. *$\mathbb{B}^\omega$ is uncountable.*

Proof. We give a variant of **Cantor's diagonal argument**, which is a proof by contradiction. Assume that $\mathbb{B}^\omega$ is countable. That is, it can be enumerated. Let the enumeration be written $b_1, b_2, b_3, \dots$. For any $n \in \mathbb{N}$, let $b_i(n)$ denote the n -th bit of $b_i \in \mathbb{B}^\omega$. Let $c \in \mathbb{B}^\omega$ be a bit sequence where

$$c(n) = \overline{b_n(n)}$$

where $\bar{b}$ is 0 if $b = 1$ and 1 if $b = 0$. That is, the n -th bit of c is the complement of the n -th bit of the n -th element of the enumeration. Note therefore that c is not in the enumeration, so the enumeration does not include all the elements of $\mathbb{B}^\omega$. Hence, $\mathbb{B}^\omega$ is not countable.

□

Proposition A.10. *A^ω is uncountable for any set A with two or more elements.*

Proof. Let $a, a' \in A$ be two distinct elements. Let $B = \{a, a'\}$. Note that $B^\omega \subseteq A^\omega$. By the same proof technique as in Proposition A.9, B^ω is uncountable. Since A^ω is at least as large, it too is uncountable.

□

An immediate consequence is that the powerset $2^\mathbb{N}$ of $\mathbb{N}$ is uncountable (because in the notation $2^\mathbb{N}$, 2 represents a set with two elements, as discussed in the box on page 308). Moreover, given any set A that is at least as large as $\mathbb{N}$ (i.e. any infinite set, countable or not, or equivalently, any set for which there is an onto function $f: A \rightarrow \mathbb{N}$), the powerset 2^A is uncountable.

A.8 Discrete Sets

Recall that a countable set is one for which there exists a one-to-one function $f: A \rightarrow \mathbb{N}$. A **discrete set** is a poset $(A, \leq_A)$ for which there exists such a function that is also an **order embedding** to $(\mathbb{N}, \leq)$, where $\leq$ is the **numeric order**. A countable set can be enumerated. A discrete set has an order relation and can be enumerated *in order*. Not all countable sets are discrete.

Proposition A.11. *The set of integers $\mathbb{Z}$ with the numeric order is countable but not discrete.*

Proof. To show this, we must show that there is no order embedding $f: \mathbb{Z} \rightarrow \mathbb{N}$. Suppose to the contrary that there is such an order embedding. The **image** $\hat{f}(\mathbb{Z})$ is a subset of $\mathbb{N}$, and hence must have a least element. Call that least element $m \in \mathbb{N}$ and let $i \in \mathbb{Z}$ be such that $f(i) = m$. Note that $i - 1 \in \mathbb{Z}$, and by (A.3), since $i - 1 \leq i$, it must also be that $f(i - 1) \leq f(i)$. Since $f(i) = m$ is the least element, it must be true that $f(i - 1) = f(i)$. But since f is required to be one-to-one, this implies that $i - 1 = i$, a contradiction.

□

The same proof method suffices to show the following:

Proposition A.12. *Every non-empty discrete poset has a least element.*

As a consequence, the set $\mathbb{Q}$ of rational numbers is not discrete. But it is not sufficient to have a least element to be discrete. The set of non-negative rational numbers, for example, has a least element, but it too is not discrete. To show that, we consider the subset containing only the *positive* rational numbers, which does not have a least element, and leverage the following proposition.

Proposition A.13. *If a poset $(A, \leq)$ is discrete and $B \subseteq A$, then $(B, \leq)$ is discrete.*

Proof. Given an order embedding $f: A \rightarrow \mathbb{N}$, it is easy to see that the [restriction](#) $f \upharpoonright B$ is also an order embedding. □

To show that the set of non-negative rationals is not discrete, knowing that the set of positive rationals is not discrete, we use the contrapositive of this proposition:

Proposition A.14. *If a poset $(B, \leq_B)$ is not discrete, then neither is any larger set $A \supseteq B$ with an order $\leq_A$ satisfying $a \leq_A a' \Leftrightarrow a \leq_B a'$ for all $a, a' \in B$.*

Specifically, let B be the set of positive rationals, which by [Proposition A.12](#) is not discrete. Then let A be the non-negative rationals, a superset. By [Proposition A.14](#), that set is not discrete.

A.9 Computability

In the classical theory of computing, a **computation** is a mapping from inputs to outputs. Given inputs, the computation produces outputs. Suppose that an input x and output y are represented as a sequence of binary digits, so $x, y \in \mathbb{B}^*$. Then it would seem that the notion of computation is the same as the mathematical notion of a function $f: \mathbb{B}^* \rightarrow \mathbb{B}^*$. But it turns out that this is not quite the case. It turns out that most such functions are not [effectively computable](#).

We can use cardinality results to study the problem of **computability**, which is about what functions can be implemented by computer programs (or more fundamentally, by algorithms that terminate with an answer).

Consider the set of all **decision problems** of the form $f: \mathbb{N} \rightarrow \mathbb{B}$. For each natural number $n \in \mathbb{N}$, the function f yields a 0 or 1.

Example A.30: Consider a function f that produces $f(n) = 0$ if n is even and $f(n) = 1$ if n is odd. We can write a computer program in any modern programming language that will implement this function. Note that this is not entirely trivial, because the input is *any* natural number, with no bound on its size. If the input is given as a bit sequence (a member of $\mathbb{B}^*$), a binary representation of the number where the first bit is the low-order bit, then the program *is* trivial. It outputs the first bit.

Since every $f \in \mathbb{B}^{\mathbb{N}}$ is a decision problem, by Proposition A.9, there are an uncountable number of such decision problems.

Consider the set C of all computer programs written in some programming language. Every such program can be represented by a finite number of bits, so $C \subseteq \mathbb{B}^*$. By Proposition A.8, $\mathbb{B}^*$ is countable, and hence C is countable. Hence, C is strictly smaller than the set of decision problems.

As a consequence, for any programming language, there are decision problems that cannot be realized by a program in that language. Such problems are said to be **undecidable**.

A.10 Conclusion

This chapter has introduced some of the mathematical tools that we will need to study concurrency. The most essential of these is the Kleene fixed-point theorem, which gives a constructive procedure (the Kleene chain) for finding the least fixed point of a continuous function. This result proves to have an astonishing variety of applications, many of which are explored in the main chapters. In particular, it gives a deterministic semantics to certain concurrent models of computation. It also suggests an execution policy for such models of computation.



Metric Spaces

B.1	Varieties of Metric Spaces	332
B.1.1	Traditional Metric Spaces	332
B.1.2	Ultrametric Spaces	333
B.1.3	Generalized Ultrametric Spaces	333
B.2	Metrics on Functions	334
B.2.1	The Cantor Metric	334
B.2.2	The Baire Distance	336
B.2.3	Contracting Functions	337
B.2.4	Lipschitz Continuous Functions	338
B.2.5	Convergence of Sequences	338
B.2.6	Complete Metric Spaces	339
B.2.7	Spherically Complete Generalized Ultrametric Spaces	339
B.3	Fixed-Point Theorems	340
B.3.1	Strictly Contracting Functions	340
B.3.2	Delta Contracting Functions	341
B.3.3	Fixed Points on Generalized Ultrametric Spaces	343
B.4	Conclusion	344

This appendix provides background material on metric spaces, which are sets together with a **distance function** that is used to determine how close members of the set are to one another. Metric spaces are used in Chapter 7 to give a fixed-point semantics to discrete-event systems. We give several variants of metric spaces because each of these variants is useful in a different context.

B.1 Varieties of Metric Spaces

B.1.1 Traditional Metric Spaces

Definition B.1. A *metric space* (X, d) is a set X together with a *metric*, which is a function

$$d: X \times X \rightarrow \mathbb{R}$$

such that for all $x, y, z \in X$,

1. $d(x, y) = 0 \iff x = y$ (identity of indiscernables)
2. $d(x, y) = d(y, x)$ (symmetry)
3. $d(x, z) \leq d(x, y) + d(y, z)$ (triangle inequality)

Intuitively, the distance between two points in a metric space should not be negative. Fortunately, this follows from the definition of a metric space.

Proposition B.1. For a metric space (X, d) , for all $x, y \in X$, $d(x, y) \geq 0$.

Proof. Letting $z = x$ in the triangle inequality, we get

$$d(x, y) + d(y, x) \geq d(x, x).$$

The identity of indiscernables now implies that

$$d(x, y) + d(y, x) \geq 0.$$

By symmetry,

$$2d(x, y) \geq 0$$

or

$$d(x, y) \geq 0.$$

□

B.1.2 Ultrametric Spaces

An ultrametric is similar, but replaces the triangle inequality with a stronger condition.

Definition B.2. An *ultrametric* on a set X is a function $d: X \times X \rightarrow \mathbb{R}_+$ such that for all $x, y, z \in X$:

1. $d(x, y) = 0 \iff x = y$ (identity of indiscernables)
2. $d(x, y) = d(y, x)$ (symmetry)
3. $d(x, z) \leq \max(d(x, y), d(y, z))$ (ultrametric inequality)

Proposition B.2. An ultrametric is a metric.

Proof. The first two properties of an ultrametric are identical to the first two properties of a metric. So we only need to show that the third property of an ultrametric implies the third property of a metric. Since all ultrametric takes values in $\mathbb{R}_+$, all distances are non-negative. Since all the distances are non-negative,

$$d(x, y) + d(y, z) \geq \max(d(x, y), d(y, z)).$$

Combining this with the third ultrametric property, we get

$$d(x, y) + d(y, z) \geq d(x, z).$$

Hence, an ultrametric is also a metric. □

B.1.3 Generalized Ultrametric Spaces

Metrics and ultrametrics measure distances using non-negative real numbers. An alternative explored by [Priess-Crampe and Ribenboim \(1999\)](#) is to use any [pointed poset](#) for the measurement.

Definition B.3. A *generalized ultrametric space* (X, D, Δ) is a set X with a *generalized ultrametric distance function* $D: X \times X \rightarrow \Delta$, where Δ is *pointed poset*, such that for all $x, y, z \in X$,

1. $D(x, y) = \perp_\Delta \iff x = y$ (identity of indiscernables)
2. $D(x, y) = D(y, x)$ (symmetry)
3. $D(x, z) \leq D(x, y) \vee D(y, z)$ (generalized ultrametric inequality)

where $\perp_\Delta$ is least element in Δ , and $\vee$ denotes the *join* or *least upper bound* in the poset.

B.2 Metrics on Functions

In this book, we will need metrics that measure the similarity of *functions* rather than arbitrary members of a set. For example, two *sequences* that have a long common *prefix* are more similar than two sequences with a short common prefix. Likewise, two functions of time that are identical for longer are more similar than two functions of time that differ earlier.

B.2.1 The Cantor Metric

Let $X = (\mathbb{R}_+ \rightarrow V)$ be the set of functions from the non-negative reals to some value set. This set might be used for functions of time.

Definition B.4. For any two functions $s_1, s_2 \in (\mathbb{R}_+ \rightarrow V)$, let $P \subseteq \mathbb{R}_+$ be defined by

$$P = \{t \in \mathbb{R}_+ \mid s_1(t) \neq s_2(t)\}.$$

That is, P is the set of times at which s_1 and s_2 differ. Let τ be the greatest lower bound of P , $\tau = \bigwedge P$, where we define $\bigwedge \emptyset = \infty$. Then the **Cantor metric** is

$$d(s_1, s_2) = 2^{-\tau}.$$

If the domain $\mathbb{R}_+$ represents time, then τ is either the earliest time at which s_1 and s_2 differ, or, if there is no such earliest time, the latest time at which they are continuously identical from time zero.

Example B.1: Let $V = \mathbb{R}$ and let s_1 and s_2 be the functions

$$s_1(t) = \begin{cases} 1 & \text{if } t < 1 \\ 0 & \text{if } t \geq 1 \end{cases}$$

$$s_2(t) = 1$$

The earliest time at which these differ is $\tau = 1$, so $d(s_1, s_2) = 1/2$.

Example B.2: Let $V = \mathbb{R}$ and let s_1 and s_2 be the functions

$$s_1(t) = \begin{cases} 1 & \text{if } t \leq 1 \\ 0 & \text{if } t > 1 \end{cases}$$

$$s_2(t) = 1$$

There is no earliest time at which they differ, but they are continuously identical from time zero until time 1. Hence, $\tau = 1$, so $d(s_1, s_2) = 1/2$.

Proposition B.3. *The Cantor metric is an ultrametric.*

Proof. First, note that if $s_1 = s_2$, then $\tau = \infty$, and hence $d(s_1, s_2) = 0$, so we have [identity of indiscernibles](#). It is also obvious that $d(s_1, s_2) = d(s_2, s_1)$, so we have [symmetry](#).

To show that the [ultrametric inequality](#) is satisfied, we need to show that for all $s_1, s_2, s_3 \in (\mathbb{R}_+ \rightarrow V)$,

$$d(s_1, s_3) \leq \max(d(s_1, s_2), d(s_2, s_3)).$$

Assume that τ is such that $2^{-\tau} = \max(d(s_1, s_2), d(s_2, s_3))$. This means that s_1 and s_2 are identical up to time τ (maybe including time τ), and s_2 and s_3 are also identical up to time τ , at least. Hence, it follows that s_1 and s_3 are identical up to time τ , so $d(s_1, s_3) \leq 2^{-\tau}$, from which the result follows. $\square$

Note that if $s_1(0) \neq s_2(0)$, then $\tau = \bigwedge P = 0$, so $d(s_1, s_2) = 1$. Hence, by the Cantor metric always lies between zero and one for functions defined on the non-negative reals.

Consider functions defined over all the reals. I.e., let $X = (\mathbb{R} \rightarrow V)$. For such functions, the distance between two functions can vary from zero to infinity.

Example B.3: Let $X = (\mathbb{R} \rightarrow \mathbb{R})$ and let s_1 and s_2 be the functions

$$s_1(t) = t$$

$$s_2(t) = t + 1$$

The two functions are never equal, so $P = \mathbb{R}$ and $\tau = \bigwedge P = -\infty$. Hence, $d(s_1, s_2) = 2^\infty = \infty$. With the Cantor metric, these two functions are infinitely far apart.

Technically, this X with the Cantor metric is not a metric space because metric spaces require the distance to be a real number. This distinction is important because attempting to apply the fixed-point theorems below to functions on this space will not work. See [Liu et al. \(2006a\)](#) for a discussion of this issue.

B.2.2 The Baire Distance

An untimed version of the [Cantor metric](#) for sequences is known as the **Baire distance**.

Definition B.5. With $A = D^{**}$, where D^{**} is the set of all finite or infinite sequences of elements of D , the **Baire distance** is a function $d: A \times A \rightarrow \mathbb{R}_+$ be such that for all $a, b \in A$,

$$d(a, b) = 1/(n + 1)$$

where n is the index of the first position at which a and b differ, and $n = \infty$ if $a = b$ (giving $d(a, b) = 0$ if $a = b$). The first element of a sequence $a \in A$ has index 0, the second has index 1, etc.

Example B.4: For example, $d((0), (1)) = 1$ and $d((0), (0, 1)) = 1/2$.

Proposition B.4. *The Baire distance is an **ultrametric**.*

Proof. The first two conditions of an ultrametric (Definition B.2) are trivially satisfied. We can show that the third condition is satisfied, and therefore d is an ultrametric.

Let n_1 be such that $\max(d(a, b), d(b, c)) = 1/(n_1 + 1)$. This means that a and b are identical at least up to index $n_1 - 1$, and b and c are also identical at least up to index $n_1 - 1$. Hence it must be true that a and c are identical at least up to index $n_1 - 1$.

Let n_2 be such that $d(a, c) = 1/(n_2 + 1)$. Hence n_2 is the earliest index at which a and c differ. Since a and c are identical at least up to index $n_1 - 1$, it must be true that

$$n_2 \geq n_1.$$

Hence

$$1/(n_1 + 1) \geq 1/(n_2 + 1),$$

which is what we needed to show. □

B.2.3 Contracting Functions

A certain class of functions on a metric space have useful fixed-point properties.

Definition B.6. A function $F: X \rightarrow X$ on a metric space (X, d) is a **contracting function** or a **contraction** if for all $x, y \in X$,

$$d(F(x), F(y)) \leq d(x, y).$$

The function is **strictly contracting** if for all $x, y \in X, x \neq y$,

$$d(F(x), F(y)) < d(x, y).$$

The function is **delta contracting** if there exists a non-negative real number $\delta < 1$ such that for all $x, y \in X, x \neq y$,

$$d(F(x), F(y)) \leq \delta d(x, y).$$

The first two definitions also apply to [generalized ultrametric spaces](#) (see Section [B.1.3](#)). However, there is no analog of the delta contracting definition for generalized ultrametric spaces.

B.2.4 Lipschitz Continuous Functions

Definition B.7. A function $F: X \rightarrow Y$ between metric spaces (X, d_X) and (Y, d_Y) is **Lipschitz continuous** if there exists a non-negative real number K such that for all $x, x' \in X$,

$$d_Y(F(x), F(x')) \leq K d_X(x, x').$$

The following proposition follows trivially by letting $K = \delta$ in Definition [B.7](#).

Proposition B.5. A delta contracting function is Lipschitz continuous.

B.2.5 Convergence of Sequences

The metric in a metric space makes it easy to formalize the notion of **convergence**. A sequence of points in a metric space converges to a point if the distance between the points and the limit point becomes arbitrarily small.

Definition B.8. A sequence of points $x_1, x_2, \dots \in X$ in a metric space (X, d) **converges** to a point $x \in X$ if for all $\epsilon > 0$, there exists an integer N such that for all $n \geq N$,

$$d(x_n, x) < \epsilon.$$

B.2.6 Complete Metric Spaces

A **complete metric space** is one that contains all its limit points.

Definition B.9. A sequence of points $x_1, x_2, \dots \in X$ in a metric space (X, d) is **Cauchy** if for all $\epsilon > 0$, there exists an integer N such that for all $n, m \geq N$,

$$d(x_n, x_m) < \epsilon.$$

Definition B.10. A metric space is **complete** if every Cauchy sequence converges to a point in the space.

B.2.7 Spherically Complete Generalized Ultrametric Spaces

For the [generalized ultrametric spaces](#) defined in Section B.1.3, we can define a stronger notion of completeness. The most interesting fixed-point theorems rely on this stronger notion of completeness.

Definition B.11. Given a generalized ultrametric space (X, D, Δ) , the **spherically closed ball** of radius $\delta \in \Delta, \delta \neq \perp_\Delta$, is the set

$$B_\delta(x) = \{y \in X \mid D(x, y) \leq \delta\}.$$

Proposition B.6. For a generalized ultrametric space (X, D, Δ) , consider $x, y \in X$ where $y \in B_\delta(x)$ for some $\delta \in \Delta, \delta \neq \perp_\Delta$. For any $\alpha \in \Delta$ such that $\perp_\Delta < \alpha \leq \delta$, we have $B_\alpha(y) \subseteq B_\delta(x)$.

Proof. Because $y \in B_\delta(x)$, we have $D(x, y) \leq \delta$. Consider an $\alpha \in \Delta$ such that $\perp_\Delta < \alpha \leq \delta$. Consider any $z \in B_\alpha(y)$. Then $D(y, z) \leq \alpha$. By the generalized

ultrametric inequality,

$$D(x, z) \leq D(x, y) \vee D(y, z) \leq \delta \vee \alpha \leq \delta.$$

Hence, $z \in B_\delta(x)$. Therefore, $B_\alpha(y) \subseteq B_\delta(x)$. □

If we choose $\alpha = \delta$, then we draw the (rather odd) conclusion that every point in a ball is also the center of the ball.

Definition B.12. A *chain of closed balls* is a set of *spherically closed balls* $\{B_{\delta_i}(x_i) \mid i \in I\}$, for some totally ordered index set I , such that for all $i, j \in I$, $i \leq j$ implies $B_{\delta_i}(x_i) \supseteq B_{\delta_j}(x_j)$.

Definition B.13. A generalized ultrametric space is *spherically complete* if every chain of closed balls has a non-empty intersection.

B.3 Fixed-Point Theorems

There are several useful fixed-point theorems for contracting functions on metric spaces.

B.3.1 Strictly Contracting Functions

Theorem B.1. For a *strictly contracting* function $F: X \rightarrow X$ on a metric space (X, d) , if F has a fixed point, then it is unique.

Proof. Suppose that x and y , $x \neq y$, are two fixed points of F . Because F is strictly contracting,

$$d(F(x), F(y)) < d(x, y).$$

Because x and y are fixed points of F ,

$$d(F(x), F(y)) = d(x, y).$$

Hence,

$$d(x, y) < d(x, y).$$

This is a contradiction, so F cannot have two distinct fixed points. □

B.3.2 Delta Contracting Functions

The previous theorem is not constructive, in the sense that it does not provide a procedure to find the fixed point. The following theorem is constructive, but it only applies to delta contracting functions.

Theorem B.2. (Banach fixed-point theorem) *A delta contracting function $F: X \rightarrow X$ on a complete metric space (X, d) has a unique fixed point, and the following sequence converges to that fixed point:*

$$\begin{aligned} x_1 &= F(x_0) \\ x_2 &= F(x_1) \\ &\vdots \\ x_{n+1} &= F(x_n), \quad n \geq 1 \end{aligned}$$

for any starting point $x_0 \in X$.

Proof. Because F is delta contracting,

$$d(F(x_{n+1}), F(x_n)) \leq \delta d(x_{n+1}, x_n).$$

By induction, we can show that for all $n \in \mathbb{N}$,

$$d(x_{n+1}, x_n) \leq \delta^n d(x_1, x_0).$$

We can use this to show that the sequence is Cauchy. Consider two $n, m \geq N$, $n > m$. By the triangle inequality,

$$d(x_n, x_m) \leq d(x_n, x_{n-1}) + d(x_{n-1}, x_{n-2}) + \cdots + d(x_{m+1}, x_m).$$

Combining these,

$$\begin{aligned}
 d(x_n, x_m) &\leq \delta^{n-1}d(x_1, x_0) + \delta^{n-2}d(x_1, x_0) + \cdots + \delta^m d(x_1, x_0) \\
 &= \delta^m d(x_1, x_0) \sum_{i=0}^{n-m-1} \delta^i \\
 &\leq \delta^m d(x_1, x_0) \sum_{i=0}^{\infty} \delta^i \\
 &= \delta^m d(x_1, x_0) \frac{1}{1 - \delta} \\
 &= \frac{\delta^m}{1 - \delta} d(x_1, x_0).
 \end{aligned}$$

Let $\epsilon > 0$ be given. Let N be such that for all $m \geq N$,

$$\frac{\delta^m}{1 - \delta} d(x_1, x_0) < \epsilon.$$

Then for all $n, m \geq N$,

$$d(x_n, x_m) < \epsilon.$$

Hence, the sequence is Cauchy. Because the space is complete, the sequence converges to a point $x \in X$. Let

$$x = \lim_{n \rightarrow \infty} x_n.$$

Then,

$$x = \lim_{n \rightarrow \infty} F(x_n) = F(\lim_{n \rightarrow \infty} x_n) = F(x),$$

where interchanging the limit and the function is justified because F is Lipschitz continuous (a consequence of the fact that F is delta contracting and Proposition B.5). Hence, x is a fixed point of F .

Moreover, because $\delta < 1$, F is strictly contracting, and hence the fixed point x is unique by Theorem B.1. □

B.3.3 Fixed Points on Generalized Ultrametric Spaces

For a [generalized ultrametric space](#) (see Section B.1.3), we cannot use the Banach fixed-point theorem (Theorem B.2) because there is no analog of a delta contracting function.

Theorem B.3. (*Priess-Crampe and Ribenboim (1999)*) *A strictly contracting function $F: X \rightarrow X$ on a spherically complete generalized ultrametric space (X, D, Δ) has a unique fixed point.*

Proof. This proof is adapted from Section 5.2 of [Priess-Crampe and Ribenboim \(1999\)](#).

Existence. For each $x \in X$, write $\pi_x = D(x, F(x))$, and suppose, for the sake of contradiction, that F has no fixed point. Then $F(x) \neq x$, and hence $\pi_x \neq \perp_\Delta$, for every $x \in X$. For each $x \in X$ we can therefore form the [spherically closed ball](#) $B_x = B_{\pi_x}(x)$, and we let

$$\mathcal{B} = \{B_x \mid x \in X\}$$

be the collection of all such balls, partially ordered by inclusion. By Zorn's lemma, $\mathcal{B}$ contains a maximal chain C (a maximal totally ordered subset). Since C is a [chain of closed balls](#) and X is spherically complete, the intersection $\bigcap_{B_x \in C} B_x$ is non-empty; choose a point z in it.

We first show that $B_z \subseteq B_x$ for every $B_x \in C$. Fix $B_x \in C$. Because $z \in B_x$, we have $D(z, x) \leq \pi_x$. Because F is contracting (a strictly contracting function is contracting), $D(F(x), F(z)) \leq D(x, z) \leq \pi_x$. Applying the generalized ultrametric inequality twice,

$$\pi_z = D(z, F(z)) \leq D(z, x) \vee D(x, F(x)) \vee D(F(x), F(z)) \leq \pi_x,$$

since each of the three terms is at most π_x (recall $D(x, F(x)) = \pi_x$). In particular $\pi_z \neq \perp_\Delta$, and $z \in B_{\pi_x}(x)$ with $\perp_\Delta < \pi_z \leq \pi_x$, so Proposition B.6 gives

$$B_z = B_{\pi_z}(z) \subseteq B_{\pi_x}(x) = B_x.$$

Since B_z is comparable with (indeed contained in) every element of C , the maximality of C forces $B_z \in C$; thus B_z is the smallest element of C .

Now consider the point $F(z)$. Since $D(z, F(z)) = \pi_z$, we have $F(z) \in B_z$. Because F is strictly contracting and $F(z) \neq z$,

$$\pi_{F(z)} = D(F(z), F(F(z))) < D(z, F(z)) = \pi_z.$$

As $F(z) \in B_{\pi_z}(z)$ and $\perp_\Delta < \pi_{F(z)} < \pi_z$, Proposition B.6 gives $B_{F(z)} \subseteq B_z$. This inclusion is proper, because $z \notin B_{F(z)}$: indeed $D(F(z), z) = \pi_z > \pi_{F(z)}$. But then $B_{F(z)}$ is comparable with every element of C and is strictly smaller than the least element B_z of C , so $C \cup \{B_{F(z)}\}$ is a strictly larger chain, contradicting the maximality of C . This contradiction shows that F must have a fixed point.

Uniqueness. This follows exactly as in Theorem B.1. If z and y were two distinct fixed points, then strict contraction would give

$$D(z, y) = D(F(z), F(y)) < D(z, y),$$

which is impossible. Hence the fixed point is unique. □

Unfortunately, unlike the Banach fixed-point theorem, this theorem is not constructive. It does not provide a procedure to find the fixed point. In Chapter 7, we define a particular generalized ultrametric space of discrete-event signals, and we show that a fixed point of a strictly contracting function on this space can be found using a constructive procedure.

B.4 Conclusion

This chapter has introduced the metric-space machinery used in Chapter 7 to give fixed-point semantics to discrete-event systems. We began with ordinary metrics and ultrametrics, then generalized distances to values in a [pointed poset](#). For functions and sequences, the [Cantor metric](#) and [Baire distance](#) measure similarity by how long two behaviors agree before they differ. Both are ultrametrics, a form particularly well suited to reasoning about prefix agreement and causality.

The central semantic objects are [contracting functions](#)—maps that bring points closer together. We distinguished contracting, [strictly contracting](#), and [delta contracting](#) maps, and related delta contraction to [Lipschitz continuous](#) functions. Completeness conditions—ordinary completeness for metric spaces and [spherically complete](#) generalized ultrametric spaces—ensure that limiting behaviors remain in the space and support existence proofs for fixed points.

The fixed-point theorems of Section B.3 parallel the order-theoretic results of Appendix A. Strict contraction yields uniqueness of any fixed point; delta contraction together with completeness yields existence and a constructive iteration scheme (the [Banach fixed-point theorem](#)); and strictly contracting maps on spherically complete generalized ultrametric spaces always have a unique fixed point, though without a general constructive procedure. In Chapter 7, these tools are instantiated for discrete-event signals, where a specialized construction provides a constructive procedure even in the generalized setting.

Bibliography

- Abadi, D., 2012: Consistency tradeoffs in modern distributed database system design: CAP is only part of the story. *Computer*, **45**(2), 37–42. doi:[10.1109/MC.2012.33](https://doi.org/10.1109/MC.2012.33).
- Allen, G. E., P. E. Zucknick, and B. L. Evans, 2007: A distributed deadlock detection and resolution algorithm for process networks. In *Int. Conf. on Acoustics, Speech, and Signal Proc. (ICASSP)*, IEEE, vol. II, pp. 33–36.
- Andalam, S., P. S. Roop, and A. Girault, 2010: Predictable multithreading of embedded applications using PRET-C. In *Formal Methods and Models for Codesign (MEMOCODE)*, IEEE/ACM, pp. 159–168. doi:[10.1109/MEMCOD.2010.5558636](https://doi.org/10.1109/MEMCOD.2010.5558636).
- André, C., 1996: SyncCharts: a visual representation of reactive behaviors. Report RR 95–52, revision: RR (96–56), University of Sophia-Antipolis. Available from: <http://www-sop.inria.fr/members/Charles.Andre/CA%20Publis/SYNCCHARTS/overview.html>.
- Arvind, L. Bic, and T. Ungerer, 1991: Evolution of data-flow computers. In Gaudiot, J.-L. and L. Bic, eds., *Advanced Topics in Data-Flow Computing*, Prentice-Hall.
- Bacelli, F., G. Cohen, G. J. Olster, and J. P. Quadrat, 1992: *Synchronization and Linearity, An Algebra for Discrete Event Systems*. Wiley, New York.

- Bailis, P., A. Ghodsi, J. M. Hellerstein, and I. Stoica, 2013: Bolt-on causal consistency. In *International Conference on Management of Data (SIGMOD)*, ACM, pp. 761–772. doi:[10.1145/2463676.2465279](https://doi.org/10.1145/2463676.2465279).
- Becker, A., 2018: *What Is Real? The Unfinished Quest for the Meaning of Quantum Physics*. Basic Books.
- Bell, J. S., 1964: On the Einstein Podolsky Rosen paradox. *Physica*, **1(3)**, 195–200.
- Benveniste, A. and G. Berry, 1991: The synchronous approach to reactive and real-time systems. *Proceedings of the IEEE*, **79(9)**, 1270–1282.
- Benveniste, A. and P. Le Guernic, 1990: Hybrid dynamical systems theory and the SIGNAL language. *IEEE Tr. on Automatic Control*, **35(5)**, 525–546.
- Berry, G., 1976: Bottom-up computation of recursive programs. *Revue Française d'Automatique, Informatique et Recherche Opérationnelle*, **10(3)**, 47–82. Available from: <http://eudml.org/doc/92028>.
- , 1999: *The Constructive Semantics of Pure Esterel*. Unpublished, draft version 3 ed. Available from: <http://www-sop.inria.fr/meije/esterel/doc/main-papers.html>.
- , 2003: The effectiveness of synchronous languages for the development of safety-critical systems. White paper, Esterel Technologies. Available from: <http://www.esterel-technologies.com>.
- Berry, G. and G. Gonthier, 1992: The Esterel synchronous programming language: design, semantics, implementation. *Science of Computer Programming*, **19(2)**, 87–152. doi:[10.1016/0167-6423\(92\)90005-V](https://doi.org/10.1016/0167-6423(92)90005-V).
- Berry, G. and E. Sentovich, 2001: Multiclock esterel. In *Correct Hardware Design and Verification Methods (CHARME)*, Springer-Verlag, vol. LNCS 2144. doi:[10.1007/3-540-44798-9_10](https://doi.org/10.1007/3-540-44798-9_10).
- Bhattacharya, B. and S. S. Bhattacharyya, 2000: Parameterized dataflow modeling of dsp systems. In *International Conference on Acoustics, Speech, and Signal Processing (ICASSP)*, pp. 1948–1951.
- Bhattacharyya, S. S., J. T. Buck, S. Ha, and E. Lee, 1995: Generating compact code from dataflow specifications of multirate signal processing algorithms. *IEEE*

- Transactions on Circuits and Systems I: Fundamental Theory and Applications*, **42(3)**, 138–150. doi:[10.1109/81.376876](https://doi.org/10.1109/81.376876).
- Bhattacharyya, S. S., J. T. Buck, S. Ha, and E. A. Lee, 1993: A scheduling framework for minimizing memory requirements of multirate dsp systems represented as dataflow graphs. In *VLSI Signal Processing VI*, IEEE, pp. 188–196. doi:[10.1109/VLSISP.1993.404488](https://doi.org/10.1109/VLSISP.1993.404488).
- Bhattacharyya, S. S. and E. A. Lee, 1993: Scheduling synchronous dataflow graphs for efficient looping. *Journal of VLSI Signal Processing Systems*, **6(3)**, 271–288. doi:[10.1007/BF01608539](https://doi.org/10.1007/BF01608539).
- Bhattacharyya, S. S., P. Murthy, and E. A. Lee, 1996a: APGAN and RPMC: Complementary heuristics for translating dsp block diagrams into efficient software implementations. *Journal of Design Automation for Embedded Systems*, **2(1)**, 33–60. doi:[10.1023/A:1008806425898](https://doi.org/10.1023/A:1008806425898).
- Bhattacharyya, S. S., P. K. Murthy, and E. A. Lee, 1996b: *Software Synthesis from Dataflow Graphs*. Kluwer Academic Publishers, Norwell, Mass.
- Bilsen, G., M. Engels, R. Lauwereins, and J. A. Peperstraete, 1996: Cyclo-static dataflow. *IEEE Transactions on Signal Processing*, **44(2)**, 397–408. doi:[10.1109/78.485935](https://doi.org/10.1109/78.485935).
- Boussinot, F., 1991: Reactive c: An extension to c to program reactive systems. *Software Practice and Experience*, **21(4)**, 401–428. doi:[10.1002/spe.4380210406](https://doi.org/10.1002/spe.4380210406).
- Box, G. E. P. and N. R. Draper, 1987: *Empirical Model-Building and Response Surfaces*. Wiley Series in Probability and Statistics, Wiley, ISBN 0471-81033-9.
- Brewer, E., 2000: Towards robust distributed system. In *Symposium on Principles of Distributed Computing (PODC)*.
- , 2012: CAP twelve years later: How the “rules” have changed. *IEEE Computer*, **45(2)**, 23–29. doi:[10.1109/MC.2012.37](https://doi.org/10.1109/MC.2012.37).
- Brock, J. D. and W. B. Ackerman, 1981: Scenarios, a model of non-determinate computation. In *Conference on Formal Definition of Programming Concepts*, Springer-Verlag, vol. LNCS 107, pp. 252–259. doi:[10.1007/3-540-10699-5_102](https://doi.org/10.1007/3-540-10699-5_102).

- Broman, D., L. Greenberg, E. A. Lee, M. Masin, S. Tripakis, and M. Wetter, 2015: Requirements for hybrid cosimulation standards. In *Hybrid Systems: Computation and Control (HSCC)*. doi:[10.1145/2728606.2728629](https://doi.org/10.1145/2728606.2728629).
- Broy, M. and G. Stefanescu, 2001: The algebra of stream processing functions. *Theoretical Computer Science*, **258**, 99–129.
- Buck, J. T., 1993: Scheduling dynamic dataflow graphs with bounded memory using the token flow model. Ph.D. Thesis Tech. Report UCB/ERL 93/69, University of California, Berkeley. Available from: <http://ptolemy.eecs.berkeley.edu/publications/papers/93/jbuckThesis/>.
- , 1994: Static scheduling and code generation from dynamic dataflow graphs with integer-valued control systems. In *IEEE Asilomar Conf. on Signals, Systems, and Computers*.
- Buck, J. T. and E. A. Lee, 1992: The token flow model. In *Data Flow Workshop*. Available from: <https://ptolemy.berkeley.edu/publications/papers/92/tokenFlow/tokenFlowAus.pdf>.
- , 1993: Scheduling dynamic dataflow graphs with bounded memory using the token flow model. In *IEEE Int. Conf. on Acoustics, Speech, and Signal Processing (ICASSP)*, vol. 1, pp. 429–432.
- Cardoso, J., E. A. Lee, J. Liu, and H. Zheng, 2014: Continuous-time models. In Ptolemaeus, C., ed., *System Design, Modeling, and Simulation using Ptolemy II*, Ptolemy.org, Berkeley, CA, ISBN 978-1-304-42106-7. Available from: <http://ptolemy.org/books/Systems>.
- Caspi, P., 1992: Clocks in dataflow languages. *Theoretical Computer Science*, **94**(1).
- Cataldo, A., E. A. Lee, X. Liu, E. Matsikoudis, and H. Zheng, 2006: A constructive fixed-point theorem and the feedback semantics of timed systems. In *Workshop on Discrete Event Systems (WODES)*. Available from: <http://ptolemy.eecs.berkeley.edu/publications/papers/06/constructive/>.
- Charron-Bost, B., 1991: Concerning the size of logical clocks in distributed systems. *Information Processing Letters*, **39**(1), 11–16. doi:[10.1016/0020-0190\(91\)90055-M](https://doi.org/10.1016/0020-0190(91)90055-M).

- Church, A., 1932: A set of postulates for the foundation of logic. *Annals of Mathematics*, **32(2)**, 346–366. Available from: <https://www.jstor.org/stable/1968337>.
- Church, A. and J. B. Rosser, 1936: Some properties of conversion. *Transactions of the American Mathematical Society*, **39(3)**, 472–482. doi:10.2307/1989762.
- Coffman, E. G., Jr. (Ed), 1976: *Computer and Job Scheduling Theory*. Wiley.
- Colaço, J.-L. and M. Pouzet, 2003: Clocks as first class abstract types. In Alur, R. and I. Lee, eds., *EMSOFT*, Springer-Verlag, vol. LNCS 2855, pp. 134–155. doi:10.1007/978-3-540-45212-6_10.
- Conway, M. E., 1963: Design of a separable transition-diagram compiler. *Communications of the ACM*, **6(7)**, 396–408.
- Copley, B. and F. Martin, 2014: *Causation in Grammatical Structures*. Oxford University Press, Oxford, England.
- Corbett, J. C., J. Dean, M. Epstein, A. Fikes, C. Frost, J. Furman, S. Ghemawat, A. Gubarev, C. Heiser, P. Hochschild, W. Hsieh, S. Kanthak, E. Kogan, H. Li, A. Lloyd, S. Melnik, D. Mwaura, D. Nagle, S. Quinlan, R. Rao, L. Rolig, Y. Saito, M. Szymaniak, C. Taylor, R. Wang, and D. Woodford, 2012: Spanner: Google’s globally-distributed database. In *OSDI*. doi:10.1145/2491245.
- , 2013: Spanner: Google’s globally-distributed database. *ACM Transactions on Computer Systems (TOCS)*, **31(8)**. doi:10.1145/2491245.
- Cousot, P. and R. Cousot, 1977: Abstract interpretation: A unified lattice model for static analysis of programs by construction or approximation of fixpoints. In *Symposium on Principles of Programming Languages (POPL)*, ACM Press, pp. 238–252.
- Creeger, M., 2005: Multicore CPUs for the masses. *ACM Queue*, **3(7)**, 63–64.
- Cremona, F., M. Lohstroh, D. Broman, E. A. Lee, M. Masin, and S. Tripakis, 2017: Hybrid co-simulation: it’s about time. *Software and Systems Modeling*, **18**, 1655–1679. doi:10.1007/s10270-017-0633-6.
- Davey, B. A. and H. A. Priestly, 2002: *Introduction to Lattices and Order*. Cambridge University Press, second edition ed.

- Dennis, J. B., 1974: First version data flow procedure language. Report MAC TM61, MIT Laboratory for Computer Science.
- Dhar, A., 1993: Nonuniqueness in the solutions of Newton's equation of motion. *American Journal of Physics*, **61**(1), 58–61. doi:[10.1119/1.17411](https://doi.org/10.1119/1.17411).
- Dijkstra, E. W., 1976: *A Discipline of Programming*. Prentice Hall.
- Ditlevsen, S. and A. Samson, 2013: Introduction to stochastic models in biology. In Bachar, M., J. Batzel, and S. Ditlevsen, eds., *Stochastic Biomathematical Models: with Applications to Neuronal Modeling*, Springer, pp. 3–35. doi:[10.1007/978-3-642-32157-3_1](https://doi.org/10.1007/978-3-642-32157-3_1).
- Dulloo, J. and P. Marquet, 2004: Design of a real-time scheduler for kahn process networks on multiprocessor systems. In *International Conference on Parallel and Distributed Processing Techniques and Applications (PDPTA)*.
- Earman, J., 1986: *A Primer on Determinism*, vol. 32 of *The University of Ontario Series in Philosophy of Science*. D. Reidel Publishing Company.
- Edwards, S., 2018: On determinism. In Lohstroh, M., P. Derler, and M. Sirjani, eds., *Principles of Modeling*, Springer, Cham, Switzerland, vol. LNCS 10760, pp. 240–253. doi:[10.1007/978-3-319-95246-8_14](https://doi.org/10.1007/978-3-319-95246-8_14).
- Edwards, S. A. and J. Hui, 2020: The sparse synchronous model. In *Forum on Specification and Design Languages (FDL)*. doi:[10.1109/FDL50818.2020.9232938](https://doi.org/10.1109/FDL50818.2020.9232938).
- Edwards, S. A. and E. A. Lee, 2003a: The semantics and execution of a synchronous block-diagram language. *Science of Computer Programming*, **48**(1), 21–42. doi:[10.1016/S0167-6423\(02\)00096-5](https://doi.org/10.1016/S0167-6423(02)00096-5).
- , 2003b: The semantics and execution of a synchronous block-diagram language. *Science of Computer Programming*, **48**(1), 21–42. doi:[10.1016/S0167-6423\(02\)00096-5](https://doi.org/10.1016/S0167-6423(02)00096-5).
- Edwards, S. A., E. A. Lee, S. Tripakis, and P. Whitaker, 2014: Synchronous-reactive models. In Ptolemaeus, C., ed., *System Design, Modeling, and Simulation using Ptolemy II*, Ptolemy.org, Berkeley, CA, ISBN 978-1-304-42106-7. Available from: <http://ptolemy.org/books/Systems>.
- Eidson, J. C., 2006: *Measurement, Control, and Communication Using IEEE 1588*. Springer.

- Eidson, J. C. and K. B. Stanton, 2015: Timing in cyber-physical systems: the last inch problem. In *IEEE International Symposium on Precision Clock Synchronization for Measurement, Control, and Communication (ISPCS)*, IEEE, pp. 19–24. doi: [10.1109/ISPCS.2015.7324674](https://doi.org/10.1109/ISPCS.2015.7324674).
- Falk, J., J. Keiner, C. Haubelt, J. Teich, and S. S. Bhattacharyya, 2008: A generalized static data flow clustering algorithm for mp soc scheduling of multimedia applications. In *Embedded Software (EMSOFT)*, ACM.
- Faustini, A. A., 1982: An operational semantics for pure dataflow. In *Proceedings of the 9th Colloquium on Automata, Languages and Programming (ICALP)*, Springer-Verlag, vol. Lecture Notes in Computer Science (LNCS) Vol. 140, pp. 212–224.
- Feng, T. H., E. A. Lee, X. Liu, C. Motika, R. v. Hanxleden, and H. Zheng, 2014a: Finite state machines. In Ptolemaeus, C., ed., *System Design, Modeling, and Simulation using Ptolemy II*, Ptolemy.org, Berkeley, CA, ISBN 978-1-304-42106-7. Available from: <http://ptolemy.org/books/Systems>.
- Feng, T. H., E. A. Lee, X. Liu, S. Tripakis, H. Zheng, and Y. Zhou, 2014b: Modal models. In Ptolemaeus, C., ed., *System Design, Modeling, and Simulation using Ptolemy II*, Ptolemy.org, Berkeley, CA, ISBN 978-1-304-42106-7. Available from: <http://ptolemy.org/books/Systems>.
- Fletcher, S. C., 2012: What counts as a newtonian system? the view from norton’s dome. *European Journal for Philosophy of Science*, **2**, 275–297. doi: [10.1007/s13194-011-0040-8](https://doi.org/10.1007/s13194-011-0040-8).
- Friedman, D. P. and D. S. Wise, 1976: CONS should not evaluate its arguments. In *Third Int. Colloquium on Automata, Languages, and Programming*, Edinburg University Press.
- Gaderer, G., P. Loschmidt, E. G. Cota, J. H. Lewis, J. Serrano, M. Cattin, P. Alvarez, P. M. Oliveira Fernandes Moreira, T. Wlostowski, J. Dedic, C. Prados, M. Kreider, R. Baer, S. Rauch, and T. Fleck, 2009: The white rabbit project. In *Int. Conf. on Accelerator and Large Experimental Physics Control Systems*.
- Gale, R., 1966: McTaggart’s analysis of time. *American Philosophical Quarterly*, **3**(2), 145–152. Available from: <https://www.jstor.org/stable/20009201>.
- Geilen, M. and T. Basten, 2003: Requirements on the execution of Kahn process networks. In *European Symposium on Programming Languages and Systems*,

- Springer, LNCS, pp. 319–334. Available from: <http://www.ics.ele.tue.nl/~tbasten/papers/esop03.pdf>.
- Geilen, M., T. Basten, and S. Stuijk, 2005: Minimising buffer requirements of synchronous dataflow graphs with model checking. In *Design Automation Conference (DAC)*, ACM, pp. 819–824. doi:[10.1145/1065579.1065796](https://doi.org/10.1145/1065579.1065796).
- Geilen, M. and S. Stuijk, 2010: Worst-case performance analysis of synchronous dataflow scenarios. In *CODES+ISSS*, ACM, pp. 125–134.
- Gilbert, S. and N. Lynch, 2002: Brewer’s conjecture and the feasibility of consistent, available, partition-tolerant web services. *ACM SIGACT News*, 33(2). doi:[10.1145/564585.564601](https://doi.org/10.1145/564585.564601).
- Girault, A., B. Lee, and E. A. Lee, 1999: Hierarchical finite state machines with multiple concurrency models. *IEEE Transactions On Computer-aided Design Of Integrated Circuits And Systems*, **18**(6), 742–760.
- Gordon, M. I., W. Thies, M. Karczmarek, J. Lin, A. S. Meli, A. A. Lamb, C. Leger, J. Wong, H. Hoffmann, D. Maze, and S. Amarasinghe, 2002: A stream compiler for communication-exposed architectures. In *10th international conference on Architectural support for programming languages and operating systems*, ACM Press, 605428, ISBN 1-58113-574-2, pp. 291–303.
- Ha, S., 1992: Compile-time scheduling of dataflow program graphs with dynamic constructs. Ph.D. Thesis Tech. Report UCB/ERL 92/43, UC Berkeley. Available from: <https://www2.eecs.berkeley.edu/Pubs/TechRpts/1992/2024.html>.
- Ha, S. and E. A. Lee, 1991: Compile-time scheduling and assignment of dataflow program graphs with data-dependent iteration. *IEEE Transactions on Computers*, **40**(11), 1225–1238. doi:[10.1109/12.102826](https://doi.org/10.1109/12.102826).
- Halbwachs, N., P. Caspi, P. Raymond, and D. Pilaud, 1991: The synchronous data flow programming language LUSTRE. *Proceedings of the IEEE*, **79**(9), 1305–1319.
- Hanxleden, R. v., B. Duderstadt, C. Motika, S. Smyth, M. Mendler, J. Aguado, S. Mercer, and O. O’Brien, 2014a: Sccharts: Sequentially constructive statecharts for safety-critical applications. In *Programming Language Design and Implementation (PLDI)*, pp. 372–383.

- Hanxleden, R. V., M. Mendler, J. Aguado, B. Duderstadt, I. Fuhrmann, C. Motika, S. Mercer, O. O'Brien, and P. Roop, 2014b: Sequentially constructive concurrency—a conservative extension of the synchronous model of computation. *ACM Transactions on Embedded Computing Systems (TECS)*, **13(4s)**, 1–26. doi:[10.1145/2627350](https://doi.org/10.1145/2627350).
- Harel, D., 1987: Statecharts: A visual formalism for complex systems. *Science of Computer Programming*, **8(3)**, 231–274.
- Harel, D., H. Lachover, A. Naamad, A. Pnueli, M. Politi, R. Sherman, A. Shtull-Trauring, and M. Trakhtenbrot, 1990: STATEMATE: A working environment for the development of complex reactive systems. *IEEE Transactions on Software Engineering*, **16(4)**, 403 – 414. doi:[10.1109/32.54292](https://doi.org/10.1109/32.54292).
- Harel, D. and A. Pnueli, 1985: On the development of reactive systems. In Apt, K. R., ed., *Logic and Models for Verification and Specification of Concurrent Systems*, Springer-Verlag, vol. F13 of NATO ASI Series, pp. 477–498.
- Hawking, S., 2002: Godel and the end of the universe. *Stephen Hawking Public Lectures*. Available from: <http://www.hawking.org.uk/godel-and-the-end-of-physics.html>.
- Helland, P., 2021: Don't get stuck in the “con” game: Consistency, convergence and confluence are not the same! *Queue*, **19(330)**, 16–35. doi:[10.1145/3475965.3480470](https://doi.org/10.1145/3475965.3480470).
- Helland, P. and D. Campbell, 2009: Building on quicksand. In *Conference on Innovative Data Systems Research (CIDR)*, ACM. Available from: <https://arxiv.org/abs/0909.1788>.
- Hennessy, M. and R. Milner, 1980: On observing nondeterminism and concurrency. In de Bakker, J. and J. van Leeuwen, eds., *Automata, Languages and Programming (ICALP)*, Springer, vol. Lecture Notes in Computer Science, LNCS 85, pp. 299–309. doi:[10.1007/3-540-10003-2_79](https://doi.org/10.1007/3-540-10003-2_79).
- Herlihy, M. P. and J. M. Wing, 1990: Linearizability: A correctness condition for concurrent objects. *ACM Transactions on Programming Languages and Systems*, **12(3)**, 463–492. doi:[10.1145/78969.78972](https://doi.org/10.1145/78969.78972).
- Hoare, C. A. R., 1978: Communicating sequential processes. *Communications of the ACM*, **21(8)**, 666–677.

- Hoefer, C., 2016: Causal determinism. *The Stanford Encyclopedia of Philosophy*, **Spring 2016 Edition**. Available from: <http://plato.stanford.edu/archives/spr2016/entries/determinism-causal/>.
- Hopcroft, J. and J. Ullman, 1979: *Introduction to Automata Theory, Languages, and Computation*. Addison-Wesley, Reading, MA, ISBN ISBN 0-201-02988-X.
- Hsu, C.-J., F. Keceli, M.-Y. Ko, S. Shahparnia, and S. S. Bhattacharyya, 2004: DIF: An interchange format for dataflow-based design tools. In *International Workshop on Systems, Architectures, Modeling, and Simulation*.
- Hu, T. C., 1961: Parallel sequencing and assembly line problems. *Operations Research*, **9(6)**, 841–848.
- Jantsch, A. and I. Sander, 2005: Models of computation and languages for embedded system design. *IEE Proceedings on Computers and Digital Techniques*, **152(2)**, 114–129.
- Jefferson, D. C., S. M. Lichten, and L. E. Young, 1996: A test of precision GPS clock synchronization. In *International Frequency Control Symposium*, IEEE. doi: [10.1109/FREQ.1996.560314](https://doi.org/10.1109/FREQ.1996.560314).
- Johnston, W. M., J. R. P. Hanna, and R. J. Millar, 2004: Advances in dataflow programming languages. *ACM Computing Surveys*, **36(1)**, 1–34.
- Kahn, G., 1974: The semantics of a simple language for parallel programming. In *Proc. of the IFIP Congress 74*, North-Holland Publishing Co., pp. 471–475.
- Kahn, G. and D. B. MacQueen, 1977: Coroutines and networks of parallel processes. In Gilchrist, B., ed., *Information Processing*, North-Holland Publishing Co., Amsterdam, pp. 993–998.
- Khomenko, V., M. Schaefer, and W. Vogler, 2008: Output-determinacy and asynchronous circuit synthesis. *Fundamenta Informaticae*, **88(4)**, 541–579.
- Kinniment, D. J., 2007: *Synchronization and arbitration in digital systems*. John Wiley & Sons, Ltd., New York.
- Kleppmann, M., 2015: A critique of the CAP theorem. ArXiv:1509.05393 [cs.DC]. Available from: <https://arxiv.org/abs/1509.05393>.

- Kopetz, H., 1997: *Real-Time Systems : Design Principles for Distributed Embedded Applications*. Springer.
- Kopetz, H. and G. Bauer, 2003: The time-triggered architecture. *Proceedings of the IEEE*, **91(1)**, 112–126.
- Lamport, L., 1978: Time, clocks, and the ordering of events in a distributed system. *Communications of the ACM*, **21(7)**, 558–565. doi:[10.1145/359545.359563](https://doi.org/10.1145/359545.359563).
- , 1984: Using time instead of timeout for fault-tolerant distributed systems. *ACM Transactions on Programming Languages and Systems*, **6(2)**, 254–280.
- , 1986: On interprocess communication - parts I and II. *Distributed Computing*, **1(2)**, 77–101. doi:[10.1007/BF01786228](https://doi.org/10.1007/BF01786228).
- Landin, P. J., 1965: A correspondence between Algol 60 and Church’s lambda notation. *Communications of the ACM*, **8(2)**, 89–101.
- Laplace, P.-S., 1901: *A Philosophical Essay on Probabilities*. John Wiley and Sons, Hoboken, NJ, translated from the sixth French edition by F. W. Truscott and F. L. Emory.
- Le Guernic, P., T. Gauthier, M. Le Borgne, and C. Le Maire, 1991: Programming real-time applications with SIGNAL. *Proceedings of the IEEE*, **79(9)**, 1321 – 1336. doi:[10.1109/5.97301](https://doi.org/10.1109/5.97301).
- Lee, E. A., 1991: Consistency in dataflow graphs. *IEEE Trans. on Parallel and Distributed Systems*, **2(2)**, 223–235.
- , 1999: Modeling concurrent real-time processes using discrete events. *Annals of Software Engineering*, **7**, 25–45. doi:[10.1023/A:1018998524196](https://doi.org/10.1023/A:1018998524196).
- , 2006a: Concurrent semantics without the notions of state or state transitions. In Asarin, E. and P. Bouyer, eds., *International Conference on Formal Modelling and Analysis of Timed Systems (FORMATS)*, Springer-Verlag, vol. LNCS 4202. doi:[10.1007/11867340_2](https://doi.org/10.1007/11867340_2).
- , 2006b: The problem with threads. *Computer*, **39(5)**, 33–42. doi:[10.1109/MC.2006.180](https://doi.org/10.1109/MC.2006.180).
- , 2008: Cyber physical systems: Design challenges. In *International Symposium on Object/Component/Service-Oriented Real-Time Distributed Computing (ISORC)*, IEEE, pp. 363 – 369. doi:[10.1109/ISORC.2008.25](https://doi.org/10.1109/ISORC.2008.25).

- , 2009: Computing needs time. *Communications of the ACM*, **52(5)**, 70–79. doi:[10.1145/1506409.1506426](https://doi.org/10.1145/1506409.1506426).
- , 2014: Constructive models of discrete and continuous physical phenomena. *IEEE Access*, **2(1)**, 1–25. doi:[10.1109/ACCESS.2014.2345759](https://doi.org/10.1109/ACCESS.2014.2345759).
- , 2016: Fundamental limits of cyber-physical systems modeling. *ACM Transactions on Cyber-Physical Systems*, **1(1)**. doi:[10.1145/2912149](https://doi.org/10.1145/2912149).
- , 2017: *Plato and the Nerd — The Creative Partnership of Humans and Technology*. MIT Press. Available from: <https://ptolemy.berkeley.edu/~eal/pnerd/>.
- , 2020: *The Coevolution: The Entwined Futures of Humans and Machines*. MIT Press, Cambridge, MA. Available from: <https://ptolemy.berkeley.edu/~eal/books/TheCoevolution.pdf>.
- , 2021: Determinism. *ACM Transactions on Embedded Computing Systems (TECS)*, **20(5)**, 1–34. doi:[10.1145/3453652](https://doi.org/10.1145/3453652).
- , 2022: What can deep neural networks teach us about embodied bounded rationality. *Frontiers in Psychology*, **25**. doi:[10.3389/fpsyg.2022.761808](https://doi.org/10.3389/fpsyg.2022.761808).
- Lee, E. A., R. Akella, S. Bateni, S. Lin, M. Lohstroh, and C. Menard, 2023a: Consistency vs. availability in distributed cyber-physical systems. *ACM Transactions on Embedded Computing Systems (TECS)*, **22(5s)**, 1–24. doi:[10.1145/3609119](https://doi.org/10.1145/3609119).
- Lee, E. A., S. Bateni, S. Lin, M. Lohstroh, and C. Menard, 2021: Quantifying and generalizing the CAP theorem. *arXiv:2109.07771 [cs.DC]*. Available from: <https://arxiv.org/abs/2109.07771>.
- , 2023b: Trading off consistency and availability in tiered heterogeneous distributed systems. *Intelligent Computing*, **2(Article 0013)**, 1–23. doi:[10.34133/icomputing.0013](https://doi.org/10.34133/icomputing.0013).
- Lee, E. A. and S. Ha, 1989: Scheduling strategies for multiprocessor real-time DSP. In *Global Telecommunications Conference (GLOBECOM)*, vol. 2, pp. 1279–1283. doi:[10.1109/GLOCOM.1989.64160](https://doi.org/10.1109/GLOCOM.1989.64160).
- Lee, E. A., J. Liu, L. Muliadi, and H. Zheng, 2014a: Discrete-event models. In Ptolemaeus, C., ed., *System Design, Modeling, and Simulation using Ptolemy II*, Ptolemy.org, Berkeley, CA, ISBN 978-1-304-42106-7. Available from: <http://ptolemy.org/books/Systems>.

- Lee, E. A. and E. Matsikoudis, 2009: The semantics of dataflow with firing. In Huet, G., G. Plotkin, J.-J. Lévy, and Y. Bertot, eds., *From Semantics to Computer Science: Essays in memory of Gilles Kahn*, Cambridge University Press. Available from: <http://ptolemy.eecs.berkeley.edu/publications/papers/08/DataflowWithFiring/>.
- Lee, E. A. and D. G. Messerschmitt, 1987a: Static scheduling of synchronous data flow programs for digital signal processing. *IEEE Transactions on Computers*, **C-36(1)**, 24–35. doi:10.1109/TC.1987.5009446.
- , 1987b: Synchronous data flow. *Proceedings of the IEEE*, **75(9)**, 1235–1245. doi:10.1109/PROC.1987.13876.
- Lee, E. A., S. Neuendorffer, and G. Zhou, 2014b: Dataflow. In Ptolemaeus, C., ed., *System Design, Modeling, and Simulation using Ptolemy II*, Ptolemy.org, Berkeley, CA, ISBN 978-1-304-42106-7. Available from: <http://ptolemy.org/books/Systems>.
- Lee, E. A. and T. M. Parks, 1995: Dataflow process networks. *Proceedings of the IEEE*, **83(5)**, 773–801. doi:10.1109/5.381846.
- Lee, E. A. and A. Sangiovanni-Vincentelli, 1998: A framework for comparing models of computation. *IEEE Transactions on Computer-Aided Design of Circuits and Systems*, **17(12)**, 1217–1229. Available from: <http://ptolemy.eecs.berkeley.edu/publications/papers/98/framework/>.
- Lee, E. A. and S. A. Seshia, 2017: *Introduction to Embedded Systems - A Cyber-Physical Systems Approach*. MIT Press, Cambridge, MA, USA, 2nd ed. Available from: <http://LeeSeshia.org>.
- Lee, E. A. and M. Sirjani, 2018: What good are models? In *Formal Aspects of Component Software (FACS)*, Springer, vol. LNCS 11222.
- Lee, E. A. and H. Zheng, 2007: Leveraging synchronous language principles for heterogeneous modeling and design of embedded systems. In *EMSOFT*, ACM, pp. 114 – 123. doi:10.1145/1289927.1289949.
- Lewis, E. R. and R. J. MacGregor, 2006: On indeterminism, chaos, and small number particle systems in the brain. *Journal of Integrative Neuroscience*, **5(2)**, 223–247. Available from: <https://people.eecs.berkeley.edu/~lewis/LewisMacGregor.pdf>.

- Lin, Y., R. Mullenix, M. Woh, S. Mahlke, T. Mudge, A. Reid, and K. Flautner, 2006: SPEX: A programming language for software defined radio. In *Software Defined Radio Technical Conference and Product Exposition*. Available from: <http://www.eecs.umich.edu/~sdrgrp/publications.php>.
- Liu, J. and E. A. Lee, 2003: On the causality of mixed-signal and hybrid models. In Maler, O. and A. Pnueli, eds., *6th International Workshop on Hybrid Systems: Computation and Control (HSCC '03)*, Springer-Verlag, vol. LNCS 2623, pp. 328–342. Available from: <http://ptolemy.eecs.berkeley.edu/papers/03/hybridCausality/>, doi:10.1007/3-540-36580-X_25.
- Liu, X. and E. A. Lee, 2008: CPO semantics of timed interactive actor networks. *Theoretical Computer Science*, **409**(1), 110–125. doi:10.1016/j.tcs.2008.08.044.
- Liu, X., E. Matsikoudis, and E. A. Lee, 2006a: Modeling timed concurrent systems. In *CONCUR 2006 - Concurrency Theory*, Springer, vol. LNCS 4137, pp. 1–15. doi:10.1007/11817949_1.
- , 2006b: Modeling timed concurrent systems using generalized ultrametrics. Report UCB/EECS-2006-45, EECS Department, University of California, Berkeley. Available from: <http://www.eecs.berkeley.edu/Pubs/TechRpts/2006/EECS-2006-45.html>.
- Lo Bello, L. and W. Steiner, 2019: A perspective on iee time-sensitive networking for industrial communication and automation systems. *Proceedings of the IEEE*, **107**(6), 1094–1120. doi:10.1109/JPROC.2019.2905334.
- Lohstroh, M., Í. Íncier Romeo, A. Goens, P. Derler, J. Castrillon, E. A. Lee, and A. Sangiovanni-Vincentelli, 2019: Reactors: A deterministic model for composable reactive systems. In *8th International Workshop on Model-Based Design of Cyber Physical Systems (CyPhy'19)*, Springer-Verlag, vol. LNCS 11971. doi:10.1007/978-3-030-41131-2_4.
- Lohstroh, M. and E. A. Lee, 2019: Deterministic actors. In *Forum on Specification and Design Languages (FDL)*. Available from: https://ptolemy.berkeley.edu/publications/papers/19/Lohstroh_Lee_DeterministicActors_FDL_2019.pdf.

- Lohstroh, M., C. Menard, S. Bateni, and E. A. Lee, 2021: Toward a lingua franca for deterministic concurrent systems. *ACM Transactions on Embedded Computing Systems (TECS)*, **20(4)**, Article 36. doi:[10.1145/3448128](https://doi.org/10.1145/3448128).
- Lohstroh, M., C. Menard, A. Schulz-Rosengarten, M. Weber, J. Castrillon, and E. A. Lee, 2020: A language for deterministic coordination across multiple timelines. In *Forum for Specification and Design Languages (FDL)*, IEEE. doi:[10.1109/FDL50818.2020.9232939](https://doi.org/10.1109/FDL50818.2020.9232939).
- Lorenz, E. N., 1963: Deterministic nonperiodic flow. *Journal of the Atmospheric Sciences*, **20**, 130–141. doi:[10.1175/1520-0469\(1963\)020](https://doi.org/10.1175/1520-0469(1963)020).
- Lynch, N., 1996: *Distributed Algorithms*. Morgan Kaufmann.
- Lázaro Cuadrado, D., A. P. Ravn, and P. Koch, 2007: Automated distributed simulation in Ptolemy II. In *Parallel and Distributed Computing and Networks (PDCN)*, Acta Press.
- Malament, D. B., 2008: Norton’s slippery slope. *Philosophy of Science*, **75**, 799–816. doi:[10.1086/594525](https://doi.org/10.1086/594525).
- Maler, O., Z. Manna, and A. Pnueli, 1992: From timed to hybrid systems. In *Real-Time: Theory and Practice, REX Workshop*, Springer-Verlag, pp. 447–484.
- Malik, S., 1994: Analysis of cyclic combinational circuits. *IEEE Transactions on Computer Aided Design of Integrated Circuits and Systems*, **13(7)**, 950–956.
- Manna, Z. and A. Pnueli, 1992: *The Temporal Logic of Reactive and Concurrent Systems*. Springer, Berlin.
- Marino, L. R., 1981: General theory of metastable operation. *IEEE Transactions on Computers*, **C-30(2)**, 107–115.
- Matloff, N., 2008: Introduction to discrete-event simulation and the SimPy language. Report, University of California, Davis. Available from: <http://heather.cs.ucdavis.edu/~matloff/156/PLN/DESimIntro.pdf>.
- Matsikoudis, E. and E. A. Lee, 2013: On fixed points of strictly causal functions. In *International Conference on Formal Modeling and Analysis of Timed Systems (FORMATS)*, Springer-Verlag, vol. LNCS 8053, pp. 183–197. doi:[10.1007/978-3-642-40229-6_13](https://doi.org/10.1007/978-3-642-40229-6_13).

- , 2015: The fixed-point theory of strictly causal functions. *Theoretical Computer Science*, **574**, 39–77.
- Mattern, F., 1988: Virtual time and global states of distributed systems. In Cosnard, M., P. Quinton, M. Raynal, and Y. Robert, eds., *Parallel and Distributed Algorithms*, North-Holland, pp. 215–226. Available from: http://www.coned.utcluj.ro/~salomie/DS_Mas/5-Additional_Materials/GlobalStates/VirtTimeGlobStates.pdf.
- Mendler, M., T. R. Shiple, and G. Berry, 2012: Constructive Boolean circuits and the exactness of timed ternary simulation. *Formal Methods in System Design*, **40(3)**, 283–329. doi:10.1007/s10703-012-0144-6.
- Mills, D. L., 2003: A brief history of NTP time: confessions of an internet time-keeper. *ACM Computer Communications Review*, **33**.
- , 2006: *Computer Network Time Synchronization — The Network Time Protocol*. CRC Press, Boca Raton, FL.
- Milner, R., 1980: *A Calculus of Communicating Systems*, vol. 92 of *Lecture Notes in Computer Science*. Springer, Berlin Heidelberg.
- , 1989: *Communication and Concurrency*. Prentice Hall, Englewood Cliffs, NJ, USA.
- Moreira, O., T. Basten, M. Geilen, and S. Stuijk, 2010: Buffer sizing for rate-optimal single-rate dataflow scheduling revisited. *IEEE Transactions on Computers*, **59(2)**, 188–201. doi:10.1109/TC.2009.155.
- Morris, J. H. and P. Henderson, 1976: A lazy evaluator. In *Conference on the Principles of Programming Languages (POPL)*, ACM.
- Müller, O. and P. Scholz, 1997: Functional specification of real-time and hybrid systems. In Maler, O., ed., *Hybrid and Real-Time Systems*, Springer, Berlin/Heidelberg, vol. Lecture Notes in Computer Science (LNCS), vol.1201, pp. 273–285. doi:10.1007/BFb0014732.
- Muller, R. A., 2016: *Now — The Physics of Time*. W. W. Norton and Company.
- Murata, T., 1989: Petri nets: Properties, analysis and applications. *Proceedings of IEEE*, **77(4)**, 541–580. doi:10.1109/5.24143.

- Murthy, P. K. and S. S. Bhattacharyya, 2006: *Memory Management for Synthesis of DSP Software*. CRC Press.
- Murthy, P. K. and E. A. Lee, 2002: Multidimensional synchronous dataflow. *IEEE Transactions on Signal Processing*, **50(8)**, 2064–2079. doi:[10.1109/TSP.2002.800830](https://doi.org/10.1109/TSP.2002.800830).
- Naundorf, H., 2000: Strictly causal functions have a unique fixed point. *Theoretical Computer Science*, **238(1-2)**, 483–488. Available from: <http://www.sciencedirect.com/science/article/B6V1G-43FWY1C-K/2/d64dbca6a23fee5596d3b12189ccf34d>.
- Norton, J. D., 2007: Causation as folk science. In Price, H. and R. Corry, eds., *Causation, Physics, and the Constitution of Reality*, Clarendon Press, Oxford, book section 2, pp. 11–44.
- NS-3 Consortium, 2015: ns-3. Available from: <http://www.nsnam.org/>.
- Object Management Group (OMG), 2007: A UML profile for MARTE, beta 1. OMG Adopted Specification ptc/07-08-04, OMG. Available from: <http://www.omg.org/omgmarte/>.
- , 2010: System modeling language specification v1.2. Standard specification, OMG. Available from: <http://www.sysmlforum.com>.
- Olson, A. G. and B. L. Evans, 2005: Deadlock detection for distributed process networks. In *ICASSP*.
- Paladino, F., S. Li, and E. A. Lee, 2026: Maxwait: A generalized mechanism for distributed time-sensitive systems. *arXiv*, **2601.21146v1 [cs.DC]**. doi:[10.48550/arXiv.2601.21146](https://doi.org/10.48550/arXiv.2601.21146).
- Park, D., 1980: Concurrency and automata on infinite sequences. In (eds), D. P., ed., *Theoretical Computer Science*, Springer, Berlin, Heidelberg, vol. LNCS 104. doi:[10.1007/BFb0017309](https://doi.org/10.1007/BFb0017309).
- Parks, T. M., 1995: Bounded scheduling of process networks. Ph.D. Thesis Tech. Report UCB/ERL M95/105, UC Berkeley. Available from: <http://ptolemy.eecs.berkeley.edu/papers/95/parksThesis>.
- Parks, T. M. and D. Roberts, 2003: Distributed process networks in Java. In *International Parallel and Distributed Processing Symposium*, pp. 138–146.

- Patterson, D. A., G. Gibson, and R. Katz, 1988: A case for redundant arrays of inexpensive disks (RAID). In *International Conference on Management of Data (SIGMOD)*. doi:10.1145/50202.50214.
- Pearl, J. and D. Mackenzie, 2018: *The Book of Why: The New Science of Cause and Effect*. Basic Books, New York.
- Pino, J. L., T. M. Parks, and E. A. Lee, 1994: Automatic code generation for heterogeneous multiprocessors. In *International Conference on Acoustics, Speech, and Signal Processing (ICASSP)*, pp. 445–448. doi:10.1109/ICASSP.1994.389626.
- Plotkin, G., 1976: A powerdomain construction. *SIAM Journal on Computing*, **5**(3), 452–487.
- Popper, K., 1959: *The Logic of Scientific Discovery*. Hutchinson & Co., London and New York.
- Priess-Crampe, S. and P. Ribenboim, 1996: Generalized ultrametric spaces I. *Abhandlungen aus dem Mathematischen Seminar der Universität Hamburg*, **66**, 55–73.
- , 1999: Logic programming and ultrametric spaces. *Rendiconti di Matematica, Serie VII*, **19**, 155–176.
- Ptolemaeus, C., 2014: *System Design, Modeling, and Simulation using Ptolemy II*. Ptolemy.org, Berkeley, CA, ISBN 978-1-304-42106-7. Available from: <http://ptolemy.org/books/Systems>.
- Reed, G. M. and A. W. Roscoe, 1988a: Metric spaces as models for real-time concurrency. In *3rd Workshop on Mathematical Foundations of Programming Language Semantics*, pp. 331–343.
- , 1988b: A timed model for communicating sequential processes. *Theoretical Computer Science*, **58**(1-3), 249–261.
- Ritchie, D. M. and K. L. Thompson, 1974: The UNIX time-sharing system. *Communications of the ACM*, **17**(7), 365 – 375.
- Rovelli, C., 2017: *Reality is Not What It Seems: The Journey to Quantum Gravity*. Riverhead Books, New York.
- , 2018: *The Order of Time*. Riverhead Books, New York.

- Russell, B., 1913: On the notion of cause. *Proceedings of the Aristotelian Society*, **13**, 1–26.
- Sander, I. and A. Jantsch, 2004: System modeling and transformational design refinement in ForSyDe. *IEEE Transactions on Computer-Aided Design of Circuits and Systems*, **23**(1), 17–32.
- Sangiorgi, D., 2009: On the origins of bisimulation and coinduction. *ACM Transactions on Programming Languages and Systems*, Vol. 31, No. 4, Article 15, Pub. date: May 2009., **31**(4), 15:1–15:41. doi:[10.1145/1516507.1516510](https://doi.org/10.1145/1516507.1516510).
- Schulz-Rosengarten, A., R. v. Hanxleden, M. Lohstroh, E. A. Lee, and S. Bateni, 2023: Polyglot modal models through lingua franca. In *Cyber-Physical Systems and Internet of Things Week (CPS-IoT)*, pp. 337–242. doi:[10.1145/3576914.3587498](https://doi.org/10.1145/3576914.3587498).
- Schulz-Rosengarten, A., R. von Hanxleden, F. Mallet, R. De Simone, and J. Dean-toni, 2018: Time in SCCharts. In *2018 Forum on Specification Design Languages (FDL)*, pp. 5–16.
- Schwarz, R. and F. Mattern, 1994: Detecting causal relationships in distributed computations: in search of the holy grail. *Distributed Computing*, **7**, 149–174. doi:[10.1007/BF02277859](https://doi.org/10.1007/BF02277859).
- Shiple, T. R., G. Berry, and H. Touati, 1996: Constructive analysis of cyclic circuits. In *European conference on Design and Test (DATE)*, IEEE Computer Society, pp. 328–333.
- Sih, G. C. and E. A. Lee, 1993a: A compile-time scheduling heuristic for interconnection-constrained heterogeneous processor architectures. *IEEE Transactions on Parallel and Distributed Systems*, **4**(2), 175–187. doi:[10.1109/71.207593](https://doi.org/10.1109/71.207593).
- , 1993b: Declustering : A new multiprocessor scheduling technique. *IEEE Transactions on Parallel and Distributed Systems*, **4**(6), 625–637. doi:[10.1109/71.242160](https://doi.org/10.1109/71.242160).
- Srini, V., 1986: An architectural comparison of dataflow systems. *Computer*, **19**(3).
- Sriram, S. and S. S. Bhattacharyya, 2009: *Embedded Multiprocessors: Scheduling and Synchronization*. CRC press, 2nd ed.

- Stark, E. W., 1995: An algebra of dataflow networks. *Fundamenta Informaticae*, **22(1-2)**, 167–185.
- Stephens, R., 1997: A survey of stream processing. *Acta Informatica*, **34(7)**.
- Stuijk, S., M. C. Geilen, and T. Basten, 2008: Throughput-buffering trade-off exploration for cyclo-static and synchronous dataflow graphs. *IEEE Transactions on Computers*, **57(10)**, 1331–1345. doi:[10.1109/TC.2008.58](https://doi.org/10.1109/TC.2008.58).
- Talcott, C. L., 1996: Interaction semantics for components of distributed systems. In *Formal Methods for Open Object-Based Distributed Systems (FMOODS)*, pp. 154–169. doi:[10.1007/978-0-387-35082-0_11](https://doi.org/10.1007/978-0-387-35082-0_11).
- Thies, W., M. Karczmarek, and S. Amarasinghe, 2002: StreamIt: A language for streaming applications. In *11th International Conference on Compiler Construction*, Springer-Verlag, vol. LNCS 2304. doi:[10.1007/3-540-45937-5_14](https://doi.org/10.1007/3-540-45937-5_14).
- Thies, W., M. Karczmarek, J. Sermulins, R. Rabbah, and S. Amarasinghe, 2005: Teleport messaging for distributed stream programs. In *Principles and Practice of Parallel Programming (PPoPP)*, ACM. doi:[10.1145/1065944.1065975](https://doi.org/10.1145/1065944.1065975).
- Turing, A. M., 1936: On computable numbers with an application to the Entscheidungsproblem. *Proceedings of the London Mathematical Society*, **42**, 230–265.
- Turjan, A., B. Kienhuis, and E. Deprettere, 2003: Solving out-of-order communication in Kahn process networks. *Journal on VLSI Signal Processing-Systems for Signal, Image, and Video Technology*, **40(1)**, 7–18. doi:[10.1007/s11265-005-4935-5](https://doi.org/10.1007/s11265-005-4935-5).
- von Hanxleden, R., 2009: SyncCharts in C - A proposal for light-weight deterministic concurrency. In *ACM Embedded Software Conference (EMSOFT)*, pp. 11–16. doi:[10.1145/1629335.1629366](https://doi.org/10.1145/1629335.1629366).
- von Hanxleden, R., T. Bourke, and A. Girault, 2017: Real-time ticks for synchronous programming. In *2017 Forum on Specification and Design Languages (FDL)*, IEEE, pp. 1–8.
- von Hanxleden, R. et al., 2014: SCCharts: Sequentially constructive Statecharts for safety-critical applications. In *ACM SIGPLAN Conf. on Programming Language Design and Implementation*, ACM, New York, NY, USA, PLDI '14, ISBN 978-1-4503-2784-8, pp. 372–383. doi:[10.1145/2594291.2594310](https://doi.org/10.1145/2594291.2594310).

- Vuillemin, J., 1973: *Proof Techniques for Recursive Programs*. Ph. d. thesis.
- Ward, B. C., 2026: Proofreader: Informally auditing formal proofs with LLMs. Available from: <https://my.vanderbilt.edu/bryancward/projects/proofreader/>.
- Wegner, P., 1998: Interactive foundations of computing. *Theoretical Computer Science*, **192(2)**, 315–351.
- Winskel, G., 1993: *The Formal Semantics of Programming Languages*. MIT Press, Cambridge, MA, USA.
- Wisniewski, R., I. Grobelna, and A. Karatkevich, 2020: Determinism in cyber-physical systems specified by interpreted petri nets. *Sensors*, **20(5565)**, 22. doi: [10.3390/s20195565](https://doi.org/10.3390/s20195565).
- Wolfram, S., 2002: *A New Kind of Science*. Wolfram Media, Inc.
- Wolpert, D. H., 2008: Physical limits of inference. *Physica*, **237(9)**, 1257–1281. doi: [10.1016/j.physd.2008.03.040](https://doi.org/10.1016/j.physd.2008.03.040).
- Yates, R. K., 1993: Networks of real-time processes. In Best, E., ed., *Proc. of the 4th Int. Conf. on Concurrency Theory (CONCUR)*, Springer-Verlag, vol. LNCS 715.
- Yates, R. K. and G. R. Gao, 1993: A Kahn principle for networks of nonmonotonic real-time processes. In *Parallel Architectures and Languages, Europe*, pp. 209–227.
- Yu, H. and A. Vahdat, 2006: The costs and limits of availability for replicated services. *ACM Transactions on Computer Systems*, **2(24)**, 70–113. doi: [10.1145/1124153.1124156](https://doi.org/10.1145/1124153.1124156).
- Zeigler, B., 1976: *Theory of Modeling and Simulation*. Wiley Interscience, New York.
- Zeigler, B. P., H. Praehofer, and T. G. Kim, 2000: *Theory of Modeling and Simulation*. Academic Press, 2nd ed.
- Zhao, Y., E. A. Lee, and J. Liu, 2007: A programming model for time-synchronized distributed real-time systems. In *Real-Time and Embedded Technology and Applications Symposium (RTAS)*, IEEE, pp. 259 – 268. doi: [10.1109/RTAS.2007.5](https://doi.org/10.1109/RTAS.2007.5).
- Zinkernagel, H., 2010: Causal fundamentalism in physics. In Suárez, M., M. Dorato, and M. Rédei, eds., *EPSA Philosophical Issues in the Sciences: Launch of the European Philosophy of Science Association*, Springer, Dordrecht.

Index

- n*-tuple, 303
- A-series time, 36
- absent, xviii, 212, 257
- abstraction, 10
- ACID 2.0, 62
- actor, 101, 146, 154
- acyclic precedence graph, xv, 185, 186, 199, 203
- agreement, 61
- Airbus, 210
- algorithm, 54
- alphabetical order, 311
- Android, 55
- Ansys, 210
- antisymmetric, 309
- APG, xv, 185, 186
- API, xv, 71
- apparent latency, 78, 80, 99
- artificial deadlock, 132
- assignment rule, 303
- associative, 62
- asynchronous networks, 59
- asynchronous procedure call, 101
- atomic, 61
- availability, 77, 86
- axiom of choice, 314, 322
- B-series time, 36
- Baire distance, 138, 336, 344
- balance equation, 174, 203, 248
- Banach fixed-point theorem, 264, 269, 270, 272, 275, 278, 294, 341, 345
- Bayesian, 8
- Bayesian probability, 38
- BDF, xv, 192, 196, 203
- behavior, 13
- Bell's theorem, 41
- Bell, John Stewart, 41
- Bernoulli actor, 244
- bijection, 305, 319, 325
- binary digits, xvii, 302
- bisimulation, 21, 149
- bistable, 33

- black hole, 34
- blocking read, 18, 107, 140, 162
- blocking reads, 111
- blocking write, 130, 172
- Bohr, Niels, 38
- boolean dataflow, xv, 192
- bottom, 22, 310
- bottom element, xviii, 310, 313
- bounded buffers, 130, 174
- bounded execution, 124, 129, 132
- Brewer, Eric, 59
- Brock-Ackerman anomaly, 147
- buffer, 107
- bulletin board, 64
- cache memory, 55
- CAL, xv, 60
- CAL theorem, 83, 84, 87, 100
- CAL Theorem, max-plus formulation, 89
- calculus of communicating systems, xv, 150
- Cantor metric, 263, 265, 267, 278, 294, 334, 336, 344
- Cantor's diagonal argument, 327
- CAP, xv, 59, 100
- CAP theorem, 59
- cardinality, 323
- cartesian product, xvii, 114, 303, 308
- Cauchy, 268, 339
- Cauchy sequence, 49
- Cauchy surface, 34
- causal, 264, 274
- causal consistency, 63
- causality, 7, 31, 293
- causality relation, 64
- causation, 7, 31
- CCS, xv, 150
- CD, 184
- cellular automaton, 24
- CERN, 71
- CET, xv, 67, 68, 73, 74, 99
- chain, 118, 222, 309, 313
- chain of closed balls, 340, 343
- channel, 107
- chaos, 24
- Church, Alonzo, 21
- Church-Rosser theorem, 21
- circuit design, 209
- clock, 70, 208, 209
- clock calculus, 246
- clock domain, 239, 256
- clock master, 69
- clock signal, 258, 272
- closed system, 262, 267, 274, 293
- code generation, 186
- codomain, 303
- collapse of the wave function, 38
- Collatz conjecture, 279, 295
- colored token, 197
- common prefix, xviii, 282
- communicating sequential processes, 150, 271
- communication graph, 89, 100
- communication latency, 67, 79
- commutative, 62
- commutative dataflow actor, 168, 169, 171, 203
- Commutator actor, 198
- compact disk, 184
- comparable, 309
- complete, 339
- complete iteration, 173, 174, 188, 189, 192, 194, 195, 199, 200, 202
- complete lattice, 318, 324

- complete lower semilattice, 142, 318
- complete metric space, 268, 294, 339
- complete partial order, 112, 118, 313, 314
- complete upper semilattice, 318
- composability, 10
- compositionality, 168
- computability, 330
- computation, 22, 329
- concatenation, xviii, 108, 166
- concatenation of empty sets, 160
- concurrency, 54, 106
- confluence, 17
- connected graph, 180
- connectivity matrix, 176, 180, 181
- conservation of energy, 42
- conservation of momentum, 42
- consistency, x, 9, 57, 76
- consistent, 176, 180–182, 200, 203
- consistent cut, 64
- consistent global snapshot, 64
- constant function, 115
- constant network approximation, 82, 84, 85, 100
- constructive, 224, 227, 233, 251, 252, 341
- constructive procedure, 158, 323
- continuous, 111, 112, 118, 140, 142, 143, 156, 158, 160, 166, 202, 222, 292, 319
- continuous-time system, 263
- contracting function, 267, 287, 288, 338, 344
- contraction, 267, 294, 338
- convergence in a metric space, 339
- convergence of a sequence of sequences, 139
- convergent execution, 139
- coordinated universal time, xvi, 67
- coordination, 208
- coordination language, 210, 260
- Copenhagen interpretation, 38
- core, 55
- coroutines, 111
- correct execution, 122
- cosmic censorship, 34
- countable, 325, 328
- countably infinite, 325
- CountDown actor, 239
- counter, 215
- CPO, 111, 138, 157, 221, 280, 313, 314, 319, 320
- CRC, xv, 25
- CSDF, xv, 196, 201, 203
- CSP, 150, 263
- cyclic redundancy check, xv, 25
- cyclo-static dataflow, xv, 196
- DAT, 184
- data flow, 154
- data type, 110, 112, 212, 221, 257
- data-driven execution, 126, 172
- dataflow, 154, 211, 229
- dataflow actor, 154, 155, 156, 163, 165, 167, 168, 171, 202
- dataflow actor functional, xviii, 157, 160
- dataflow architecture, 154
- dataflow process, 156
- DDF, xv, 195, 196
- DE signal, 258
- deadlock, 110, 127
- decidable, 185
- decision problem, 330
- Default actor, 241
- delay component, 269
- delta causal, 264, 269, 274, 294

- delta contracting, 263, 267, 269, 294, 338, 344
- delta cycles, 260
- demand-driven execution, 111, 125, 172
- Dennis dataflow, 154
- denotational semantics, 116, 122, 123, 163, 171, 202
- dependency matrix, 88, 93, 96
- determinate vs. deterministic, 20
- determinism, 13
- deterministic, 13, 106, 155, 211
- dictionary order, 311
- digital audio tape, 184
- digital physics, 47
- directed set, 137, 158, 160, 313, 314
- discrete set, 258, 259, 267, 328
- discrete-event, 229, 260
- discrete-event component, 261, 261, 264
- discrete-event signal, 258, 262, 265, 268, 280, 281
- discrete-event system, 20, 211, 262
- discrete-event tuple, 259, 261
- distance function, 332
- distributed architecture, 186
- distributed system, 55
- domain, 303
- downset, 64, 281, 282, 285, 295
- dynamic dataflow, xv, 195
- effective, 124, 132, 134, 139
- effective execution, 125, 127, 130, 134, 140, 172
- effectively computable, 101, 110, 329
- Einstein train, 35
- Einstein, Albert, 40
- elastic collision, 42
- embedded processor, 186
- empty function, xviii, 156
- empty sequence, xviii, 307
- empty set, xvii, 302
- empty signal, 257
- enabled actor, 171, 182, 185
- EnabledComposite actor, 239
- endofunction, 222, 225, 236, 267, 304, 322
- engineering model, 6, 25
- enumerate, 326
- epoch, 72, 72
- Esterel, 208, 210, 211, 229
- Esterel Technologies, 210
- event, 257
- event horizon, 34
- event types, 63, 75
- eventual consistency, x, 66, 76
- eventually delta contracting, 263, 272, 274–276, 278, 279, 294
- eventually strictly contracting function, 287, 292
- Everett III, Hugh, 39
- execution time, 79
- existence of fixed points, 322
- extended data type, 212, 221
- Facebook Timeline, 64
- fair execution, 123, 130, 131
- fault, 24, 25, 211
- fault detection, 9
- feedback, 149, 164
- FIFO, xv, 107
- filter, 190
- finite sequences, xviii, 307, 314
- finite state machine, 199, 229
- firing, 208
- firing function, 155, 163, 196
- firing rule, 155, 163, 182, 196
- first in first out, xv, 107

- fixed point, [112](#), [166](#), [216](#), [322](#)
- fixed point theorem, [322](#)
- flat order relation, [316](#)
- flat partial order, [142](#), [221](#), [316](#)
- flight control, [210](#)
- floating-point, [73](#)
- floor on unavailability, [81](#), [83](#), [85](#), [87](#), [89](#),
[100](#)
- fork, [125](#), [133](#)
- frequentist, [8](#)
- FSM, [199](#), [229](#)
- full abstraction, [117](#)
- fully dynamic scheduling, [186](#)
- fully-static scheduling, [186](#)
- function, [xvii](#), [117](#), [303](#)
- function composition, [xvii](#), [304](#)
- function to a power, [xvii](#)
- functional, [157](#), [215](#), [237](#)
- future, [101](#)

- Galois field, [245](#)
- generalized distance, [281](#), [287](#), [295](#)
- generalized ultrametric distance function, [334](#)
- generalized ultrametric inequality, [290](#),
[334](#)
- generalized ultrametric space, [281](#), [285](#),
[295](#), [334](#), [338](#), [339](#), [343](#)
- GLB, [xv](#), [xviii](#), [312](#)
- global clock, [208](#), [245](#), [256](#)
- global positioning system, [71](#)
- Gödel, Kurt, [26](#)
- Google Spanner, [25](#), [66](#), [68](#), [71](#)
- GPS, [71](#)
- graph, [308](#)
- graph of a function, [303](#)
- greatest lower bound, [xv](#), [xviii](#), [312](#)
- guarded count, [239](#)

- Gustave function, [142](#), [162](#)
- gyroscope, [36](#)

- halting problem, [23](#), [127](#)
- Hamming numbers, [132](#), [193](#)
- happened before, [64](#)
- hardware description language, [260](#)
- Hasse diagram, [223](#), [310](#), [316](#)
- Hausdorff space, [137](#), [280](#)
- Hawking, Stephen, [37](#)
- HDF, [xv](#), [199](#)
- Heisenberg uncertainty principle, [46](#)
- Heisenberg, Werner, [38](#)
- heterochronous dataflow, [xv](#), [199](#)
- heterogeneous multiprocessor, [186](#)
- hidden variables, [40](#)
- Hoare, Tony, [16](#)
- homogeneous SDF, [173](#), [197](#), [211](#)
- Hu level scheduling, [188](#)

- idempotent, [62](#)
- identity function, [xviii](#), [304](#)
- identity of indiscernables, [332](#), [335](#)
- identity process, [109](#), [116](#)
- IDF, [xv](#), [192](#)
- IEEE 1588, [71](#)
- IEEE 802.1, [71](#)
- ill-formed composition, [232](#)
- image, [304](#), [328](#)
- image function, [xviii](#), [304](#)
- image processing, [190](#)
- imperative, [101](#), [107](#), [154](#)
- incomparable, [64](#), [309](#)
- incompleteness, [26](#)
- incompleteness of determinism, [41](#), [48](#)
- inconsistency, [xviii](#), [67](#), [74](#), [76](#), [78](#), [79](#), [82](#),
[99](#)
- inconsistent, [176](#), [178](#)

- infinite execution, [123](#)
- infinite sequences, [xviii](#), [307](#), [308](#)
- infix notation, [309](#)
- information order, [223](#)
- initial token, [164](#), [173](#)
- injective, [305](#)
- inputs, [13](#)
- instruction set architecture, [154](#)
- integer dataflow, [xv](#), [192](#)
- integers, [xvii](#), [302](#)
- interaction, [149](#)
- Internet protocol, [8](#)
- internet protocol, [xv](#)
- interrupt service routine, [86](#)
- intuitionistic logic, [228](#)
- inverse, [305](#)
- inverter, [248](#)
- iOS, [55](#)
- IP, [xv](#), [8](#)
- ISA, [154](#)
- job-shop scheduling, [173](#), [185–187](#), [199](#)
- join, [xviii](#), [281](#), [312](#), [334](#)
- join semilattice, [318](#)
- joinable, [142](#), [312](#), [313](#)
- Kahn principle, [111](#), [123](#), [134](#)
- Kahn process network, [xv](#), [18](#), [106](#), [185](#), [195](#), [263](#)
- Kahn, Gilles, [18](#), [111](#)
- Kahn-MacQueen process, [107](#), [118](#), [140](#), [154](#), [162](#)
- Kant, [14](#)
- Kleene chain, [222](#), [225](#), [323](#)
- Kleene closure, [307](#)
- Kleene fixed-point theorem, [120](#), [158](#), [159](#), [221](#), [222](#), [224](#), [282](#), [322](#), [324](#)
- Kleene operator, [307](#)
- Kleene star, [307](#)
- Kleene, Stephen Cole, [307](#), [324](#)
- Knaster, Bronislaw, [324](#)
- Knaster-Tarski fixed-point theorem, [324](#)
- Kolmogorov space, [137](#)
- KPN, [xv](#), [18](#), [106](#)
- lag, [78](#), [78](#), [85](#), [99](#)
- lambda calculus, [21](#), [23](#)
- language, [16](#)
- Laplace’s demon, [26](#)
- Laplace, Pierre-Simon, [26](#), [38](#)
- last writer wins, [66](#)
- latency, [78](#)
- lattice, [223](#), [317](#)
- law of the excluded middle, [228](#)
- lazy evaluators, [111](#), [125](#)
- LCM, [178](#)
- leap second, [71](#)
- leap seconds, [72](#)
- least common multiple, [178](#)
- least fixed point, [163](#), [322](#)
- least fixed-point semantics, [111](#), [116](#), [144](#), [159](#), [168](#)
- least upper bound, [xv](#), [xviii](#), [311](#), [314](#), [334](#)
- LET, [xv](#), [93](#)
- level, [188](#)
- LEX, [xv](#), [79](#), [80](#), [88](#), [89](#), [96](#), [100](#)
- lexicographic order, [257](#), [311](#), [319](#)
- lifted, [118](#), [158](#), [166](#), [304](#), [306](#), [319](#)
- lifted projection, [xviii](#), [307](#)
- limit, [313](#), [314](#)
- linearizable, [61](#)
- linearly dependent, [179](#)
- linearly independent, [179](#), [181](#)
- Lingua Franca, [20](#), [26](#), [72–75](#), [86](#), [92](#), [297](#)
- Linux, [55](#)

- Lipschitz continuous, [280](#), [338](#), [344](#)
- local artificial deadlock, [132](#)
- local deadlock, [110](#), [131](#)
- locality, [40](#)
- logical execution time, [xv](#), [93](#)
- logical time, [67](#), [73](#), [74](#), [256](#), [260](#), [293](#)
- logical time chases physical time, [81](#)
- logically simultaneous, [68](#), [73](#), [256](#)
- lower bound, [312](#)
- lower semilattice, [282](#), [317](#)
- LUB, [xv](#), [xviii](#), [138](#), [311](#)
- Lustre, [208](#), [210](#)
- macOS, [55](#)
- makespan, [187](#)
- many worlds, [40](#)
- marking, [197](#)
- master clock, [245](#)
- max-plus identity matrix, [90](#)
- maximal execution, [122](#), [139](#)
- MDSDF, [190](#)
- meet, [xviii](#), [282](#), [312](#)
- meet semilattice, [317](#)
- merge operation, [62](#), [64](#), [75](#)
- message passing, [106](#)
- metastable state, [33](#)
- metric, [332](#)
- metric space, [49](#), [332](#)
- microstep, [73](#)
- Milner, Robin, [16](#)
- minimum inconsistency, [76](#)
- MinMax actor, [193](#)
- modal dataflow, [196](#), [203](#)
- model, [48](#)
- model checking, [186](#)
- monotonic, [117](#), [157](#), [163](#), [282](#), [283](#), [292](#), [295](#), [318](#)
- multicore, [186](#)
- multidimensional SDF, [190](#)
- multimedia, [186](#), [190](#)
- Multiplexor actor, [198](#), [200](#)
- multitasking, [56](#)
- Murphy's law, [24](#)
- must-may analysis, [234](#)
- mutex, [101](#), [146](#)
- naked singularities, [34](#)
- natural numbers, [xvii](#), [302](#)
- network latency, [xviii](#), [79](#)
- network penalty, [xv](#), [79](#)
- network time protocol, [xvi](#), [67](#)
- Newton's first law, [30](#)
- Newton's hypothesis, [44](#)
- Newton's second law, [13](#), [27](#)
- Newton's third law, [32](#)
- Newtonian time, [62](#), [68](#)
- no cloning theorem, [38](#)
- non-negative real numbers, [xvii](#), [302](#)
- non-strict, [214](#), [219](#), [241](#)
- non-Zeno component, [261](#)
- non-Zeno tuple, [259](#), [261](#)
- nonblocking write, [107](#), [140](#)
- nonblocking writes, [111](#)
- nondeterminism, [197](#)
- nondeterministic merge, [146](#), [162](#)
- NonStrictDelay actor, [239](#), [251](#)
- Norton's dome, [28](#)
- Norton, John, [27](#)
- NP-complete, [187](#)
- NTP, [xvi](#), [67](#), [71](#), [75](#)
- numeric order, [310](#), [313](#), [319](#), [320](#), [326](#), [328](#)
- observation, [149](#)
- observational equivalence, [22](#)
- observer, [17](#)

- one-hot, [181](#), [182](#)
- one-to-one, [305](#), [319](#), [325](#)
- onto, [305](#), [319](#), [325](#)
- open neighborhood, [135](#), [136–138](#), [280](#), [285](#)
- open set, [135](#)
- operational semantics, [116](#), [159](#), [163](#), [171](#)
- order, [307](#)
- order embedding, [262](#), [319](#), [328](#)
- order isomorphic, [319](#)
- order isomorphism, [319](#)
- order preserving, [318](#)
- OrderedMerge actor, [133](#), [193](#)
- ordinals, [314](#)
- OS, [xvi](#), [55](#)
- output, [15](#)
- parallel, [54](#)
- parallel AND, [214](#), [222](#), [249](#)
- parallel OR, [213](#), [221](#), [237](#), [248](#)
- parallel or, [170](#)
- parallel scheduling, [185](#), [186](#)
- parameter, [108](#)
- parameterized SDF, [xvi](#), [202](#)
- Park, David, [22](#)
- Parks' algorithm, [130](#), [140](#), [195](#)
- partial execution, [316](#)
- partial function, [xvii](#), [303](#)
- partial order, [xviii](#), [309](#)
- partially ordered set, [309](#)
- partially synchronous networks, [59](#)
- partitioning, [59](#)
- PASS, [xvi](#), [182](#), [183](#)
- Pearl, Judea, [32](#)
- Penrose, Roger, [34](#)
- periodic admissible sequential schedule, [xvi](#), [182](#), [203](#)
- permuting elements of a tuple, [306](#)
- Petri nets, [197](#)
- Petri, Carl Adam, [197](#)
- physical time, [xviii](#), [70](#), [73](#)
- physical time interpretation, [xviii](#), [73](#), [257](#)
- pipes, [111](#)
- place, [197](#)
- PN, [xvi](#), [106](#)
- pointed, [156](#), [281](#), [282](#), [310](#), [313](#), [333](#), [334](#), [344](#)
- pointwise, [311](#)
- pointwise order, [156](#), [221](#), [223](#), [311](#), [320](#)
- pointwise prefix order, [118](#), [316](#)
- Poisson's hypothesis, [44](#)
- Popper, Karl, [26](#)
- port, [106](#)
- poset, [xviii](#), [156](#), [282](#), [309](#), [311](#), [328](#), [333](#), [334](#), [344](#)
- post-fixed point, [288](#)
- power systems engineering, [209](#)
- powerdomain, [147](#)
- powerdomain construction, [231](#)
- powerset, [xvii](#), [302](#), [304](#), [308](#), [310](#), [317](#), [328](#)
- powerset lattice, [317](#), [324](#)
- precision time protocol, [xvi](#), [71](#)
- predictability, [9](#), [22](#)
- prefix, [108](#), [309](#), [334](#)
- prefix order, [xviii](#), [112](#), [116](#), [309](#), [314](#), [318](#)
- prefix order on functions, [156](#)
- prefix order on signals, [282](#)
- PRET-C, [229](#)
- probability, [7](#)
- process, [55](#), [106](#), [209](#)
- process network, [xvi](#), [106](#), [211](#)
- processing offset, [53](#), [74](#)
- program, [54](#)

- projection, [xviii](#), [306](#)
- Proofreader, [xiii](#)
- PSDF, [xvi](#), [202](#)
- pseudo-random number, [23](#)
- Ptides, [25](#)
- Ptolemy II, [45](#), [72–74](#), [229](#), [260](#)
- PTP, [xvi](#), [71](#)
- publish-and-subscribe, [101](#)
- punctuated chaos, [19](#)
- pure, [228](#)

- quantum physics, [27](#), [37](#)
- quasi-static schedule, [192](#), [196](#)
- quasi-static scheduling, [186](#)

- radar, [190](#)
- radius, [280](#)
- Ramp actor, [244](#)
- randomness, [7](#)
- range, [304](#)
- rank, [180](#)
- rational numbers, [xviii](#), [312](#)
- RC, [229](#)
- reaction, [86](#), [238](#)
- Reactive C, [229](#)
- reactive system, [210](#)
- reactor, [20](#)
- reactor model, [242](#), [297](#)
- read event, [63](#), [75](#)
- real numbers, [xvii](#), [302](#)
- real-time constraint, [81](#), [85](#), [87](#)
- real-valued time, [257](#)
- reality, [40](#)
- receive event, [63](#), [75](#)
- reflexive, [309](#)
- relation, [303](#), [309](#)
- relational interpretation, [41](#)
- relativity, [34](#)

- remote procedure call, [101](#)
- rendezvous, [150](#), [209](#), [211](#)
- repeatability, [9](#)
- resolution, [73](#)
- restriction, [xviii](#), [282](#), [284](#), [306](#), [329](#)
- ring laser gyroscope, [36](#)
- Rovelli, Carlo, [41](#)
- Russell, Bertrand, [31](#)

- safety, [210](#)
- Sagnac effect, [36](#)
- SC, [229](#)
- SCCharts, [229](#)
- scenario, [202](#)
- Schrödinger equation, [37](#)
- scientific model, [6](#), [25](#)
- scientific positivism, [26](#)
- Scott open, [137](#)
- Scott topology, [135](#), [137](#), [138](#), [280](#)
- Scott, Dana, [137](#)
- SDF, [xvi](#), [172](#), [196](#)
- security, [21](#)
- Select, [191](#)
- Select actor, [124](#), [125](#), [127](#), [129](#), [191](#), [195](#),
[250](#)
- self-timed scheduling, [186](#)
- semantics, [112](#), [315](#)
- send event, [63](#), [75](#)
- sentence, [16](#)
- sequence, [108](#), [256](#), [307](#), [334](#)
- sequences, [xviii](#), [110](#), [307](#), [309](#)
- sequential, [140](#), [162](#), [170](#)
- sequential process, [63](#), [75](#)
- sequential program, [54](#)
- sequential schedule, [181](#)
- set, [302](#)
- set of all continuous functions, [xviii](#), [320](#),
[320](#)

- set of all functions, [304](#)
- set of events, [257](#)
- set of signals, [257](#)
- set subtraction, [xvii](#), [302](#)
- setup time, [86](#)
- shared variable, [56](#)
- sieve of Eratosthenes, [144](#), [145](#)
- Signal, [208](#), [210](#)
- signal, [257](#), [293](#)
- signal processing, [190](#), [209](#)
- simple path, [274](#)
- simulation, [22](#)
- simultaneous and instantaneous, [210](#)
- simultaneously and instantaneously, [209](#)
- singleton set, [xvii](#), [156](#), [159](#), [213](#), [265](#),
[303](#), [307](#), [308](#)
- sink, [126](#)
- sonar, [190](#)
- source, [126](#)
- Spanner, [25](#), [66](#), [68](#), [71](#)
- sparse synchronous model, [208](#), [242](#)
- speculative execution, [54](#)
- spherically closed ball, [285](#), [339](#), [340](#), [343](#)
- spherically complete, [285](#), [295](#), [340](#), [344](#)
- sporadic, [92](#), [96](#)
- SR, [172](#), [207](#)
- SR component, [208](#)
- stable, [142](#), [162](#), [170](#)
- stable functions, [111](#)
- standard topology, [136](#)
- state, [14](#), [164](#), [165](#), [187](#), [214](#)
- state transition function, [214](#)
- Statecharts, [210](#), [229](#)
- static assignment scheduling, [186](#)
- static dataflow, [172](#)
- stream, [107](#), [112](#), [256](#)
- StreamIt, [190](#)
- strict partial order, [64](#), [311](#)
- strict poset, [xviii](#), [311](#)
- strictly causal, [264](#), [265](#), [283](#), [294](#)
- strictly contracting, [263](#), [267](#), [287](#), [294](#),
[338](#), [340](#), [344](#)
- strong consistency, [76](#)
- strongly consistent, [67](#)
- subclocks, [239](#)
- subset, [xvii](#), [302](#)
- subset order, [309](#), [312](#), [317](#), [318](#)
- superdense time, [73](#), [257](#), [278](#), [281](#), [295](#)
- surjective, [305](#)
- Switch actor, [129](#), [195](#), [198](#), [250](#)
- symbolic execution, [182](#), [185](#), [199](#)
- symmetric multicore, [55](#)
- symmetry, [332](#), [335](#)
- SyncCharts, [229](#)
- synchronized keyword in Java, [209](#)
- synchronous, [209](#), [260](#), [271](#)
- synchronous block diagrams, [211](#)
- Synchronous C, [229](#)
- synchronous circuits, [208](#)
- synchronous dataflow, [xvi](#), [172](#), [203](#), [209](#)
- synchronous dataflow scenarios, [202](#)
- synchronous digital circuit, [211](#)
- synchronous language, [209](#), [210](#), [229](#)
- synchronous languages, [211](#)
- synchronous procedure call, [101](#)
- synchronous-reactive, [207](#), [210](#)
- synchroty hypothesis, [209](#), [252](#)
- Synopsys, [210](#)
- tag, [xviii](#), [73](#), [99](#), [257](#), [293](#)
- tag set, [257](#)
- Tarski's fixed-point theorem, [324](#)
- Tarski, Alfred, [324](#)
- Tarskian, [324](#)
- TCP, [xvi](#), [8](#), [12](#), [18](#), [25](#), [62](#), [107](#)

- teleport message, [190](#)
- thread, [56](#), [63](#), [146](#), [209](#)
- threat model, [21](#)
- throughput, [189](#)
- tick, [208](#)
- time, [68](#)
- time interval, [xviii](#), [70](#)
- time sensitive network, [xvi](#), [71](#)
- timespec, [72](#)
- timestamp, [66](#), [73](#), [73](#), [86](#), [257](#), [263](#)
- timestamp-ordered, [66](#)
- token, [107](#), [112](#), [155](#), [197](#), [256](#)
- token ring, [219](#)
- top, [317](#)
- topology, [135](#), [136](#)
- total function, [303](#)
- total order, [310](#), [311](#)
- totally ordered set, [73](#), [314](#)
- trace, [63](#), [75](#), [83](#), [99](#)
- transformational system, [210](#)
- transition in Petri nets, [197](#)
- transitive, [309](#)
- transitive closure, [90](#)
- transmission control protocol, [xvi](#), [8](#)
- triangle inequality, [332](#)
- truth values, [xvii](#), [302](#)
- TSN, [xvi](#), [71](#)
- tuple, [xvii](#), [303](#)
- tuple of signals, [259](#)
- Turing, Alan, [26](#)
- Turing complete, [23](#), [110](#), [185](#)
- Turing machine, [22](#), [127](#)
- Turing, Alan, [22](#)
- type, [110](#)
- UDP, [xvi](#), [12](#)
- ultrametric, [333](#), [337](#)
- ultrametric inequality, [333](#), [335](#)
- unavailability, [xviii](#), [67](#), [77](#), [80](#), [85](#), [99](#)
- unbounded execution, [174](#)
- unbounded list, [111](#)
- unbounded memory, [110](#)
- uncaused action, [7](#)
- uncountable, [327](#)
- undecidability, [26](#)
- undecidable, [110](#), [127](#), [195](#), [330](#)
- undirected cycle, [132](#)
- uniqueness of fixed points, [322](#)
- unit delay, [108](#), [110](#), [117](#), [121](#), [129](#), [133](#),
[163](#), [164](#), [270](#), [271](#)
- UnitDelay actor, [129](#), [133](#)
- universal Turing machine, [127](#)
- Unix time, [72](#)
- unknown, [xviii](#), [212](#)
- upper bound, [311](#)
- upper semilattice, [318](#)
- upper set, [135](#), [135](#), [137](#)
- upward complete, [135](#)
- user datagram protocol, [xvi](#), [12](#)
- UTC, [xvi](#), [67](#), [71](#)
- vector clocks, [66](#)
- vectorization factor, [188](#), [189](#)
- video processing, [190](#)
- von Neumann, [308](#)
- von Neumann numbers, [xviii](#), [308](#), [313](#)
- wave function, [37](#)
- WCET, [xvi](#), [86](#)
- well-formed composition, [232](#)
- When actor, [241](#), [244](#), [246](#), [248](#)
- White Rabbit, [71](#)
- wildcard, [162](#)
- Windows, [55](#)
- Wolpert, David, [26](#)
- worker thread, [172](#)

worst-case dependency, [82](#), [83–85](#), [88](#),
[100](#)

worst-case execution time, [xvi](#), [86](#)

write event, [63](#), [75](#), [80](#)

Xor actor, [193](#)

year 2038 problem, [72](#)

Zeno signal, [259](#), [276](#), [279](#), [292](#)

